
Release Information

Alation User Documentation

Oct 23, 2024

TABLE OF CONTENTS

1	Release Notes	1
1.1	Release Notes 2024.3 (General Availability)	1
1.2	Release Notes 2024.1 (General Availability)	18
1.3	Release Notes 2023.3 (General Availability)	66
1.4	Release Notes 2023.1 (General Availability)	102
1.5	Release Notes 2022.4 (General Availability)	143
1.6	Release Notes 2022.3 (General Availability)	159
1.7	Release Notes 2022.2 (General Availability)	177
1.8	Release Notes 2022.1 (General Availability)	190
1.9	Release Notes 2021.4 (General Availability)	204
1.10	Release Notes 2021.3 (General Availability)	223
1.11	Release Notes 2021.2 (General Availability)	236
1.12	Release Notes 2021.1 (General Availability)	252
1.13	Release Notes 2020.4 (General Availability)	263
1.14	Release Notes 2020.3 (General Availability)	279
2	Support Matrices	291
3	End of Support	293
3.1	Alation End of Support Matrix	293
4	Snowflake Column-Level Lineage Parser Add-On	299
4.1	Snowflake Column-Level Lineage 2021.4	299
5	Important Notices and Security Alerts	301

RELEASE NOTES

Release Notes prior to Alation version 2020.3 have been moved to the Documentation Archive. For Alation version 4, see [Archived Version 4 Release Notes](#). For Alation version V, see [Archived Version V Release Notes](#).

- [Release Notes 2024.3 \(General Availability\)](#)
- [Release Notes 2024.1 \(General Availability\)](#)
- [Release Notes 2023.3 \(General Availability\)](#)
- [Release Notes 2023.1 \(General Availability\)](#)
- [Release Notes 2022.4 \(General Availability\)](#)
- [Release Notes 2022.3 \(General Availability\)](#)
- [Release Notes 2022.2 \(General Availability\)](#)
- [Release Notes 2022.1 \(General Availability\)](#)
- [Release Notes 2021.4 \(General Availability\)](#)
- [Release Notes 2021.3 \(General Availability\)](#)
- [Release Notes 2021.2 \(General Availability\)](#)
- [Release Notes 2021.1 \(General Availability\)](#)
- [Release Notes 2020.4 \(General Availability\)](#)
- [Release Notes 2020.3 \(General Availability\)](#)

1.1 Release Notes 2024.3 (General Availability)

Note: The release notes on this page reflect changes to the Alation application. If you're looking for OCF connector release notes, jump to </sources/OpenConnectorFramework/ReleaseNotes>.

1.1.1 PATCH RELEASE 2024.3.1.3

build 19.1.3.90272

Released October 18, 2024

Note: This patch release is available for Alation Cloud Service instances on the cloud-native architecture. This patch release doesn't apply to customer-managed (on-premise) instances and Alation Cloud Service instances that haven't been migrated to the cloud-native architecture.

- Fixed an issue where Datadog metrics were not being sent or received across multiple clusters and tenants. The update ensures that metrics and logs are now properly handled. (AL-176268)

1.1.2 PATCH RELEASE 2024.3.1.2

build 19.1.2.89681

Released October 16, 2024

Note: This patch release is available for customer-managed (on-premise) instances only. This patch release doesn't apply to Alation Cloud Service instances.

- Fixed an issue where users experienced performance degradation during metadata extraction and query log ingestion (QLI), characterized by backed-up task queues, QLI executions missing the schedule, and high CPU usage by Celery processes. The performance issues were traced to a specific method that was called frequently, each time taking 0.1-1 second per call and slowing the system. We implemented optimizations to decrease the frequency and impact of database calls made by this method, enhancing the QLI performance. (AL-174546)

1.1.3 PATCH RELEASE 2024.3.1.1

build 19.1.1.89401

Released October 11, 2024

Note: This patch release is available for Alation Cloud Service instances on the cloud-native architecture. This patch release doesn't apply to customer-managed (on-premise) instances and Alation Cloud Service instances that haven't been migrated to the cloud-native architecture.

- Addressed an issue where Server Admins of Alation Cloud Service instances were receiving excessive email notifications for missed Celery jobs, occurring every 30 minutes. We've introduced a feature flag **Enable Email Alerts for Celery Jobs** under Admin Settings > **Feature Configuration** to manage such email notifications. This feature flag is disabled by default and can be enabled as needed. When enabled, it triggers hourly and daily checks, notifying admins of critical issues like missed jobs and task queue backlogs in the Celery queues ['default', 'ingestion', 'run_scheduling', 'stewardship']. (AL-174801)
- Fixed an issue where the watchdog process responsible for backing up site data files failed to start during the provisioning of new tenants. (AL-174982)

1.1.4 RELEASE 2024.3.1

build 19.1.0.88821

Released October 03, 2024

Note: This release is available for all Alation instance types: customer-managed (on-premise) and Alation Cloud Service.

In this release:

- [New User Experience Updates](#)
- [Lineage](#)
- [Compose](#)
- [Search](#)
- [Platform](#)
- [Improvements](#)
- [Bug Fixes](#)

New User Experience Updates

Alation Cloud Service instances on the cloud-native architecture

- Additional customization options for Server Admins:
 - Ability to configure primary and secondary application colors in Admin Settings > **Customize Branding**
 - Ability to hide the **Switch to New Experience/Switch to Classic Experience** button using the new feature configuration flag **Remove the Back to Classic button in the new UI** in Admin Settings > **Feature Configuration**
- More enhancements to the *wayfinder* (the left-side navigation panel):
 - Automatically hiding the objects that haven't been fully configured for non-administrator users to prevent confusion
 - NoSQL objects have been incorporated under the Data Sources hierarchy
 - Stability improvements across several object types
- Miscellaneous other updates:
 - ALLIE-generated titles can be reviewed from the parent object pages
 - Mentions within rich text fields are now easier to trigger
 - Object set field values in the main panel now support pagination
 - The Filters tab now indicates if a filter is required or not
 - The Source Comments field formatting is now in parity with the Classic Experience
 - The data type STRUCT on the column catalog page is now included
 - The **About this Instance** page now displays the Frontend build number reflecting the current version of the New User Experience

Find more details in [New User Experience](#).

Lineage

Custom Field Metadata Layers (Beta)

Server Admins and Catalog Admins can now configure custom fields to appear as metadata layers in lineage charts. Users can select custom field metadata layers on the Compound Layout lineage chart view from the **Metadata Layer** list. When applied to the chart, the layer displays the corresponding custom field value for the lineage nodes that have this field. Custom field metadata layers can be created in Admin Settings > **Lineage Settings** > **Custom Field Label**. (AL-165274)

Note: The feature is controlled by the **alation_conf** parameter `alation.feature_flags.dev_enable_lineage_custom_field_metadata` and is enabled by default.

Bulk-Import Lineage Information (Beta)

The issues with this beta feature that were identified in version 2024.3 have now been addressed:

- The upload template can now be downloaded on Alation Cloud Service instances.
- On newly created instances, there are no more issues with uploading the template.

The feature is enabled by default on the Alation Cloud Service instances. On customer-managed instances, the feature can be enabled by setting the **alation_conf** parameter `alation_conf alation.feature_flags.dev_enable_lineage_bulk_upload` to true.

Lineage between RDBMS, BI, and file system objects can be uploaded using the template. External objects are currently not supported. The row limit for uploading is 200.

Deprecation Warning Propagation Worker Functionality

Alation Cloud Service instances on the cloud-native architecture

We've refactored the legacy deprecation warning propagation service for lineage actions into a new service. This update enhances the handling of deprecation-specific propagation events. Key improvements include:

- Automatic adjustment of deprecation propagation when view definitions change. If a view definition changes for a particular downstream object that was along the deprecation path from an upstream object, then deprecation propagation will adjust to not highlight the view.
- Consistent addition and removal of propagated deprecation warnings across table, column, data source, and schema-level objects.
- Dynamic updates to the lineage chart when dataflows are added or removed, reflecting accurate deprecation warning paths. (AL-154463)

Compose

Queuing in Compose with Query Service

For data sources where Query Service is enabled, we've implemented a concurrency control queue for query executions done within the same connection. This ensures that for a particular connection, we can queue multiple query executions without the *Concurrent queries cannot be executed* error. Currently, this feature is disabled by default, and will be enabled by default in a later release. To activate it now, append the JVM flag `-Dconnector.use_concurrent_execution_per_connection=true` preceded by a space to the existing values of the parameter `connector.extra_flags` in **alation_conf**. Previously, users would encounter an error in the results pane when attempting to run multiple queries on the same connection while another was still processing. With the new concurrency control, users can submit several queries simultaneously. These queries will be queued and executed serially in the background, similar to how queries are managed when exceeding the maximum concurrent query limit. (AL-166099)

Note: Alation Cloud Service customers can request server configuration changes through Alation Support.

Search

Exact Search Without Double Quotes

For search strings without double quotes, whether they contain underscores or not, titles or names that exactly match the search query will now appear at the top of the search results. Previously, search strings with underscores were broken into tokens and ranked based on the scores of each token, which did not prioritize exact matches. (AL-171164)

Platform

PostgreSQL Upgrade to 16.4

Customer-managed instances

The internal PostgreSQL databases (the internal Alation database and the Alation Analytics database) have been upgraded from version **16.2** to **16.4** to enhance application security. The databases will be upgraded automatically during the Alation update to 2024.3.1. (AL-169856)

Ability to Enable SSL for the Kafka and Zookeeper Components

Customer-managed instances

Alation now supports enabling SSL/TLS encryption for Kafka and SASL over TLS authentication for Kafka clients and Zookeeper. Even though all services, producers, and consumers run on the Alation server and use localhost for connections, this enhancement addresses potential security concerns. (AL-122489)

Note: The feature is disabled by default and can be enabled by the instance administrator on the server backend. Steps to enable:

```
sudo su alation
rm -rf /var/lib/event_bus/kafka-ssl/*
alation_conf event_bus.kafka.ssl.enabled -s True
alation_action kafka-generate-ssl
alation_action deploy_conf_kafka_ssl
alation_supervisor restart event-bus:*
alation_supervisor restart lineage
```

Java Component Upgrade

The internal Java component has been upgraded to a newer version (Corretto 11.0.24.x.x) to incorporate the latest security fixes and enhancements. (AL-164803)

Enhanced Audit Logging for Database Connectors, File System Connectors, and Custom Field Permissions

The user actions on the data source settings page are now audited and the corresponding logs are stored in **alation-event.log** in the **/opt/alation/site/logs** directory. The log captures information about the user who triggered the action, when it was triggered, what was the outcome of the action, what was the previous value of the field before the user acted on it, and what is the new value changed by the user. For example, if a user changes the URI for a source, this action will be captured by the log. (AL-123286, AL-168788, AL-168620)

Improvements

- Optimized the lineage compound layout to eliminate redundant calculations. This enhancement will significantly improve chart load times. (AL-170941)
- Implemented support for non-conforming session IDs in the QLI pipeline service, which now strips extraneous details such as usernames and timestamps from the session ID in custom MDE queries. This correction was implemented for the Snowflake OCF connector and will be extended to all other OCF connectors in the future. (AL-168154, *Alation Cloud Service instances on the cloud-native architecture*)
- The **alation_conf** flag `enable_obfuscating_view_sql` is now set to `true` by default. This ensures that when the **Obfuscate Literals** option is enabled for a data source, literals in the extracted view SQL will be automatically obfuscated. (AL-163780)
- Updated the `scan_postgres` utility:
 - Integrated `scan_postgres` with `pg_amcheck` for improved performance and reliability
 - Added progress indicators
 - Enhanced efficiency for larger databases (AL-164467, *Customer-managed instances*)
- Previously, some customers experienced congestion in Celery Worker resources due to a large number of catalog set-related tasks getting queued. We've implemented a rate limit on the affected tasks to deduplicate tasks with identical arguments (such as the catalog set ID). Due to rate limiting, only one unique task will be queued and executed within a configured interval. This approach ensures more efficient use of resources and reduces task backlog. Rate limiting was added for the `sync_members_from_catalog_set` and `update_dynamic_sets` tasks by debouncing them within a specified interval. The interval is controlled by the **alation_conf** flags `stewardship.sync_members_from_catalog_set.debounce_interval` and `stewardship.update_dynamic_sets.debounce_interval` respectively, with

a default interval of 600 seconds. During the debounce period, updates to catalog set members will not reflect in these features if the task is rescheduled within this time.

Impact on features:

- The `sync_members_from_catalog_set` debounce period: Delays the synchronization for the Catalog Activity and Mentioned On features for Documents and Folders by the debounce interval.
- The `update_dynamic_sets` debounce period: Delays the synchronization for Catalog Sets Membership and updates to Catalog Sets Shared Fields by the debounce interval. (AL-171500)

Bug Fixes

- Fixed an issue where the bulk upload of lineage with an Excel template failed on newly installed instances. (AL-172593)
- Fixed an issue where objects that were @-mentioned in rich text fields displayed only icons, without titles. After this fix, the titles of referenced objects are displayed as expected. (AL-172065)
- Fixed an issue where characters like the ampersand & appeared in their HTML-escaped form in catalog set titles on catalog pages. (AL-128763)
- Previously, any query execution error incorrectly included an additional message: *If unable to reconnect, please validate if your Hive Server is up and running properly.* This issue has been resolved. Now, only the syntax error message relevant to the current database type is displayed. (AL-169746)
- *Alation Cloud Service instances on the cloud-native architecture:* Fixed a bug where published queries from Compose were sometimes not shown in the catalog under the table catalog page > **Queries**. (AL-152405)
- Fixed an issue where some data dictionary-related tasks were stuck in *processing* or *queued* states due to concurrent invocations. To address this, concurrent API requests on the same resource will now be blocked and return a 400 error status, preventing tasks from getting stuck. Additionally, a daily scheduled task has been implemented to scan the database for any data dictionary tasks running for over 24 hours. The expiration threshold can be adjusted with the **alation_conf** flag `alation.data_dictionary.task.pending.seconds_to_keep`. (AL-168954)
- *Alation Cloud Service instances on the cloud-native architecture:* In the New User Experience, fixed an issue where a generic error message was shown when a data dictionary export was already in progress. Implemented status code checks to improve error messaging. (AL-168794)
- *Alation Cloud Service instances on the cloud-native architecture:* Fixed a pagination issue in the New User Experience, ensuring that object set values within template main sections now match the pagination behavior of the Classic Experience. (AL-165421)
- *Customer-managed instances:* Fixed an issue where the health check monitor URL (**<Base_URL>/monitor/health_check/**) previously took about 2 minutes to display results. Health checks now load in less than 3 seconds on average. (AL-168548)
- Fixed an issue with the Kafka processes using plain text listeners. We've introduced the ability to enable encryption for both Kafka and Zookeeper processes, enhancing security. (AL-143490)
- Fixed an issue with the Snowflake OCF connector with the dbt add-on connector where metadata extraction (MDE) failed with the error *Unrecognized token 'version': was expecting (JSON String, Number, Array, Object or token 'null', 'true' or 'false')* when the manifest file was large (>100MB). Added a fix to get large manifest files from the correct raw GitHub URL. (AL-167304)

- Fixed a bug where using the public API to add or remove domain memberships (**integration/v2/domain/membership/**) with 100 or more objects recursively resulted in a 500 error and a code exception. Updated the API exception handling logic to include the limit argument, ensuring that a correct 400 response with a detailed message is returned. (AL-165401)
- Fixed an issue where `max(last_processed_record_ts)` for `bi_report` and `documenthub` in Alation Analytics showed a future date due to an ETL bug that incorrectly set the timestamp when the ETL checkpoint was reset. This fix ensures accurate timestamp setting during the ETL process. (AL-164137)
- Fixed an issue where a user who was provided access to the Alation Analytics data source, was suspended before they could create their password from the **Account Settings > DB Connections** tab. With this fix, only those users whose credentials already exist can be suspended. (AL-162736)
- Fixed an issue where the `ts_updated` value for custom fields wasn't reflected in the Alation Analytics database. Now, the `ts_updated` field for all the object types will be updated in case a custom field associated with an object is updated in the catalog. (AL-161824)
- Fixed an issue where Alation Analytics wasn't capturing deleted fields for data sources, schemas, tables, and columns. Updated the query used by the Analytics ETL to ensure we capture records which have been deleted. (AL-155423)
- Fixed the Analytics ETL to ensure `ts_deleted` is reflected correctly for Policy tables. (AL-170039)
- Fixed a bug in the loading of the `curation_history` data into Alation Analytics, which caused some records to be missed. (AL-165866)

1.1.5 PATCH RELEASE 2024.3.0.1

build 19.0.1.87345

Released September 20, 2024

Note: This patch release is available for Alation Cloud Service instances on the cloud-native architecture. This patch release doesn't apply to customer-managed (on-premise) instances and Alation Cloud Service instances that haven't been migrated to the cloud-native architecture.

- Fixed an issue where objects that were @-mentioned in rich text fields displayed only icons, without titles. After this fix, the titles of referenced objects are displayed as expected. (AL-172065)

1.1.6 RELEASE 2024.3.0

build 19.0.0.85267

Released September 3, 2024

Note: This release is available for all Alation instance types: customer-managed (on-premise) and Alation Cloud Service.

In this release:

- [New User Experience Updates](#)
- [Alation Anywhere](#)
- [Homepage](#)

- [Alation Analytics](#)
- [Compose](#)
- [Lineage](#)
- [Search](#)
- [Authentication and User Management](#)
- [Governance](#)
- [Open Connector Framework](#)
- [Bug Fixes](#)
- [Known Issues](#)

New User Experience Updates

Alation Cloud Service instances on the cloud-native architecture

Ability to Remove the Switch to New Experience/Switch to Classic Experience Button

We've introduced a new feature flag, **Remove the Back to Classic button in the new UI**, which allows Server Admins to hide the **Switch to New Experience/Switch to Classic Experience** button from the user interface. Enabling this feature establishes the New User Experience as the user interface without the ability to go back to the Classic Experience. Pages that haven't yet been converted to the New Experience will still be presented in the Classic Experience. (AL-164824)

Alation Anywhere

Alation Cloud Service instances on the cloud-native architecture

Public Preview of Alation Chrome Extension (Beta)

The Alation Chrome Extension enhances browser functionality by allowing Alation users to search and preview data catalog objects directly in Chrome. After installing the extension from the [Chrome Web Store](#), users can log into their Alation cloud instance, conduct quick searches, and view metadata previews. Snowflake users get additional benefits from this integration. After signing into Snowflake Snowsight, users can access detailed metadata about schemas and tables directly through the extension, enriching their data context with metadata from Alation, without leaving Snowflake. (AL-159629, AL-159628)

Homepage

Dynamically Rearrange Custom Blocks on the Homepage

Alation Cloud Service instances on the cloud-native architecture

We've introduced a dynamic user interface enhancement, enabling the flexible positioning of custom blocks on the homepage. Server Admins and Catalog Admins can now rearrange homepage blocks by dragging-and-dropping them. Save the changes for the updates to be applied to the homepage. (AL-153691)

Alation Analytics

Group Access Control for Alation Analytics

Data Source Admins now have the capability to manage access to the Alation Analytics data source through groups. On the **Access** tab of the Alation Analytics data source settings, you can add or remove entire groups and adjust their access levels. Previously, access had to be granted to users individually, which could become inefficient with the increasing number of Alation Analytics users. (AL-82666)

Compose

Scheduled Query Execution Notifications for Allowed Users and Groups

Alation now sends notifications on scheduled query execution to the users and groups with whom the results have been previously shared. To opt out, users can disable all scheduled query alerts via their profile preferences. The feature is controlled with two **alation_conf** flags:

- **notify_allowed_users_and_groups**: Enables or disables sending notifications to users and/or groups listed in the **Scheduled Query Setting > Sharing Settings** dialog, with whom the results have been shared previously. Enabled by default.
- **notify_users_in_groups**: Enables notifications directly to individual users in a group specified in the **Scheduled Query Settings > Sharing Settings** dialog. Disabled by default, and the specified group is notified through the group email. (AL-168820)

The Sapphire Compose Theme

The Sapphire Compose theme is designed to complement the New User Experience. It will be automatically applied for all users after an admin enables the New User Experience feature flag. This ensures a consistent aesthetic across all Alation interfaces. If users opt to switch from the Sapphire theme to another, the auto-application of the theme will be disabled, and subsequent themes will require manual selection. (AL-159192)

Note: The Sapphire theme is also available on customer-managed (on-premise) Alation instances.

Differentiate Between Launching and Running a Query

We've introduced a new status indicator for running queries to differentiate between launching and running a query. Previously, the status message *Launching query* was displayed for the duration of the query execution. Now, it'll change to *Executing query* once the database connection is established. This enhancement is intended to help discern the stage of query processing and quickly spot any connection delays. (AL-21736)

Change to Default Setting for Automatic Connections

The default value for the **alation_conf** parameter `allow_automatic_connection_on_tab_open` has been changed to `false`. This update disables the automatic creation of connections to the datasource when users open a new tab or switch tabs in Compose, causing database account lockout problems when passwords have expired. (AL-165067)

Change to Default Setting for DDL Query Processing

We've updated the default behavior for processing of Data Definition Language (DDL) queries from Compose. Previously, the incremental extraction and lineage generation for DDL queries ran synchronously. Now, the **alation_conf** parameter `alation.query_exec.ddl.run_callbacks_async` setting is enabled by default, and these processes will run asynchronously. This update aims to enhance performance and reduce the time it takes for queries to complete. (AL-136855)

Deprecation of Desktop Compose

The Desktop Compose installer is no longer available for download from Alation, and the option to select Desktop or Web as the default Compose user experience has been removed from admin and user settings. Existing Desktop Compose installations will continue to function, but users are encouraged to transition to the Compose Web interface that is accessible via a regular browser. All queries and settings like theme and font size are preserved in the Web interface, but open tabs and connections will need to be re-established. Critical support for Desktop Compose will cease by August 2025. (AL-154706)

Lineage

Lineage Support for SQL Server Stored Procedures

Alation Cloud Service instances on the cloud-native architecture

We've enhanced lineage support for cataloged stored procedures and functions for SQL Server OCF connectors. It is available for systems on the cloud-native architecture and the `mt_lineage` service enabled. To activate this feature, the **alation_conf** parameter `alation.ocf.mde.rdbms.parse_function_ddl.enabled` should be set to `true`. This enables the detailed parsing and lineage tracing of function and procedure definitions. (AL-156708)

Known Issue with SQL Server Lineage from Stored Procedures

The lineage chart may not be accurate when queries include columns with unqualified names. (AL-170059, AL-170264)

Note: Alation Cloud Service customers can request server configuration changes through Alation Support.

Lineage Minimap Available by Default

Alation Cloud Service instances on the cloud-native architecture

The lineage minimap feature in lineage charts no longer requires enablement through a feature flag and is available by default. (AL-168144)

Bulk-Import Lineage Information (Beta)

Users with the Data Stewards, Data Source Admins, Catalog Admins, and Server Admins roles can now bulk-import lineage information via an Excel template. This feature allows users to easily identify the success or failure of each submission through a detailed upload report. The feature is currently in beta and can be enabled with the **alation_conf** parameter `alation.feature_flags.dev_enable_lineage_bulk_upload`. (AL-159481)

Note: UPDATE

September 20, 2024

We have identified a few issues affecting the functionality of this beta feature in version **2024.3**:

- The Excel template cannot currently be downloaded on Alation Cloud Service instances.
- On newly created instances, there are issues with uploading the template.

The Engineering team is actively working on resolving these issues to ensure full functionality in the next version, **2024.3.1**, planned for early October 2024. In the meantime, we advise refraining from using this beta feature. For any concerns or immediate needs, please contact the Alation Support team.

Lineage Settings User Interface for RDBMS OCF Connectors (Beta)

Data Source Admins can now use a separate user interface for configuring cross-system lineage. The feature is currently in beta and can be enabled with the **alation_conf** parameter `alation.feature_flags.DEV_enable_etl_elt_lineage_settings`. (AL-152599)

Note: Alation Cloud Service customers can request server configuration changes through Alation Support.

Search

Struct Column Visibility in Search

For data sources containing STRUCT data types, the visibility of STRUCT columns can now be managed via **Per-Object Parameter** settings for schemas and tables. When a schema or table is configured as non-browsable, all STRUCT columns within that object will automatically be hidden. This ensures that these columns are no longer searchable in Alation when their parent objects are set to non-browsable, providing better access control. (AL-152087)

Authentication and User Management

User Profile Upload Limited to 25 Updates

We've introduced new limits for uploading user profiles. Each CSV can now contain no more than 25 user creations or updates. If a CSV contains more than 25 rows but the total number of user creations and updates does not exceed 25, the upload will proceed. Details about this restriction are available in the info panel in **Admin Settings > User Profiles**. Uploads exceeding this limit will show a preview but disable the **Confirm** button, with an error message displayed below the preview area. (AL-166371)

Bypassing SAML Authentication Is Disabled

The mechanism to bypass SAML authentication has been disabled for administrators of both Alation Cloud Service on the cloud-native architecture and on-premise deployments. In emergency situations, instance admins will need to contact Alation Support to request a temporary bypass. (AL-127742)

Governance

Improvements to Domain Membership Rules and Performance

Catalog Pages will now always accurately reflect the object's Domain membership, without the lag time that was previously seen. The rules for domain propagation have changed. Previously, exclusions of child objects from Domains were very persistent:

- An exclusion rule was permanent unless an object was added to the Domain directly.
- When a parent object was included recursively, then child objects were only added to the Domain if no parent objects at any level were excluded.

Now, the most recent action takes precedence. For example, if a user adds a Schema recursively this week, it will override a previous Domain exclusion on a child Table from last month. This change is designed to align more closely with user expectations, ensuring that catalog pages immediately and accurately reflect an object's Domain membership. (AL-166059, AL-166060)

Handling of People Set Reviewer Deletion in Workflows

We've updated how existing workflows handle the removal of People Sets designated as reviewers. If a People Set set as a Reviewer is removed from the system, it will now be marked as **[Gone]** in the Reviewer field to indicate its absence. (AL-164891)

De-duplication of Snowflake Tags Enabled by Default

The **alation_conf** parameter `alation.feature_flags.enable_centralized_tags_deduplication` is now enabled by default. Snowflake tags from a common database or schema that are used across other databases or schemas will now be de-duplicated by default during extraction. (AL-166048)

Deprecation of the Suggested Stewards Feature

The Steward Suggestions feature is being deprecated. While customer instances with this feature currently enabled can still view, approve, or reject steward suggestions, the option to enable it through the feature flag in **Admin Settings > Feature Configuration** has been removed. However, the corresponding **alation_conf** parameter (`alation.feature_flags.enable_allie_steward_suggestions`) remains available, allowing for the feature to be disabled if necessary. (AL-168008)

Open Connector Framework

OCF Connectors

Updates to the Open Connector Framework (OCF) connectors are released every week. See </sources/OpenConnectorFramework/ReleaseNotes> for information on the updates in the month of **August 2024**. A few August highlights are:

Databricks Unity Catalog OCF Connector Improvements

We have released version **3.0.0** of the Databricks Unity Catalog OCF connector that enhances its metadata extraction capabilities. This update includes switching from a driver-based extraction method to a more robust query-based extraction from the `information_schema`. (AL-157004)

Enable Token Refreshing for Google BigQuery Compose OAuth

We've added the capability to auto-refresh OAuth tokens via the Compose authentication flow for Google BigQuery data sources. Now, when Data Source Admins enable the **Request Refresh Tokens** checkbox on the **Compose** tab of the settings page, Alation will auto-refresh tokens for connections established in Compose. (AL-162370)

Improved Workspace Extraction from Azure Power BI

Version **2.4.0** of the Power BI Scanner connector introduces an option to Extract all workspaces from Azure Power BI within the Customize additional extraction scope section on the **Metadata Extraction** tab.

- **Default Behavior:** With the **Extract all workspaces from Azure Power BI** toggle turned off, only workspaces accessible to the configured service principal will be extracted.
- **Enhanced Extraction:** When the toggle is turned on, all workspaces from Azure Power BI will be extracted, regardless of the service principal's access permissions.

Additionally, this version supports lineage from Redshift native query-based datasets and Snowflake-based datasets. (AL-168903)

Bug Fixes

New User Experience

Alation Cloud Service instances on the cloud-native architecture

- Previously, if a policy title consisted only of multiple spaces, it would display as blank and would be non-editable. Now, titles composed of multiple spaces will be displayed as *Untitled <object type>*. (AL-163248)
- In the Policy Center, the policy group membership dialog previously showed the **Save** button enabled by default, even if no groups were selected. We've updated this behavior so that the **Save** button will only become enabled when a user selects or deselects any policy group. (AL-162835)
- Previously, when the **Enable New UI - Default and UI Promotion access** flag was enabled, links on some pages under **Govern** would still redirect to pages on the Classic Experience. Now, all pages will correctly redirect to New User Experience when it is set as the default behavior. (AL-163421)
- Fixed an issue with the Title field validation when creating new workflows. (AL-163548)
- Addressed an issue that occurred when attempting to edit workflows associated with custom fields that had been deleted from the system. Previously, trying to edit such workflows would result in an error in both Classic and Neo Experiences. Now, instead of blocking the user, the system will display the label **[Gone]** for the deleted custom field within the workflow. (AL-166130)
- Previously, when users were assigning Reviewers to workflows, the Reviewer list failed to display all existing People Sets. This fix ensures that all People Sets are now visible when adding a Reviewer to workflows. Additionally, this issue was fixed in the bulk actions and catalog activity features, ensuring consistent functionality. (AL-166710)
- Addressed an issue where changes to the order of columns in the Document table were not persisting after users attempted to swap them and then save. (AL-163769)
- Fixed an issue where a homepage name disappeared from the list of homepages in **Admin Settings > Homepage Customization** when the homepage was disabled and then enabled and saved or enabled and then disabled enabled and saved. Now, the homepage name remains visible in the list after enabling or disabling. (AL-167711)
- Fixed an issue with adding terms to custom blocks on the homepage where adding a term, folder, or document to a custom block on any homepage resulted in an unexpected error. These objects can now be successfully added to custom blocks without errors. (AL-167993)
- Fixed an issue where logo updates done on the **Admin Settings > Homepage Customization** page weren't reflected on the homepage despite a message indicating a successful save. The issue has been resolved, and the logo now updates correctly in the top left corner after saving. (AL-168163)
- Fixed an issue where the Column Organization widget was missing from the Policy Center in New User Experience. The Column Organization feature is now available for policy pages. (AL-162838)

OCF Platform and Connectors

- Fixed a bug in the OCF downstream job for tags that previously led to errors when handling missing tags. The issue caused deleted tag links in the source to remain uncleared in Alation. With this fix, the system now correctly cleans up missing tag links. (AL-166880)
- Enhanced the ingestion process for Snowflake tags to address an issue where tag values containing newline characters were not being handled properly. This fix ensures that all options for the Snowflake tag picker field are populated correctly. (AL-164875)
- Updated the SQL Parser grammar type for the Athena OCF connector to address issues with query parsing. This enhancement improves the accuracy of parsing and ingesting queries from Athena into Alation. (AL-167009)
- Resolved an issue where the `GET_DDL()` function was called indiscriminately for procedures and functions in Snowflake, leading to extraction query failures and multiple errors. This update enhances the stability and efficiency of metadata extraction from Snowflake data sources. (AL-163439)
- Fixed an issue in the Snowflake OCF connector data sampling process. Previously, when the maximum rows to store was set to 0, the table row count was incorrectly displayed as 0, rather than reflecting the actual row count. This has been corrected in connector version **2.4.4** to ensure that the row count accurately represents the data in the table, regardless of the max rows setting. (AL-162276)
- Resolved a limitation in the Tableau OCF connector where Excel sheets embedded from OneDrive weren't appearing on the **Connections** tab in Alation. We have now added support for the creation of connections for cloud file types with connector version **1.9.5**. (AL-161089)
- Fixed an issue with storing `STRUCT` columns in the Alation internal database as NoSQL attribute objects where some of them were duplicated as could be discovered through Search. Duplicating is the result of the internal logical model for NoSQL data types. After this fix, duplicate objects will no longer be discoverable by users, as they only service the internal model. (AL-147379)
- Fixed an issue where the impact analysis report for a Snowflake data source resulted in empty Path values for BI objects. With this fix, the correct path information will be displayed for BI objects in both Impact Analysis and Upstream Audit reports, ensuring lineage traceability. (AL-164388)
- Resolved an issue after upgrading to version **2023.3.5.1**, where the **connector.log** file was missing and older connector log files were not being compressed using gzip. This was due to an internal process incorrectly removing the connector from the log path. A conditional check has now been implemented to prevent this from happening. (AL-161279)

Governance

- Resolved an issue in the Stewardship Workbench **Bulk Actions** feature, where the **Select All** option incorrectly applied actions to more objects than displayed in the search results. Now, when using the **Select All** option, only the objects shown in the search results will be updated. (AL-163762)
- Optimized the internal APIs behind the workflow step execution after observing slow response times in logs. This enhancement reduces the overall execution time of the APIs. (AL-161158)
- Corrected a typo in the spelling of **Catalog Activity** on the Curate and Govern landing page in the Classic Experience. (AL-164324)
- Resolved an issue that prevented users from editing existing workflows, which was caused by the system attempting to fetch all users and groups. The solution now bypasses the user or group fetch API if no unique users or group profiles are added as reviewers. This change eliminates unnecessary data retrieval, ensuring that workflow edits can be made efficiently. (AL-164472)

- Fixed an issue causing data dictionary uploads to fail with a 500 internal server error. The problem was traced back to extra double quote characters that were introduced by a third-party CSV editor tool. The fix handles these characters properly, ensuring successful uploads. (AL-165445)
- Resolved an issue in the Bulk Utility where uploading XLSX files with terms resulted in a 504 time-out error. (AL-161711)
- We've introduced the capability to schedule a backend job that recomputes the members of all conditional catalog sets, which is useful for addressing an issue where previously deleted objects that were restored via an MDE job, did not appear in conditional sets with matching rules. This job is disabled by default and can be enabled with the feature flag `stewardship.DEV_recompute_conditional_catalog_sets_members.enabled`. You can customize the schedule by modifying the following feature flags according to Celery crontab standards:
 - Day of week: `stewardship.DEV_recompute_conditional_catalog_sets_members.schedule.day`
 - Hour: `stewardship.DEV_recompute_conditional_catalog_sets_members.schedule.hour`
 - Minute: `stewardship.DEV_recompute_conditional_catalog_sets_members.schedule.minute`

Changes to these scheduling flags also require a restart of the `celery-beat` service to take effect. We recommend contacting Alation Support if you observe this issue and need to adjust the flags. (AL-168175)

Alation Analytics and Analytics Experience

- Resolved a synchronization issue between the catalog and the Alation Analytics database regarding deleted fields. Although fields marked as deleted in Rosemeta were accurately updated, this change was not reflected in Alation Analytics. This was corrected by updating the ETL query to ensure it now properly includes records marked as deleted. (*Alation Cloud Service instances on the cloud-native architecture*, AL-155423)
- Fixed an issue where Steward updates in term objects made through **Bulk Actions** were not appearing in Alation Analytics. The solution includes a migration script in the upgrade package that resets the ETL checkpoint for terms. Consequently, the ETL process for terms will restart from the beginning after the update. Given Alation Analytics' ETL capacity to process only six months of data per ETL run and depending on the volume of terms data, multiple ETL cycles may be necessary to fully update the terms data. This could result in a delay in reflecting the latest terms changes. (AL-158292)
- Fixed an issue in Alation Analytics that caused errors during data insertion into Snowflake when non-ASCII characters were present in custom fields. (*Alation Cloud Service instances on the cloud-native architecture, Analytics Experience*, AL-147037)

Other Fixes

- Resolved an issue where Search results filters failed to populate additional filter data when any documentation object type was used as the object filter. Additionally, deleted folders will no longer appear in full search results. (AL-160582)
- Addressed an issue in the `LMSTitleDescriptionSerializer`, where each call to a specific method triggered a separate query execution. This problem was present when the Annotation Value Service was enabled. To optimize performance, unnecessary calls within the `LMSTitleDescriptionSerializer` have now been removed. (*Alation Cloud Service instances on the cloud-native architecture*, AL-164402)

- To enhance application security we have now restricted the uploading of SVG files for articles. (AL-135045)
- Fixed an issue with sending alerts to admin users. The system now performs status checks on Celery jobs every 30 minutes. If a job missed its scheduled run or if any Celery queues become excessively backed up, an email notification will be sent out. (*Alation Cloud Service instances on the cloud-native architecture*, AL-156446)
- Fixed an issue where emails were not being sent to group emails when posts were made in conversations. This update ensures that all conversation assignees, including groups, now receive email notifications. (AL-159349)
- Previously, users had no option to return to the original object page from the data dictionary upload page except by using the browser's Back button. Now, users can easily navigate back to the data source from the upload dictionary page using a breadcrumb trail. (AL-152150)
- Added support for displaying partition key information in the table view of Alation Search. The fix applies to any data source that supports partition keys and includes partition key information in its extracted metadata. (AL-159752)

Known Issues

- *Alation Cloud Service instances on the cloud-native architecture*: In version **2024.3**, objects that are @-mentioned in rich text fields display only icons, without titles. Such objects as articles, terms, documents, and domains are affected. Please note, the titles are intact in the internal database but are not appearing in the user interface. A fix has been provided with [PATCH RELEASE 2024.3.0.1](#). (AL-172065)
- Conversations that are deleted and subsequently restored are currently not appearing in the Inbox. (AL-169183)

1.2 Release Notes 2024.1 (General Availability)

Note: The release notes on this page reflect changes to the Alation application. If you're looking for OCF connector release notes, jump to </sources/OpenConnectorFramework/ReleaseNotes>.

1.2.1 PATCH RELEASE 2024.1.5.5

build 18.5.5.89300

Released October 14, 2024

Note: This patch release is available for customer-managed (on-premise) instances only. This patch release doesn't apply to Alation Cloud Service instances.

- Fixed an issue where users experienced performance degradation during metadata extraction and query log ingestion (QLI), characterized by backed-up task queues, QLI executions missing the schedule, and high CPU usage by Celery processes. The performance issues were traced to a specific method that was called frequently, each time taking 0.1-1 second per call and slowing the system. We implemented optimizations to decrease the frequency and impact of database calls made by this method, enhancing the QLI performance. (AL-174546)

1.2.2 PATCH RELEASE 2024.1.5.4

build 18.5.4.85494

Released September 6, 2024

Note: This patch release is available for customer-managed (on-premise) instances only. This patch release doesn't apply to Alation Cloud Service instances.

- Optimized several internal APIs to reduce loading times of tasks generated by workflows. (AL-161158)

1.2.3 PATCH RELEASE 2024.1.5.3

build 18.5.3.84439

Released August 22, 2024

Note: This patch release is available for Alation Cloud Service instances on the cloud-native architecture. This patch release doesn't apply to customer-managed (on-premise) instances and Alation Cloud Service instances that haven't been migrated to the cloud-native architecture.

- Platform improvements for utilization. (AL-168402)

1.2.4 PATCH RELEASE 2024.1.5.2

build 18.5.2.83447

Released August 14, 2024

Note: This patch release is available for customer-managed (on-premise) instances only. This patch release doesn't apply to Alation Cloud Service instances.

- This patch release addresses a security vulnerability. (AL-164410)

1.2.5 PATCH RELEASE 2024.1.5.1

build 18.5.1.82675

Released August 2, 2024

Note: This patch release is available for Alation Cloud Service instances on the cloud-native architecture. This patch release doesn't apply to customer-managed (on-premise) instances and Alation Cloud Service instances that haven't been migrated to the cloud-native architecture.

- Updated the helm chart to address an issue where non-string custom label values were not handled properly. The issue previously hindered the upgrade of cloud instances. With this fix, the helm chart now supports non-string values for custom labels. (AL-166922)

1.2.6 RELEASE 2024.1.5

build 18.5.0.82175

Released August 1, 2024

Note: This release is available for all Alation instance types: customer-managed (on-premise) and Alation Cloud Service.

- *New User Experience*
- *Authentication & User Management*
- *Search*
- *Data Dictionary*
- *Document Hubs*
- *Lineage*
- *Open Connector Framework (OCF)*
- *Bug Fixes*
- *Known Issue*

New User Experience

- **Curate & Govern**
 - Bulk operations on related data assets with Catalog Sets (AL-153502, AL-153503)
 - Bulk Actions in the full-search experience to bulk-manage Data Policies, Business Policies, Domains, custom fields, Stewards, and People Sets.(AL-156924)
 - The Catalog Activity and Curation Progress reports (AL-153501, AL-153500)
 - Creation and configuration of Workflows (AL-151062)
 - Creation and management of Data and Business Policies (AL-137933)
- **Homepage**
 - Dynamic layout: Admin users can change the layout of the homepage through a drag-and-drop experience. Items can be added in the page preview mode and moved around the screen. Page elements can be reorganized within a block. Admin users can also change the width of custom blocks to occupy either the full screen or half of it. (AL-153691)
- Migration of RDBMS data source pages to OCF is available in New User Experience (AL-156576)
- File system objects are integrated into the New User Experience (AL-160203, AL-160171)
- Agent and connector dashboards are integrated into the New User Experience (AL-151872)
- Multiple improvements to My Inbox, including accurate conversation counts and displaying resolved conversations by default (AL-158174)

Authentication & User Management

Auth Service on Agent: IAM User Profile Support

As part of the **2024.1.5** release, we've added the IAM user profile-based authentication in Auth Service on Agent, where customers would store their credentials as profiles in a credentials file on their instance. This authentication mechanism is compatible with agents operating in an AWS environment, other platforms, and on-premise servers. (AL-163038)

Restriction on the User Profile Upload

For every CSV file with user profiles that is uploaded, the total number of user updates and creations is now allowed to be no greater than 50. Whenever a Server Admin uploads a CSV file, irrespective of this threshold being exceeded or not, the preview will always be shown. When the 50 limit is exceeded, then users won't be able to confirm the changes. The **Confirm** button appears disabled and an error message will be shown right below the preview table. (AL-158623)

Search

Alation Search now supports the AND operator (all caps) to specify the logical "and" condition among search terms. For example, if a user enters a search string like *revenue AND costs*, then only search results that contain both "revenue" and "costs" will be retrieved. Results that contain only "revenue" or only "costs" but not both will not show up. (AL-158501)

Data Dictionary

Support for Custom Field IDs in the Dictionary Upload

We've improved the handling of custom fields in the user interface to address issues arising from duplicate custom field names. Alation users can create multiple custom fields of different types with the same name, previously leading to conflicts during the data dictionary upload where the system could not resolve a custom field name to a specific field record. To prevent such issues, we've now enabled the inclusion of custom field IDs alongside names in the dictionary file headers. The allowed formats for entering custom fields IDs are:

- `field_id`
- `field_id | field_name`

For example, if there are two custom fields named `status`, one with ID `10001` and another with ID `10002`, users can specify the ID in the header to distinguish between them, like `10001|status` or `10002|status`. This update ensures that custom fields are correctly identified during the upload. (AL-160329)

Document Hubs

Customize Document Tables

We have enhanced the column organization experience of Document tables on Document Hub and Folder pages. Document tables now have a standard (out-of-the-box) column configuration with capabilities for Admin users to assemble and save a custom configuration for specific Document Hubs. Admins can also apply the custom configuration to all Folder pages in the Document Hub. The **Reset** behavior has been enhanced to restore settings to the standard configuration or the Document-hub specific configuration. (AL-141394)

Lineage

Business Lineage: General Availability

This release marks a significant milestone in our multi-month effort to enhance, declutter, and simplify lineage charts. Alation Lineage now provides a comprehensive set of features that deliver a balanced and engaging chart experience for all users, including data analysts, data engineers, business users, and viewers. Starting with version **2024.1.5**, features that were previously in Public Preview and labeled as *Beta* are now generally available. Users have multiple interaction options with the chart, allowing them to adjust the view and layout to meet their specific analytical needs. Below are the main highlights of the Lineage improvements you can now take advantage of. To access these features below, the Compound Layout must be enabled.

- **All instance types**
 - The *Compound Layout lineage view* (released in **2024.1.0**)
 - Availability of *manual lineage in Compound Layout* (released in **2024.1.0**)
 - *Business titles* in Compound Layout (released in **2024.1.1**)
 - *Data Trust Metadata Layer* (released in **2024.1.0**)
 - *Filtering of dataflow objects* (released in **2024.1.2**)
 - *Column-level lineage tracing* (released in **2023.3.0**)
 - *Alphabetical sorting of columns in Compound Layout* (released in **2023.3.4**)
 - *Data source logos* (released in **2023.3.2**)
 - *Enhanced Edge Highlighting* (released in **2023.3.3**)
 - *Data Quality Metadata Layer* (a **2024.1.5** feature)
 - *Lineage Chart Position Focus Improvement* (a **2024.1.5** feature)
- **Alation Cloud Service on the cloud-native architecture**
 - Availability of the *Lineage functionality in the New User Experience*
 - *Lineage Minimap* (a **2024.1.5** feature)
 - *Lineage Chart Legend* (a **2024.1.5** feature)
 - *Deprecation Warning Icon Change* (a **2024.1.5** feature)

Data Quality Metadata Layer

Added support to viewing Data Quality flags on the lineage graph and to view additional information in a popup pane for a specific table or column. (AL-155339)

Lineage Chart Position Focus Improvement

This improvement ensures that the cursor focus remains consistent within the lineage chart when users toggle display options, like business titles, or navigate between tabs, such as switching from full lineage to column-level lineage tracing. This ensures a seamless and consistent viewing experience across different aspects of lineage. (AL-145290)

Lineage Minimap

Alation Cloud Service on the cloud-native architecture, Classic UI and New User Experience

The lineage chart will now show a small minimap with bi-directional drag for easy navigation to help users orient on large lineage charts. Currently, the minimap is behind a feature flag (`alation.feature_flags.dev_enable_lineage_minimap`) and requires enabling. (AL-142151)

Note: On customer-managed instances, you can modify the values of backend feature flags [using the `alation_conf` tool](#).

Alation Cloud Service customers can request server configuration changes through Alation Support.

Lineage Chart Legend

Alation Cloud Service on the cloud-native architecture, Classic UI and New User Experience

Added a legend layer to lineage charts that shows all the icons on the chart with clarifications. The legend layer will be updated based on icons present at the time of viewing. If users filter dataflows or data sources the legend will also be filtered. This feature is currently behind a feature flag (`alation.feature_flags.dev_enable_lineage_legend`) and requires enabling. (AL-142152)

Note: On customer-managed instances, you can modify the values of backend feature flags [using the `alation_conf` tool](#).

Alation Cloud Service customers can request server configuration changes through Alation Support.

Deprecation Warning Icon Change

Alation Cloud Service on the cloud-native architecture, New User Experience

The approach to Trust Flags in the New User Interface (Endorse - Add Important Info - Deprecate) is different from Classic (Endorse - Warn - Deprecate). Instead of the warning (yellow) icon, the propagated deprecation alerts will appear with the blue Important info label on lineage charts to be consistent with the new Trust Flag paradigm. (AL-142158)

Remove the Dataflow for BI Sources Feature Flag

The configuration flag `alation.feature_flags.enable_dataflow_for_bi_lineage` has been removed. By default, dataflow objects will be created for all BI sources and displayed on lineage charts. (AL-163393)

“Disable automatic lineage generation” Renamed to “Enable automatic lineage generation”

The **Disable automatic lineage generation** checkbox on the **General Settings** tab of data sources has been renamed to **Enable automatic lineage generation**. The default behavior of this setting remains unchanged (automatic lineage generation is enabled). (AL-152080)

The Compound Layout Excludes Dataflows for Viewer Users

The Dataflow filter in Compound Layout is now set to *Hide* by default for users with the Viewer role. (AL-162532)

GSP Upgrade

Upgraded the General SQL Parser (GSP) jar to version **2.9.x.x**. The upgrade enables support for stored procedure lineage and for specific naming conventions, such as the use of the at sign @ in Db2 table names and numeric values at the end of catalog names for Google BigQuery (GBQ) target tables. (AL-159884)

Open Connector Framework (OCF)

OCF Connectors

Updates to the Open Connector Framework (OCF) connectors are released every week. See </sources/OpenConnectorFramework/ReleaseNotes> for information on the updates in the month of **July 2024**. A few July highlights are:

- **Certification of query log ingestion (QLI) from Redshift Serverless** (AL-156721)
 - Certified the Redshift OCF connector for Redshift Serverless instances on customer-managed and Alation Cloud service instances. Provided support for QLI using the SYS tables for Serverless, Provisioned, and Spectrum.
- **Column-level lineage from SAP HANA data sources** (AL-143424)
 - The OCF Connector for SAP HANA now supports column-level lineage. The General SQL Parser (GSP) will use the HANA grammar to generate lineage at the column level. This feature is behind a flag **Automatically extracts Column Level Lineage for SAP data source** in Admin Settings > **Feature Configuration**.
- **New user experience for Power BI Scanner OCF connector** (AL-136714)
 - The Power BI Scanner OCF connector now has an improved user experience with better navigation, multiple visual cues, and step-by-step guidance making it easier to configure and manage Power BI sources. The user interface now includes connection and access checks to help prevent configuration errors. Troubleshooting is made simpler as the error messages include actionable recommendations.

OCF Platform

Removed Feature Flags for Add-on Connectors

Alation now allows installation and usage of add-on connectors without the requirement to enable any feature flags. This change removes the flags for Fivetran (**Enable Fivetran Connector for OCF RDBMS sources (Paid Feature)**) and dbt (**Enable DBT Connector for OCF RDBMS sources (Paid Feature)**) from the user interface in Admin Settings > **Feature Configuration**. With this change, Data Source Admins will always see the add-on connector configuration tabs in the settings of the main connector. (AL-157919)

Removed Feature Flags for BI Column-Level Lineage

The configuration flags `alation.feature_flags.enable_auto_cll_tableau` and `alation.feature_flags.enable_auto_cll_powerbi_scanner` and the respective feature flags **Automatically extracts Column Level Lineage for Tableau BI data source** and **Automatically extracts Column Level Lineage from Power BI Scanner data source** are removed. By default all BI connectors will generate column-level lineage along with the table-level lineage for all the applicable data, if the respective RDBMS data source has its column-level lineage flag enabled. (AL-159784)

Bug Fixes

API

- For the Relational Integration API, we have fixed the documentation on the response type for the GET methods. They are now properly indicated as arrays of objects. (AL-153910)
- For both the Native Data Sources API and OCF Data Sources API, the `unresolved_mention_fingerprint_method` parameter is now correctly identified as an integer rather than a string. (AL-153910)
- Fixed a bug that was preventing the BI Server API from appearing in the Swagger UI in an Alation instance. (AL-161394)
- Fixed an issue with the Document API (GET `integration/v2/document`), where it might return a 504 timeout error for certain documents. The performance of the GET endpoint has been optimized by limiting the number of rows scanned to fetch the templates for objects. (AL-163608)
- Fixed an issue where the GET all articles API returned a 504 timeout error after one minute, even with a limit set to 1. The query process for fetching article visibility information has been optimized to retrieve visibility information in a single query. (AL-162485)

Data Dictionary

- Fixed an issue in the data dictionary download where the downloaded file included an extra description column, leading to a header error during the upload process (*Error at header 13, repeated header found: 'description'.*) (AL-160318)

Governance

- Fixed an issue with extracting Dynamic Masking policies from Snowflake that affected instances with multiple Snowflake data sources using a connection to the same database. Previously, when running metadata extraction (MDE) from one data source, the policy links in the **Data Managed by Policy** field were incorrectly deleted for other Snowflake data sources. To apply the fix, metadata extraction should be rerun on affected Snowflake data sources after updating Alation to **2024.1.5**. (AL-161825)
- *Classic UI*: Previously, the flag score values for child RDBMS objects (schemas of a data source, tables of a schema, columns of a table) reflected only flags that are directly set on the child objects. Now, they will also reflect the flags set via catalog sets if a child object is a member of those catalog sets. Previously, the flag score column in the user interface was sortable, but now the sorting is disabled due to performance considerations. Users can navigate through all pages for viewing the flag score of child objects but cannot sort by this column anymore. (AL-147639)

Note: The flag score is the sum of all the trust flags set applied to an object. For example, if a schema has two endorsements and one deprecation, the flag score is +1 (+2-1). Warnings don't have any impact on the score.

- Resolved a synchronization issue in manual catalog sets that occurred during the data dictionary upload process. Previously, changes made through uploads did not reflect in the manual catalog sets. To address this, we've introduced support for synchronizing the metadata of manual catalog sets' children during the upload flow. This functionality is controlled by a new feature flag, `alation.feature_flags.DEV_update_equivalent_objects_in_upload_data_dictionary`, which is disabled by default. Users looking to test or utilize this feature must explicitly enable the flag. (AL-142458)
- Fixed an issue with conditional catalog sets where changes to **Shared Fields** were not reflecting on newly added objects that met the membership criteria. Previously, intermittently, updates to descriptions or other shared values in a catalog set were not automatically applied to existing or new members unless the membership expression was manually refreshed. The problem has been fixed, ensuring that any updates to catalog sets are now correctly and immediately applied to all relevant objects. (AL-115297)

Document Hubs, Glossaries, & Articles

- Fixed an issue in the New UI where changes to the names of custom fields were not reflected in the table with customized Column Organization on the Folder page displaying documents with those altered fields. After this fix, if an admin has saved a table configuration and subsequently updates a custom field's name, the table configuration will unselect this field but will display the updated name correctly in the Column Organization widget. An admin will need to re-select the updated field. (AL-160664)
- Resolved an issue occurring when the feature flag **Document Hubs 3: Enable Hub Creation (Public Preview)** was activated. Previously, the **Edit Columns** widget did not display custom fields. This issue has now been fixed. (AL-154994)
- Resolved an issue in the Bulk Utility where the download process would become unresponsive if a user attempted to download terms assigned to a titled domain and an untitled subdomain. This fix ensures that downloads proceed smoothly, regardless of the titling status of the domain or subdomain involved. (AL-160170)

OCF Connectors

- Resolved an issue reported for the SQL Server OCF data sources where profiling failed for tables in schemas containing a period in their names. For example, while profiling was successful for tables in the ABC schema, it failed for tables in the ABC.ABC schema. The logic was modified to pass the schema name and catalog name separately in table profiling requests, which ensures successful profiling regardless of schema naming conventions. The changes have been done on both the OCF SDK **4.8.7** and the Alation application side. Specific connector updates will follow. (AL-141614)
- Resolved an issue with the Elasticsearch OCF connector not functioning in Compose with Query Service disabled. Users previously encountered the error *Credentials incorrect. Error connecting to the database* when attempting to add credentials through the **Connection Settings** dialog. This has been fixed, ensuring the connector now operates correctly in Compose. (AL-157169)
- Addressed security concerns reported for the OCF Oracle connector using JRE **1.8.0.x**. In response, we have upgraded the JDK to a newer version where the identified vulnerabilities were fixed to ensure enhanced security and compliance of the Oracle connector. The fix is applied to connector version **1.2.2** and requires Alation version **2024.1.5**. (AL-150899)
- Fixed an issue with timezone normalization in the QLI results, which previously led to the *partial success* errors. The problem stemmed from the user-selected date range being converted to UTC for database queries, but not being similarly converted when checking if results fell outside this range. This normalization inconsistency has now been corrected, and the checking of dates in both the fetching and condition align correctly. (AL-152646, AL-154254)

Alation Analytics

- Addressed a discrepancy between Alation Analytics on PostgreSQL and the Alation-managed Analytics Experience on Snowflake, regarding mismatched FP values in BI-related tables. The issue was traced to a loss of precision in larger integer values during the data transfer from PostgreSQL to Snowflake in the data copying framework. The problem has been corrected, ensuring that the FP values now consistently match across Alation Analytics and the Analytics Experience on Snowflake. (AL-157364)

Other Fixes

- Resolved an issue where modifying OAuth Client Applications through the UI resulted in an “IAMException AlationIAMRequest - 400” error. This error occurred during attempts to create or update OAuth clients, particularly when adjusting user information such as username and activity status. The fix ensures that users can now successfully manage OAuth Client Applications under Authentication settings. (AL-163384)
- Upgraded the Java component of the Alation server (Amazon Corretto) to version **11.0.23.x.x** to take advantage of bug fixes and enhanced functionality. (AL-125882)
- A recent penetration test identified a potential vulnerability related to file deletion. When a user uploaded and subsequently deleted an attachment, the file was not properly removed from the server. We have addressed this issue to ensure that attachments are securely removed from the server upon deletion. (AL-135046)
- Addressed a discrepancy between the Classic UI and New User Experience: the **Last Updated** and **Created On** columns, previously only available in the Column Organization widget on Folder pages in the Classic UI, are now also included in the New UI. (AL-161955)

Known Issue

The data dictionary preview user interface issues an internal API error when a custom field name included into the data dictionary file contains the colon special character :. In this case, users can validate the preview by downloading the preview report. The upload to the catalog still works as expected and the updates from the dictionary will be saved accordingly. (AL-137006)

1.2.7 PATCH RELEASE 2024.1.4.2

build 18.4.2.82056

Released July 26, 2024

Note: This patch release is available for Alation Cloud Service instances on the cloud-native architecture. This patch release doesn't apply to customer-managed (on-premise) instances and Alation Cloud Service instances that haven't been migrated to the cloud-native architecture.

- Fixed an issue in the Celery task logic, originally designed to fix a race condition affecting the updating of task statuses in job tables. Previously, if a task was not found in the database, an exception would trigger a retry of the update. However, if retries were initiated for other reasons, they would block the task. Additionally, a call within the metrics module occasionally triggered the same exception for different tasks, leading to the blocking of the current one. This issue caused some disruptions, such as scheduled queries being randomly skipped, with a message in the user interface that they were skipped *due to technical difficulties*. This patch resolves the problem and ensures more reliable task management. (AL-165523)

1.2.8 PATCH RELEASE 2024.1.4.1

build 18.4.1.80840

Released July 11, 2024

Note: This patch release is available for all Alation instance types: customer-managed (on-premise) and Alation Cloud Service.

- Fixed an issue where users encountered a pop-up message *Sorry, there was an error in Alation* in the Classic user interface on various pages after a recent update. While users could close this message to continue working, it was associated with a *500 internal server error* in the internal API. After the fix, the error no longer occurs, ensuring uninterrupted access to Alation. (AL-163785)

1.2.9 RELEASE 2024.1.4

build 18.4.0.80242

Released July 4, 2024

Note: This release is available for all Alation instance types: customer-managed (on-premise) and Alation Cloud Service.

- [Compose](#)

- [Search](#)
- [Document Hubs](#)
- [Homepage](#)
- [Core Platform](#)
- [Alation Analytics](#)
- [Open Connector Framework](#)
- [Alation Anywhere for Tableau](#)
- [Authentication](#)
- [Lineage](#)
- [Governance](#)
- [Alation API](#)
- [Improvements](#)
- [Bug Fixes](#)

Compose

Configurable Timestamp Format for Query Results

Introduced a new configuration option, **Timestamp Format**, for selecting a timestamp format in Compose Settings > **Results**. Users can now select a preferred format for displaying timestamp types: either **Standardized** (RFC 3339) or **Legacy**. This configuration applies to only newly stored data. Previously stored result data will not be affected by changes to this setting. (AL-157186)

Compose SSO with AWS Databricks Unity Catalog

Alation now supports an additional configuration of Compose OAuth for Databricks Unity Catalog on AWS. Users can connect to Databricks Unity Catalog data sources in Compose via partner applications, no longer relying on personal access tokens. To configure, you'll need to create a new app connection in Databricks under Account Console > Settings > App Connections and provide its properties in the data source settings. (Admin privileges are required to create a new app connection.) (AL-155129)

Load Query

Compose users now have the ability to load a query from a file by clicking **Load Query** on the Compose dropdown menu (top left of the page), selecting a text file from the local machine, and loading the content of the file to the Compose query editor to create a new query. (AL-158717)

Disable Excel Live Reports by Default

New Alation installations and instances upgraded to **2024.1.4** will have their default values set to *off* for Excel Live Report. This affects both Cloud Service and customer-managed instances. (AL-158713)

Query Forms Support Object Names with Double Quotes

For Query Forms, the WHERE clause will now allow column names with double quotes for non-ASCII characters. (AL-116577)

Search

Alation Cloud Service instances on the cloud-native architecture

We're beginning to deploy *Intelligent Search* (or *Semantic Search*) to qualified Alation Cloud Service instances after their search index is migrated to Elastic Cloud Service. [Information about Intelligent Search](#) is available on Alation Community (requires a login). (AL-119678)

Document Hubs

Delete Document Hubs

Users now have the ability to delete a document hub, which removes it from the user interface and prevents the creation of objects within it. Deleting is only allowed for hubs that contain soft-deleted folders and documents. (AL-152300)

Homepage

Documents and Folders on the Homepage

Alation Cloud Service instances on the cloud-native architecture

Documents and folders can now be embedded on a custom homepage. For documents, the **Description** field will be displayed. For folders, the **Description** field and a list of up to five of its documents will be displayed. Documents and folders can only be embedded using the homepage editor in the New User Experience, but they will be visible when viewing the homepage in both the New and Classic User Experiences. (AL-156127, AL-161045)

Free-Text Field on the Homepage

Alation Cloud Service instances on the cloud-native architecture

You now have the option to add a free-text RTF field to your custom homepages. RTF fields can only be embedded using the homepage editor in the New User Experience, but the RTF field will be visible when viewing the homepage in both the New and Classic User Experiences. (AL-161045)

Core Platform

Postgres Upgrade to 16.2

Customer-managed instances

The internal Alation database (Postgres) was upgraded to version **16.2** to take advantage of new features and security enhancements. (AL-151431)

Support for TLS 1.3 Connections

Alation Cloud Service instances on the cloud-native architecture

Added support for TLS 1.3 connections to all external non-FIPS/commercial Application Load Balancers. Customers will now see TLS 1.3 as the primary connection type, with fallback to TLS 1.2 for clients that do not support TLS 1.3. (AL-122273)

Alation Analytics

Alation Analytics Postgres Upgrade to 16.2

Customer-managed instances

The Alation Analytics PostgreSQL database was upgraded to version **16.2** to take advantage of new features and security enhancements. (AL-151431)

Custom Field Scrubbing on Alation Analytics

A new configuration flag, `alation_analytics-v2.extract.enable.custom_field_formatter`, is now available for enabling custom field scrubbing on the Alation Analytics V2 database. This feature converts all empty or placeholder values (such as "", {""}, and { }) to NULL. To activate this functionality, enable the flag using **alation_conf** (requires a restart of the Alation Analytics Celery service (`alation_supervisor restart celery:celery-alationanalytics_0`)). By default, this flag is disabled. (AL-143652)

Note: Alation Cloud Service customers can request server configuration changes through Alation Support.

Open Connector Framework

Simplified Installation and Upgrade for Connectors

Alation Cloud Service instances on the cloud-native architecture

We've made changes to the user experience of administrators managing connectors, helping them streamline their actions. The experience is enabled by default for all Cloud instances on the cloud-native architecture.

- **Simplified installation**—Server admins can now search for and install connectors directly from the Connector Dashboard with a few clicks.
- **Connector upgrade indicator**—The Connector Dashboard now displays an *Upgrade available* indicator when a new version of the connector that you have installed becomes available.

- **Simplified connector upgrade**—Server admins can upgrade a connector with a few clicks, without the need to download a newer connector locally first. (AL-146347)

Manage Access to BI Folders and Reports

Previously, access permissions could only be set at the BI server level, with all the child BI objects inheriting these settings. With this update, permissions can also be set individually on BI folder and report objects. If no explicit permissions are defined, folders and reports will inherit settings from their parent object (BI server or BI folder). BI Admins can enable granular permissions by selecting **Enable explicit permission** checkbox on a folder or report, allowing for specific user access control. This feature is disabled by default and can be activated by a Server Admin using **alation_conf** with the parameter `alation.granular_object_privacy.enabled`.

Note: Alation Cloud Service customers can request server configuration changes through Alation Support.

Alation Anywhere for Tableau

Upload Tag Metadata to Tableau

For Tableau BI sources, the Metadata Upload job now transfers tag information from Alation to Tableau for workbooks, dashboards, and view objects. (AL-158283)

Authentication

Retrieve Data Source Credentials from AWS Secrets Manager via Alation Agent

Alation Cloud Service instances on the cloud-native architecture

This feature allows customers to keep their data source credentials within their Virtual Private Cloud so that they never traverse the Alation Cloud Service infrastructure. To achieve this, the [/cloud/AlationAgent/AuthServicePlugin](#) can now be installed on an Alation Agent. This release offers support for AWS Secret Manager vaults with instance-based IAM roles. With Auth Service installed on the Agent, credentials stored in the vault won't leave the perimeter of the company network when being retrieved by Alation for authenticating with a data source for extraction or query log ingestion. Moreover, these credentials will not be persisted in any of the Alation components including Alation Agent. For more information, see [/sources/OpenConnectorFramework/ConfigureSecretsforOCFConnectors/ConfigureAWS SecretsManager/index](#). (AL-149546)

Note: This feature requires customers to upgrade to the new [Alation Agent \(1.7.2.4360\)](#) and a connector version that supports the feature. To find out if a specific connector supports this feature, see the documentation for the specific connector in the [/sources/OpenConnectorFramework/index](#) section.

Lineage

New Lineage Chart Rendering Algorithm

With this release, we introduced a new chart rendering algorithm for the Compound Layout in lineage charts, enhancing both performance and user experience:

- **Node and Path Limits**—The system now loads the entire graph until it reaches a hard limit of 2,000 nodes and 50,000 paths. If these limits are reached:
 - Use the Impact Analysis reports to export a complete list
 - Apply the Dataflow filter to hide certain elements
 - Navigate to another object within the same lineage graph to view its lineage
- **Deprecated Features**—The features known as *horizontal truncation* (or *node pagination*) and *vertical truncation* (or *collection nodes*) have been deprecated. The collection node dialog has been removed from the user interface.
- **Bug Fixes**—Ensured each data source group in the Compound Layout has one cluster on the chart and resolved an issue with lineage path highlighting in looping graphs. (AL-150882)

Governance

Enable De-duplication of Snowflake Tags

By default, Snowflake tags from a common database or schema that are used across other databases or schemas are duplicated in Alation for each database or schema where those tags are applied. This creates duplicate tag custom fields and necessitates admins to add all the duplicated fields to templates for usage. We've added a new feature flag `alation.feature_flags.enable_centralized_tags_deduplication` that enables de-duplication of common tags by creating only a single copy of the custom field for each, even though they are applied across different databases or schemas. The flag is disabled by default. (AL-157018)

Clear Values Using the Data Dictionary

While uploading a data dictionary in the user interface, users can now choose between replacing existing values with skipping empty values or with applying empty values from the dictionary. (AL-156670)

Alation API

Find a short summary of the Alation API 2024.1.4 changes below. More details are available in the [2024.1.4 API Release Notes](#) on Alation's Developer Portal.

Allie SDK

Alation has released an open-source Python SDK—Allie SDK—to better enable working with Alation APIs. Check out the [SDK documentation](#) and the [Allie SDK public GitHub repository](#).

Documents API

You can now specify which fields should be returned when [getting documents](#) by providing the query parameter values and supplying a comma-delimited list of fields that should be returned. (AL-138823)

Folders API

You can now specify which fields should be returned when [getting folders](#) by providing the query parameter values and supplying a comma-delimited list of fields that should be returned. (AL-156335, AL-162492)

Members Permission and Privacy Settings APIs

We've added BI folders and BI reports as supported object types for the [Members Permission API](#) and [Privacy Settings API](#), so you can now control access to these object types using Alation APIs. (AL-159617)

Improvements

- The feature flag **3rd Party Metrics** has been removed from the Admin Settings > **Feature Configuration** page. (AL-149013)
- Upgraded the Java component of the Alation server to a newer version (**amazon-corretto-11.0.23.x.x**) to take advantage of bug fixes and improved security features. (AL-125882)
- Upgraded the Redis Python package from version **3.x** to **5.x.x**. This update provides enhancements and new features that improve performance and compatibility. Due to compatibility issues, the `django-redis-cache` library, which requires a Redis version less than **4.0**, has been replaced with `django-redis`. It offers similar functionality as `django-redis-cache` and ensures better support and updates. (AL-157423)
- Addressed performance bottlenecks in the `LMSTitleDescription` serializer, which previously accounted for about 50% of the processing time when creating terms with a description that includes @-mentions. This optimization significantly speeds up the links endpoint of the internal server API, used for hydrating the Document tables on Folder pages. (AL-155303)

Bug Fixes

Articles, Terms, and Documents

- Addressed an issue where users experienced a *504 Gateway Timeout* error when accessing articles from the Classic user interface. Fixed by improving the performance of the internal article API, resulting in faster response times and enhanced efficiency when retrieving articles. (AL-160021)
- Resolved an issue in the Classic user interface where attempting to remove or a selected template on the term or document catalog page (by selecting the topmost option in the template selector dropdown) previously resulted in a 500 error. Although the error page was displayed, if the user refreshed the

page, the template change was correctly persisted. This fix ensures that the template can be updated without the need to refresh the page for the change to take effect. (AL-158942)

- Fixed an issue where article picker and multi-picker fields were not displayed on the catalog pages of article objects after the **Document Hubs 1: Setup** and **Document Hubs 2: Show Glossary Hub** feature flags were disabled. These fields are now consistently visible regardless of the status of these feature flags. (AL-160185)
- Fixed an issue in the Classic user interface where Stewards and Composers could edit picker fields for terms, despite editing restrictions set at the glossary level. This fix ensures that editing restrictions are properly enforced. (AL-159727)
- Addressed an issue where the download of glossary terms would stall when using the Bulk Utility if the terms were associated with a domain hierarchy. For instance, downloads previously halted if a term was linked to both a parent Domain and its child subdomain. This fix ensures that downloads proceed without interruption. (AL-152126)
- Fixed a bug where articles failed to load properly due to an issue with the URI transformation in version 2023.3.4 (the / symbol was incorrectly converted into ASCII.) This fix ensures that all special characters in article titles are handled correctly. (AL-149344)

Bulk Utility

- Fixed an issue where tags were not uploaded when using the Bulk Utility to add multiple tags to a term. This problem persisted regardless of whether users tried to upload a single tag or multiple tags, or whether they were creating or updating tags. After this fix, new tags can be added via the Bulk Utility. (AL-155181)

Governance Fixes

- Fixed an issue where the option **Reviewers of this step can bypass this entire workflow in workflow** settings no longer worked. Now, a workflow can be bypassed if the user is a reviewer, and the workflow has this reviewer bypass flag enabled. (AL-156067)
- Resolved a display issue on the main Policy Center page where suspended users appeared in the **Stewards** column. The Stewards column no longer shows suspended stewards, only displaying active stewards associated with policies and policy groups. (AL-153864)
- Fixed an issue where users without the necessary permissions could view private data sources in the **Titles & Descriptions Chart**. Private data sources are now excluded from the list in this view, reinforcing data privacy and access controls. (AL-158729)

Other Fixes

- Updated the New User Experience to match Classic, allowing all users to resolve or reopen conversations, not just the creators. This fix ensures the **Resolve conversation** or **Reopen conversation** button is available to everyone. (AL-157619)
- Fixed a bug in the homepage editor (Admin Settings > **Homepage Customization**) where users could save the homepage after selecting **Image** as the banner background mode without actually uploading an image. Now, if a user chooses to set a banner image but does not upload one, an error message will appear, and saving the homepage will be blocked until an image is uploaded. (AL-141037)

1.2.10 PATCH RELEASE 2024.1.3.2

build 18.3.2.79136

Released June 20, 2024

Note: This patch release is available for customer-managed (on-premise) instances and Alation Cloud Service instances on the cloud-native architecture.

This patch isn't available for Alation Cloud Service instances that haven't been migrated to the cloud-native architecture.

- Fixed the issue of slow response times when checking job statuses. Previously, viewing the data source settings page was delayed by one to two seconds for each job in the Job History table due to checks on the status of Celery processes. To address this, we have removed the check for the Celery process from internal Celery tasks. This update is intended to improve response times and overall performance when accessing job statuses. Users should now experience faster loading times on the data source settings page. (AL-161486)

1.2.11 PATCH RELEASE 2024.1.3.1

build 18.3.1.78990

Released June 18, 2024

Note: This patch release is available for customer-managed (on-premise) instances and Alation Cloud Service instances on the cloud-native architecture.

This patch isn't available for Alation Cloud Service instances that haven't been migrated to the cloud-native architecture.

- Reverted the [Django library update](#) introduced in 2024.1.3 that had altered the handling of SAML attributes without an explicit NameFormat. The change in 2024.1.3 had caused login failures for customers with Microsoft Entra ID or Active Directory Federation Service (AD FS) as identity providers, where the authentication application settings did not include the NameFormat property (undefined by default). **We have reverted this change to restore the previous behavior.** This gives customers time to update their Entra ID or AD FS configurations accordingly. If your IdP is Entra ID or AD FS, please ensure that your IdP configurations are [updated as recommended](#). (AL-161311)
- Implemented a new response message for SAML replay attacks to enhance security. Any repeated use of a SAML response assertion will now trigger the message: *Received SAMLResponse assertion has already been used*. Additionally, following the upgrade to version 2023.1.3.1, all active user sessions that were established before the upgrade will be logged out upon the first attempt to access the application after the upgrade. This measure ensures that all sessions start under the latest security and application updates. (AL-161311)
- Fixed a bug where the **Custom Field Permissions** page experienced significant performance issues, with the page becoming unresponsive or taking more than ten seconds to respond when attempting to edit permissions. The fix optimizes the loading process for the Custom Field Permissions page, reducing response times and resolving issues of unresponsiveness. (AL-157895)
- Fixed an issue where Compose queries would continue running indefinitely in cases where the target table didn't exist. Implemented a fix to terminate query service connections immediately if a table does not exist. This fix also resolves the issue of queries with typos continuing indefinitely. (AL-157061)

1.2.12 RELEASE 2024.1.3

build 18.3.0.76841

Released June 3, 2024

Note: This release is available for all Alation instance types: customer-managed (on-premise) and Alation Cloud Service.

- [Core Platform](#)
- [New User Experience](#)
- [Governance](#)
- [Lineage](#)
- [Alation Public APIs](#)
- [OCF Connectors](#)
- [Known Issues](#)
- [Bug Fixes](#)

Core Platform

The Django Component Upgrade

We have upgraded our software infrastructure to Django **4.2**, replacing the end-of-support Django 2.x. This new version contains enhancements that improve the robustness, security, and performance of the Alation application:

- Improved WebSocket support—Django **4.2** enhances real-time communication capabilities for SaaS applications through improved WebSocket support.
- Async views and tasks—This version allows for more efficient management of long-running operations through asynchronous views and tasks.
- Async StreamingHttpResponse—With ASGI (Asynchronous Server Gateway Interface), this version enables the StreamingHttpResponse class to work with asynchronous iterators, improving the handling of large data streams. (AL-141005)

Important: ACTION REQUIRED

This upgrade impacts the SAML authentication configuration in Microsoft Entra ID (former Azure Active Directory) and Active Directory Federation Service (AD FS) as identity providers (IdPs). The default value of the Name Format property of SAML attributes on the Alation server has been changed to `unspecified` from the previous value of `URI`. This creates a requirement to explicitly set the Name Format property to `URI` or `Basic` in the SAML application in Microsoft Entra ID, as it also defaults to `Omitted` (default).

Customers who use Entra ID or AD FS as IdPs are asked to perform an additional configuration in their SAML application by **July 30, 2024**. Refer to the advisory [Important Advisory - SAML authentication failure with Azure/Microsoft IDP and Microsoft ADFS after upgrade to 2024.1.3](#) on Alation Community for specific details.

New User Experience

Alation Cloud Service on Cloud-Native Architecture

- In New User Experience, the column organization widget was added to the **Folder** and **Documents** tables on catalog pages of Document Hub objects and to the **Documents** table on catalog pages of Folder objects. (AL-141392)
- The **ALLIE AI Suggested Stewards** feature is now available in New User Experience. (AL-132059)
- Added more functionality to the **alation_conf** flag `alation.feature_flags.hide_articles` that can be used by customers who have completed the transition from articles to Document Hubs. Now, it also hides articles from the **Recently Visited** section of the homepage in New User Experience. The default value is `False` (no change in behavior). Alation Cloud Customers can request this change through Alation Support. (AL-146401)

Governance

Folder Object Type as a Field Option

Added the folder object type as a selectable option for Object Set and Reference custom field types. Note that it will appear as “Glossary” if the Document Hubs 3 feature flag is disabled. (AL-152523)

The Mentioned On Field

Added a new built-in field **Mentions** that lists all Documents and Folders that mention a specific catalog object. On the sidebar of a catalog page or when being added via Customize Catalog, it appears with a title of **Mentions**. When the field is located on the main panel of a template, it will display as two separate tables, one titled **Mentioned on Folders** and the other **Mentioned on Documents**. (AL-152316)

Data Dictionary

In release **2024.1.3**, uploading a data dictionary file with empty values via the user interface will not reset the existing catalog field values. This is a change from the previous release (**2024.1.2**), where the catalog fields would be reset if the corresponding field in the uploaded data dictionary file was blank. (AL-155567)

Lineage

Support for Cross-Source Lineage with Direct Lineage Extraction

Cross-source lineage configuration now supports data sources with *direct lineage extraction* (Databricks Unity Catalog, SAP HANA). This addition follows the capabilities introduced in the **2023.1.6** release, which enabled cross-source lineage for data sources connected via the same host and port and extends these capabilities to include the direct lineage pipeline. (AL-148534)

Alation Public APIs

Find a short summary of the Alation API changes below. More details are available in the [2024.1.3 API Release Notes](#) on Alation's Dev Portal.

Folder API

We've added a new [Folders API](#) for creating, retrieving, updating, and deleting folder object types. (AL-152751)

Custom Field Values API

The Custom Field Values API now supports the `glossary_v3` object type, which refers to both glossary and folder objects when the Document Hubs feature is enabled. (AL-155886)

Search API

The Search API now supports the `glossary_v3` object type (can refer to both glossary and folder object types), and the `glossary_term` object type (can refer to both document and term object types). (AL-155886)

Data Dictionary API

The Data Dictionary API's [Upload a data dictionary endpoint](#) now has parameters that control the behavior of blank values in the data dictionary. You can ignore blank values or cause them to reset the value of the field in the catalog. (AL-155567)

Documents API and Terms API

- We updated the sample response for the GET method to be more accurate. The sample response is now an array of documents, and the `otype` and `is_public` items have been removed, as they are not returned when getting documents. (AL-149027)
- Added throttling configurability for the Document API at the Celery component level to prevent API performance issues. The following **`alation_conf`** parameters were added:
 - `alation.api.throttling.document_public_apis.read` (defaults to null, which means “no rate limit”)
 - `alation.api.throttling.document_public_apis.write` (defaults to 100/minute) (AL-151897)

OCF Connectors

Updates to the Open Connector Framework (OCF) connectors are released every week. See [/sources/OpenConnectorFramework/ReleaseNotes](#) for information on the updates in the month of **May 2024**. A few May highlights are:

The Talend OCF Connector Public Preview

The [OCF connector for Talend](#) is now available on Connector Hub. The connector is compatible with Alation version **2023.3** and higher. (AL-143292)

Enhance the Parser to Handle Parameterized Values

Previously, users reported issues with extracting the connection and lineage information for Power BI datasets due to limitations in the connector logic in handling API responses that included parameterized values within expressions. We made enhancements to the connector to ensure it accurately handles parameterized values. Additionally, a fix was implemented to address an out-of-memory error caused by the repeated replacement of commonly used parameter names leading to excessive string expansion. Another fix ensures that table names in queries with special characters are correctly quoted, improving the resolution of nodes for cross-source lineage and enhancing the parser's accuracy and reliability. (AL-153882)

Resolving Unqualified Column Names in Query Log Ingestion

Customer-managed (on-premise) instances and Alation Cloud Service instances that were not migrated to the cloud-native architecture

Query log ingestion (QLI) for Snowflake OCF data sources now supports the resolution of unqualified column names from query history. An unqualified column name is a reference to a column that does not include parent objects. Previously, joins and filters were not captured if the column names weren't qualified with table aliases in the join and filter conditions. The change applies to customer-managed (on-premise) instances and Alation Cloud Service instances that were not migrated to the cloud-native architecture. It will be implemented on Alation Cloud Service instances on the cloud-native architecture in a future release. (AL-154123)

Known Issues

- The **2024.1.3** release raised the default iteration count of the password hasher for user CSV files from 390,000 to 600,000. This adjustment has extended the processing time for password hashes in CSV files, resulting in UI request timeouts when bulk-uploading users from CSV. Despite the UI timeout, the process continues in the backend, and if the data is correct, the updates will still be completed. For information on the workaround, please refer to this [Support Knowledge Base article](#). (AL-158623)
- In Classic UI, removing the selected template by selecting the topmost option in the template selector dropdown on a term or document catalog pages results in the 500 error in the user interface. If the user refreshes the page, the change is saved and the error is resolved. (AL-158942)
- Creating or updating an OAuth client application by Server Admins through the user interface in Admin Settings > **Server Admin** > **Authentication** > **OAuth Client Applications** fails with the error *IllegalArgumentException AlationIAMRequest - 400*, preventing the creation or modification of any app used to generate authentication tokens. If you need to create or update an app, please contact Alation Support. The authentication tokens generated by existing applications remain unaffected. (AL-163384) (FIXED in version **2024.1.5**)

Bug Fixes

Document Hubs

- Previously, all custom fields from document templates, irrespective of the Hub, were displayed in the Column Organization widget for the Document table on a Document Hub's catalog page. After this fix, the list of columns in Column Organization becomes specific to each Hub. (AL-148730)
- Previously, documents and folders from unpublished Document Hubs appeared in Alation Search results. After this fix, documents and folders from unpublished Document Hubs won't be discoverable through Search. (AL-154062)

Governance

- Previously, the stewards listed on the Policy Center page didn't match the actual stewards on policy group objects and suspended stewards were displayed. After this fix, suspended users are no longer visible in the list of policies, the row access policy links table, and the dynamic data masking policy links table. (AL-157793)
- Fixed an issue where users with roles other than Server Admin couldn't edit the Picker custom fields which had no value set. The issue occurred when updating a custom field on an object that a user had access to, but this object had a corresponding synced object established via a manual catalog set that the current user couldn't access. As a result of this fix, when a user updates any field value on a catalog object that falls under manual catalog sets metadata sync settings, the system automatically updates the same field value on equivalent objects. This system action overrides any object-level permissions on these equivalent objects, ensuring the field value is updated across all related items. Additionally, the field edit history will record the details of the actual user who made the update, both on the edited object and its equivalent objects. (AL-142427)
- Previously, the Bulk Utility encountered a 500 error and failed to open if there was an object template with the same name as one of the article templates. To fix the issue, we implemented an object type filter in Bulk Utility to ensure that only article templates are considered. (AL-155973)
- When using Bulk Utility to download glossary terms, the download became unavailable if any of the terms had a domain hierarchy associated with them. For example, if a parent domain and a child subdomain were both associated with the term, then the download was unsuccessful. The issue has been fixed to ensure reliable functionality in downloading terms. (AL-152126)

Alation Analytics

- An ETL sub-job for Alation Analytics, DatasourcesTransformJob, failed in version **2024.1.1** due to an error *Current transaction is aborted. ...* This issue was traced back to handling null values for URIs of OCF data sources. The fix ensures that the URI field adheres to the NOT-NULL constraint, stabilizing the ETL process. (AL-155188)
- Extended the consumer_timeout setting for RabbitMQ to six hours. This change ensures that longer-running tasks can complete without interruption due to premature timeouts. (AL-154677)
- Fixed the error *Duplicate key value violates unique constraint 'bi_connection_column_fp_key'* reported in the Alation Analytics ETL process. The resolution involves removing the unique key constraint for the fp column in the bi_connection and bi_connection_column tables to prevent such conflicts during ETL operations. (AL-150501)

Other Fixes

- Fixed an issue where for data sources supporting direct lineage extraction, the cross-source lineage configuration resulted in temporary objects on lineage charts. Extended the support for cross-source lineage to data sources with direct lineage extraction. (AL-148534)
- For the Power BI OCF connector, fixed an issue where semantic models (renamed datasets) were missing connections after extraction. Fixed in connector version **2.2.1** (AL-147313)
- Fixed an issue where query log ingestion (QLI) didn't retrieve the join and filter information if column names were not qualified with parent table names. Fixed by extending QLI support to unqualified column names for on-premise instances and Alation Cloud Service instances that haven't been migrated to the cloud-native architecture. (AL-141950)

1.2.13 PATCH RELEASE 2024.1.2.2

build 18.2.2.78391

Released June 17, 2024

Note: This patch release is available for customer-managed (on-premise) instances and Alation Cloud Service instances on the cloud-native architecture. The patch will be deployed for specific customers as required.

(This patch release isn't available for Alation Cloud Service instances that haven't been migrated to the cloud-native architecture.)

- Fixed a bug where the **Custom Field Permissions** page experienced significant performance issues, with the page becoming unresponsive or taking more than ten seconds to respond when attempting to edit permissions. The fix optimizes the loading process for the Custom Field Permissions page, reducing response times and resolving issues of unresponsiveness. (AL-157895)
- Fixed an issue where Compose queries would continue running indefinitely in cases where the target table didn't exist. Implemented a fix to terminate query service connections immediately if a table does not exist. This fix also resolves the issue of queries with typos continuing indefinitely. (AL-157061)
- Introduced a new configuration option, **Timestamp Format**, for timestamp formatting in Compose settings. Users can now select the preferred format for displaying timestamp types: either **Standardized** (RFC 3339) or **Legacy**. This configuration applies to only newly stored data. Previously stored result data will not be affected by changes to this setting. (AL-153948, AL-153947, AL-157186)

1.2.14 PATCH RELEASE 2024.1.2.1

build 18.2.0.73651

Alation Analytics for Alation Cloud build 1.39.1.2516

Released May 9, 2024

Note: This patch release is available for Alation Cloud Service instances on cloud-native architecture.

- Extended the `consumer_timeout` setting for RabbitMQ to six hours. This change ensures that longer-running tasks can complete without interruption due to premature timeouts. (AL-154677)

1.2.15 RELEASE 2024.1.2

build 18.2.0.73651

Released May 1, 2024

Note: This release is available for all Alation instance types: customer-managed (on-premise) and Alation Cloud Service.

- *New User Experience (General Availability)*
- *Alation Agent*
- *Authentication*
- *Lineage*
- *Governance*
- *Compose*
- *Alation Analytics*
- *Logical Metadata Services*
- *Alation Public APIs*
- *OCF Connectors*
- *Bug Fixes*

New User Experience (General Availability)

Alation Cloud Service on Cloud-Native Architecture

We've added more functionality to the **General Availability** (GA) release of the [New User Experience](#) arriving just prior to the upcoming **1.0** release. The following list captures a summary of additions since the Public Preview IV release. All additions are available on Alation Cloud Service on Cloud-Native Architecture, unless also specified for other instance types.

- Query page: the Run Query feature
- Policy object pages
- Lineage: compound layout view
- Manual lineage curation
- Governance landing page: the Curate section (in Classic but linked from the new user interface)
- My Inbox
- User page
- Alation add-ons: added Data Catalog Maturity Assessment (DCMA)
- BI pages
 - Customizable BI layouts (*also available on customer-managed (on-premise) instances and Alation Cloud Service instances that haven't been migrated to the cloud-native architecture*)
 - Granular labels on BI report and folder catalog object types (site, project, report, dashboard, page, view, and others)

- Customizable table layouts for data sources

Note: As we move into GA, you no longer need beta language in your contract or an addition to the whitelist to enable the New User Experience.

Alation Agent

Notifications for Alation Agent's loss of connection and certificate expirations are now enabled by default for Server Admins. Existing preferences for Server Admins who have previously configured these notifications will remain unchanged. These notifications were previously opt-in by default. The change to opt-out by default aims to improve awareness and response to these important system events. Server Admins can adjust the notification settings under the **Notifications** tab on the **Account Settings** page. (AL-153656)

See </cloud/AlationAgent/AgentReleaseNotes> for the full release notes for Alation Agent.

Authentication

Inactive User Suspension

Users who have not logged into the Alation catalog for an extended period can now be automatically suspended:

- **Non-FedRAMP instances**—By default, this setting is disabled. To enable, navigate to **Admin Settings > Server Admin > Authentication**, find the **Inactive User Suspension Setting** section, and set a value between 30 and 365 days. To deactivate, reset the value to 0.
- **FedRAMP instances**—Suspension is enabled by default at 90 days and cannot be altered or disabled. (AL-124754)

Lineage

Column-Level Lineage for Temporary Tables

Alation Cloud Service on Cloud-Native Architecture

Added support of column-level lineage for temporary (temp) tables. Temp tables are usually created with `CREATE TABLE` queries using either the `TRANSIENT` or `TEMPORARY` keyword. Such queries can also be generated by tools like dbt. Tables that are created and dropped in the same session are classified as temp tables as well. On lineage charts, column-level lineage will be available for a temp table if column names are explicitly mentioned in the query. If columns are not explicitly mentioned (queries like `SELECT *`), then column-level lineage will only be available for the `CREATE TABLE AS SELECT *` queries.

Column-level lineage from the `ALTER TABLE` queries is supported if these queries are executed during the same session as the `CREATE TABLE` queries. (AL-126794).

Column-Level Lineage for Custom SQL

Alation now generates column-level lineage between an RDBMS and a BI source when data was imported into the BI server with custom SQL queries. The feature applies to RDBMS sources with table and column-level lineage support. (Information about column-level lineage support is available in the *Support Matrix*.)

- Custom SQL query content can be found on the **DataSource Connection** tab of BI objects. The feature has been currently certified for Tableau and Power BI OCF BI sources.
- A limitation: column-level lineage is not displayed correctly for custom SQL in scenarios involving column name changes on the BI server side. (AL-147324)

Filtering Dataflows (Public Preview)

You now have the ability to filter dataflow objects out of lineage charts. This feature is intended to reduce clutter that may be caused by recurring dataflows due to automated intake from query log ingestion. To access dataflows when they are filtered out of the chart, you can trace the line between nodes on the chart and click to view dataflow content on the bottom and side panes. To access all dataflows for a particular node connection, you can alternatively use the Impact Analysis lineage reports filtering them by dataflow objects. This feature is currently in Public Preview with a Beta “chip” in the user interface. (AL-140791)

Lineage Settings Tab for BI Sources (Public Preview)

Introduced a new **Lineage Settings** tab in the settings of OCF BI sources. BI Source Admins can now configure cross-system lineage directly within BI source settings, improving usability and streamlining the setup process. The configuration doesn't require manual entry of the `host:port` information. Instead, admins can select upstream RDBMS sources by name. This feature is currently in Public Preview with a corresponding Beta “chip”. The enhanced lineage settings have been certified with Tableau, Power BI, and MicroStrategy OCF BI sources in this release. (AL-145817)

Governance

Introduced a new feature flag `alation.datalake_tags.centralized_tag_schemas` to manage an allowed list of repositories for tag definitions in data lakes. Acceptable values for this list are the Fully Qualified Domain Names (FQDNs) of schemas that organize tags, such as `SNOWFLAKE.CORE` (default). When enabled in conjunction with the feature flag `alation.feature_flags.enable_snowflake_tag_filtering_by_schema`, this feature flag allows tags from specific schemas to be visible on catalog pages, regardless of whether or not the tag's schema matches the object's schema. (AL-145407)

Compose

Query Result Streaming

It is now possible to execute queries in Compose without storing query results on the Alation server or other Alation-managed architectures. This feature is intended to facilitate cloud migration strategies for customers with strict security policies that prohibit storage of query result data on Alation systems although the data does transit through these systems to be displayed in the user interface.

Enablement

This feature cannot be activated via the user interface and requires involvement from Alation Support for Alation Cloud Service instances or backend access for on-premise instances. Additionally, scripts must be executed to erase any previously stored query data.

With Query Results Streaming enabled:

- Certain functionality will be disabled, such as query scheduling, result sharing and downloads, and preserving results indefinitely.

Note: In this release, the Run & Export feature for queries is also disabled, but support for enabling exports independently of enabling result streaming is planned for a future release.

- Query results are accessible only within the Compose user interface or on query form pages immediately after query execution. They will be lost after refreshing the browser or after a re-login. (AL-148201)

Alation Analytics

Alation Cloud Service on Cloud-Native Architecture

Data Quality Improvements

- As part of the `/installconfig/AlationAnalyticsV2/Connect3rdPartyBIToolsToAAV2` feature, the ordering of custom fields for all object types will now be in ascending order of the `field_id` field to ensure that custom field creation happens in the same order on PostgreSQL and Snowflake. Previously, when multiple custom fields had the same name but different data types, they were created in different order on PostgreSQL and Snowflake, causing some data inconsistency issues. (AL-145364)

ETL Stability Improvements

The enhanced Alation Analytics pipeline now utilizes the system disk (STORAGE) in conjunction with System Memory (RAM) for efficient ETL operations. This reduces the likelihood of Out Of Memory (OOM) errors, particularly when handling vast amounts of data. As a result, you can expect significant improvement of ETL stability, especially when regularly ingesting large volumes of objects. (AL-152601)

Logical Metadata Services

Alation Cloud Service on Cloud-Native Architecture

For Alation Cloud Service customers in Asia Pacific, Canada, and Europe, we are upgrading the underlying infrastructure for Logical Metadata Services (LMS) to improve performance and reliability. LMS is the service that powers titles, descriptions, and custom field values in the Alation catalog. The upgrade will make Custom Field Value API response times six times faster, reduce page load times for glossary terms by over 50%, and decrease the occurrence of 504 gateway timeouts when loading glossary terms and catalog pages. This change also lays the groundwork for future performance and scalability improvements across the Alation platform. We will upgrade customers in the United States in the near future.

As part of this improvement, the Upload Logical Metadata API will be deprecated in the 2024.3 (September) release of Alation. After that, metadata updated with this API will not be displayed correctly in the catalog. Please visit the [Alation Developer Portal for a detailed list of replacement APIs](#). (AL-148679)

Alation Public APIs

Find a short summary of the Alation API changes below. More details are available in the [2024.1.2 API Release Notes](#) on Alation's Dev Portal.

NoSQL API

The new [NoSQL API](#) allows creating, updating, and deleting NoSQL virtual data sources. The API relies on the optimized ingestion pipeline for NoSQL objects that enables uploading metadata of complex nested data types. (AL-148931)

- **Known Issue:** Using object name, path, or ID as request parameters in the Delete endpoint incorrectly deletes data from attributes with the same name under different schemas. The deletion should only affect the specified object and its child objects. (AL-155752)

Homepage Preferences API

The new [Homepage Preferences API](#) lets Catalog Admins and Server Admins set the homepage preference for multiple groups that are already assigned to a specific homepage. (AL-143659)

BI Source API

Added a new endpoint for [migrating curation data](#) after migrating a BI server from a native connector to an OCF connector. (AL-146464)

Custom Field Values API

The way you can get the Job ID of the PUT endpoint of the Custom Field Values Async API has changed. When [updating custom field values](#), this API now runs two background jobs. The API response returns the ID of the first job. Use that ID to query the Jobs API. The response from the Jobs API returns the ID of the second job. (AL-147696)

Queries API

When [creating a query via API](#), you can now set additional fields:

- Description
- Domains
- Tags (AL-139445)

Disable Public Unauthenticated APIs

Alation Cloud Service on Cloud-Native Architecture

Addressed a potential vulnerability in seven Alation Cloud Service monitoring APIs that allowed unauthenticated read access of the status information from the public internet. The following insecure monitoring endpoints are no longer publicly accessible:

- `monitor/num_active_users`
- `monitor/backup_status`
- `monitor/hostname`
- `monitor/replication`
- `monitor/aav2/status`
- `integration/v1/monitor/services_healthcheck`
- `integration/v1/monitor/admin_alert_checks` (AL-148305)

OCF Connectors

Updates to the Open Connector Framework (OCF) connectors are released every week. See </sources/OpenConnectorFramework/ReleaseNotes> for information on the updates in the month of April 2024. A few April highlights are:

New Connectors

- </sources/OpenConnectorFramework/CosmosDBOCFConnector/index> (AL-144006)
- </sources/OpenConnectorFramework/ServiceNow/index> (AL-144008)
- </sources/OpenConnectorFramework/SSASOCFConnector/index> (AL-146586)

Connector Settings User Experience Redesign

We've completely redesigned the user experience for Alation admins setting up OCF sources for the connectors listed below. Now, the configuration parameters are organized by steps, with new validation features added, such as testing connection and permissions to access the metadata. The extraction process has been fortified, and more clarifications and troubleshooting tips are now available in the user interface in case of extraction or ingestion errors.

- </sources/OpenConnectorFramework/SQLServer/index> (AL-139399)
- </sources/OpenConnectorFramework/PostgreSQLOCFConnector/index> (AL-139424)
- </sources/OpenConnectorFramework/TableauOCFConnector/index> (AL-137578)

Simplified Migration from Native to OCF Tableau BI Sources

For Tableau sources, the native to OCF migration will run as a downstream job for the first MDE after migration. Migration can also be done using API ([BI Source API](#)). (AL-146464)

Report Fields Extraction and Column-Level Lineage For Power BI Scanner OCF BI Sources

Starting in **2024.1.2** and connector version **2.2.0**, the OCF connector for Power BI supports extraction and cataloging of report fields (dimensions) and column-level lineage. A Server Admin can enable this functionality with a dedicated feature flag **Automatically extracts Column Level Lineage from Power BI Scanner data source** in **Admin Settings > Server Admin > Feature Configuration**. (AL-132853)

Bug Fixes

- Fixed an issue where exported term search results were missing custom fields. Now, exports will include all columns displayed in the search result table view. (AL-154073)
- Resolved an intermittent issue with the Search functionality on the **Queries > History** tab for schema objects' catalog pages. (AL-149787)
- Previously, the Explain feature in Compose wasn't functioning properly for multiple OCF connectors after enabling the Query Service feature flag. The issue was fixed and connector certification was performed to ensure that the Explain feature now functions as expected. (AL-143362)
- Resolved an issue where data dictionary uploads were failing after upgrading to version 2024.1. A content validation error in the data dictionary preview UI was corrected, allowing the upload of files that contain specific Unicode characters. (AL-152621)
- Resolved a bug in workflows where approving multiple changes for the same object led to only one change being applied. The system now supports multiple parallel updates to custom fields through workflow approvals. (AL-143259)
- Resolved an issue where People Set custom field values were not included in files downloaded using the Bulk Utility. (AL-150153)
- Resolved a Bulk Utility issue where a term's Last Updated property was not updating correctly on upload. The `ts_updated` property for terms will now be correctly updated after the Bulk Utility upload. (AL-148881)
- Resolved a bug affecting automated backups after upgrading to version **2023.3**. Previously, backups did not execute when unfinished jobs persisted in `jobs_job`. Added logic now identifies and invalidates stuck jobs; if a job is active for over a day without a corresponding Linux process, it's marked as failed, allowing new backup jobs to proceed. If these conditions aren't met, existing jobs continue, and new jobs are postponed. (AL-149526)
- Resolved a problem where notifications for failed jobs due to service or server outage were significantly delayed. Notifications will now be issued within 30 minutes of a job failure. (AL-148721)
- Resolved a problem with the dbt add-on connector where the absence of the `compiled_code` field in the dbt manifest file caused lineage streaming failures for a Snowflake data source. The connector now handles cases where this field is empty. (AL-147412)
- Resolved a bug in the Databricks Unity Catalog OCF connector where extraction failed due to the error *A name's original value can't be null or the empty string*. The lineage extraction queries have been updated to exclude empty column and table names from the result set. Fixed in connector version **2.2.2** (AL-141050, AL-140061)

- Fixed an issue with the 3rd party BI access to Alation Analytics on Snowflake, where counts for schemas in PostgreSQL and Snowflake for the `rdbms_datasources`, `rdbms_schemas`, `rdbms_tables`, `rdbms_columns` tables were different. The bug was caused by the NULL records on join columns. We are now using the COALESCE function along with default values on join columns in Snowflake to prevent the NULL values. (AL-145893)
- Now, during the installation of Alation Analytics, after installing the required Python packages using the Python Package Installer (pip), it's removed from the base image to prevent a potential security vulnerability. (AL-145515)
- Resolved a replication issue for the Alation Analytics table `rdbms_schemas` to Snowflake. This problem had multiple causes, including NULL values coexisting with other values in array column types and NULL records in join columns. All identified issues have been successfully addressed. (AL-140093)
- Resolved an issue where the user's display name during the LDAP login flow was not updated when `first_name` or `last_name` gets updated in the LDAP server (AL-87975).

1.2.16 RELEASE 2024.1.1

build 18.1.0.70857

Released April 4, 2024

Note: This release is available for all Alation instance types: customer-managed (on-premise) and Alation Cloud Service.

- [New User Experience \(Public Preview IV\)](#)
- [Alation Agent](#)
- [Domains](#)
- [Governance](#)
- [Lineage](#)
- [Alation Analytics](#)
- [Alation Public APIs](#)
- [Bug Fixes](#)

New User Experience (Public Preview IV)

Alation Cloud Service on Cloud-Native Architecture

The [New User Experience](#) has now moved into Public Preview IV. The following list captures a summary of the additional features and pages currently available in Public Preview IV:

- Data source catalog details page
- Schema catalog details page
- Table catalog details page
- Column catalog details page
- Catalog navigation
- Data source hub

- Domains Hub and details
- Homepage basic features and Admin guided user interface (GUI)
- Query page (with functionality to run the query in the catalog page)
- BI hub and all detail pages
- The dynamic template layout in the **Admin Settings > Customize Catalog > Custom Templates** for the BI object types:
 - BI datasource column
 - BI report column
 - BI datasource
 - BI server
 - BI report

Note: The dynamic template layout will be available for the Classic User Experience and the New User Experience after the update to the **2024.1.1 release**. All cloud-native environments (New User Experience environments) should be updated to this release by May 3, 2024.

- The ability for Server Admins to make the New User Experience default if desired
- Search (list and table view, saved search, export results)

Note: The **More** filter experience and other improvements are coming in the 1.0 release of the New User Experience.

- Conversations
- Tags object page
- Backreference group move-ability in templates
- Lineage (simple layout view)
 - The compound layout view and manual lineage are coming soon
- Dataflow object page
- Document hub basics
- Alation add-ons

Note: To participate in the New User Experience Public Preview, you must have Beta language contained in your contract or have a Beta addendum signed. Reach out to your Alation account manager for more information if you would like to access the Preview.

Alation Agent

Alation Cloud Service on Cloud-Native Architecture

For BI sources using Alation Agent in polling mode, an informational message has been added to the **Test Connection** dialog to inform admin users that the operation may require additional time to complete. (AL-141782)

Domains

Filtering Catalog Content Based on Domain: Alation now allows for a more personalized catalog experience with filtering the left-hand navigation based on users' domains. The feature is intended to provide quicker access to important information by filtering out data unrelated to domains. (AL-121118)

Governance

Create Multiple Table Templates

We have introduced the capability to create multiple templates for the table object type in **Customize Catalog > Custom Templates**. By default, this feature is disabled but can be activated with the feature flag **Enable creation of multiple templates for table catalog pages** in Admin Settings > **Feature Configuration**. Once enabled, both Catalog Admins and Server Admins can create multiple templates on the Table tab under **Customize Catalog > Custom Templates** and associate them with specific data sources. While a single template can be associated with multiple data sources, each data source is limited to one template. (AL-137912)

Redesign of the Govern Landing Page

In the **Apps** menu, the **Govern** tile was renamed to **Curate & Govern**. The landing page for Curate & Govern was redesigned and now includes tiles for each part of the Governance App. This change has a dedicated flag in **alation_conf**: `enable_governance_home_page_redesign` (enabled by default). (AL-147712)

Data Dictionary

Handling of blanks and invalid values: In a downloaded data dictionary, invalid cells are now represented by underscores, replacing the “N/A” and “null” values from the legacy process. In an uploaded data dictionary, blank values now trigger the “unset” operation for existing entries. The **Replace Existing Values** option also supports the “unset” operation—blank values in the upload will overwrite existing ones if selected; if not selected, existing values will be preserved. (AL-132472)

Lineage

Business Titles in the Compound Layout View: You can now choose to display business titles inside nodes on lineage charts instead of technical names. The option is located in the Filters dialog accessible on the top left of the chart. Business titles are sourced from the Title catalog field of an object. For objects with no Title, their technical names will be displayed. (AL-146600)

Alation Analytics

- Added the table `documenthubs` to the Alation Analytics database. Users can now retrieve document hub information and categorize and analyze their document and folder objects by specific document hubs. (AL-150170)
- Added a new `document_hub_id` property to the `terms` and `glossaries` tables in the Alation Analytics database. Alation Analytics users can leverage this field when retrieving information on a document hub (number of folders, documents, document templates, etc.) (AL-142164, AL-147417)
- The column `rdbms_datasources.uri` in the Alation Analytics database will now be populated for OCF data sources. This change specifically accommodates the URI for the GBQ and Hive data sources. (AL-125234, AL-151676)

Alation Public APIs

Find a short summary of the Alation API changes in **2024.1.1** below. More details are available in the [2024.1.1 API Release Notes](#) on Alation's Dev Portal.

New My Domains API

The new [My Domains API](#) allows Catalog and Server Admins to assign users and groups to domains and more. The API is part of the Federated Catalog public preview. To use it, enable the Enable Federated Catalog Features flag on the **Feature Configuration** page in Admin Settings. (AL-121118)

Note: To participate in the Federated Catalog public preview, you must have Beta language contained in your contract or have a Beta addendum signed.

Domains API

- Added endpoints to the [Domains API](#) to move a domain to a different parent and retrieve a single domain. (AL-121319)
- The Domains API can now recursively assign domain membership to an object and all its children. (AL-141426)

Cross-System Lineage API

The Cross-System Lineage API endpoints are being merged into the existing [Lineage API](#). The separate Cross-System Lineage API is being deprecated. (AL-145817)

OCF Data Sources API

- You can now [get](#) and [update](#) all metadata extraction options for an OCF data source. The additional options are accessible through a new `configuration_parameters` body parameter. (AL-148936)
- When updating metadata extraction options, you can now Base64 encode parameter values, enabling you to send multi-line SQL queries. (AL-149793, AL-148936)
- For the [Trigger metadata extraction job](#) endpoint of the OCF Data Sources API, we will be deprecating the three body parameters `limit_schemas`, `exclude_schemas`, and `remove_filtered_schemas`. (AL-148936)

Relational Integration API

- The [Relational Integration API](#) can now get and update STRUCT columns and child columns. (AL-152225)
- The Relational Integration API documentation was updated to indicate that the POST methods should not be used to update objects. (AL-146270)

Data Dictionary API

- Added an endpoint to [get data dictionary task errors](#) so you can see what might have caused a data dictionary upload to fail. (AL-150783)

Policy API

The [List Policy Groups](#) endpoint now works with a trailing slash. It will also continue to work without a trailing slash. (AL-145478)

Bug Fixes

Curate & Govern

- Resolved a problem where removing values from Object Sets bypassed the workflow process. Object Set and People Set fields now fully adhere to workflows. (AL-144777)
- Resolved a discrepancy in the display of Snowflake tag names between **Customize Catalog > Custom Fields** and **Customize Catalog > Custom Field Permissions**. Previously, tags shown with fully qualified names in the **Custom Fields** section were displayed with partial names in the **Custom Field Permissions** section. Now, tags are consistently displayed with their full names in both sections. (AL-144307)
- Fixed an issue where email notifications for custom field modifications were sent without considering the custom field's permission settings. Now, email notifications undergo a permission check so that only users with custom field access are notified of changes. (AL-147500)

Data Dictionary

- Fixed an issue where updating a custom field through data dictionary uploads wasn't reflected in field history. Field history will now be correctly updated when a custom field update is done through data dictionary upload. (AL-149122)
- Fixed an issue preventing the user interface from accepting "N/A" as a valid option for Picker custom fields during data dictionary uploads. "N/A" values are now supported for uploading in Picker fields. (AL-146164)
- Fixed an issue with the data dictionary upload where blank values led to inconsistent updates: Titles were updated while Descriptions were ignored when the **Keep Existing Values** option was selected. Now, both the Title and Description fields are updated. (AL-146220)

Alation Analytics

- Fixed a bug causing the ETL process for Alation Analytics to fail with an error *Must be a superuser to terminate superuser process*. Now, Alation no longer aborts queries run by super users in Postgres, as it's not allowed since version **13.13**. (AL-148020)
- Fixed an issue where the deletion of multiple terms using Bulk Utility wasn't reflected in the Alation Analytics database by subsequent ETL jobs, as the ETL checkpoint failed to recognize the deletions. Now, objects deleted through the Term and Document public API will not alter the `ts_updated` field, ensuring accurate tracking of deletions. (AL-134334)
- Resolved a bug where term IDs remained in the `glossaries.glossary_links` column of the Alation Analytics database after their removal from a glossary. This issue was addressed by modifying the term extraction query to consider the `ts_deleted` timestamp along with `ts_updated`. (AL-134089)

Terms, Glossaries, and Documents

- Resolved a bug with the Terms API where using the PUT method removed all glossary associations within a glossary instead of just the specified terms. Now, only the intended terms are removed. (AL-147499)
- Resolved an issue causing search facets to still display folders and document templates after disabling the Document Hubs 3 feature flag. With this fix, document hub-related search facets are no longer shown in search filters after disabling the Document Hubs feature. (AL-151800)
- Fixed an issue causing the document templates to still be displayed in the list of term types in **Customize Catalog > Custom Templates** disabling the Document Hubs feature. Now, only term types for terms will be displayed in **Customize Catalog > Custom Templates** if the Document Hubs feature is disabled. (AL-149245)

Rich Text Fields Formatting

- Fixed an issue that caused the loss of formatting when indented text, rows, or cells copied from MS Excel were pasted into the Rich Text editor for articles. Formatting is now preserved as expected. (AL-145495)
- Fixed an issue where during the editing of a lengthy article, adding a carriage return caused the cursor to jump to the top of the article. (AL-129702)

- Fixed an issue where formatting in Rich Text fields updated through a catalog set wasn't rendered correctly on catalog pages of objects in the set. Resolved by adding the necessary classes to properly render the formatting. (AL-103510)
- Resolved a rendering issue where the ampersand symbol (&) in object titles was not displayed correctly. (AL-71755)

The Connect Platform

- Resolved a bug causing cross-system lineage between two SQL Server data sources to link to soft-deleted records after metadata extraction. Cross-system lineage will now create correct associations with existing objects in the catalog. (AL-148764)
- Fixed an issue where partition keys (part of table metadata) were not extracted by the OCF connector for Impala on CDP. Updated table extraction logic in the connector code to include extraction of partition keys. (AL-147592)
- Fixed an issue with extraction from DB2 sources. Previously, metadata extraction failed with an error *MissingFormatArgumentException: Format specifier '%s'*. To fix the issue, the OCF SDK for the DB2 OCF connector was upgraded to version **4.4.0**. (AL-147543)
- Fixed an issue where metadata extraction failed for an SSRS OCF BI source with the *ArrayOutOfBoundException* exception. (AL-146517)
- Resolved an issue in the OCF SQL Server connector where the value title lookup was not functioning with dynamic profiling. The catalog attribute page's title lookup feature has been updated to support value title lookup when dynamic profiling is enabled. (AL-131988)

Other Fixes

- Fixed an issue where administrator users were unable to check the status of Alation services with the `alation_action status_alation` command after updating to version **2024.1**. This command now requires execution as the alation user. (AL-150411)
- Fixed an issue where after upgrading to **2023.3.5**, search export containing glossary term objects failed with an error *Failed to prepare download: 'NoneType' object is not iterable in the logs*. (AL-148667)
- Fixed an issue where switching object types while in the saved search view changed the column selection to default for the saved search. (AL-127725)
- Resolved an issue where canceling a query for a Hive data source prevented subsequent queries from being executed due to concurrent query checks. This was fixed by adding a cancellation check alongside the concurrent queries check. (AL-142466)
- Made a change to the internal server API forbidding direct POST requests to the URL `/data/0/upload_docs`. Now, valid otype and oid are required for this URL, which then enforces restrictions on file types. (AL-137774)

1.2.17 PATCH RELEASE 2024.1.0.2

Alation Cloud installer build 9.0.0.20165

Released April 5, 2024

Note: This release is available for Alation Cloud Service instances that haven't been migrated to the cloud-native architecture.

- Corrected the Datadog volume name which results from chroot jail removal to ensure the upgrade process functions correctly (internal architecture fix). (AL-152926, AL-152845)
- Added exception handling if Amazon RDS Custom Resource Definition (CRD) is missing (internal architecture fix). (AL-151792)

1.2.18 PATCH RELEASE 2024.1.0.1

Alation Cloud installer build 9.0.0.19815

Released March 26, 2024

Note: This release is available for Alation Cloud Service instances that haven't been migrated to the cloud-native architecture.

- Resolved a bug in the Open Connector Framework affecting the **Test access** operation for metadata extraction and query log ingestion. This issue impacted OCF connectors for Snowflake, Oracle, and MicroStrategy, where Test access failed displaying a `KeyError['podIP']` error in the user interface and an error *Stream closed with status: Status{code=UNIMPLEMENTED, description=Method not found: grpc.health.v1.Health/Check, cause=null}* in the logs. The problem was caused by an internal race condition when an operation was requested from OCF connectors. (AL-151828)

1.2.19 RELEASE 2024.1.0

build 18.0.0.66893

Released March 4, 2024

Note: This release is available for all Alation instance types: customer-managed (on-premise) and Alation Cloud Service.

- *New User Experience (Public Preview III)*
- *Document Hubs (Public Preview)*
- *Lineage*
- *Enterprise Readiness*
- *Authentication*
- *Alation Analytics*
- *Governance*
- *OCF Connectors*

- [Bug Fixes](#)

New User Experience (Public Preview III)

Alation Cloud Service on Cloud-Native Architecture

The [New User Experience](#) has now moved into Public Preview III. The list below captures a summary of the major features and pages currently available in the New User Experience:

- Data Source Catalog Details page
- Schema Catalog Details page
- Table Catalog Details page
- Column Catalog Details page
- Catalog Navigation
- Search List View
- DataSource Hub
- Domains Hub and Details
- Homepage basic features + Admin Guided User Interface (GUI)
- Query Page (without functionality to run the query in the catalog page)
- BI Hub and all detail pages

Note: To participate in the New User Experience Public Preview, you must have Beta language contained in your contract or have a Beta addendum signed. Reach out to your Alation account manager for more information if you would like to access the Preview.

Document Hubs (Public Preview)

This release introduces **document hubs**, designed to enhance your organization's ability to categorize and manage extensive, written content directly related to your data. Document hubs offer the flexibility to create customized documents, tailored to meet your organization's data documentation needs.

Document Hubs are set to eventually replace the current articles functionality.

Check out the [documentation for Document Hubs](#).

Key Features

- **Centralized data documentation:** Each document hub serves as a dedicated space within the catalog for organizing your written content.
- **Structured organization:** Document hubs consist of a hierarchical structure with hubs at the top level, followed by folders, and documents within those folders. This structure facilitates easy navigation and management of content.
- **Management of content:** Administrators have the capability to create and publish document hubs. Users can configure folders, add or remove documents, modify titles and descriptions, and utilize search and sorting functionalities. (AL-143825, AL-142897)

Note: Document hubs are available as part of Alation's public preview program. To participate in the Document Hubs Public Preview, you must have Beta language contained in your contract or have a Beta addendum signed. Reach out to your Alation account manager for more information if you would like to access the Preview.

Lineage

Lineage Compound Layout Is General Availability

The Compound Layout lineage view, initially launched into public preview in version 2022.3.0, is now a General Availability offering for all customers in all environments. The "Beta" label was removed from the Compound Layout view switch on lineage graphs.

Note: The Compound Layout view organizes objects on lineage graphs into groups under data sources, such as Snowflake or Tableau, enabling focused analysis of data movement across different platforms. Users have the capability to drill down from the data source to column level, tracing objects across data source groups and achieving a comprehensive end-to-end lineage view.

Limitation & Workarounds

There is still a performance limitation in the Compound Layout view feature, where lineage graphs with a large number of objects may experience slow initial loading times. We've prioritized efforts to resolve this issue in the upcoming releases.

If you run into this issue, you can use the following workarounds until the issue is resolved:

- **Adjust graph limits:** Ask your Alation Server or Catalog Admin to adjust graph limits and try reloading the graph with the Compound Layout disabled.
- **Clean up the lineage graph:** Disable the Compound Layout, and in the default legacy view, remove unnecessary lineage and dataflows either manually or via APIs. Enable the Compound Layout view and reload the graph.
- **Request a targeted node cleanup:** Reach out to Alation Support for targeted node cleanup.
- **Revert to the default view:** If a graph's loading time exceeds 30 seconds or fails, refresh the page, revert to the default view, and attempt to reload with the Compound Layout view disabled.
- **Try out optimized graph display:** Alation Cloud Service customers can request evaluation of a new beta feature for optimized graph displays. Contact Alation Support to have it enabled on your Alation instance.

Beta Features Within Compound Layout

Some elements of the Compound Layout view are still marked with a “Beta” label, as they remain in public preview:

- *Manual Lineage in Compound Layout (Public Preview)*
- *Data Trust Metadata Layer (Public Preview)*

These features have recently transitioned from private to public preview as we seek broader feedback before finalizing their release.

Manual Lineage in Compound Layout (Public Preview)

We’ve added an Edit button directly in the Compound Layout view of lineage graphs. The button enables users to manually create, publish, and, if necessary, delete lineage. The feature is still in public preview, as the user interface for previewing manual lineage and data source grouping is still under development. As a temporary measure, upon publishing manual lineage, users are redirected back to the compound layout to review their modifications within the established data source groups. Users also have the flexibility to refine their manual lineage as needed after reviewing the graph. (AL-145095)

Data Trust Metadata Layer (Public Preview)

We’ve enriched the lineage diagrams with the Data Trust Metadata Summary designed to provide users with quick access to key information about trust flags. By clicking the overlay icons like depreciation, endorsement, or warning symbols, users can view a pop-up summary containing asset descriptions, steward contacts, trust flag messages, and impact analysis. The feature is still in public preview (beta), with more details and additional metadata layers to come later based on user feedback. (AL-141586)

General SQL Parser Upgrade

The General SQL Parser’s JAR file was upgraded to version **2.8.5.8** to take advantage of the bug fixes and improvements in the newer version. (AL-143872)

Enterprise Readiness

Postgres Upgrade to 13.13

Upgraded Alation’s internal database (PostgreSQL) to version **13.13** to take advantage of new features and to fix vulnerabilities. (AL-143302)

Apache Kafka Upgraded to 3.6.1

Upgraded Alation's internal event streaming platform (Apache Kafka) to version **3.6.1** to take advantage of new features and to fix vulnerabilities. (AL-143661)

Chroot Removal

Alation Cloud Service on Cloud-Native Architecture

The Alation Cloud Service instances no longer use the Chroot-restricted environment in their builds. (AL-105007)

Note: This change does not impact customer-managed installations, which will remain unchanged.

Authentication

Alation Is Available Through the Microsoft Entra App Gallery

Alation is now included in the Microsoft Entra application gallery, enabling straightforward integrations with Microsoft Entra ID (formerly Azure AD) for SAML-based single sign-on (SSO). (AL-143594)

Reauthentication Due to Inactivity

Alation Cloud Service on Cloud-Native Architecture

Starting with this release, when a user session remains inactive for a predetermined period, it will be automatically terminated. Users receive a warning two minutes prior to session expiration. For FedRAMP deployments, the inactive timeout is fixed at 15 minutes and cannot be adjusted. For non-FedRAMP deployments, the timeout interval can be configured as follows:

- Default duration: 60 minutes
- Minimum duration: 5 minutes
- Maximum duration: 1440 minutes (24 hours) (AL-125476)

Cross-Account IAM Support for AWS Secrets Manager

The authentication support for AWS Secrets Manager has been expanded to include cross-account IAM roles, in addition to the existing methods of IAM user and IAM user+IAM role. The **Authentication Configuration Methods for External Systems** section of the user interface (**Admin Settings > Server Admin > Authentication**) has undergone some enhancements for AWS Secrets Manager. Only the relevant configuration fields are now displayed, ensuring a more intuitive user experience. (AL-119588)

Alation Analytics

Alation Analytics PostgreSQL Upgrade to 13.13

The Alation Analytics PostgreSQL database was upgraded to version **13.13** to take advantage of new features and to fix vulnerabilities. (AL-143302)

Information About Tables and Columns Linked to Data Policies Added to Alation Analytics

As of version **2023.1.7**, Alation Analytics has supported querying data policies directly from the `data_policy` table. This enhancement introduces a new table, `data_policy_link`, which provides detailed associations between data policies and specific table or column objects. This table can be particularly useful for administrators and analysts looking to understand the application of data policies at the granular level of tables and columns.

Columns

- `dim_load_id`, `dim_ts_created`, `dim_checksum`—ETL bookkeeping columns.
- `object_id`, `object_uuid`—Object ID in numeric and UUID formats respectively.
- `object_type`, `object_type_id`—Object type in the text (table, attribute) and number (27, 1) formats respectively.
- `id`—Primary Key for the table in Alation Analytics.
- `policy_link_id`—Internal ID for each policy link.
- `policy_id`—Policy ID.
- `ts_created`, `ts_updated`, `ts_deleted`—The created, updated, and deleted timestamps.
- `deleted`—Boolean to denote if a link is soft-deleted. (AL-139218)

Expire Old WAL Archives

A new scheduler will now automatically remove WAL archives of the Alation Analytics database that are older than 48 hours. This change is designed to save disk space by ensuring that WAL archives are regularly rotated and outdated files are efficiently purged. (AL-139378)

Transition from Redis to Amazon ElastiCache

Alation Cloud Service on Cloud-Native Architecture

Alation Analytics now utilizes cache from Amazon ElastiCache instead of Redis within the Alation core container. This is a significant enhancement in the Alation Analytics architecture, aiming to provide a better and more reliable service to Alation Cloud customers. (AL-141201)

Governance

Resolution of Duplicate Tag Names

The display of Snowflake tags on catalog pages (including schema, table, and column pages) has been enhanced for better handling of tags with duplicate names across different schemas:

- A new feature flag, `alation.feature_flags.enable_snowflake_tag_filtering_by_schema`, has been introduced to allow optional filtering of tags by schema within catalog pages (disabled by default).
- The **Custom Fields** section under **Customize Catalog** now displays schema information for each Snowflake tag within picker fields. This addition is intended to help users differentiate between tags with identical names but belonging to different schemas. (AL-138710)

Auto-Enablement of Policy Center and Snowflake Tags Backend Parameters

Starting with this release, when the feature flags for Policy Center and Snowflake Tags are enabled in **Feature Configuration**, the corresponding set of backend flags is also automatically enabled. This enhancement eliminates the need for additional manual backend configuration when enabling the Policy Center and Snowflake Tag features in the Alation user interface. The following backend flags are auto-enabled:

- `alation.ocf.mde.policy.enable_extraction`
- `alation.ocf.mde.policy.enable_ingestion`
- `alation.ocf.mde.policy.enable_search`
- `alation.ocf.mde.policy.enable_sync`
- `alation.ocf.mde.custom_field.enable_sync`
- `alation.ocf.mde.custom_field.enable_extraction`
- `alation.ocf.mde.custom_field.enable_ingestion` (AL-143847)

OCF Connectors

Updates to the Open Connector Framework (OCF) connectors are released every week. See </sources/OpenConnectorFramework/ReleaseNotes> for information on the updates in the month of February 2024.

A few February highlights are:

- </sources/OpenConnectorFramework/DynamoDBOCFConnector/index> (AL-143589)
- </sources/OpenConnectorFramework/Elasticsearch/index> (AL-127896)
- </sources/OpenConnectorFramework/TibcoDV/index> (AL-127897)
- </sources/OpenConnectorFramework/TIBCOspotfire/index> (AL-139985)

Bug Fixes

Governance & Glossaries

- Fixed an issue where uploading a data dictionary with custom field names in Japanese failed after upgrading to version **2023.3.3.1**. The original parsing limitation to ASCII characters has been expanded to accept Unicode characters, ensuring successful uploads for Japanese and other Unicode characters. (AL-141572)
- Fixed an issue where titles remained unchanged during a data dictionary file upload when the **Keep Existing Values** option was selected and the initial title and description values were blank. The fix ensures that titles are now updated correctly when the Keep Existing Values option is selected when the current titles are blank. (AL-146220)
- Removal of field from an object set is now honored via workflows. (AL-144777)
- Fixed an issue where manual catalog set synchronization failed with the error *Job is not running. It might have been killed or crashed unexpectedly or there might have been a system outage*. The fix is expected to lead to reduced job failures in catalog sets synchronization (AL-140909).
- Resolved an issue where changing the term template via the Bulk Utility did not update the template. The Bulk Utility now supports template updates effectively. (AL-143246)
- Fixed an issue in the Bulk Utility where the downloaded terms showed deleted objects in Object Sets. Specifically, when columns related to a term's Object Set custom field Related Objects were dropped and subsequently soft-deleted during metadata extraction, they would still appear as *[GONE]* in the field history, despite being absent in the user interface. The Bulk Utility has been updated to exclude deleted objects from Object Set values. (AL-136475)
- Resolved an issue introduced in version **2023.3.1.1** where HTML tags were visibly displayed in the glossary term pages within the **Related Terms** section. This fix ensures that HTML tags no longer appear on catalog pages for Object Sets. (AL-137238)
- Resolved an access issue where users with the Steward roles were unable to use the data dictionary public API. Now, when the `alation.feature_flags.disable_perm_check_on_upload_dd` feature flag is set to `False`, users in the Steward, Composer, Source Admin, Catalog Admin, and Server Admin roles are allowed to make data dictionary public API calls, including uploading and retrieving task details. (AL-144523)
- Resolved an issue where workflows reviewer selection was failing on specific instances when user tried to select group as reviewers (AL-145175)

Alation Analytics and Analytics Experience

- Resolved a synchronization issue affecting column objects' data between Alation and Alation Analytics, specifically for column titles and descriptions updated through data dictionary upload. To reload the title and description missing data, [/admins/Troubleshooting/ResetETLCheckpointsForRDBMS](#). (AL-143810, AL-146271)
- Fixed an issue in the `rosemeta_allvisits` table where updates for conversation visits were not accurately reflected. Enhanced the logging mechanism for tracking visits to the conversation page in the `public.visits` table within Alation Analytics. This fix ensures that all user visits to conversations are properly logged and updated. (AL-142986)
- Resolved ETL process failures that began following an upgrade to version 2023.3.2.1 and were manifested through the error *TransformManager.transform(): Column name not supported in Snowflake: ACCOUNT*. To address the incompatibility with Snowflake's reserved keywords used as custom field names (ACCOUNT), we've implemented a solution to rename the affected columns in tables to the

format <CUSTOM_FIELD_NAME>_<ID>, ensuring successful data loading into Snowflake without errors. (AL-142557)

- Fixed an issue where transform containers failed to restart in a healthy state after an unattended server reboot. This fix ensures a coherent and reliable startup of all necessary containers post-reboot. (AL-136558)

Lineage

- Resolved an issue where upstream column-level lineage links (CLL) failed to generate for certain tables with quoted column names, although downstream lineage remained unaffected. With this bug fix, we're ensuring upstream lineage visibility for columns with quoted names. (AL-140987)

General SQL Parser

- Fixed lineage parsing issues between dbt and Snowflake. The fix ensures lineage accuracy from constructs like the CONNECT BY subclause. (AL-144356)
- Fixed an issue where Google BigQuery published queries were not associating with their respective tables, resulting in their absence from the **Queries** tab. This fix ensures that Google BigQuery published queries correctly link to tables. (AL-141819)
- Resolved a parser error affecting Snowflake DDL views, where lineage for views defined and ingested by MDE was not displayed, although connections to related Tableau objects were visible. This update addresses parsing issues, including those related to the CONNECT BY subclause, ensuring accurate lineage representation for Snowflake DDL views. (AL-140906)

Connected Sheets

- Addressed the Excel login error encountered on Chrome, which was caused by changes to third-party cookie handling. (ACS-562)
- This fixes the issue where the connected sheets hub page redirects in loops because of chrome's third-party cookie changes (ACS-589)

Connect Platform

- Fixed an issue with Hive data sources where, after canceling a query within a session, subsequent queries failed to execute due to concurrent query checks. This fix introduces a cancellation check alongside the concurrent queries check, ensuring smooth execution of subsequent queries after a cancellation event. (AL-142466)
- Fixed an issue where profiling failed for an AWS Databricks data source with the following error: *The connector was provided with a list of columns to profile from a table, however the schema of the returned result set did not match those columns.* The solution skips tables that cause failures and continues with the sampling process, ensuring profiling tasks complete successfully. (AL-141595)
- Addressed an issue where multiple installed versions of a connector on an Alation instance were indistinguishable in the user interface due to the absence of version numbers. This led to confusion over which version was being selected. Now, in case multiple connector versions are installed, the dropdown shows the connector version next to its name. (AL-146195)

- Resolved an issue where uploading data via XLSX files to native and OCF Redshift data sources failed with the error `AttributeError: 'NoneType' object has no attribute 'columns'`. Users can now successfully use XLSX files to upload data to both new and existing tables. (AL-138753)

Other Fixes

- Fixed an issue where content such as @-references, links, and photos included in user profile descriptions disappeared upon saving. The fix ensures that all content added to profile descriptions is saved. (AL-143451)
- Resolved a deadlock issue in uWSGI worker threads introduced by the upgrade from Python **3.6** to **3.9**. The deadlock occurred when threads loaded a Django application alongside the Python module traceback. From version 2024.1 onwards, this is addressed by disabling traceback by default. (AL-141594, AL-145955)

1.3 Release Notes 2023.3 (General Availability)

Note: The release notes on this page reflect changes to the Alation application. If you're looking for OCF connector release notes, jump to </sources/OpenConnectorFramework/ReleaseNotes>.

1.3.1 RELEASE 2023.3.5.2

build 17.5.2.73574

Released April 26, 2024

Note: This release is available for customer-managed (on-premise) Alation instances.

- Addressed an issue of performance degradation during search data ingestion. This issue was caused by a recently added attribute in one of the functions that led to a large number of unnecessary queries against the Alation database. The issue was addressed by query optimization and enhanced caching. (AL-154844)

1.3.2 RELEASE 2023.3.5.1

build 17.5.1.68029

Released March 11, 2024

Note: This release is available for customer-managed (on-premise) Alation instances.

- Resolved a deadlock issue in uWSGI worker threads introduced by the upgrade from Python 3.6 to Python 3.9. The deadlock occurred when threads loaded a Django application alongside the Python module traceback. The issue was addressed by disabling traceback by default. (AL-145955)
- Fixed an issue in the `rosemeta_allvisits` table where updates for conversation visits were not accurately reflected. Enhanced the logging mechanism for tracking visits to the conversation page in

the `public.visits` table within Alation Analytics. This fix ensures that all user visits to conversations are properly logged and updated. (AL-142986)

1.3.3 RELEASE 2023.3.5

build 17.5.0.64295

Released February 6, 2024

Note: This release is available for all Alation instance types: customer-managed (on-premise) and Alation Cloud Service.

- *The Connect Platform*
- *Alation Public APIs*
- *Search*
- *Navigation*
- *User Management*
- *Renaming of Glossary Terms Feature Flags*
- *OCF Connectors*
- *Bug Fixes*

The Connect Platform

Catalog Visibility Management for BI Sources

To provide better control over who can view BI sources in the catalog, we've introduced the ability to manage catalog visibility for BI sources and a function of the *BI Server Admin*, similar to that of the *Data Source Admin* for RDBMS data sources. Users with the Server Admin role can grant the *BI Server Admin* access to themselves and other users. A BI Server Admin can access the settings of a BI source and configure its visibility on the new **Access** tab:

- **Public BI Server**—Visible to all catalog users
- **Private BI Server**—Visible to specific catalog users or groups designated by a BI Server Admin

By default, a newly created BI source is public. (AL-124106)

AWS PrivateLink Support for Amazon S3 OCF Sources

Alation Cloud Service on Cloud-Native Architecture

AWS PrivateLink connections are now supported from the Amazon S3 OCF connector to an S3 bucket located in the customer's VPC for VPC's in the same AWS region. (AL-118070)

Alation Agent Polling Mode

Alation Cloud Service on Cloud-Native Architecture

Alation Agent can now be configured to use polling. In the polling mode, users will see an informational message that some operations may take longer time to finish. The Agent dashboard in **Admin Settings > Manage Connectors** will indicate whether the Agent is configured to use the standard or the polling mode. (AL-138169)

Alation Public APIs

Access to the Custom Field Value API for Catalog Admins

Alation users with the Catalog Admin role can now use the [Custom Field Value API](#) endpoints. (AL-140940)

Ability to Manage BI Source Visibility via API

You can now make BI servers private or public when [creating a BI server](#) and [updating a BI server](#) with the API. When [getting a BI server](#), the response indicates whether the server is private. This API accompanies the new feature *Catalog Visibility Management for BI Sources*. (AL-142964)

New API to Manage Data Source Connections for a BI Server Source

Added [new API endpoints](#) allowing users to get, create, update, and delete the data source connections information for a BI server source. (AL-141305)

Terms API Can Filter by Glossary and Provides More Info

When [retrieving terms via API](#), you can now filter terms by which glossary they belong to. The response includes which glossary the terms belong to, any custom fields associated with the term, and the ID of the term's template. (AL-111871)

Fixes to Data Dictionary API and Articles API

- Fixed a bug in the [Data Dictionary API](#) where a trailing slash was included in the documentation for the GET endpoint, which was causing API calls to fail. (AL-142026)
- The [Articles API](#) now returns the correct error code (401) when a call is made with an invalid token. Previously it returned the 403 error code. (AL-134789)

Search

Filtering by Catalogs

We've introduced a new search facet **Catalog** that can be used to filter metadata by the catalog object. The Catalog facet is applicable to schemas with multipart names ("catalog.schema"), such as Google BigQuery. Users with the Server Admin and Catalog Admin roles can add the Catalog facet to the left-hand filter panel in **Admin Settings > Catalog Customization > Customize Search Filters**. (AL-122598)

Navigation

Control Users' Ability to Select Preferred Homepages

Server Admins now have the capability to control whether or not catalog users can select their preferred homepages. We've added a new feature flag `alation.feature_flags.homepage.enable_admin_assigned_homepage` to **alation_conf** to control the ability for users to set a preferred homepage from the **My Account** menu.

- By default, the flag is set to `True`, which means only administrators assign homepages to users.
- When the flag is set to `False`, users will have the ability to designate a preferred homepage.

The feature flag provides organizations with flexibility to determine if they want users to select from and work with multiple homepages. (AL-145020)

User Management

User Sign-Up Moderation Enabled by Default

Previously, user sign-up moderation was disabled by default on customer-managed and Alation Cloud Service instances that are not on the cloud-native architecture. From this release onward, user sign-up moderation is enabled by default on all new installations, both customer-managed and Alation Cloud Service, and on the cloud instances that are newly migrated to Alation Cloud Service on the cloud-native architecture.

On subsequent upgrades, the user sign-up moderation setting will be preserved.

Important: If your current instance is customer-managed or an Alation Cloud Service instance that hasn't been migrated to the cloud-native architecture yet and you have never changed your user sign-up moderation preferences, then expect this setting to become enabled by default after updating to 2023.3.5.

If you've changed this setting before, your previous configuration will be preserved.

We recommend that user sign-up moderation should be enabled by default on all instance types to adhere to security best practices. Please review your user sign-up moderation settings after updating to 2023.3.5 and update the setting as per your organization's policies. You can find the user sign-up moderation settings under **Admin Settings > Authentication > User Signup Moderation Preference**. (AL-141813)

SCIM Integration Enhancements

We introduced a new read-only attribute groups in the GET request response to the SCIM user detail request. The SCIM user endpoints for retrieving either a list of users or a specific user will now return the group information in the response as a read-only attribute. This makes Alation's SCIM API compliant with RFC 7643 and allows supporting custom-developed workflows that rely on this attribute. (AL-140681)

Renaming of Glossary Terms Feature Flags

In preparation for the public preview of the Documentation Hub feature that is coming up in the major 2024.1 release (March 2024), the feature flags enabling the Glossary Terms feature have been renamed in **Admin Settings > Feature Configuration**:

- **Glossary Terms Setup** has been renamed to **Document Hubs 1: Setup** (disabled by default)
- **Glossary Terms GA** has been renamed to **Document Hubs 2: Show Glossary Hub** (disabled by default)
- **Glossary Terms Performance** has been renamed to **Improvement Document Hubs 4: Performance Improvement** (enabled by default)
- **Glossary Terms Total Count** has been renamed to **Document Hubs 5: Document Total Count** (disabled by default)

In addition, there is a new feature flag **Document Hubs 3: Enable Hub Creation (Public Preview)** (disabled by default). Customers are not advised to enable this feature in 2023.3.5. The enablement guidelines will be available with the Document Hub Public Preview information. (AL-133300)

OCF Connectors

Updates to the Open Connector Framework (OCF) connectors are released every week. See </sources/OpenConnectorFramework/ReleaseNotes> for detailed information on the connector updates in the month of January 2024.

Some January highlights include:

New Connector Releases

- [Fivetran add-on OCF connector \(GA\)](#)
- [Kafka OCF connector \(public preview\)](#)
- [Mongo DB OCF connector \(public preview\)](#)

Connector Settings User Experience Redesign

We've completely redesigned the user experience for Alation admins setting up OCF sources for the connectors listed below. Now, the configuration parameters are organized by steps, with new validation features added, such as testing connection and permissions to access the metadata. The extraction process has been fortified, and more clarifications and troubleshooting tips are now available in the user interface in case of extraction or ingestion errors.

- [Google BigQuery OCF connector](#)
- [Oracle OCF connector](#)

- Databricks Unity Catalog OCF connector
- Redshift OCF connector

Other Updates

Cross-Data-Source Lineage for MicroStrategy BI Sources

Previously, only BI lineage was supported by the MicroStrategy OCF connector. Cross-system table-level lineage was not supported due to unavailability of the host and port information via the MicroStrategy REST API. As a result, Alation couldn't trace a report or document back to the upstream data source.

Now, we leverage the Alation API to input host and port for the data sources where it is not available via the MicroStrategy API, providing end-to-end table-level lineage for in-memory cubes.

OCF Tableau Enhancements

Native to OCF migration is available for Tableau sources:

- The migration is enabled through both the user interface and a curation migration script.
- The script now has improved logging.

OCF Oracle E-Business Suite 12.2.12 Certification

The OCF connector for Oracle has been quality-certified for Oracle E-Business Suite version 12.2.12. (AL-133313)

Alation Agent Certification with Compose for More Connectors

Compose is now available for the following connectors installed on Alation Agent (AL-139438):

- OCF connector for Impala on CDH
- OCF connector for Impala CDP

Bug Fixes

Data Dictionary

- Fixed an issue where the field editing history is not updated when a user updates fields using a data dictionary upload. Now, the field editing history for a field appears in the catalog after uploading the data dictionary. (AL-141644)
- Previously, while parsing the data dictionary header, the grammar was restricted to only ASCII characters, so Unicode characters were rejected by throwing an error message. We have fixed this behavior by modifying the grammar to include Unicode characters too. For example, now the upload data dictionary feature works for Japanese characters and other Unicode characters. (AL-141572)
- Fixed an issue where the data dictionary upload using a CSV file failed in the Firefox browser with the error Invalid file format. Only CSV or TSV files are supported. Now, users can upload both TSV and CSV data dictionary files on any version of Firefox, Chrome, and other supported browsers. (AL-141001)

Alation Analytics

- Increased the consumer timeout for RabbitMQ to two hours to ensure that ETL does not truncate large tables. (AL-135579)
- Any flag object extracted into Alation Analytics from the Alation database with the `is_for_temp_object` column set, will be skipped by the queries. Going forward, we are excluding temporary objects from being stored in Alation Analytics as they do not reference objects stored in the Alation database. (AL-135191)
- Ensured that Alation Analytics now does not use the Docker 4.7 update instructions alone in Dockerfiles by updating the base image to use a newer version of RHEL 8 which doesn't have the update command alone as instruction. (AL-133173)
- Added a new column `groupprofile_id` to the `alation_group` table. This table will hold the ID of the group profile related to the authentication group. This will ensure consistency between Alation Analytics and the custom fields in the catalog UI, like the Steward field. (AL-128206)
- Fixed the issue that caused the following error during ETL: *Failed to insert records into stg_bi_datasource_columns due to psycopg2 error: null value in column "bi_ds_fp" of relation "stg_bi_datasource_columns" violates not-null constraint*. Added a not-null constraint to the `bi_ds_fp` column in the `(stg_)*bi_datasource_columns` table. (AL-127789)
- Added support for null values in the `ts_started`, `ts_updated`, and other fields to prevent failures of the ETL job `StewardshipBulkActionTransformJob`. (AL-137837)

Governance

- Fixed an issue causing pod eviction. Optimized memory usage by the Celery service of the downstream MDE jobs for OCF Snowflake data sources that handles ingestion of Snowflake tags and tag links. (AL-140666)
- Fixed an issue where workflow reviewers were unable to suggest changes, receiving an error message: *Your suggestion couldn't be created at this time. The feature is disabled or you do not have sufficient permissions*. (AL-133561)
- Fixed an issue where an initial extraction from a BI source did not display some of the extracted reports in the user interface, with those reports appearing in the catalog automatically a few days later. The issue was caused by a Celery queue for long-running tasks. Addressed by optimizing how processing queues are organized for long-running stewardship tasks. (AL-127200)
- Optimized the downstream jobs that run during MDE for OCF Snowflake data sources. In addition, the job that handles the clean-up of deleted Snowflake tags will now skip the deletion of the Snowflake tag custom field if the new **alation_conf** flag `alation.feature_flags.disable_snowflake_customfield_deletion` is enabled (disabled by default). This change decreases memory usage during MDE and shortens MDE duration. (AL-140666)

Other Fixes

- Fixed an issue that prevented the installation and upgrade on FIPS-enabled operating systems. (AL-135158)
- Fixed an issue where the user moderation setting became enabled after each Alation upgrade since migrating the instance to the cloud-native architecture. Now, the preference set by a Server Admin will be preserved. (AL-137401)
- Addressed a security vulnerability in the Golang packages by upgrading to their latest version. (AL-143662)

1.3.4 PATCH RELEASE 2023.3.4.1

build 17.4.1.63336

Released January 24, 2024

Note: This patch release is for Alation Cloud Service instances on the cloud-native architecture.

This patch includes enhancements to the Alation Analytics offering on the cloud-native architecture. It isn't universally available in all AWS regions. If your Alation Cloud Service instance is not planned for this update, the patch may not be applicable to your instance. (AL-139775)

1.3.5 RELEASE 2023.3.4

build 17.4.0.60909

Released January 4, 2024

Note: This release is available for all Alation instance types: customer-managed (on-premise) and Alation Cloud Service.

Alation Anywhere

Alation Anywhere for Microsoft Teams (Cloud Native Architecture only)

Microsoft Teams application users can now unfurl a conversation link from Alation. Users can quickly start a new conversation from Microsoft Teams per a specific object, reply to a conversation, or mark the conversation as resolved. (AL-124015)

Alation Analytics

Alation Analytics Dashboard Improvements

- On the ETL Status Dashboard an admin can mark all the ETL failure messages as read using the Mark all as read button. In addition, an admin can refresh the ETL Status Dashboard without reloading the tab using the Refresh Table button. (AL-138679)
- On the Top Contributors Dashboard, if a user no longer exists in Alation but does exist in the Alation Analytics data source, then an asterisk will be appended to the user's name in the Top Contributors Dashboard. (AL-138679)

Enable Third-Party Metrics in Alation

Naming and default values for feature flags that enable outgoing requests to Pendo and other third-party analytics tools have been made more consistent. This is enabled by default. Each third-party analytics tool has a dedicated feature flag that customer admins can enable or disable from the feature flags admin page. (AL-140552)

Alation Analytics Backup and Restore Improvements

Improved overall performance of Alation Analytics backup and restore. (AL-133755)

Improvements to Alation Analytics for BI Access

- Fixed an issue where the initial Alation Analytics data replication to Snowflake failed for certain tables which had the `NULL` value among the values in the array-type columns. (AL-133055)
- The script `copy_data_to_snowflake.py` now accepts the arguments which can be used to copy only the required tables. The `time_period` table is now also part of the initial data copy process to Snowflake. (AL-141482)

Compose Improvements

Ensured that the Compose Desktop menu item **Visit Help Center** Cloud Authorization redirects to the help center page as intended. (AL-137205)

Lineage: Alphabetical Sorting of Columns in Compound View

The column objects in the Compound Layout of lineage charts are now sorted alphabetically. This sorting facilitates tracking of data movement between objects in both full and column-level lineage views. (AL-137732)

Cloud Authorization

SCIM Token Generation Warning

Users will now receive a warning before generating a new SCIM token when a previous token already exists (expired or not). The warning will indicate that the previous token will be invalidated. (AL-138329)

Data Dictionary Improvements

Upload Data Dictionary improvements include: (AL-138675) - You can now link Terms and Policies to data objects using Upload Dictionary. Terms and Policies links will also be downloaded in the object set field as part of the data dictionary file.

- Period characters are now supported. The downloaded dictionary will have key_name quoted as required by the upload feature when names contain a period character.
- Non-admin users can now access upload dictionary without a page refresh.
- Enhanced object reference presentation in the data dictionary preview table. Previously, the title appeared as the text of the link within the object reference column. Now, when the key is available, it displays as link text; however, if the key is unavailable, the object's title will be shown as text in the link.
- Data Dictionary preview is now able to render fields with a comma in names.

Glossary and Articles

Users will now see the **Stewards** field on all Glossary details pages. (AL-133299)

Governance Improvements

Policy ingested timestamp for policies that were ingested using the Native Snowflake Connector will be back-filled with the created timestamp. (AL-138443)

Group-based Homepages

Catalog admins can assign a homepage to one or more user groups. (AL-134390)

- Users that are part of the assigned group(s) that have a homepage will be able to view this homepage when they log in to the application.
- Users that are already using role-based homepage associations will automatically be migrated to the group-based scheme.
- If a user is part of multiple groups that have different homepages assigned, the user will be able to choose a preferred homepage.
- When a user is not part of any group that has a homepage assigned or they haven't chosen a preferred homepage, they will be shown the Alation default homepage.

Search Enhancements

Automated the task to analyze and update search index data. The task is configurable using `alation_conf` (to enable, disable, and schedule the task). (AL-129324)

APIs

OCF Data Sources API Compose Support

Enhanced the API to be able to set up compose for data sources using OCF connectors. (AL-130098)

Bug Fixes

Cloud Authorization: Bypass User Email Confirmation for SSO-based Authentication Flows

A recent update from Python 3.6 to 3.9 caused empty passwords to be handled differently; as a result, users or requests provisioned in the system without a password and not a confirmed email address were redirected to the verify email page even if the password is not required for those users (like in the case of SCIM). Now the system will set non-empty passwords for all the users during the user creation flows to ensure that users won't get redirected to verify their email if they are not password users. (AL-139068).

Glossary and Articles

Removed an errant "b" that was prepended to the title slug in Article URLs. (AL-138356)

KVstore Improvements

Some users are affected by a memory hike on the KVstore and need to regularly reboot KVstore. The KVstore has been upgraded to 8.3.2 and will now consume less memory. (AL-142108)

1.3.6 PATCH RELEASE 2023.3.3.2

build 17.3.2.58486

Released December 13, 2023

Note: This release is available for customer-managed (on-premise) Alation instances.

- Fixed an issue that prevented users from connecting to their OCF Impala data sources in Compose after transitioning from the Hive driver to the Cloudera Impala driver within the CDP connector. The corresponding driver in Compose was updated accordingly. Users can now connect and run queries as previously experienced. (AL-138914)

1.3.7 PATCH RELEASE 2023.3.3.1

build 17.3.1.58212

Released December 5, 2023

Note: This release is available for Alation Cloud Service instances on the cloud-native architecture.

- Fixed a bug where a new Python-based wrapper script that was implemented to support additional Datadog features caused a bug which resulted in no logs in Datadog from any of the affected clusters. Now, the log availability in Datadog has been restored. (AL-140466)

1.3.8 RELEASE 2023.3.3

build 17.3.0.57463

Released December 4, 2023

Note: This release is available for all Alation instance types: customer-managed (on-premise) and Alation Cloud Service.

Compose

Past Scheduled Executions Dashboard for Admins

A new dashboard is now available as a second tab next to the existing scheduled query admin dashboard that provides a tabular view of past scheduled executions. This feature aims at helping Server Admins understand the patterns of scheduled queries on their instance. They will be able to monitor if any schedules fail, take significant execution time, or were canceled. (AL-128832)

Improved Query Export Scalability

Previously, the Run and Export actions sometimes failed for smaller exports and consistently failed for exports of over 1M records. The Run and Export functionality should now be stable for both small and large exports. (AL-125680)

Data Dictionary Improvements

- Introduced the feature flag `alation.feature_flags.disable_perm_check_on_download_dd` to disable permission checks when downloading data dictionaries. Permission checks may take a significant amount of time, which in some cases can delay the delivery of the download URL to users by a few hours. This flag allows disabling the permission checks from the backend. By default, the flag is set to `False`, which means the **Download Dictionary** feature will check the permissions and the download option will be available for users with the Composer role and above. Setting this flag to `True` will disable the permission checks and the **Download Dictionary** option will only be visible to Catalog Admins and Server Admins. (AL-138107)

- Data Dictionary download now includes a new column, **al_datadict_column_data_type**, which contains the data type for each column. This information was previously contained in the **al_datadict_item_properties** column, but has now been broken out on its own. This column should be considered read-only.
- Added support for custom field names in upper or mixed case in the data dictionary file. (AL-137003)
- Added support for custom field names with a comma in headers. (AL-138310)
- A summary is now displayed for the uploaded dictionary on the Upload Dictionary screen after the upload dictionary process is completed. (AL-136514)
- The data dictionary upload now accepts the TSV file format. (AL-136379)
- Previously, we relied on object keys (names) to link objects using object set fields in the upload dictionary flow, which caused some performance challenges and issues identifying the right objects in case of duplicates. To address this, we added support for the object ID in the format `objectId | ObjectKey(name)`. You can now provide the object ID or the name or both to link objects via the dictionary upload. When downloading the data dictionary file, we will include both the object id and the object key for each object set field value. (AL-127462)

Alation Analytics

- Added a “no special characters” password policy for the Alation Analytics database password: the password can't contain any special characters. (AL-138780)
- Now you can skip load tables from the subsequent ETL runs using a new **alation_conf** flag `alation_analytics-v2.etl.skip_load_tables`. This setting can unblock all other ETL jobs if some jobs are failing. To add tables to this flag, use the following syntax: `alation_conf alation_analytics-v2.etl.skip_load_tables -s '["users", "query"]'`. To reset the flag, use the following syntax: `alation_conf alation_analytics-v2.etl.skip_load_tables -s '[]'`. (AL-128081)
- **Cloud-native architecture** The `alationanalyticsv2_entitlement_counts` table has been renamed to `license_u`
 - `num_licenses`—The number of licenses available.
 - `pct_utilization`—The number of users as a percentage of available licenses.
 - `pct_adoption`—The number of active users as a percentage of available licenses.

This table is also now replicated to Alation Analytics on Snowflake. (AL-136878)

Note: The `license_usage` table is populated on Alation instances on Cloud-native architecture only. On customer-managed instances, this table is present but not populated.

Glossary Improvements

The `enable_glossary_terms_performance` feature flag, which controls the performance enhancement for Glossaries, now defaults to enabled. The flag still remains in the flag list on the **Feature Configuration** page. If you had previously toggled it on and then off again, it will remain disabled after the upgrade to 2023.3.3. You can manually toggle it back on again to experience a performance improvement. (AL-136812)

Absolute User Session Timeout (Cloud-Native Architecture only)

For security reasons, users will no longer be able to stay logged in for an indefinite period of time. A user session will be terminated forcing the user to re-login once the time limit for session duration is reached. The absolute user session timeout is configurable on the Authentication page of Admin Settings, in the **Session Timeout Settings**, using the field **Maximum Session Duration (in hours)**. This feature applies only to customers on the cloud-native architecture. (AL-123328)

Lineage

Enhanced Edge Highlighting

Improved user focus and navigation within lineage charts by fading out nodes and paths not directly connected to a selected object's upstream or downstream paths. This visual enhancement helps users concentrate on the most relevant parts of the lineage. (AL-134668)

Alation API

Custom Field Value APIs

- Added support for `api_resource`, `api_resource_field`, and `api_resource_folder` object types in the Public Custom Field Value API and Public Custom Field Value Async API (Refer to [Custom Field Value Overview](#) on Alation Developer Portal). (AL-134107)
- Fixed an issue with the Public Custom Field Value API where validation errors were not thrown correctly when using an unsupported object type (AL-134176).

Public API for SCIM Token Generation

- Added support for the API endpoint for SCIM token operations (Refer to [SCIM 2.0 API](#) on Alation Developer Portal). (AL-136805)

Cloud Infrastructure

Cloud-native architecture

- Re-enabled the admin UI for viewing active tasks which was previously disabled when this service launched due to some issues. (AL-124848)
- The License dashboard page now shows the contract expiration date. Previously, the license expiration date was displayed. In Alation, licenses are renewed every year until the contract expires. (AL-137415)
- ElastiCache is now in use on instances on the cloud-native architecture. Alation components leveraging ElastiCache have better performance by storing and retrieving frequently used data within ElastiCache. (AL-107135)

Bug Fixes

- Fixed an issue where after the upgrade to 2023.3.1, the query forms user experience for the Viewer role was hampered when the Obfuscate Literals setting was enabled (all literals became unavailable in the query form result). This fix returns query forms behavior to align with the behavior in Compose where a user with run (or edit) permissions on a query would see literals, even when literal obfuscation is turned on, as opposed to a viewer who would not see the literals. A user with only view permissions on a query would not see literals on any Alation page when obfuscation is turned on. (AL-136260)
- Fixed an issue when the dynamic multi-select widget did not show any error messages on query forms. (AL-137032).

1.3.9 PATCH RELEASE 2023.3.2.2

build 17.2.2.66531

Released February 28, 2024

Note: This release is available for customer-managed (on-premise) Alation instances.

- Resolved a deadlock issue in uWSGI worker threads introduced by the upgrade from Python 3.6 to Python 3.9. The deadlock occurs when a thread attempts to load a Django application alongside the Python module traceback. The issue has been addressed by disabling traceback by default. (AL-145955)

1.3.10 PATCH RELEASE 2023.3.2.1

build 17.2.1.55993

Released November 8, 2023

Note: This release is available for all Alation instance types: customer-managed (on-premise) and Alation Cloud Service.

- This patch release includes platform updates to support the enhanced user experience for configuring OCF connector settings, which will be available in upcoming connector releases, starting with the Snowflake OCF connector. (AL-137533)
- Fixed a Bulk Utility issue where an internal .csv file would become corrupted when users clicked the **Download** button on the Bulk Utility page. This issue occurred where the initial use of the Bulk Utility involved clicking **Download** as opposed to selecting a specific saved search, resulting in the generated .csv file missing the Search column and leading to a 500 error. (AL-136141)

1.3.11 RELEASE 2023.3.2

build 17.2.0.55403

Released November 1, 2023

Note: This release is available for all Alation instance types: customer-managed (on-premise) and Alation Cloud Service.

Alation Analytics

Basic Authentication for RabbitMQ

The RabbitMQ component of Alation Analytics now requires basic authentication for incoming connections from other Alation components. The RabbitMQ password is auto-generated during the installation or upgrade of Alation Analytics. The password needs to be set in **alation_conf** via the `alation_analytics-v2.rmq.config.password` parameter before initializing the Alation Analytics database.

This is a one-time setup for customer-managed instances of Alation. Alation will perform this setup for Alation Cloud Service customers.

Object and User Counts

Object and user counts are now stored in Alation Analytics in the `alation_analytics_v2_entitlementcounts` table. The table is updated daily with a snapshot of all object usage across an Alation instance.

Note: For Alation Analytics Snowflake users: This data is not replicated into Snowflake in this release, but it will be synced starting from the next release (the November release).

Improvements for Third-Party BI Access via Snowflake

Cloud-native architecture

Snowflake Password Requirements Enforced

If you have a connected Alation Analytics database in Snowflake, your Alation Analytics password must meet Snowflake's password requirements. This is now enforced in the Alation user interface under **Account Settings > DB Connections**. To create or update the password for the Alation Analytics data source, users will need to follow certain password policies imposed by Snowflake (a lowercase alphabet, an uppercase, a digit, and the length of eight characters).

Clear Analytics Data - Fixed

The Snowflake database was not getting cleared when clearing up Alation Analytics using the **Clear Analytics Data** button. Now, the cleaning up of the database will be done asynchronously as part of a Celery task, and the button click will post a message to Celery and return success. (AL-125158)

Alation Anywhere

Alation Anywhere for Slack now shows rich previews for glossaries and terms.

Authentication

We have updated and expanded audit logging for important security-related events, such as updates to authentication settings, authorization changes, and user and group management updates. Authentication logs are written to **alation-events.log**.

Compose Enhancements

- You can now sort queries alphabetically by name in ascending and descending order.
- Compose now deletes blank queries after closing the query tab. In addition, on Sunday, Compose automatically deletes any blank queries from the past week.
- Embedded queries can now be executed in query forms.
- When query results are exported or downloaded in any way, Alation logs the event for auditing purposes. The logs are recorded to the `alation-events.log` file.
- Fixed an issue where the **Connection Settings** dialog didn't open when attempting to run a query with a connection error.
- The Query API now respects configuration settings that prevent query results from being exported or downloaded.
- Multi-select parameters in query forms can now have their value options defined with a pre-query. Once the pre-query has been run, query form users can select values based on the results of the pre-query. For usability purposes, pre-query results have an upper limit of 100 result items.

Connected Sheets OAuth Enhancements

- Roles are kept separate for the username and SSO credentials so users will not see roles auto-populated inconsistently between the different credential forms.
- Error messages have been moved to the top so that they are visible without having to scroll.
- A “Connected” badge was added showing the credential that is being used.
- The credentials form now automatically expands and collapses.

Governance

Data Dictionary upload has a new, streamlined interface and now runs asynchronously. Users should avoid quoting header names containing space characters to avoid issues with well-known CSV editors.

Data Dictionary download now includes a new first column, **al_datadict_item_properties**, which contains object types and object IDs for each object in the data dictionary. When updating an RDBMS object with a data dictionary, this column should not be modified. When migrating curation information from one instance to another, this column should be removed, as the object IDs may not be correct for the target instance.

Install and Upgrade

- On customer-managed (on-premise) instances of Alation, you can now place the internal Alation database on the Amazon Relational Database Service (Amazon RDS) instance in your own AWS account. This capability provides the benefits of RDS to self-managed customers.
- Added support for Red Hat Enterprise Linux (RHEL) 9.x, Ubuntu Linux 22, and Debian Linux 11 as base operating systems for installation of the on-premise Alation software package.
- The legacy tab—**Software Updates**—has been removed from Admin Settings.

Lineage

Asset Grouping Icons for Datasource Cluster Nodes

Data source logos are now shown in the Compound Layout lineage view for each data source group. Generic data source logos are used for virtual data sources and Custom DBs. (AL-128656)

Search

Synonym Search On Cloud-Native Architecture

Synonym search is now supported in the cloud-native architecture.

New Public API for the Management of Synonyms

A new public API has been added for the management of synonyms so that customers can customize synonym terms in their catalog searches. For more information, refer to [Search Overview](#) on Alation's Developer Portal.

User Management: Generate SCIM Tokens in the User Interface

Server Admins can now generate SCIM tokens on the **Admin Settings > Authentication** page, making it easier to maintain the connection between your identity provider and Alation. Server Admins will get alerts about SCIM token expiration. Email alerts will occur 15 days before expiration, 1 day before, on the day of expiration, and the day after expiration. Server Admins will also see a banner appear in the Alation interface 15 days before the token expires.

Connect Platform

Zero Touch Agent - Downloadable Installers through Public API

Installer packages for the Alation Agent can now be downloaded programmatically through the public Alation Agent API.

OCF Connectors

HDFS OCF Connector—New!

The HDFS OCF connector is now generally available (GA).

MicroStrategy OCF Connector—New!

MicroStrategy OCF connector is now generally available (GA). The MicroStrategy OCF connector supports TLS and other features based on MicroStrategy REST APIs.

Salesforce OCF Connector—Limited Availability

Limited availability release of the Salesforce OCF connector which would be replacing the custom DB connector.

- Known issue: The OAuth support for MDE requires manual token refresh upon expiration.

Bug Fixes for Multiple OCF Connectors

- Column profiling: Fixed a profiling error raised when profiling a column with more than 2,147,483,647 rows (AL-130007). Applies to OCF connectors:
 - **AWS Databricks OCF connector**
 - **Azure Databricks OCF connector**
 - **Redshift OCF connector**
- The JDBC URI is auto-updated to have `useNativeQuery=0` if not given in the URI. This fixes custom query sampling errors (AL-129772). Applies to OCF connectors:
 - **AWS Databricks OCF connector**
 - **Azure Databricks OCF connector**
 - **Databricks Unity Catalog OCF connector**
 - **Databricks Unity Catalog OCF connector**
- Removed the `jtds` prefix appended in `sqlserverazure` in the connector code (AL-134059). Applies to OCF connectors:
 - **Azure SQL DB OCF connector**
 - **SQL Server OCF connector**

Athena OCF Connector

- Getting query string from Batch API using query execution ID. (AL-127679)
- Removed the ORDER BY clause from the column profile queries (stats). (AL-133322)

AWS Glue OCF Connector

Upgraded to the latest SDK and passing the Alation compatible version. (AL-132242)

Azure Blob Storage OCF Connector

Azure Service Principal-based authentication feature is supported from Alation version 2023.3.1. (AL-119081, AL-133846)

Denodo OCF Connector

Fixed selective extraction and now the connector filters tables and columns based on the selected schema. (AL-126061)

EMR Presto OCF Connector

Fixed an issue with dynamic sampling taking user credentials from the **General Settings** page. (AL-131731)

Google BigQuery OCF Connector

Fixed an issue with child statements of create procedure query not being parsed. (AL-128150)

Redshift OCF Connector

- We need to check the limit while accumulating/appending the string instead of at the end of the sequence to avoid unnecessary processing. (AL-133320)
- Removed temporary instance profile option from the **General Settings** page. (AL-130294)

Snowflake OCF Connector

- Fixed extraction queries which caused an issue when we don't have a catalog name in the JDBC URI and the selected Include/Exclude filter option was **all schemas except**. (AL-133694)
- Snowflake has custom streamer class for filter extraction and that doesn't handle validation exception similar to SDK's streamer class. As part of the fix we added catch block to handle validation exception. (AL-134421)

Starburst Trino OCF Connector

Fixed an issue with dynamic sampling taking user credentials from general settings page. (AL-134314)

Teradata OCF Connector

- Introduced the **Extract Schemas From DBASE** setting, enabled by default, allowing schema extraction from `DBC.Dbase` or `DBC.DatabasesV` based on permissions, providing a mechanism for scenarios with limited table access. In the absence of this support, MDE was failing if the service account did not have permissions for `DBC.Databasesv`. (AL-133434)

Tableau OCF Connector

Popularity for the Workbook and Report/View objects are now supported in Alation when using the Tableau OCF Connector.

Vertica OCF Connector

Made a modification of the grammar type from ANSI to VERTICA. Previously, when the ANSI grammar type was in use, certain SQL queries, such as `SELECT * FROM TABLE LIMIT 10;`, didn't undergo proper parsing by Alation during QLI. With this fix, these types of queries now undergo successful parsing without encountering any errors. (AL-129496)

Bug Fixes

Governance

Revamped Dictionary Upload Flow: More performant and robust data dictionary upload flow is now enabled by default on all environments. (AL-134121)

Snowflake Tag Fixes for Policy Center

- Snowflake tags are removed from data dictionary during download as we do not support persisting tag value updates via data dictionary. (AL-131699)
- For the native Snowflake connector, the tooltip text for custom fields wasn't getting populated correctly. This has now been fixed for OCF as well. (AL-130336)
- In Snowflake, tags can have `allowed_values` set for them (see [Object Tagging](#) in Snowflake documentation). Previously, during MDE, the values in a tag's `allowed_values` field were not populated as Custom Fields. However, when tags were attached to objects with a value, those values got populated. Now, MDE is fixed to get the `allowed_values` of Snowflake tags as Custom Fields values. (AL-130309)

Term & Glossaries

- Fixed an issue causing corruption of article snapshots when articles are edited using Bulk Utility. (AL-133010)
- Fixed an issue that resulted in a new row when copying and pasting into a table within an article. (AL-133281)

User Management

- Suspending users flow is optimized to clean up the active session data in a background task. In some situations, there are 3-4M session entries that previously caused the 504 timeouts. While the background task sweeps and cleans up sessions, a middleware change further blocks these users and prevents their logins for security. (AL-122130)

1.3.12 PATCH RELEASE 2023.3.1.2

build 17.1.2.54864

Released October 25, 2023

Note: This release is available for all Alation instance types: customer-managed (on-premise) and Alation Cloud Service.

- In a recent update of Datadog Agent, a configuration change was omitted. This resulted in email alerts being sent every 15 seconds, which led to an excessive volume of notifications for Alation administrators. With this fix, email alerts will be sent based on your configuration in **alation_conf**, ensuring that their frequency aligns with your expectations. (AL-134151)

1.3.13 PATCH RELEASE 2023.3.1.1

build 17.1.1.54419

Released October 20, 2023

Note: This release is available for Alation Cloud Service instances on the cloud-native architecture.

- Fixed an issue related to the Okta SCIM integration with Alation. Previously, if a user was deleted from the Okta SCIM app and subsequently re-added, their reactivation in Alation did not occur as expected. The issue was indicated by the following error in the `scim-debug.log` file: *Exception - InternalServer: Unexpected error in updating SCIM user with id <...>*. With this fix, users who are removed and then re-added to the SCIM app in Okta will now be automatically re-provisioned and re-activated in Alation. (AL-134886)

1.3.14 RELEASE 2023.3.1

build 17.1.0.52371

Released October 5, 2023

Note: This release is available for all Alation instance types: customer-managed (on-premise) and Alation Cloud Service.

Compose

Granular Download and Export Permissions

Introduced user and group-level permissions for downloading and exporting query results to ensure that only authorized users and groups can export and download the data. Previously, Data Source Admins could enable or disable the Export and Download feature for all users of a data source. Now, they can select individual users and/or groups and grant them the ability to export and download. The access settings are available:

- On the **Compose** tab of the Settings page for an OCF data source
- On the **General Settings** tab of the Settings page for a native data source or Custom DB.

Query Export Alerts

Added status alerts for the Export and Download feature in Compose. Compose will show an alert if a download gets stuck or cannot be completed. You will also be informed about the progress status (number of rows downloaded so far). This feature can be enabled from the Compose preferences.

Ability to Disable Automatic Connections to Data Sources

Server Admins can now disable automatic connections to data sources in Compose. With automatic connections disabled, Compose won't use the expired credentials to make connections automatically and cause the data source to lock up the account.

Sort Query Groups

Users can now apply sorting to query groups.

Multiple Compose Bug Fixes

See [Compose Bug Fixes](#) below.

Search

Alation now searches for matches across all text fields (picker, multi-picker, rich text custom fields) on term tables.

Governance

Data Dictionary download now includes a new first column, **al_datadict_item_properties**, which contains object types and object IDs for each object in the data dictionary. When updating an RDBMS object with a data dictionary, this column should not be modified. When migrating curation information from one instance to another, this column should be removed, as the object IDs may not be correct for the target instance.

Connect Platform

Query Log Ingestion Pipeline Improvements

Stabilized the performance of query log ingestion (QLI) by restricting a single QLI job to run for seven days or less. Multiple QLI jobs can still be submitted to account for more days as an alternative. If a QLI job is submitted for more than seven days, an error message will be displayed in the QLI job submission UI and referenced in the QLI job history page and the job will not run until resolved. If any exceptions are required to loosen these restrictions and submitting multiple QLI jobs is not an option, please request an exception via a Support ticket.

OCF Connectors

- [Tableau OCF Connector](#)
- [Databricks on AWS OCF Connector](#)
- [Databricks on Google Cloud OCF Connector](#)
- [Databricks Unity Catalog OCF Connector](#)
- [Amazon S3 OCF Connector](#)
- [Azure Blob Storage OCF Connector](#)
- [Teradata OCF Connector](#)
- [SQL Server OCF Connector](#)

Tableau OCF Connector

Version 1.6.7.6780

- Removed the settings **Disable auto pagination**, **Number of workbooks to return per request**, **Number of datasources returned per request**, **Number of connections to return per request**, and **Number of fields/columns to return per request** from the connector settings page. Pagination is now always enabled. (AL-131539, AL-129442)
- Added the ability to enable or disable the extraction of auto-generated embedded data sources (**Enable MDE for auto-generated embedded data sources**). (AL-132136, AL-125129)
- The REST API calls will now all have a timeout to prevent latency issues leading to extremely long extraction times. (AL-129663)

- Fixed an issue with extraction where users observed duplicate and stale objects in the catalog. Duplicates or stale objects will no longer be created. (AL-129311)
- Fixed an issue where a failure in the popularity API request led to an error during extraction. Added logic to correctly handle empty popularity API responses. (AL-130126)

Databricks on AWS OCF Connector

Version 2.1.3.5925

Upgraded the Databricks driver in the connector to the the latest version. In this driver version, the JDBC URI protocol was changed from `spark://` to `databricks://`. (AL-123135)

Databricks on Google Cloud OCF Connector

Version 2.0.0.5925

Upgraded the Databricks driver in the connector to the the latest version. In this driver version, the JDBC URI protocol was changed from `spark://` to `databricks://`. (AL-123141)

Databricks Unity Catalog OCF Connector

Version 2.0.4.6984

Fixed an issue where a table that was previously extracted and later excluded from extraction, was still displayed, but with the *Deleted from the data source* message. Selective extraction now properly filters catalogs, schemas, tables, and columns, and no longer contains objects from catalogs and schemas that are not included in extraction. (AL-126886)

Amazon S3 OCF Connector

Version 3.8.6.6821

Fixed an issue with extracting folders that don't have a name, for example, folders that are stored with a / and no explicit name. Added the ability to extract the S3 folders with no name but having contents. Previously, we didn't show such folders on Alation. (AL-128411)

Azure Blob Storage OCF Connector

Version 2.4.1.7072

Added support for authentication with an Azure service principal (Azure Data Lake Storage Gen 2). (AL-119081)

Teradata OCF Connector

Version 1.0.9.6817

Fixed an issue with dynamic sampling and profiling where it didn't work when the **Use LDAP Server** checkbox was selected in **General Settings**. Dynamic sampling and profiling now works with LDAP authentication. (AL-131715)

SQL Server OCF Connector

Version 1.3.5.6961

Fixed an issue with the Upload Data feature where it failed with the error *The uploaded data does not seem to match the selected table*. Added support for the data types `nchar`, `nvarchar`, and `timestamp`, which had not been available previously. All built-in data types of SQL server that are normalized to boolean, string, integer, float, date, time, and timestamp are now supported. (AL-129508)

Enterprise Readiness

Elasticsearch Upgrade

Upgraded the Elasticsearch component to a newer version to leverage support for semantic search.

Python Upgrade

Upgraded the internal Python language libraries to a newer version to take advantage of the new, more secure, functionality.

Bug Fixes

Compose

- The parameter **Always show column names** in the **Results** section of Compose Settings previously only applied to query runs. Now, it has been extended to query exports. (AL-127968)
- Previously, changing the query owner of a scheduled query didn't change the owner of the schedule. After this fix, a change to query ownership will also change the ownership of the schedule. (AL-130402)
- Fixed an issue where **My Queries**, **Drafts**, and **By Query Group** lists on the **Open Query** screen did not include the unpublished queries that a user had access to. Now, users will see all of their queries, including unpublished ones in these query lists. (AL-124378)
- Fixed an issue where the **Password** field in the **Connection Settings** dialog was not clear when users tried to create a second new connection. (AL-101879)
- Fixed an issue where the **Format** action could break the formatting for some queries. In query variables, both keys and values are now permitted to be case-insensitive. (AL-94131)
- Fixed an issue with the OAuth authentication for OCF Snowflake data sources where the Connection Settings dialog showed the *Authentication Successful* message when the access and refresh tokens were expired. Now, the *Authentication Failure* message will be correctly displayed when token expiration happens. (AL-100851)

- Removed the OAuth authentication status notifications that were shown in a slide-out in the bottom right corner of the Compose user interface. OAuth authentication status messages will be now displayed in the Connection Settings dialog. (AL-126545)
- Fixed an issue where HTML sanitation incorrectly recognized some HTML tags within SQL. Only script tags will now be sanitized. (AL-131716)
- Fixed an issue where query form results did not obfuscate literals while the data source had the **Obfuscate Literals** setting enabled. (AL-115548)
- Fixed a bug in the Compose user interface where un-publishing and then re-publishing a query caused the corresponding query version to appear as a parentless version in **Version History**. (AL-36005)
- Fixed an issue where a query schedule could be saved without providing the connection information. (AL-121813)
- Fixed an issue where the Conversation pull-out icon did not appear on the query form catalog page. (AL-133248)
- Fixed an issue with the **Load More** button for queries on the **Queries** tab under schema and table catalog objects where pressing this button did not upload more queries into the interface. (AL-131497)
- When the service account is configured with key pair authentication for Snowflake, the data associated with it is being used for Compose connections. The key pair parameters should not be used for Compose connections and the change in this ticket does an extra check to ensure this. (AL-128080)

Alation Analytics

- Fixed an issue where the `ts_updated` column was not updated when custom templates were modified. (AL-114792)
- Fixed an issue where the current timezone of a query schedule was not extracted into Alation Analytics. Now you'll be able to query the timezone of the scheduled query from the Alation Analytics data source. This can help in estimating the exact runs of the cron jobs handling the queries. To retrieve this information, make use of the `schedule_cron_tz` column in the query table. (AL-119622)

Glossaries and Articles

- Fixed a performance issue where the loading of linked terms on the glossary details page took excessive time. The fix ensures much quicker responses to hydrate both Glossary Hub and Glossary Term tables. It is controlled by the **alation_conf** parameter `alation.feature_flags.enable_glossary_terms_performance` that is disabled by default and needs to be enabled if you are affected by this issue. (AL-115695)
- Fixed an issue where the object titles mentioned in the description were missing on the hover-over the **Columns** field on catalog pages of table objects. Now, the full description will be displayed. (AL-129977)
- Previously, there was a page length control on the **Term Glossary** membership dialog. It defaulted to 5, and could be adjusted to show up to 25 glossaries per page. From Alation version 2023.1.6, this page size control was no longer visible, but the default page size of five could not be changed. Now, fixed: the page length can again be controlled by users in the user interface. (AL-132113)
- Fixed an issue with Bulk Utility where domains with an ampersand (&) in the name were not migrated. Alation now allows the domains with an ampersand to be successfully applied to terms during migration from articles using Bulk Utility. (AL-132214)

- Fixed an issue where the rich text editor didn't retain formatting. Users can now choose if they want to keep the formatting from outlook or discard it. (AL-132418)
- Fixed an issue where suspended users were not displayed as GONE on the **Access** dialog for object-level permissions-enabled objects. (AL-133111)

Data Dictionary

- Users will now be able to upload a data dictionary file containing rich text fields without experiencing the *Unexpected Error* caused by the presence of rich text fields in the data dictionary file while generating a preview. Additionally, we introduced the **alation_conf** parameter `hydrate_rich_text_field_values_on_data_dictionary_preview_ui` (disabled by default). It controls whether the existing values of rich text fields should be hydrated by showing the corresponding display text for each of the mentions on the preview screen. (AL-129392)

Other Bug Fixes

- Addressed an issue where the **Subdomains** field was not available for addition to the domain template after it was deleted. Alation now prevents users from removing of the Subdomains and Domain Users components from the domain template. (AL-127364)
- Made a performance enhancement to speed up retrieval of object sets. You can still return to the previous way of retrieving object sets by disabling the **alation_conf** flag `DEV_enable_object_set_value_visible_v2` (enabled by default). (AL-112701)
- Fixed an issue where file system object count reporting query used up lot of disk space. Optimized the query to avoid excessive disk usage. (AL-132803)
- Fixed a behavior where LMS history files were deleted when any error happened during processing. (AL-132602)

1.3.15 RELEASE 2023.3.0 - General Availability

build 17.0.0.49181

Released September 5, 2023

Alation Agent

Alation Agent Notifications

Once a user subscribes to Alation Agent notifications, the user will receive email alerts 60, 45, 30, 20, 10, 5 and 1 day before the certificate expires. If the certificate expires due inaction, daily reminders will be sent to ensure action is taken.

Alation Agent Diagnostic Tool

The Agent diagnostic tool allows users to troubleshoot connectivity issues between the Alation Agent and Alation. It runs a series of diagnostic steps, and is installed on the Agent node. The tool provides an easy way to gather all the connector logs with a simple command.

Alation Agent Docker Version Update

Upgraded the version of Docker bundled with the `alation_container_service` to 24.0.5.

Alation Analytics

Stewardship Data Is Available in Alation Analytics

Stewardship Workbench (Bulk Actions) data is now available in Alation Analytics. With this change, customers can query Alation Analytics to get metrics around bulk actions.

Third-Party BI Access to Alation Analytics

Cloud-native architecture only

Alation Cloud Service customers on the cloud-native architecture can now connect third-party BI tools to Alation Analytics. This is done by replicating your Alation Analytics database to Snowflake, where you can connect your BI tools to access Alation Analytics data.

Alation Anywhere

Slack Conversations

Added support for previewing Conversation object types in Slack, and added the conversation option for all supported object types. A user is able to quickly start new conversation from Slack per specific object, reply to conversation, or mark one as resolved.

Preview Dashboard Object Types in Slack and Microsoft Teams

Support of enhanced Dashboard preview in Alation Anywhere for Slack and Microsoft Teams integration.

Alation Database Schema Enhancements

Modified the Alation database schema to support larger volumes of metadata. Users can now ingest more metadata on a frequent basis without running into issues with integer overflow. Note that migrations during an upgrade can add an additional 30 minutes depending on the volume of the metadata.

Compose

Flexible Query Result Protection

To ensure Alation instances comply with customer security policies, this feature provides more control and flexibility to datasource admins on how query results can be shared. Datasource admins have three levels of control on whether the sharing the query result is allowed or not: No sharing is permitted, only private sharing is allowed, or allowing private and public sharing. Datasource admins can also control if users have the ability to preserve query results without automatic deletion

Compose Compatibility with Query Service-enabled OCF Connectors

Compose is now compatible with Query Service enabled Starburst Trino and SybaseIQ OCF Connectors, so that clients using these data sources can migrate to these OCF Connectors.

Compose Export Notifications

When Exporting a query execution, Compose will show an alert message containing the title of the query and the number of result rows exported upon each completed query statement, so that users are informed if the Export is successful.

Connected Sheets Improvements

- Users can now see Column Titles were available in search left navigation bar
- Search results auto expand in left navigation

Conversations

Conversations are now Reflected in Domains

Conversations about catalog objects, which are associated with data domains, are now transitively associated to the same domains.

Multiple Users and Groups Assigned to Conversations

- Handling multiple assignees for a single conversation (before the change only one user could be assigned to a conversation).
- Added possibility to assign group profiles to conversations.

Authentication

Metadata Extraction based on Cross-Account IAM Role-based Authentication

Support for cross-account authentication is needed for Metadata Extraction (MDE) between the Alation instance and AWS resources that might reside in different AWS accounts. With IAM role-based authentication, customers do not have to enter IAM user credentials in the Alation Catalog. Authentication is performed based on trust set up between an IAM role in the Alation account and an IAM role set up in the customer's AWS account where the AWS resource is running. Metadata extraction and query log ingestion from the following OCF AWS connectors: S3, Glue, Athena, and Redshift, are supported by this feature.

API Authentication using JWT - *Cloud-native architecture only*

APIs can now be secured with JWT token-based authentication in addition to using opaque tokens. With this new authentication, customers can register an external application with Alation and use the JWT tokens assigned to the application for authentication and authorization.

Cloud Configuration - *Cloud-native architecture only*

New Cloud-Native Architecture Service: Entitlement and License Management Service

Cloud Entitlement and License service displays customer licenses in the license user interface page based on entitlement grants/revocations and track usage of features. The license screen can be viewed only by Server Admins.

Connectors

Azure Blob Storage OCF Connector, version 2.3.4

Improved the Azure Blob Storage OCF Connector performance by streaming directories on the fly instead of storing it in memory.

Azure Databricks OCF Connector

Azure Databricks OCF Connector Driver has been upgraded to 2.6.33. (AL-97592)

Databricks Unity Catalog OCF Connector, version 2.0.3

- Added QLI support and upgraded the Simba Spark Driver 2.6.21 to the Databricks JDBC driver 2.6.33. (AL-118979)
- Disabled QLI Lineage by default. The datasource admins no longer need to disable the automatic lineage generation flag on the datasource settings page manually to turn off QLI Lineage. (AL-130658)

Google BigQuery OCF Connector, version 1.1.8

Improved the memory management and performance of the Google BigQuery OCF Connector. (AL-125140)

Power BI OCF Connector

Power BI Dataflow features introduced for Power BI Scanner Connector. (AL-111227)

- PowerBI Dataflow objects are now cataloged within Alation for the PowerBI OCF Scanner Connector.
- Dataflow objects provide reusable transformation logic that can be shared by Datasets and Reports within PowerBI and the Power Platform.
- Dataflow objects strengthen security around underlying datasources by exposing required data within it for reports instead of from the datasource.
- Dataflows are cataloged alongside Datasets in Alation.
- Users can now view the datasource referred by the Dataflow along with Datasets and Reports that are created from it and the fields that comprise the dataflow object.
- Hybrid sources that power a Dataset, for example, created from Dataflow(s) and another datasource are now visible.
- Dataflow objects can be nested i.e. created from other Dataflows.

Snowflake OCF Connector

- Added check for supported Alation versions during installation. (AL-129242)
- Installation will fail on unsupported Alation versions. (AL-129242)

SQL OCF Connector

- Added check for supported Alation versions during installation. (AL-129243)
- Installation will fail on unsupported Alation versions. (AL-129243)

Starburst Trino OCF Connector, version 1.2.8

- Fixed the issue when a customer adds catalog name in JDBC URI, GET list of schemas picks up schemas for all catalogs. (AL-127362)
- Changed Parser Grammar type from ANSI to PRESTO to align with native connector implementation. (AL-127352)

Tableau OCF Connector, version 1.6.1

- Resolved out of memory issues and added support for report fields value sampling. (AL-127096) (AL-120383)
- Fixed report columns value sampling feature. (AL-114383)
- Added owner property for projects and sub projects. (AL-127028)
- Added lineage support for tableau dashboards with hidden sheets. (AL-125127)

Enhanced User Experience for the OCF Connector Authentication Process (Private Preview)

OCF Connector Authentication improvements include: (AL-125067)

- Enlarged the URI field to improve user readability and editing.
- Clear help text for URI and QLI fields.
- Streamlined the interface to display relevant fields based on the selected authentication method.
- Rearranged the authentication process to follow a logical order: Connection, Service Account, Permissions, and QLI.

Governance – Machine Learning

Allie Steward Suggestions (Public Preview – Cloud-native architecture only)

Allie Suggested Stewards helps users automatically find the right people as stewards for each data object. This feature leverages a machine learning model that suggests stewards based on interactions, and also improves the suggestions over time by learning from user actions. This is applicable for cloud-native architecture customers only.

Lineage

Column-Level Lineage Tracing (Beta)

The compound layout has been enhanced to drill down from data sources to the column level with the ability to trace columns upstream and downstream across data source groups to capture a holistic end-to-end lineage view. (AL-126630)

Support for Rename and Deletion of RDBMS data sources in Lineage Source Groups

- Deleting an RDBMS datasource deletes the corresponding lineage source group.
- Renaming a RDBMS source renames the corresponding lineage source group.

Improved Lineage Performance

Created scripts to clean up duplicate temp objects and dataflows from lineage to address performance issues with large graphs. (AL-107764)

Native to OCF Migration

Improvements made related to Bulk API for Native to OCF Migration. (AL-12868)

PostgreSQL Upgrade

PostgreSQL is upgraded to 13.11 from 13.6.

Search

Save Column Selections in Saved Search

Saved Search will automatically save the column selections and their sequence in the Search Results Table view, so that users will see their prior column selections when they retrieve their Saved Searches.

Terms and Glossaries Improvements

- Customers will now see all valid custom fields, fields present on one or more term type templates, as selectable options on Term Tables by default. (AL-127525).
- Verbiage on Bulk Utility article body has been updated to indicate that other usages are “not supported” (previously “not recommended”). (AL-127569).
- Term detail pages now show the “Description” field name. (AL-128745).

APIs

Public APIs for Table User Permission / Privacy Updates

Added Public APIs for updating the table permissions and privacy settings. More information is available in the swagger docs under `/openapi/members_permission/` and `/openapi/privacy_settings/`. Server admins by default have access to both APIs. Members with catalog or source admin roles can only access these APIs if they are tagged as data source admins for the table's parent data source.

Upload Data Dictionary Public API (Public Preview)

Public Preview of Upload Dictionary Public API that brings an improved, performant and asynchronous approach to update records supporting a csv as an input file.

Conversation Public API Version 2

The new API is available at the following endpoint: `/integration/v2/conversations/`. This API includes the ability to handle multiple assignees for a single conversation and added the possibility to assign group profiles to conversations.

Conversations Public API Improvements

- Conversations Public API is updated with advanced handling of deleted conversations, including new search options and possibility to restore deleted conversation.
- Updated the Conversations Public API V2 and Conversations Internal API to support new searching and filtering options. Users can search for conversations by creation date, assignee, or last response date. (AL-125306).

Workflows Public APIs

Public APIs for workflows are now available. These will allow users to list all the workflow requests, create a change request, and show workflow request details.

Bug Fixes

- **Alation Analytics:** - Fixed issue causing customer to not be able to query `rdbms_columns`. (AL-126878)
 - Fixed issue causing deprecation flag description that was updated using the API not being updated in Alation Analytics V2. (AL-123876)
- **BI Migration:**
 - Migration script validates if source and target are the same before starting migration using the script. (AL-127623)
 - Added validation for source and target when migrating servers from native to OCF in Alation. (AL-126483)
- **Bulk Import Utility:** Fixed an issue causing keys inconsistency and incomplete, missing data source ID (`ds_id`) and causing an error. Download template now has all object sets of type table, data source, column, schema prepended with `ds_id`. (AL-119516)
- **Bulk Utility:** Fixed issue to avoid Bulk Utility Download process issue on encountering a Term linked to a Glossary without title. (AL-125179)
- **Compose:**
 - Comment syntax inside query variables are now treated as plain text. (AL-128065)
 - The “Open New Query In Compose” button from AAV2 dashboards now creates queries without titles and prevents them from being automatically published to avoid duplicate queries. (AL-126656)
 - Fixed issue causing Schedule Result Sharing Modal to not follow the admin setting ‘Allow Public Sharing Of Compose Results’. After the fix, public sharing is disabled if it is disabled from admin settings. (AL-124132)
- **Conversations:** Fixed an issue causing the conversation resolution status to be empty or null. (AL-106466)

- **Google BigQuery OCF Connector:** Fixed issue causing Lineage to not display for Temporary Tables and other connected upstream and downstream tables, although the queries are parsed successfully. Added support to show the lineage via a temporary node. (AL-117243)
- **Governance/Workflow:** Fixed issues for Review Policy/Term object set and Suggest column object set. (AL-128809)
- **Lineage:**
 - Added support to create lineage for duplicate queries with different comments across data sources. (AL-125865)
 - Fixed issue causing incorrect list of data sources to display for DataFlow sources. Deleting or renaming a RDBMS source now deletes/renames its corresponding lineage source groups. (AL-120403)
 - Fixed issue to support column-level lineage (CLL) for Select * queries with SET operators. (AL-77284)
- **Power BI Scanner OCF Connector:** Fixed connector links for apps and app related objects. (AL-120876)
- **RDBMS Migration:** Migration script validates if source and target are the same before starting migration using the script. (AL-127532)
- **Search:**
 - Fixed an issue causing multi-picker custom fields in catalog search to not display an accurate not_set count for custom fields. (AL-116323) (AL-95802)
 - Fixed an issue when using Edit Columns on the Advanced Search page to not display custom fields. (AL-100937)
 - Fixed appearance of backend value for “not set” filter. (AL-130483)
- **Snowflake OCF Connector:**
 - Fixes *KeyError* that occurs while Snowflake tag ingestion when a schema/table/column associated with a Snowflake tag is not found in Alation. (AL-128018)
 - Fixed failure in policy links ingestion due to missing objects caused by inconsistent custom queries for tables and columns. (AL-125463)
 - Fix failure in tag links ingestion due to missing objects caused by inconsistent custom queries for schemas, tables and columns. (AL-125462)
 - Fixed issue causing the error when trying to establish database connection in Compose when Snowflake datasource is configured to use OAuth. (AL-124305)
 - Fixed issue when detecting role switching for Snowflake configured using External Authentication. (AL-122112)
 - Lineage is not getting created for Snowflake Create View DDL queries having “With Tag” clause. (AL-115186)
- **Tableau OCF connector:** Fixed connector links in Tableau OCF connector. (AL-118648)

1.3.16 KNOWN ISSUES

- The upgrade to 2023.3.1 or a later minor version of 2023.3 may fail due to Postgres failing to start. The issue affects some of the Alation Cloud Service instances that have not been migrated to the cloud-native architecture and the on-premise instances that used a custom approach to the installation and installed the Postgres component in a Docker container. This issue can be addressed through a workaround: </admins/Troubleshooting/20233UpgradeFailsPostgresDoesNotStart>. A fix will be provided in a future Alation version.

1.4 Release Notes 2023.1 (General Availability)

1.4.1 RELEASE 2023.1.7.6

build 16.7.6.52631

Released October 11, 2023

Note: This patch release is available for customer-managed (on-premise) instances.

- Users will now be able to upload a data dictionary file containing rich text fields without experiencing the *Unexpected Error* caused by the presence of rich text fields in the data dictionary file while generating a preview. Additionally, we introduced the **alation_conf** parameter `hydrate_rich_text_field_values_on_data_dictionary_preview_ui` (disabled by default). It controls whether the existing values of rich text fields should be hydrated by showing the corresponding display text for each of the mentions on the preview screen. (AL-129392)
- Added a scheduled job that runs daily and deletes the data dictionary input files and generated report files that are expired. The expiration period is set to five days with a feature flag which can be adjusted dynamically. (AL-124514)

1.4.2 RELEASE 2023.1.7.5

build 16.7.5.49213

Released August 30, 2023

Note: This patch release is available for customer-managed (on-premise) instances.

- Fixed an issue with the **Health Check** not showing any data. (AL-129229)
- Fixed an issue with the **Health Check** page not loading the status of services. (AL-118674)

1.4.3 RELEASE 2023.1.7.4

build 16.7.4-49180

Released August 30, 2023

Note: This patch release is available for Alation Cloud Service instances on the cloud native architecture.

- When deployed on the same node, the `alationfc` pods were sharing the process ID. This resulted in a number of issues, such as the `supervisord` crashing and containers having access to processes in other containers. The issue has been fixed. (AL-131522)

1.4.4 RELEASE 2023.1.7.3

build 16.7.3.48794

Released August 25, 2023

Note: This patch release is available for Alation Cloud Service instances on the cloud native architecture.

- This patch disables the synonym search feature on cloud native architecture, as we have identified an issue with it and are actively working on a fix. We've made a decision to disable synonym search for now to ensure your catalog experience is not affected. (AL-130610)

1.4.5 RELEASE 2023.1.7.2

build 16.7.2.48581

Released August 24, 2023

Note: This patch release is for all Alation instance types.

This patch release delivers important security fixes, addressing potential security risks in a number of components of the Alation application. Your security tools may report vulnerabilities in Alation if this patch is not applied.

The patch includes these fixes:

- Upgraded the package `golang.org/x/crypto` in the Alation Connector Manager code to a newer and more secure version. (AL-127369)
- Native connector dependencies were upgraded to address potential security issues. (AL-127903, AL-127194)
- Addressed potential security risks in `AuthServer`. (AL-128630, AL-127600, AL-129757)
- Fixed a potential security risk in `Compose` by upgrading to Log4j 1.2 bridge. (AL-127586)
- Upgraded the General SQL Parser to a newer version to take advantage of the most recent bug fixes and improvements. (AL-124451)

Additionally:

- Made changes to enable the transition to a new KVStore binary, providing the dynamic configuration on `cache_index_and_filter_blocks` and partitioned index and filter. (AL-124995)

1.4.6 RELEASE 2023.1.7.1

build 16.7.1.47079

Released August 16, 2023

All instance types

On-premise, cloud native, and instances that haven't been migrated to cloud native architecture

- Fixed the error `KeyError` that occurred during tag ingestion from a Snowflake data source when a schema, table, or column associated with a Snowflake tag was not found in Alation. (AL-128018)
- Fixed an issue where views ingested through incremental metadata extraction were incorrectly renamed and as a result, incorrect lineages were created and displayed on Lineage diagrams. Solved by enhancing the view parsing logic during the calculation of lineage. (AL-127249)

Instances on cloud native architecture

Cloud native only

- Previously, enabling the MT QLI flag on a cloud native architecture instance would result in enabling the QLI 2.0 flow for all tenants. This flow should be restricted to the Enterprise cluster only. This fix ensures that enabling the MT QLI flag only applies to tenants in the Enterprise cluster. (AL-128327)

1.4.7 RELEASE 2023.1.7

build 16.7.0.44693

Released August 1, 2023

Note: Features or updates that are “on-premise only” or “cloud native architecture only” are identified with a label in these release notes. If there is no label, it means the feature or update applies to both architectures with the exception of items listed in the Cloud section.

Alation Agent

- You can now subscribe to email notifications that alert you when any of your Alation Agents have become disconnected from your Alation Cloud Service instance for an extended period of time.

Alation Analytics

- Data for policies and policy groups are now included in Alation Analytics.

Alation Anywhere

- The OpenWith API now provides additional fields: stewards, userName, and userEmail.

Cloud Platform - *Cloud native architecture only*

- Added a new cluster in Tokyo to better support customers in that region. (Cloud native architecture only)
- Align token_buffer_time in OAuth and AWS IAM Authentication configuration: The “token_buffer_time” parameter in OAuth Authentication configuration is now allowed to have a value between 1 to 10 minutes, unlike previously which allowed a value between 1 to 60 minutes, For all existing OAuth Authentication configuration whose “token_buffer_time” is greater than 10 minutes, it must be updated to comply with the new allowed range. This change makes the allowed range of values for “token_buffer_time” consistent across OAuth and AWS IAM Authentication configurations. (Cloud native architecture only)

Compose

- Compose now provides code folding, which enables a user to selectively hide or display subsections of long SQL queries, improving readability and making it easier to focus on specific sections of code.
- When viewing the results of a SQL query in Compose, you can now click on a row number in the results pane to highlight that entire row. You can toggle highlighting off and on for multiple rows by clicking their row numbers. This makes it easier to keep track of your place when scrolling horizontally.

Connected Sheets - *Cloud native architecture only*

- Sharing Connected Sheets: Users can now share Excel spreadsheets stored in OneDrive or SharePoint, or Google Sheets, with their colleagues. This feature is not available in Excel files stored in the local file system. See the Alation Connected Sheets documentation for more information. (Cloud native architecture only)

Search

- Users can now export selected columns in the Search table view. Previously, users were only able to export four default columns.

Workflows

- Creator Optional in Change Request Workflow: While configuring a change request workflow, the user can decide whether the creator of the object for object types like Policy, Policy Group, Term, and Glossary should go under review or if they are allowed to bypass workflow.

Connectors

Snowflake OCF Connector

- Added Fields on MDE settings page to support custom tag and tag link feature. You can now curate queries to extract tags and tag links to personalize which tags have to extract.
- Updated Default Snowflake QLI query to extract only successful statements.
- Changed session column of default query in the query to take up only the session ID.

Starburst Trino OCF Connector

- Support added for SSL connection with .jks files in addition to .cer files.
- Updated current single catalog extraction logic for a generic jdbc URI.

Open Connector Framework

- Added new endpoints to the Data Sources API (OCF) to get and patch general configuration settings for data sources. The new endpoints allow all settings to be modified except certain authentication profiles like OAuth.

Alation APIs

- The Public Custom Field API now has access to Catalog Admins, in addition to Server Admins.
- New Workflow API endpoints to list active reviews, approve an active review, or reject an active review.
- New Policy API endpoint to delete policies in bulk.

Bug Fixes

Alation Analytics

- Fixed an issue where users receive an email indicating Alation Analytics setup failed when first installing it, followed by a duplicate email. (AL-121948)
- Fixed an issue where the Alation Analytics groups table was not getting updated when a name change was made to a custom group. (AL-122632)

Compose

- Fixed an issue where the Scheduled Query Sharing Settings dialog would allow users to share query results, even when the public sharing of query results were disabled in Admin Settings. (AL-124132)
- Fixed an issue that was causing Server Admins and Catalog Admins to be unable to edit trust check flags from within Compose. (AL-123960)
- Fixed an issue where query forms would still allow you to click the Run Form button even if the query was empty. (AL-82406)

Glossaries, Terms, and Articles

- Fixed glossary pagination issues. (AL-103526)
- Fixed an issue where custom fields were missing from articles after the upgrade to 2023.1.5. (AL-126298, AL-126255)
- Fixed an issue that prevented members from being added to private data sources. (AL-123954)

Governance

- Fixed an issue in Stewardship dashboard causing duplicate and incorrect schemas when using the dropdown filter. (AL-121211)

Native to OCF Connector Migration

- Fixed issue causing View Query to be missing from View Catalog page after migrating from a Native to OCF connector. This fix is for both single data source and bulk migrations of RDBMS sources. (AL-116534)

Connectors

- **Denodo OCF Connector:** Fixed an issue for column extraction for null datatype returning from driver (v1.2.1, AL-121751)
- **Oracle OCF Connector:** Fixed an issue where MDE failed as the Default Column Extract query was adding (,). (v1.4.6, AL-177498)
- **Snowflake OCF Connector:** Fixed table-based QLI query (v1.2.2, AL-124415)
- **SQL Server OCF Connector:** Fixed an issue causing last run usernames to display incorrectly. (v1.3.1, AL-121091)
- **Teradata OCF Connector:** Fixed issue with Compose query execution in script mode (v1.0.8, AL-122878)

1.4.8 PATCH RELEASE 2023.1.6.1

16.6.1.45432

Released August 3, 2023

Note: This patch release is for Alation Cloud Service instances that have not yet been migrated to cloud native architecture.

- Fixed an issue where views ingested through incremental metadata extraction were incorrectly renamed and as a result, incorrect lineages were created and displayed on Lineage diagrams. Solved by enhancing the view parsing logic during the calculation of lineage. (AL-127249)

1.4.9 RELEASE 2023.1.6

build 16.6.0.43241

Released July 10, 2023

Note: Features or updates that are “on-premise only” or “cloud native architecture only” are identified with a label throughout these 2023.1.6 release notes. If there is no label, it means the feature or update applies to both architectures with the exception of items listed in the Cloud section.

Alation Analytics

Alation Analytics Backup and Restore Support - *On-premise only*

- Alation Analytics (AA) now supports backup and restore for AA data. Using this feature, users can restore to the latest available backup instead of rehydrating ETL from start (which could take more time depending on the data size). This feature is *not* applicable to Docker rootless on-premise installations or Alation Cloud Service (ACS) installations, and there is no support for scheduled backups.

ETL Status Dashboard and Failure Notification feature

- ETL status dashboard and failure notification feature (delivered in the 2023.1.4 release) can be controlled by the feature flag. The default status is disabled/ OFF.

Alation Anywhere

Alation Anywhere for Tableau - Private SSL Support

- Implemented self-signed SSL certificates and Personal Access Token (PAT) support for Tableau Metadata Upload. This feature is behind a feature flag and provides improved handling of connection encryption during metadata upload from Alation to Tableau (use of custom self-signed SSL certificates or disabling SSL certification).

Alation Anywhere OpenWith Improvements

- Improved support for OCF data sources in OpenWith functionality.
- OpenWith is available for datasources for which it was not configured. Added change to ensure that requests filter datasource types properly.

Cloud Platform

Alation Cloud Service URLs now have subdomains

- Alation is introducing subdomains to the URL for Alation Cloud Service. This will make it more clear that your data is hosted in your geography, provide isolation between regions, and improve compliance with local regulations. It also allows Alation to better optimize the performance of each region.

AWS IAM Authentication has new `aws-global` option in the region field

- The Admin feature `aws-global` is now an option in the region field for AWS IAM Authentication configuration. This provides the ability to select global STS endpoints instead of region specific endpoints, as global endpoints might be preferred because of high availability.

PrivateLink Support for Alation on Cloud Native Architecture

- Alation Cloud Service customers using PrivateLink can connect to data sources after migrating to Alation cloud native architecture. Alation cloud native architecture provides PrivateLink connectivity to data sources in all supported regions.

New Status Page for Alation Cloud Service

- Users can directly subscribe to updates or view anytime the status of Alation Cloud Service at status.alationcloud.com. The status page displays the status by region for all tier one shared services. The status page also displays any major incidents affecting all customers in a given region.

Compose

Compose Improvements

- Compose users can select an object (schema, table, column) in the left side bar and drag it to the Edit pane. The full object name is inserted into the current cursor location in the Edit pane, by default, instead of previous hyperlink functionality.
- Added the ability to support PKCE for Snowflake integration with External Oauth with Azure AD by removing an extra = character.
- Decreased the default polling time for compose query completion to increase responsiveness. It is still configurable by the `alation_conf.alation.query_exec.polling.min_interval`.
- Added support to Compose SSO to work with Databricks URI, as per naming conventions specified for Databricks Driver 2.6.29.

- When using an Advanced schedule, users now have the option to chose a timezone whenever they are in a different timezone. Previously, the Advanced schedule could not be converted from one timezone to another.
- Users can now go directly from Compose to the Query Catalog page by selecting the “View Query Page” dropdown button.

Connected Sheets - *Cloud native architecture only*

- Support Azure SQL from Connected Sheets from both Microsoft Excel and Google Sheets.

Governance App

Policy Center

- Made minor improvements to how policies are fetched, sorted, and paginated. (AL-118350)

Workflows

- Users can configure that reviewers of a step will bypass the entire workflow while adding or updating step configuration. This feature is under a dev flag.

Lineage

Cross System Lineage

- Added support in the MDE/QLI lineage flow for cross-system lineage across data sources with the same host and port configuration.

Lineage Improvements

- Dataflow Public API will now return HTML content instead of plain text for the description field. (AL-122740)
- Updated the migration script `deferred_lineage_group_sync` to ensure temp objects appear in default groups to make it easier to differentiate them from other catalog objects. (AL-118343)
- Made changes to ensure that the expand and collapse operation of asset groups will focus the groups and the root node will be centered initially. (AL-122595)
- Support added for temp node creation in MDE lineage flow and cross-system lineage across data sources with the same host configuration. (AL-114337)

Connectors

Improvements and Updates to OCF Connectors

- Fixed issue on the Datasource page, causing the page to not render properly when using the browser back button to move from one tab to other. (AL-121951)

Amazon S3 OCF Connector Enhancements

- Added Proxy support for OCF S3 connector (Basic Proxy and Auth Proxy). Upgraded to the latest cData csvdriver to resolve a sampling issue.
- Provided a fix for failing MDE for referential directories for OCF Connector v3.5. Also added exception handling while decoding path string to prevent MDE failure. (AL-119420)
- Fixed issue causing extraction to fail due to connection timeout. (AL-117232)

Azure Blob OCF Connector Enhancements

- Upgraded the JRE version to Temurin JRE 11.0.18, and ensured that the timestamp field in the parquet file will have the correct value.
- With connector version 2.3.1 onwards timestamp field in parquet file will have correct value.

Azure Synapse OCF Connector Enhancements

- Ensured Azure Synapse OCF Connector server-less extracts views and columns for schemas as expected. (AL-121471)

Impala OCF Connector Enhancements

- Compose works with SSL enabled CDH. (AL-120004)

Looker OCF Connector Enhancements

- Significant improvements related to MDE performance and logging.

Postgres OCF Connector Enhancements

- Fixed an issue in which Postgres OCF MDE was failing to extract the definition of aggregate functions using QUERY_FUNCTION. (AL-120774)

PowerBI Scanner OCF Connector Enhancements

- Added the PowerBI field level descriptions as source descriptions in Alation. These source comments will appear on the overview page in the Catalog.
- Added source description for datasource fields.

Redshift OCF Connector Enhancements

- Upgraded the SDK 4.1.31. This SDK ensures complex column datatype, null and empty object ingestion failures no longer occur.

Snowflake OCF Connector Enhancements

- Policy extraction is available for TABLE, VIEW, COLUMN. Policy linked with a tag is not supported and removed from extraction (1.1.5).
- Improved MDE Column Query performance for selective extraction. (1.1.3) (AL-119284)
- Updated query column to exclude METADATA\$ROW_VERSION system generated column. (1.1.4) (AL-122046)

SQL Server OCF Connector Enhancements

- Fixed an issue in which usernames appeared incorrect in query history. (AL-121091)

Starburst Trino OCF Connector Enhancements

- Starburst Trino OCF Connector now picks up only schemas of the catalog specified in JDBC URI.

Tableau OCF Connector Enhancements

- Migration Script fix for Tableau Native to OCF migration of bi_folder and bi_folder_filter data which has composite_key= "postgres_id"

Query Service Support for More OCF Connectors

- Implemented Query Service in Impala OCF connector (1.1.0), Impala CDP OCF Connector (1.1.0), and Databricks Unity OCF Connector (1.2.1), Fixed bug for Query fails for certain field types in SQL server (1.3.2) and MySql (1.3.8) with Query Service enabled.

OCF SDK Enhancements

- Modified the Maven dependencies to remove log4j 1.2.17 and replace it with the log4j1.2 API bridge library. This ensures no log4j 1.x is used in the SDK.

Alation APIs

Dataflow Public API Improvements

- The Dataflow public API will return HTML content instead of plain text for description field.

Policies API

- Added an API to update business policies.

Public Bulk Create Custom Field API

- Added support for bulk creation of Custom Fields in the Public Custom Field API. Update and Delete behavior are not supported.

Bug Fixes

Alation Analytics

- Fixed issue causing tenant to fail to provision due to null credentials from vault for Snowflake. (AL-124078)
- Fixed observing error in the backup command in Docker Rootless Ubuntu16 instance. (AL-123612) - *Cloud native architecture only*
- Postgres is stopped before cleaning archive directory, allowing proper cleanup of old files. (AL-123253, AL-122786)
- Fixed issue causing an invalid failure notification email to be sent when setting up Alation Analytics in cloud native infrastructure. (AL-123094) - *Cloud native architecture only*

Glossary Page

- Fixed issue causing alert text to not display when removing term from Glossary Page. Alert text now displays properly. (AL-120139)

Compose

- Fixed an issue where Query Form Authors, Description, and Run information were not updating when modified. (AL-67395)
- Fixed issue so that OCF Teradata respects the data upload chunksize for number of rows, same as Native Teradata data source. (AL-120471)
- Fixed issue causing “Compose new Query” link from Query execution to not work properly. (AL-118756)
- Fixed issue in data upload when data source is configured to use query service. (AL-118675)
- Query permission modal was not showing the update fields without page refresh. Fixed the issue so that re-opening the modal will show the updated fields. (AL-111298)
- Fixed issue in Compose Preview, where the description of schema, table, and attribute were only showing the icon for the mentions and the name was missing. (AL-49226)

Cloud Platform

- Fixed issue causing /dev/shm mount missing inside chroot after upgrade. (AL-116942)

Search

- Fixed an issue to ensure Viewers are able to access and navigate to functions as expected. (AL-123062)

Workflows

- Fixed a blocking behavior for reviewers. Previously if a user took action on a suggestion (for example, approved a change) and the user was part of the approval workflow downstream as part of a group, the user could no longer take action on the suggestion. Now the same reviewer in multiple steps can act on all steps. (AL-121448)

1.4.10 PATCH RELEASE 2023.1.5.1

build 16.5.1.40261

Released June 16, 2023

Support for AWS PrivateLink on Cloud Native Architecture

AWS PrivateLink is now supported for Alation Cloud Service instances on cloud native architecture. Alation supports AWS PrivateLink between a customer’s Alation instance VPC and a customer’s VPC in the same AWS region. You can use AWS PrivateLink to connect to your regionally local data sources over a non-public internet connection. See [/cloud/CloudConfiguration/AWSPrivateLinkWithAlationCloudService](#) for more details.

Bug Fixes

- For Alation Cloud Service deployments that haven't been migrated to the cloud native architecture, we upgraded the RDS component to version 13.8. The previous version 13.6 is at end of life, which caused the need for this upgrade. (AL-123330)
- Fixed an issue with Lineage diagrams where cataloged objects appeared as temporary nodes that disappeared when the temporary object filter was applied. This fix ensures that cataloged objects will not disappear from the graph when the temp filter is applied. (AL-122844)
- Fixed an issue with the lineage service consuming a large amount of memory because of unbounded queries. (AL-122667)
- Enabled `cache_index_and_filter_blocks` to use the block cache for indexing and filtering. (AL-120852)
- The *connection reset by peer* error was often observed when Alation was reading configurations from CMS on Alation Cloud Service deployments that haven't been migrated to the cloud native architecture. This blocked some jobs from moving through processing. Now, the `cmsclient` `retry` has been added to reduce the possibility of failure when the *connection reset by peer* error occurs. (AL-120592)

1.4.11 RELEASE 2023.1.5

build 16.5.0.39106

Released June 1, 2023

Alation Analytics

- Glossary Terms and Glossary Data are now available in Alation Analytics (AAv2). Visualizations in Alation Analytics will not explicitly call out Glossary Terms and Glossaries, but the data is available for querying.
- Added Workflows data to Alation Analytics: Workflows data is now available in Alation Analytics, enabling customers to query Alation Analytics for Workflows usage data.

Alation Anywhere

Alation Anywhere for Microsoft Teams integration

The Alation App is now available on the Microsoft teams app store and provides customers with the following capabilities:

- Search and filter objects within Alation
- Preview objects
- Star and watch objects
- Support for the following object types: Tables, Queries, Data Source, Domains, Glossary, Terms, Articles, and Conversations.

Alation Anywhere Improvements

Resolved some notification related issues.

Cloud Platform

AWS Secrets Manager and Azure Key Vault

- Connection settings for OCF Connectors can now be pulled from user-maintained vaults, either AWS Secrets Manager or Azure Key Vault. This allows users to consolidate secrets previously stored directly in the Alation database, preventing “credential sprawl” and allowing organizations to comply with their IT security policies.
- Credentials stored in AWS Secrets Manager may be stored as plain text, JSON Key/Value pairs, or as binary certificates.
- Credentials stored in Azure Key Vault must be plain text only. In particular, it is not possible to obtain binary certificates from Azure Key Vault.

Support Migration of Alation Analytics Database from On-prem to ACS

Users have full access to Alation Analytics data after the migration from on premise to Cloud Native Architecture without having to rebuild the Alation Analytics database.

OCF Connector Updates

Google BigQuery OCF Connector

Support Auto Lineage for Non-data Manipulation Queries: Added auto-lineage for SELECT queries coming with the destination table for Google BigQuery. The destination table is considered the target table for the particular SELECT query and lineage is built accordingly.

Impala CDH OCF Connector

HTTPS, HttpFS, and keytab are all supported. Tables with space in the HDFS location are extracted without errors.

Power BI Scanner OCF Connector

- Support for Dataset Parameters: Power BI uses parameters to make reports dynamic. The connector now supports these values within Datasets making it easy to use dynamic PowerBI reports and view them within Alation.
- Enhanced MDE import: Only active workspaces are retrieved now within MDE and exclude deleted workspaces. This improves MDE performance while also making it more robust.
- Updated JRE image: JDK version updated to 11.0.18 from previous version 11.0.16.

File System Connectors

Verified folder navigation working properly for files/folders with special characters in their name.

File System Native Connectors to OCF Connector Migration

Added support for a Migrate button in UI for Amazon S3 and Azure Blob Storage native file system connectors.

Query Service Support for OCF Connectors

- Google BigQuery
- Databricks on Google Cloud
- Azure Synapse Analytics
- SAP ASE (Sybase ASE)
- SingleStore (MemSQL)
- Amazon Athena
- Denodo

RDBMS Connector SDK Metadata Extraction

The OCF Connector SDK now skips streaming and logging the objects received during RDBMS MDE that violate the following constraints for text fields within them: fields exceeding the maximum character limit, fields having a non-zero length and also a null character in payload, and fields representing the original name for objects as empty.

Compose

Compose Improvements

- Added the ability to switch to advanced scheduling for schedules created before the 2022.3 release and ensured query schedules appear correctly regardless of query owner timezone. Also resolved various timezone related issues within query scheduling.
- A Compose user can now select max number of columns to display per page for query results, making it easier to link columns far apart on the same row.
- Ensured that the time appears as expected for simple and advanced schedules in query search.

Governance App

- Support of Policy as permitted otype in Object set custom field. With this, you can now link policies to objects using object set custom fields.
- When rejecting a change request, users will be asked to provide a rejection reason.

Workflows

- In My Tasks and Workflow Configuration table, user title was taking priority over username. Fixed this to take username as higher priority while displaying.
- Added Workflows data to Alation Analytics: Workflows data is now available in Alation Analytics, enabling customers to query Alation Analytics for Workflows usage data.
- Field Level Workflows: Alation now allows creating change requests on specific fields. Change request workflow can be configured on the field level by selecting specific fields while in the change request workflow configuration. Users can still configure change request workflow on an object by selecting all fields.

Connected Sheets

- Auto refresh capability is now available in Excel add-in. This allows users to schedule automated data refresh on hourly, daily, weekly and monthly cadence and users can also opt into email notifications for successful and/or failed scheduled jobs.
- Connected Sheets now supports Databricks on AWS and Azure Databricks via Alation Agent.

Lineage

Lineage Impact Analysis Improvements

- Changes made to the Impact Analysis feature allow lineage users to first select the filters and then 'Apply Filter' to load relevant data based on user selection instead of loading all data by default.
- Added support for pagination on the lineage service for Impact Analysis requests to improve performance.
- Added the ability to cancel an impact analysis request that is being processed.
- Added support for more fine-grained incremental control on the 'Max Distance' filter option instead of the preset dropdown values.

Search

- Custom Field Search: Enable text search on custom_field values so that searching a value in custom fields will return the relevant search results (objects). To enable this setting, set the flag `alation.search.enable_custom_field_search=True`.

Alation API

Conversation API improvement

The full path of an element was not being returned by the GET request, with the same issue visible in the user interface. Improved so that GET values return full path.

Public API Pagination Support for OCF datasources

- The BIServer APIs to get folders, reports etc., supports limit and offset API parameters for pagination.
- The API to sync the schemas from datasource to Alation has been split into two APIs to improve performance. There is a new API to sync the schemas from data source to Alation. This will allow users to trigger sync job and check job status. Once the job is done, the `available_schemas` endpoint can be used to fetch a list of schemas in a paginated way.

Public API Support for Policies

- Added the ability to create Policies in bulk through public API.
- Added support to list policies through API. API also supports pagination, sorting, filtering and searching based on some parameters.

Alation Developer Portal

For a full list of Alation API updates, see the [2023.1.5 API Release Notes](#) on the [Alation Developer Portal](#).

Bug Fixes

- Alation Analytics: Fixed an issue in ETL Jobs which caused extended running and resulted in stalling the ETL pipeline. This was due to sub-optimal extraction of filesystems, files and directories. This task is now optimized to work in smaller batches so that the ETL pipeline does not stall and AAv2 data refresh stays on cadence. (AL-116308)
- Alation Analytics: Fixed an issue for a AAv2 where users could see the AAv2 settings page without enabling it by going to the URL directly. After the fix, the page is only visible if the admin enables the feature flag for AAv2 in Admin settings > Feature configuration page. (AL-73040)
- Alation Analytics: Fixed an issue affecting Custom Group names that were added and removed frequently in the Alation catalog were not reflecting correctly in AAv2 due to unique constraints implemented in the associated table. This constraint is removed and users can view how custom group names have been added and removed via AAv2. (AL-113239)
- Compose: Fixed issue causing schedule time to display incorrectly. Fix ensures the time appears as expected for simple and advanced schedules in query search. (AL-120473)
- Compose: Previously disabling/enabling a schedule from query search/scheduled query dashboard/compose scheduling modal was modifying the timezone of schedule, causing the user interface to display incorrect time and also triggering schedule runs at incorrect time. Now disabling/enabling schedule will not alter the timezone of schedule. (AL-112937)
- Compose: Fixed issue causing variables in commented lines of query display when not supposed to display. (AL-107727)

- QLI Pipeline: The task name for QLI Event Archiving was `rosemeta.tools.event_archiver_with_partitioning.archive_execution_events_and_mentions`. The fix is to update the entry for the QLI Event Archiving task field in the `PeriodicTable` entry. (AL-118821)
- Looker OCF Connector: Fixed issue for selective MDE option to handle extraction of excluded/deleted folders and associated objects appropriately. (AL-119475)
- SSRS OCF Connector: Fixed issue with MDE for ensuring consistency of folders and objects from SSRS that are extracted into Alation. (AL-118662)
- Power BI Scanner OCF Connector: Added buffer time to the retry mechanism for API retries when rate limit is reached. (AL-114578)
- Power BI Scanner OCF Connector: Fixed issue to redirect to correct URL within customer Power BI instance using the connector option within Alation's catalog page for paginated reports. (AL-107494)
- dbt OCF Connector: Fixed issue causing metadata to not ingest from tables with an alias. (AL-117369)
- EMR Presto OCF Connector: Fixed issue causing columns to be missed in Alation for delta tables after successful metadata extraction. (AL-119574)
- Amazon Redshift OCF Connector: Fixed issue to add support for complex data types having nested fields with more than 100 characters. (AL-118629)
- Amazon Redshift OCF Connector: Fixed issue causing dynamic sampling to not work with the URL provided in Compose section. (AL-117835)

Known Issues

BI Datasource catalog pages are not rendering correctly when dataflows are extracted

- When the user is navigating between tabs (Connections, Fields, Overview, etc.) in the BI Datasource catalog pages, and if they click on the browser's back button to navigate between these tabs, there is HTTP 500 error thrown. This issue is because of the async operation to pass the properties within the feature flag, but ideally, it should have been initialized before the feature flag. The user will see sporadic errors when they click the browser's navigation buttons instead of clicking the tabs to navigate from one page to another.
- **Workaround:** Use the tabs in the application to navigate between pages (BI Datasource catalog page), or reload the page when encountering this page rendering issue.

1.4.12 PATCH RELEASE 2023.1.4.2

build 16.4.2.37601

Note: This patch is for Alation Cloud Service customers who haven't moved to our cloud native architecture. It is not available for customer-managed (on-premise) instances or Alation Cloud Service instances on the cloud native architecture.

- The BI (GBMv2) public API now includes support for pagination with `limit` and `offset` parameters.

1.4.13 PATCH RELEASE 2023.1.4.1

build 16.4.1.36590

Note: This patch is only released for Alation Cloud Service instances on the cloud native architecture and is not available for customer-managed (on-premise) instances.

- Previously, when restoration of a backup of an Alation version with Postgres version 13.6 was attempted on the 13.8 instance image, updating Postgres version configuration failed, as it is not allowed on the cloud native architecture. Now, Alation will skip updating Postgres version if RDS is enabled, ensuring that the backup restoration can move forward.
- Fixed an issue with Hive Connector 1.0.0 where query log ingestion (QLI) returned the error *File /user/history/done does not exist*. The Hive connector expects query log files to use a certain format, and Hive QLI job failed when processing malformed files. After this fix, instead of failing the job, Alation will skip such files in processing.

1.4.14 RELEASE 2023.1.4

build 16.4.0.35588

Alation Analytics

- **Health Check Notifications to Server Admins**
 - Server Admins can turn off Alation Analytics ETL status alert email notifications.
 - Server Admins can see the latest ETL failure status alerts on the Homepage. Detailed ETL failure statuses along with the history will also be available on the ETL status page.
 - Created monitors on Datadog for ETL failures for the SRE team for environments on Cloud Native Architecture.

Alation Anywhere

- **Alation Anywhere for Tableau**
 - **Metadata upload to Tableau feature works with OCF Connectors.**
 - * Previously, metadata upload worked only if extraction was done using native connectors. After this change metadata upload to Tableau works if extraction was done using OCF Connector for: Snowflake, MSSQL Server, Oracle, Postgres, MySQL, Azure Databricks, AWS Databricks, Redshift and BigQuery.
 - New Tableau enhancements, such as sending the endorsement flag to Tableau objects, sending catalog set definitions and inherited tags, adding common branding for endorsement and quality flag consistency.

Connect Platform

- **OCF Connector Migration Improvements**

- Added support for bulk migration using command and Public API which can allow users to bulk migrate their sources of the same database type from Native to OCF connectors.
 - Enhanced Native to OCF migration command and UI option to migrate all the connection configurations, MDE configuration, Compose configuration, sampling configuration and QLI configuration except SSL certificate migration for all the connectors which supports SSL auth and User Auth certificate.
- Column Level Lineage for Sybase is now GA.
 - Column Level Lineage for Hive is now GA.
 - Incremental metadata extraction support for RDBMS data sources.
 - Alation RDBMS connectors now support incremental direct lineage extraction by default. This is useful for data sources such as Databricks Unity Catalog that have metadata that can be used to leverage incremental lineage data extraction.
 - OCF Alation Connector Manager in Rootless mode
 - Alation Connector Manager and Docker can now be run in rootless mode. Deploying the services requires root or sudo privileges.
 - Query Service Support
 - Added QueryService Support for a variety of OCF connectors including AWS Databricks OCF, Azure Databricks, Google BigQuery, Starburst, and Oracle.
 - Email notifications for file system sources
 - Alation now sends email notifications to Server Admins and File System Admins about extraction failures on *file system sources*. Admins will be notified if the following jobs fail: metadata extraction, schema extraction, Get List Of Schemas, or scheduled extraction. Notifications are sent for both OCF and native file system sources. Previously, email notifications were only sent for RDBMS sources (*data sources*). Admins can opt out of these notifications by disabling the toggle **a scheduled metadata extraction/query log ingestion/profiling job fails on a data source** or by opting out of all email notifications on the **Notifications** tab of their user account page. For more on email notifications from Alation, refer to </welcome/UserProfileandPreferences/EditProfile>.

Connected Sheets

- Enhanced search with ability to search on Title and Description, or use catalog object URL in search bar of Add-in.
- Improved performance on catalog load and search experience with caching. Users can also refresh the cache at any given time.
- View trust flags in Add-in for objects downloaded.
- Easily access catalog page for downloaded objects.
- Connected sheets now support Postgres, SQL server, Oracle and MySQL via Agent.

Lineage

- Added an Apply Filter button to load data relevant to the user selection.
- Instead of using a dropdown there is now a number input with increment / decrement buttons for max depth.
- Added pagination for Impact Analysis to improve performance.

Governance App

- A public API has been added to list workflows.
- Add support to provide reason for rejection during workflow add term / associate term request process.
- Data Dictionary Improvements. Improvements to the data dictionary upload API and pagination of the preview list in order to prevent failure during file upload.
- Fixed issued causing properties section was showing as duplicated in the business policy catalog page.
- Fixed issue causing Policy Center page to fail to load due to null values being present in the API response.
- Untag values for Snowflake only for the data source currently being ingested instead of all sources.

Installation and Upgrade

- Some existing jobs (Metadata Extraction / QLI / Profiling) which were in LocalTimeZone are still time-shifted 8 hours earlier after a previous Alation upgrade. A new script is available to update time zones to UTC.

Cloud Platform

- **OCF Agent Rootless mode**
 - Alation Connector Manager and Docker can now be run in rootless mode. Deploying the services requires root or sudo privileges.
- **LMS History Service: Data Archival**
 - Support archiving history data in the new, performant LMS History Service which is used on Cloud Native Architecture platform.

Alation API

- **Public APIs for Alation Agent**
 - APIs allow admins to configure the Agent certificate and resync operations programmatically. These APIs allow admin users to automate the lifecycle operations in an unattended fashion.
- **Article API**
 - Article Swagger UI is always available, as Articles API is GA.
- **Conversations API**
 - Users can now like a post in a resolved conversation.
 - Mentioned link text (URL and name) is now returned by GET request.

Bug Fixes

- Fixed formatting issue when uploading business metadata containing special characters. The solution modifies the existing process to treat all data passed as part of description as text.
- Ensured Data source, Schema, Table or Attribute name loads when reviewing suggestions on workflows opened from Task Inbox.
- Fixed issued causing properties section showing as duplicated in the business policy catalog page.
- Fixed issue causing Policy Center page to fail to load due to null values being present in the API response.
- Untag values for Snowflake only for the data source currently being ingested instead of all sources.
- Fixed issue causing search with lineage to not return the most relevant attribute on the top of search results.
- Fixed issue causing attribute objects to not be indexed after upgrading Alation to version 2023.1.
- Fixed data format issue when uploading a data dictionary. The upload process was treating numerical data without any text incorrectly, rather than considering it as text. This issue was fixed by modifying the existing process to treat all data passed as part of description as text.
- Fixed issue causing glossary pages and glossary hub not to open and showing an error message.
- Fixed issue when searching for a node outside of the returned paginated data to return no values. The search functionality on the asset grouping modal for collection nodes when searching for an object now returns search results from the entire collection for the currently opened lineage graph.
- Fixed issue to increase the limit of object set individual records from a limit of 50 characters to 256 characters.
- Fixed an issue causing Policy custom templates to create a new visual_config clone on every edit and save action.
- Fixed an issue causing ampersands (&) to show as '&,' in the left hand navigation. Ampersands now display correctly in the left hand navigation.
- Fixed issue causing the arrow icon in the right navigation pane to not work for the data source.
- Fixed issue where OCF Connector status is showing as unknown after server restart when using the Alation OCF Agent.
- Fixed issue causing QLI 2.0 to display an incorrect timestamp in a job status message.
- Fixed issue causing QLI 2.0 to trigger airflow when there are no queries to ingest and status message displays "Query Ingestion Started".
- Fixed issue causing Fivetran to display the dataflow object incorrectly. The dataflow object now displays correctly between the source and the destination.
- Added a script in the diagnose tool to solve a QLI archival issue, where event_id is null for some partitions.
- Fixed issues to support cross-system lineage.
- Fixed issue causing search with lineage to not return the most relevant attribute on the top.
- Added more object types being tracked to the script that checks the indexing status of different object types. No changes for how to use or new flows.
- Fixed issue causing an incorrect Healthcheck error message in the native cloud environment.
- Performance improvements for Upload Data Dictionary flow.

- Fixed the logic for how/when to show the “Restore draft” on Article Body and rich text fields.
- Fixed formatting issue when uploading business metadata containing special characters. The solution modifies the existing process to treat all data passed as part of description as text.

1.4.15 RELEASE 2023.1.3

build 15.3.3.33838

Note: This build is available for Alation Cloud Service customers on the cloud-native architecture.

- Fixed an issue related to the background task scheduler. Previously, background jobs running for over an hour could cause subsequent jobs to crash. In the user interface, the issue manifested itself through an error message in the **Job History** table for the jobs that had crashed: *Job is not running. It might have been killed or crashed...* With this fix, long-running Celery jobs will not impact subsequent jobs.

1.4.16 RELEASE 2023.1.2

build 15.3.2.32690

Alation Catalog and Compose Improvements

- SQL literals can now be obfuscated for all view objects in the catalog. To enable this feature, set the **alation_conf** flag `alation.feature_flags.enable_obfuscating_view_sql` to True. This will mask the view SQL query on the catalog pages where it is displayed for all views. See [/admins/AdditionalConfiguration/ObfuscateLiteralsinViewSQL](#) for more details.
- Added a new Compose setting **Always show column names** which is turned on by default. With **Always show column names** enabled, the column headers of query results tables will show physical column names even when columns also have curated titles in Alation. With **Always show column names** disabled, the column headers of query results tables will show titles or physical column names depending on whether a column has a title. Titles will be shown for columns that have them and physical column names will be shown for columns without titles.

Alation API

- Added support for Catalog Set (`dynamic_set_property`) and Catalog Set Shared Fields (`propagated_catalog_set`) in Supported Object Types in the [Public Custom Field Values API](#).

OCF Connector Updates

- [Azure Blob Storage OCF Connector](#)
- [Azure Synapse OCF Connector](#)
- [Databricks Unity Catalog OCF Connector](#)
- [dbt OCF Connector](#)
- [EMR Presto OCF Connector](#)
- [Impala on CDP OCF Connector](#)

- *Power BI OCF Connector (Non-Scanner)*
- *Amazon Redshift OCF Connector*
- *Amazon S3 OCF Connector*
- *SAP BW OCF Connector*
- *Snowflake OCF Connector*
- *Tableau OCF Connector*

Azure Blob Storage OCF Connector

*New connector version available: **2.2.0.4373***

- Fixed an issue where extraction failed with an “unexpected error” when information for some fields that were optional for extraction was missing in the extracted files. Now, if object properties that are optional for extraction, for example, size or modified time, are missing, MDE will successfully complete without ingesting these object properties.
- The rule name parameter is now a mandatory field for extraction. Extraction will fail in case of an empty rule name or empty destination container, or both, with the corresponding message in the connector log.
- Added an option to exclude individual files during schema extraction (applicable if **Schema Path Pattern** is enabled).

Azure Synapse OCF Connector

*New connector version available: **1.1.0.4029***

- Added ability to extract metadata from serverless SQL pools.

Databricks Unity Catalog OCF Connector

*New connector version available: **1.0.3.4144***

- Added support for lineage generation.

dbt OCF Connector

*New connector version available: **2.1.0.4367***

- Added support for the Google BigQuery data source.
- Added support for the Databricks on AWS, Databricks on GCP, and Azure Databricks data sources.

EMR Presto OCF Connector

New connector version available: 1.0.4.4060

- Added support for extracting Hive Delta tables.

Impala on CDP OCF Connector

New connector version available: 1.0.3.4132

- Fixed an issue with Kerberos with a keytab connection.
- Fixed an issue with SSL-enabled connection.
- Fixed an issue with extraction where it failed if external tables pointed to the HDFS file path with spaces in the file name.
- Added support for extracting Kudu tables.

Power BI OCF Connector (Non-Scanner)

New connector version available: 3.0.0.4371

- Fixed a number of metadata extraction issues.

Amazon Redshift OCF Connector

New connector version available: 1.2.0.4216

- Added support for EC2 instance profile-based authentication.

Amazon S3 OCF Connector

New connector version available: 3.5.0.4216

- Added support for extraction from the PSV and TSV files.
- Fixed an issue where extraction failed with an *unexpected error* when information for some fields that were optional to extraction was missing in the extracted files. Now, if object properties that are optional to extraction, for example, size or modified time, are missing, metadata extraction will successfully complete without ingesting these object properties.
- Added an option to exclude individual files during schema extraction (applicable if **Schema Path Pattern** is enabled).

SAP BW OCF Connector

New connector version available: 1.0.9.4297

- Fixed a number of extraction issues.
- Added support for schema prefix extraction. When we extract from JSON files with a schema prefix, metadata extraction now completes successfully and all columns are extracted. Previously, column names were not extracted from JSON files with a schema prefix.

Snowflake OCF Connector

New connector version available: 1.1.1.4303

- Added ability to reuse the QLI view created for the native Snowflake data source.
- Added support for Snowflake key pair authentication with the public and private key generated using the aes-256-cbc encryption.

Tableau OCF Connector

New connector version available: 1.5.0.4134

- Fixed a number of metadata extraction issues, including extraction failing due to an out of memory exception and other errors.
- Added ability to successfully extract metadata from a Tableau BI source when embedded data sources have a very large number of upstream columns.
- Updated the connector's JRE image.
- Fixed issues with metadata upload when the data sources are Google BigQuery and MySQL.
- Fixed an exception-handling bug where the connector froze without erroring out when an API request failed. The connector will now log a clear error message and skip the object that caused the error.
- Updated the Tableau BI source settings user interface. The fields required for Personal Access Token authentication now appear on top of the settings page under **Connector Settings > Server Connection**.
- Added ability to extract the owner information for workbooks, dashboards, reports, embedded data sources, and published data sources.

Connected Sheets

- Added ability to edit the dataset configuration including columns and filter criteria at any time. Available in both Google and Excel add-ins.
- Added the trust flags information to data objects downloaded from an add-in. You will see the relevant trust flag descriptions in the add-in with a link to view more details on the catalog page.
- Connected Sheets can now work with an on-premise PostgreSQL data source connected with Alation Agent.

Bug Fixes

- Resolved an issue where the HA pair rebuilding failed after the upgrade to version 2023.1 with cluster splitting due to a bug in the replication service. Now, the admin performing the upgrade of an HA pair with cluster splitting should be able to successfully complete it.
- Fixed a bug in Compose where large query executions on an Amazon Redshift data source were getting stuck. The issue was due to a race condition on the server created by one of the internal API calls. Now, users should be able to successfully execute large queries on Redshift.
- Fixed a bug that caused an error when accessing the Policy Center page on some Alation instances. The issue was due to the API response for the “/all_policies” GET request containing unexpected null values for the Stewards field on some of the policy objects. Fixed by adding logic to the frontend to properly handle unexpected values.
- Fixed a bug where the QLI archiving job was set to run once a year in the internal database even though users had previously added a different schedule in the settings of a data source. The QLI archiving job will now respect schedules set by users.
- Fixed an issue with the Compound Layout lineage view where the More nodes indicator disappeared from the graph after adding a node from the Search for more objects dialog. Now, when there are collection nodes present for the graph, the down arrow indicator will also be present.
- Fixed an issue where the Show Parent and Show Children actions did not have an effect after a user attempted to expand a lineage node on a large Lineage diagram.
- Fixed an issue where the Show Parent and Show Children actions only displayed a part of the available lineage nodes on large Lineage diagrams.
- Fixed a number of user interface issues to stabilize and improve the Lineage Asset Grouping feature.
- Added a warning to the Delete BI server confirmation dialog that deleting large BI sources may cause issues in the Alation user interface. Deleting a large BI source may cause a 500 error on the Alation server. We recommend contacting Alation Support for help with deleting such BI sources until this issue is resolved. Note that BI sources with a small number of objects or with no objects can be safely deleted from the user interface.
- Fixed an issue where the lineage generation job was stuck in the Running status even though the corresponding metadata extraction job was completed and assigned the status Succeeded. The job history will now reflect the correct status of all jobs subordinate to metadata extraction.
- Fixed the time zone discrepancy in the **Next Run** column in **Admin Settings > Monitor > Scheduled Query Dashboard**. Now the **Next Run** time will be shown in an admin's local time zone.
- Fixed an issue where Alation failed to start after remounting the data volumes from the old to the new Alation instance. The problem was due to a permission issue when trying to restart a process. Fixed by updating the directory permission.

Known Issues

- The Alation update of an HA pair **without splitting the cluster** will fail. We recommend upgrading HA instances using the cluster splitting upgrade method until this issue is resolved.
- Deleting a large BI source may cause a 500 error on the Alation server. We recommend contacting Alation Support for help with deleting such BI sources until this issue is resolved. Note that BI sources with a small number of objects or with no objects can be safely deleted from the user interface.

1.4.17 RELEASE 2023.1.1

build 15.3.1.30768

- Fixed an issue where tag values on multiple Snowflake data sources cataloged using the Snowflake OCF connector were reset to Not set after metadata extraction on one of them. The fix ensures that after metadata extraction on one of the existing Snowflake OCF data sources, the tag field values will not change on other Snowflake OCF data sources.
- Fixed an issue for an Oracle data source where the Compose connection kept timing out with the error *Could not start query. Check your connection and try again* in the UI and the *Connection closed* error in the log. Added a configuration option to enable a check that the Oracle database connection is open before running a query. To enable, append `-DcheckOracleConnIsOpen=true` to the **alation_conf** parameter `connector.extra_flags`.

Note: Do not override existing values. If the parameter `connector.extra_flags` already has some values set, add your new value:

```
alation_conf connector.extra_flags -s ' -Darg1=value1 -Darg2=value2 -
DcheckOracleConnIsOpen=true'
```

On how to use **alation_conf**, see /admins/AdditionalConfiguration/Usingalation_conf.

- Fixed an issue where users were unable to see the table and column object lists on the **Curate > Assign Stewards** page. The issue was due to the internal APIs used for curation dashboards performing multiple queries to the database in an inefficient manner, which resulted in the 504 error. Added a fix where for every given batch of 1000 curation object records, all related records will get pre-fetched using a single bulk query.
- Queries ingested with QLI with the only change in the comments will no longer create duplicate dataflow objects.
- Fixed a performance issue with Multi-Select Picker custom fields. With multiple populated Multi-Select Picker custom fields on a catalog page, multiple duplicate calls to the internal API were happening to refresh the Multi-Select Picker field data. These extra calls were redundant and caused performance slowness. Fixed by updating the conditional check on Multi-Select Pickers to only fetch data when the selected values have changed.
- The health check API `http(s)://<your_alation.com>/monitor/i_am_alive` now has a configurable timeout threshold. The health status check of the lineage service will be terminated if the threshold is reached. The threshold is controlled by the **alation_conf** parameter `lineage-service.health-check.timeout`, set in seconds. Previously, this threshold could not be configured. When it was reached, the `i_am_alive` API response or the corresponding email alert indicated that the lineage service health check failed but did not contain any specific exception. Now, the API response and the email alert will include a message that the request timed out based on the health check timeout threshold.
- Previously, Alation users could run into a bug where the column profile histogram on the **Sample** tab of the column catalog page only added new counts and values but never removed values that had been removed from the underlying database. This issue has now been addressed and column samples should correctly reflect the removal of a value during any corresponding profiling request (scheduled table profiling or manual column profiling).
- Fixed a SAML authentication issue on Alation Cloud Service instances where the login failed for new users with the error *RelatedObjectDoesNotExist: User has no profile*. After the fix, Alation will process SAML responses with empty attributes lists correctly. For example, if an empty list is received from the IdP for an attribute like `Title`, then an empty string will be saved instead of breaking the login flow.

- Fixed a SAML authentication issue on Alation Cloud Service instances where active users were unable to log in or suspended users could not be activated if the case in their username had changed in the IdP. Now, a username can be patched in IAM using the PATCH endpoints. If there is any change in the username in the IdP, the username will be sync'd to IAM as well.

1.4.18 RELEASE 2023.1.0 - General Availability

build 15.3.0.29706

Released February 21, 2023

Governance App

Add Object Request Workflow for Glossary Terms

This new workflow aims at governing the creation of terms in Alation. It allows you to set up a review process for any newly created glossary term before it is added to a glossary. For any new term under a workflow, there are two modes:

- **DRAFT**—A private mode for terms in a workflow. It allows the creator to work on the term in a private space and invite other users to collaborate. The term is not accessible to users unless they are granted access by the creator or are Catalog or Server Admins.
- **UNDER REVIEW**—Review mode. Once the creator has finished drafting the term, they can submit it for review. This initiates the review process, and the term enters the Under Review mode. The term is available for review by a defined set of reviewers, which includes the creator and collaborators. The term remains in the Under Review mode until it is reviewed and approved. Once it is approved, the term goes live and can be viewed by all users.

All workflow participants get real-time email notifications about the changes and reviewers get a rich experience enabling them to act on terms pending review.

Compose

Granular Access Permissions for Queries

Previously, Alation didn't have the capability to set granular access permissions for query objects that would allow only certain users to view, edit, or run them. Starting with this release, you can use the new **Sharing & Access** feature to configure access to queries in Compose and the catalog, thus controlling who can view, edit, and run them. Levels of access supported for queries are:

- *No access*
- *View Access*
- *View & Run Access*
- *View, Run, & Edit Access*
- *Owner.*

The access rules are set on individual query objects and apply to both unpublished and published queries. For new queries, the default is *View & Run Access* for everybody.

If you are using the **alation_conf** parameter `alation.catalog.unpublished_query_visibility_level` to control visibility of unpublished queries, note that this setting will take precedence over individual access permissions on unpublished query objects. If the parameter `alation.catalog.unpublished_query_visibility_level` is currently set to `not_visible` (unpublished queries can only be accessed by Owners and users with whom they were explicitly shared), then after enabling the **Sharing & Access** feature, unpublished queries will remain invisible to users even if they are given *View Access* in the user interface. We do not recommend changing the parameter `alation.catalog.unpublished_query_visibility_level` to `visible` before all protected unpublished queries have been assigned the correct access privileges.

Scheduled Query Dashboard for Admins

The **Scheduled Query Dashboard** in **Admin Settings > Monitor** empowers Server Admins to manage scheduled queries across all users and data sources in an Alation instance:

- All scheduled queries at-a-glance
 - Tabular view of all scheduled queries with the following details: Query Title, Data Source, Query Owner, Owner Status, Schedule, Next Run Time, Status, and Actions.
- Actions on specific scheduled queries
 - Change Schedule Settings
 - * Modify schedule settings to spread query loads over time, improving overall system performance.
 - Change Permission
 - * Grant or revoke View, Run, and Edit permissions to or from individuals or groups in case the query owner is not available.
 - * Change the owner of a query to avoid “ghost” scheduled queries.
 - Cancel Next Run
 - * Cancel the current execution or the next pending execution of a scheduled query to save resources.
 - Disable Schedule
 - * Disable a query from running on a schedule, so that scheduled queries with owners who are no longer employed at the company or queries that are no longer needed don’t keep running forever.

Ability to Use Hive Variables in Compose

Previously, since Hive variables were in conflict with Query Form variables due to their syntax, they could not be used in Compose. With this enhancement, Compose users can use an escape character (`${hivevar:tablename}`) to resolve syntax conflicts between Query Form variables and Hive variables, enabling the use of Hive variables in Compose queries.

Lineage

Compound Layout View

Alation now provides the ability to represent lineage objects based on their logical alignment with the data source and dataflow objects: the **Compound Layout View** of the lineage data. The view supports the **Show Parent/Show Children** functionality which will display additional upstream and downstream lineage nodes that are not directly linked to the main object.

This feature is not enabled by default. To make it available on your Alation instance, set the following **alation_conf** flags to True:

- `alation.feature_flags.DEV_enable_lineage_asset_grouping`—Enables the **Compound Layout View** toggle on the **Lineage** tab.
- `alation.feature_flags.DEV_enable_lineage_collection_nodes_modal`—Adds the Show Parent/Show Children functionality.

If a data source supports DBT connections, users will see the DBT dataflow objects grouped together.

Compound Layout View currently doesn't support column-level lineage.

Lineage Default View Configuration—No Restart Required

Previously, after the lineage settings were modified in **Admin Settings > Customize Catalog > Lineage**, the lineage service had to be restarted. Starting in 2023.1, the restart is no longer required.

Impact Analysis Performance Improvements

We implemented a number of changes to improve the general performance of the Impact Analysis reports:

- Reports now use lazy loading to load faster.
- Temp objects are excluded by default. Previously, the configuration parameter `lineage-service.filters.exclude_temp` was applied only to the graph view. Now, it's applied to both the graph and the Impact Analysis reports.

Catalog

Domain and Term Objects Now Respect Object-Level Permissions

- On the catalog page for a domain, if a user doesn't have View permission for a specific object, they will not see this object in the **Browse this Domain** list.
- Users can no longer apply changes to objects they don't have Edit permission for when adding an object to a domain.
- The **Domains** field on the catalog pages of objects will now respect object-level permissions. Only users with permissions to edit an object will be able to modify its domain membership.
- If a user does not have Edit permission for a parent glossary, they will not be able to apply membership changes to child terms. A user will need View or Edit access to a term object and Edit access to a glossary to add or remove a term to or from a glossary.
- If a term object is added to more than one glossary and is configured to inherit permissions from all of them, a user will need access to all glossaries to have access to this term.

- A user will be able to see and/or update the template of a term object based on their object-level permissions for this term.

Term and Glossary Navigation and Search Improvements

- We added the **Glossaries** column to the Terms table to indicate which glossaries the terms belong to. This applies to the **Glossaries** landing page and each individual glossary pages.
- We added the Glossaries facet to full-page Search for term objects. When searching for terms with the full-page search, you can now filter by glossary.

Data Quality Indicators Added in Search Results

Alation Search users will be shown an icon indicating a Data Quality warning that affects objects in the search results. This icon represents quality issues affecting the object or upstream objects.

Quality Propagation of Downstream Data Objects—General Availability

Alation's Data Quality service now includes quality values generated from ingestion from OCF data sources. On detection of deleted data objects, Data Quality creates a deleted quality value on that object which then propagates to the downstream nodes on the lineage graph. This information will be displayed for downstream tables, BI reports, and BI datasources and is surfaced on the new **Data Health > Upstream Issues** tab for each respective catalog page.

This functionality, introduced in the previous release as Beta, is now General Availability.

Search Results Span Title and Rich Text Fields

The Search function now enables search results that span the title as well as rich text fields. For example, an article with “new york city” in the title and “local park” in the body can be found by searching for “new york park”.

Note: There is a small accompanying change in the search spell checker behavior. When a search query consists of one word, it will be checked for spelling. When a search query consists of several words, the spell checker will be invoked less often. It will only check and suggest a correction if there are no results for the word(s) that are spelled correctly.

For example, in the query “healthcare revnue”, “revnue” is misspelled. However, if the search finds results matching “healthcare”, the spell checker will not correct the second word, returning results that match “healthcare”. If both words are misspelled, the spell checker will correct the first word. For example, if you typed “halthcare revnue”, the search would still return matching results for “healthcare”.

Search Indexing Improvements

- Implemented the force merge API for Elasticsearch for better search index storage management. Force merge executes every Sunday.
- Completed a number of modifications to reduce the search index size, including setting the current default stored values for fields to False so that Elasticsearch saves source data only.

Alation Analytics App

Supportability Improvements

- Added functionality to improve the ETL process and failure tracing for Alation Analytics. Implemented an error notification framework in order to communicate ETL errors. Server Admins will now receive email notifications if the ETL process ends in an error, reducing the time required for them to become aware of an issue.
- The health status of Alation Analytics Postgres and RabbitMQ containers can now be monitored from the user interface in **Admin Settings > Monitor > Health Check**. Note that on the instances on Cloud Native Architecture, only the RabbitMQ health is displayed, as Postgres is a managed RDS service. Failure notifications are sent as an email to Server Admins.

Data Model Additions

- Added fields `file_url`, `directory_url`, and `filesystem_url` to tables `files`, `directories`, and `filesystems` respectively in Alation Analytics to improve consistency with columns available for other data objects like `datasources`, `schemas`, `tables`, and `columns`.
- Added a new boolean column `private` to the `rdbsms_tables` table in Alation Analytics to indicate whether the table is private or public.

Alation APIs

New Conversations API

Added a new API for working with Conversations, introducing endpoints for getting, creating, updating, and deleting Conversations and Conversation replies (posts).

Updates to OCF Data Sources API

Added new endpoints to the OCF Data Sources API:

- API now allows users to update OCF data source information using the PUT endpoint `/integration/v2/datasource/{datasource_id}/`.
- Added **datasource config checks** group of endpoints that allow performing connection tests that are equivalent to the Test Connection functionality in data source settings in the user interface:
 - `/integration/v2/datasource/{datasource_id}/configuration_check/`—Check all configurations of the data source.
 - `/integration/v2/datasource/{datasource_id}/configuration_check/service_account_authentication/`—Check the service account authentication status of the data source.

- /integration/v2/datasource/{datasource_id}/configuration_check/service_account_privileges/—Check the service account privileges of the data source.
- /integration/v2/datasource/{id}/configuration_check/network_connection/—Check if the OCF connector is responsive. Returns response 200 if responsive and response 400 when a gRPC timeout occurs.

Updates to User V2 API

Enhanced the GET methods to return the date the user was created and their last login date. We also added new query parameters to enable filtering on the user's creation date and last login date. This will enable Alation program owners to drive adoption of the catalog by programmatically identifying users who haven't logged in for a given time. You can then send users a message encouraging them to take advantage of Alation.

Updates to Virtual File Systems API

Fixed a bug where POST calls were failing if the `ts_last_accessed` parameter was missing from the request. This has been fixed so that `ts_last_accessed` is no longer required for uploading virtual file systems.

Alation Integrations

Connected Sheets

Connected Sheets for Microsoft Excel

The Connected Sheets for Microsoft Excel is a new add-on that helps users discover data sources in the Alation catalog, download data using filter & sort criteria, and refresh on demand. The Connected Sheets page, accessible from the **Apps** menu in Alation, will have the MS Excel add-on installation link.

This feature is available on the Alation Cloud Service with cloud-native architecture. Users must have an Explorer or Creator license to use the add-on.

Connected Sheets for Google Sheets

Users can now edit the configuration of the downloaded sheet.

Alation Anywhere for Tableau

- To provide more visibility into automated data upload jobs for Server Admins, we added the **Upload Job Status** table to the **Metadata Upload From Alation to BI Server** section of the Tableau BI source settings showing the status of automated data uploads. The details for each upload job are available on clicking the **View Details** link in the **Details** column of the table. The table will be available if Alation Anywhere for Tableau is enabled on the Alation instance and if automated uploads are enabled for the Tableau data source.
- The @-mentions used in the **Description** field and in the trust flag comments of data source and table objects are now transferred to Tableau as part of metadata upload.

Alation Anywhere for Slack

User Experience Improvements

- Added a success indicator that an item was sent to a chat to the Alation Search dialog in Slack.
- Undefined items have been removed from preview.

Cloud Platform

New Explorer Role

Alation Cloud Native architecture customers will now be able to use a new user role Explorer with the corresponding new license type. This role offers all the capabilities of a viewer and also provides access to Connected Sheets, query forms, and the Queries API, except the CREATE endpoint.

Bring Your Own Key (BYOK)

Bring Your Own Key (BYOK) is an encryption key management system that allows customers to encrypt their data and retain control and management of their encryption keys. Available for Alation Cloud Native architecture customers.

Connect Platform

Column Level Lineage Add-Ons

- Column-level lineage for PostgreSQL will be shown on the Lineage tab along with table-level lineage after running MDE or QLI jobs.
- Column-level lineage (CLL) for Presto is GA (previously Beta).
- Column-level lineage (CLL) for Databricks is GA (previously Beta).
- Column-level lineage (CLL) for Hive is Beta.
- Column-level lineage (CLL) for Sybase is Beta.

User Interface for Migrating Native BI Sources to OCF

Admins can now migrate native BI sources to OCF in the Alation user interface. Previously, this action required access to the backend of the Alation server.

View Jinja Code on Dataflow Catalog Pages for dbt Dataflows

Added ability to view Jinja code on the dataflow catalog pages for DBT dataflows.

New OCF Connectors

Amazon Athena OCF Connector

The OCF connector for Amazon Athena is available in this release.

Databricks GCP OCF Connector

The Databricks GCP OCF connector is available in this release.

dbt Add-on OCF Connector

The dbt add-on OCF connector is available in this release. The connector supports dbt Core and Cloud, and extracts dbt descriptions for tables and columns in the catalog and dbt lineage for supported databases and data warehouses (Snowflake, Amazon Redshift, and PostgreSQL). The configuration settings for add-on connectors are available on the **General Settings** tab of OCF RDBMS connectors when the add-on is enabled on the Alation instance.

Also see [View Jinja Code on Dataflow Catalog Pages for dbt Dataflows](#).

If the Compound Layout View is enabled on Lineage diagrams, the dataflow objects from dbt will be grouped together.

SSRS OCF Connector GA

The SQL Server Reporting Services (SSRS) OCF BI connector is now General Availability.

Improvements and Updates to OCF Connectors

Power BI OCF Connector Enhancements

- Added support for commas in queries available in dataset expressions returned in Power BI REST API responses for the SQL Server data source.
- The connector can now parse queries with multiple quotes and spaces to get the connection information for the SQL Server data source.
- The connector now validates port numbers in dataset expressions returned by Power BI REST API responses for all data sources.
- The connector now gets connection information from a non-quoted connection string in the dataset expression returned by Power BI REST API responses for the Snowflake datasource.
- Enhanced the connector to work with HTTP proxies (for example, Squid).
- Added the ability to restrict access to only those workspaces to which the service principal has access.
- Enhanced the connector to support Google BigQuery data source for table-level lineage.

Tableau OCF Connector Enhancements

- Added support for column-level lineage.
- Added an error log summary report at the end of metadata extraction.
- Added extract or live indicators and extract date for Tableau datasources (both published and unpublished).
- Removed pagination fields from the Tableau OCF settings page as pagination is now handled within the system.
- Implemented a fix to remove un-extracted sites and (or) projects when selective MDE is run.
- Enhanced the connector to support BigQuery data source for table-level lineage.

Looker OCF Connector Enhancements

- From version 1.4.1, the Looker OCF connector has support for Looker 4.0 API.
- Fixed an issue where the curation data for BI reports and BI datasources was lost after extraction if a column had been added or removed to a report in Looker. Alation provides a migration script to update Looker BI objects in Alation to restore missing curation data. The script should be run before the connector upgrade.
- Enhanced the connector to support BigQuery data sources for table-level lineage.

Azure Blob Storage OCF Connector Enhancements

- Added support for schema extraction for CSV, TSV, PSV, and Parquet files.
- Added support for on-demand end-user driven Sampling for a CSV, TSV, PSV, and Parquet file. Access keys and Shared Access Signatures are supported authentication mechanisms. Azure Active Directory authentication is also supported via AuthService. Available with Azure Blob Storage OCF Connector version 2.0.0 and Alation version 2023.1.

AWS Databricks OCF Connector Enhancements

- AWS Databricks OCF connector version 2.0.0 now uses Databricks driver version 2.6.29.
- AWS Databricks OCF connector version 2.0.0 is certified to be used with Databricks runtime 11.x.

Databricks Unity Catalog OCF Connector Enhancements

Databricks Unity Catalog OCF connector version 1.0.2 supports Sampling, profiling and Compose. Azure Active Directory auth is supported for Compose and Dynamic Sampling for Databricks on Azure.

Query Service Support for More OCF Connectors

Query Service enables the use of Compose for OCF data sources that are added through Alation Agent. In this release, Query Service support was added for DB2, Teradata, Greenplum, SAS Base, SAP HANA, Vertica, SAP IQ, and EMR Presto.

QLI Support for RDS-Based Data Sources

Added QLI support for MySQL, Oracle, and SQL Server on RDS. Previously, Alation did not support QLI for these RDS-based data sources.

Direct Lineage Extraction Is General Availability

The direct lineage extraction feature is now available to all users as the feature flag `alation.ocf.mde.direct_lineage.enable_extraction` has been removed. OCF connectors that support direct lineage will extract direct lineage metadata from data sources as part of metadata extraction.

Alation Agent

For Alation Cloud Service instances, we've certified the Amazon S3 OCF file system source with Alation Agent running in the customer VPC.

Enterprise Readiness

Postgres Upgrade to 13.8

The Alation internal PostgreSQL database is upgraded to version 13.8 for customer-managed (on-premise) Alation instances. PostgreSQL version 13.8 includes a security vulnerability fix and provides performance gains over previous versions.

Default Backup Tool Is pgBackRest

The default backup tool is now **pgBackRest**, which has improved backup reliability and performance capabilities over the previous tool, **pg_probackup**. The `pg_probackup` tool remains as an option to restore older backups.

Bug Fixes

- Fixed an issue where if term objects were added as Object Set field values, their descriptions did not appear on the table view of the Object Set field in the main panel of the catalog template. Terms will now have their descriptions displayed in the Object Set field tabular view, as expected.
- Fixed a bug where the **Details** column of the **Job History** table on the **Query Log Ingestion** tab of data source settings did not display the success message for the Succeeded status. Now, the Details information will be available for every QLI status.
- Fixed an issue where back references from term objects did not appear in the **See more** dialog. The **See more** link is displayed when there are more than five back references in the **Related Articles** field. Now, the dialog will show all available back references.

- Fixed an issue where permission to share a query from a catalog page was not available to users with the Viewer role. Now, Viewers can share queries they have access to. The **Share** button will be available to them under the **More** menu on catalog pages of queries and query forms.
- Fixed an issue where Server Admins could not see the Access dialog for queries if they were not the owner or editor. Now, Server Admins can change query permissions on the query catalog page or in Compose, if necessary.
- Fixed an issue where the workflow lookup did not include the workflow after removing an excluded object from the workflow configuration. Former known issue in LA.
- For Alation Anywhere for Slack, fixed an issue where the search results were incorrect when the search string contained spaces. Users can now use search strings with spaces.
- Fixed an issue where the **Impact Analysis** or **Upstream Audit** reports were sending duplicate API requests to the database. For large lineage graphs, this could noticeably affect the lineage data loading time. After the fix, the loading time for large graphs has been significantly improved.
- Optimized the migrations of the database which were slowing down the upgrade of the Alation Analytics app. After the installation or upgrade of 23.1, the flags will need to be reconciled between Alation Analytics and the internal database. To do that, execute the one-off script `reconcile_flag` that is located in the one-off-script folder under Alation Analytics after updating Alation to 23.1. (Included into the 23.1 upgrade steps).
- Fixed an issue with tags, which were not updated in Alation Analytics after being removed from the catalog. After the installation or upgrade, the tags will need to be reconciled between Alation Analytics and the internal database. To do that, execute the `reconcile_tags` script that is located in the one-off script folder of Alation Analytics. (Included into the 23.1 upgrade steps).
- Alation Cloud Service instances will no longer send a redundant email notification that was sent to Server Admins after the Alation update that extraction from Alation Analytics failed due to the outdated Alation Analytics version.
- Alation now respects custom field permissions during field editing set through People Sets when users use Bulk Actions. Previously, bulk updates of custom fields of picker type by users given edit permissions through a People Set did not work from Stewardship Workbench.
- Fixed an issue where users observed a discrepancy between lineage graphs on Lineage V2 and V3 when temp tables were excluded and when the main node was connected to all temp nodes on the first lineage level. Lineage V3 rendered an empty graph, while Lineage V2 displayed the main node. Now, the main node will be displayed even if temp tables are excluded.
- Fixed a bug where the column width of the results table in Compose did not scale correctly and clipped the last column, adding horizontal scrolling. The fix also addresses other display issues where content was not visible when the browser window (Chrome) was set to 90% Zoom.
- Fixed an issue where a user with the Catalog Admin role, assigned as a Data Source Admin to a data source, could not use the Test Connection feature on the data source settings page. The connection test resulted in an *unexpected error*. The fix allows Catalog Admins to use the test connection feature.
- Fixed a bug where for an Amazon Redshift data source, query snippets for queries containing a regex expression produced result sets different to query results for a query without the snippet. The issue occurred due to a discrepancy in the expansion of the regex in the snippet and the query. The issue has been resolved.
- Fixed an issue in the **Account Settings** section where users were unable to navigate to other tabs from the **DB Connections** tab. The fix ensures that users can click between tabs of the Account Settings page.
- Fixed an issue where the Compose auto-complete feature suggested column name in lowercase for columns where the column name had uppercase letters. The attribute model has two fields for storing

the table column name. The `name` field is case-insensitive, while the `original_name` field maintains the case of the original column name. In order to maintain the case of the column during auto-suggestion, Alation now returns the value of the `original_name` field.

- Fixed an issue where the GBM V2 GET folders bulk API threw an internal server error for a non-existing parent folder. Now, if the parent folder does not exist, the API will return null for the `parent_folder` parameter.

Known Issues

2023.1 Update Known Issue

If you have a password set on the internal PostgreSQL database (Rosemeta), the update will result in an error similar to the following:

```
WARN: unable to check pg-1: [DbConnectError] unable to connect to 'dbname=
'postgres' port=5432':
FATAL: password authentication failed for user "postgres"
password retrieved from file "/home/postgres/.pgpass"
ERROR: [056]: unable to find primary cluster - cannot proceed
```

The update process requires that the password on the internal PostgreSQL database should be cleared for the time of the update. The password can be set again after the update is completed. On how to set the Postgres password, see [/admins/AdditionalConfiguration/SetPasswordforPostgreSQLInstance](#).

Lineage Compound Layout View Known Issues

- Searching for a node that is not part of the currently opened pagination page does not return it. Searching for any available node in the cluster should return it even if it's not part of the current page.
- The **Show Parent** and **Show Children** actions do not have an effect after a user attempts to expand a lineage node on a large Lineage diagram. (**FIXED IN 2023.1.2**)
- The **Show Parent** and **Show Children** actions only display a part of the available lineage nodes on large Lineage diagrams. (**FIXED IN 2023.1.2**)
- After a manual update, the Lineage graph displays the incomplete lineage data with some disconnected lineage nodes.
- After a manual update to the Lineage graph and adding new nodes to existing nodes, the updated part of the graph is not displayed and the node that was updated is removed.
- Cyclic Lineage graph between multiple nodes results in duplicated nodes and incorrect lineage links.
- On the Compound Layout view, the **More nodes** indicator disappears from the graph after adding a node from the search for more objects dialog, although the nodes are listed in the dialog itself and can be added to the graph. (**FIXED IN 2023.1.2**)

Scheduled Query Dashboard Known Issue

On the Scheduled Query Dashboard, the Next Run column is displayed in UTC while the Schedule column is displayed in the user's time zone.

Term and Glossary Known Issues

- Sometimes a term in Draft or Under Review mode shows up in search. Terms should show up in search only after they were reviewed and approved.
- Selection of reviewers from People Sets is not available for the Add Object Workflow, as glossaries do not support People Sets yet.
- The first click within the **Add / Remove Glossaries** dialog seems to get “swallowed” and has no effect, but further clicks succeed. The same behavior exists on the **Add / Remove Policy Group** dialog.

Workflows Known Issue

Accessing the Workflows page for the first time after enabling it in **Admin Settings > Feature Configurations** shows the *permission denied* error.

Catalog Known Issue

After switching a data source from Public to Private, all data source objects are still showing in catalog searches.

Known Issues in 2023.1 Limited Availability Release (*build 15.2.0.27498*)

- Server Admin cannot see the permission dialog if they are not the owner/collaborator: In 2023.1 LA, a Server Admin who is neither the owner nor the author cannot see the permission dialog for a query and make permission changes to it. This is a capability that used to be available in 2022.4 and earlier releases. **FIXED IN GA**
- Viewers do not have permission to share a query: In 2023.1 LA build, Viewers do not have the permission to share a query with another Alation user. This is a capability that used to be available in 2022.4 and earlier releases. **FIXED IN GA**
- Workflow lookup not including workflow after removing excluded object in workflow configuration: In certain scenarios, workflow running on a glossary does not show up on the respective Glossary page. **FIXED IN GA**

1.5 Release Notes 2022.4 (General Availability)

1.5.1 RELEASE 2022.4.11

build 14.4.11.45870

Note: This build is available for customer-managed (on-prem) instances.

- Fixed an issue that caused an empty value in the settings of some OCF connectors of newer versions. Alation version 2022.4 didn't recognize AWS_IAM as an authentication option and loaded an empty value into the corresponding dropdown list. This caused the connection test and MDE to fail due to an unknown authentication value. We added logic to handle the new value correctly. (AL-127207)

1.5.2 RELEASE 2022.4.10

build 14.4.10.43552

Note: This build is available for customer-managed (on-prem) instances.

- Added a backward compatibility fix to resolve an OCF connector upgrade failure on Alation versions that don't have the cross-account IAM authentication feature yet. The new authentication type in question was added to the OCF SDK in the latest version but is not available on versions prior to 2023.1.6. Without this fix, you would need to first update Alation to version 2023.1.6 or later to install or update some OCF connectors. The fix enables you to install or update an OCF connector without upgrading Alation to 2023.1.6. (AL-123743)
- Fixed a number of issues with the Terms Public API List endpoint:
 - Increased the API performance by re-structuring the way term objects are iterated over, eliminating duplicate permissions checks, and breaking once the limit is reached.
 - Fixed the behavior of the skip query parameter.
 - Fixed an off-by-one error with the limit query parameter, which had caused one additional term to be returned. (AL-115877)

1.5.3 RELEASE 2022.4.9

build 14.4.9.38945

Note: This build is available for customer-managed (on-prem) instances.

- Fixed an issue causing incorrect displaying of query schedule times in the user interface. Fix ensures the time appears as expected for simple and advanced schedules in query search. (AL-120473)
- Previously, existing query schedules were shown incorrectly in the user interface after the upgrade to versions newer than 2022.3. In addition, users reported other issues with query schedules, for example, the inability to switch to advanced scheduling for schedules created before the upgrade or new query owners seeing incorrect schedules if their time zone was different from that of the previous query owner. All these issues have been fixed. (AL-113141)
- Previously, disabling or enabling a schedule from the query search, Scheduled Query Dashboard, or the scheduling dialog in Compose also modified the time zone of the schedule, causing the user interface to display an incorrect time and triggering query runs at the incorrect time. Now, disabling or enabling a schedule will not alter its time zone. (AL-112937)

1.5.4 RELEASE 2022.4.8

build 14.4.8.35994

Note: This build is available for customer-managed (on-prem) instances.

- Improved the memory usage and performance for the data dictionary upload feature. Previously, in the data dictionary upload preview, we used an HTML table with all the records of the API response. Now, we have replaced the table with a client-side paginated table. In a client-side paginated table, the data is divided into smaller “pages,” with only a limited number of rows or items displayed per page.
- Some customers previously reported that displaying a preview of a data dictionary for uploading took a long time, or sometimes resulted in the *No data found* error. The root cause was the previous implementation of this feature, where a large number of database calls was made to form a preview. To resolve this, we have optimized the code to minimize the number of database calls by batching records and temporarily caching results. Additionally, pagination has been introduced on the client side, and results are now rendered in batches across pages. As a result of this optimization, Alation is expected to successfully upload a 3 MB file and its preview in approximately two minutes; and for a 25 MB file, this should take approximately 23 minutes.

1.5.5 RELEASE 2022.4.7

build 14.4.7.31281

Note: This build is available for customer-managed (on-prem) instances.

Lineage Default View Configuration—No Restart Required

Previously, after the lineage settings were modified in **Admin Settings > Customize Catalog > Lineage**, the lineage service had to be restarted. Starting in 2023.1, the restart is no longer required.

Bug Fix

- Fixed an issue for an Oracle data source where the Compose connection kept timing out with the error *Could not start query. Check your connection and try again* in the UI and the *Connection closed* error in the log. Added a configuration option to enable a check that the Oracle database connection is open before running a query. To enable, append `-DcheckOracleConnIsOpen=true` to the **alation_conf** parameter `connector.extra_flags`.

Note: Do not override existing values. If the parameter `connector.extra_flags` already has some values set, add your new value:

```
alation_conf connector.extra_flags -s ' -Darg1=value1 -Darg2=value2 -
DcheckOracleConnIsOpen=true'
```

On how to use **alation_conf**, see /admins/AdditionalConfiguration/Usingalation_conf.

1.5.6 RELEASE 2022.4.6

build 14.4.6.29505

Note: This release is available for customer-managed (on-prem) instances.

- Fixed an issue where back references from term objects did not appear in the **See more** dialog. The **See more** link is displayed when there are more than five back references in the **Related Articles** field. Now, the dialog will show all available back references.
- Fixed an issue where users were unable to see the table and column object lists on the **Curate > Assign Stewards** page. The issue was due to the internal APIs used for curation dashboards performing multiple queries to the database in an inefficient manner, which resulted in the 504 error. Added a fix where for every given batch of 1000 curation object records, all related records will get pre-fetched using a single bulk query.
- Fixed a performance issue with Multi-Select Picker custom fields. With multiple populated Multi-Select Picker custom fields on a catalog page, multiple duplicate calls to the internal API were happening to refresh the Multi-Select Picker field data. These extra calls were redundant and caused performance slowness. Fixed by updating the conditional check on Multi-Select Pickers to only fetch data when the selected values have changed.
- The health check API `http(s)://<your_alation.com>/monitor/i_am_alive` now has a configurable timeout threshold. The the health status check of the lineage service will be terminated if the threshold is reached. The threshold is controlled by the **alation_conf** parameter `lineage-service.health-check.timeout`, set in seconds. Previously, this threshold could not be configured. When it was reached, the `i_am_alive` API response or the corresponding email alert indicated that the lineage service health check failed but did not contain any specific exception. Now, the API response and the email alert will include a message that the request timed out based on the health check timeout threshold.

1.5.7 RELEASE 2022.4.5

build 14.4.5.28479

- Fixed an issue where the previews of report, document, and dossier metadata objects were not available for a MicroStrategy BI source. Fixed by improving preview extraction performance for MicroStrategy sources.
- Fixed a bug where Alation Analytics containers did not auto-start after a server reboot. Enabled a new start-up systemd service `aacontainers_startup.service` to bring the containers up after a reboot.
- Previously, due to certain race conditions on the server, scheduled query executions failed with the following error: *Could not connect to the database. Please reconnect or reload.* The issue has now been fixed.
- Fixed an issue where the **More nodes** lineage node was not displayed for target tables on Lineage charts. Now, the **More nodes** option will be available for both source and target tables.
- Fixed an issue where the **Curation Progress** report on My Stewardship Dashboard auto-refreshed a short time after loading filtered results and reset to the data source level. The issue was happening due to a race condition that triggered a re-fetch of the table data. The issue has been fixed.
- Corrected a number of typos in the User V2 API documentation.

1.5.8 RELEASE 2022.4.4

build 14.4.4.26642

Note: An Alation Cloud Service release.

This build is not available for customer-managed (on-prem) instances.

- Added a fix to ensure high availability containers are placed in different zones from those of the core `alationfc` container, which guarantees high availability.
- Fixed a bug to exclude configuration replication in cloud native architecture.
- The first user on the Enterprise Edition can now change their password on the **User Profile** page after logging in.
- The **Forgot Password** functionality can now be used by the first user on the Enterprise Edition before they log in.
- Fixed an issue where users were not able to connect to the Alation Analytics database in Compose using their own Alation Analytics account. Added the Connect privilege on the Alation Analytics database for users who create Alation Analytics accounts. This ensures users can connect from Compose.
- Added internal provisioning enhancements.

1.5.9 RELEASE 2022.4.3

build 14.4.3.25341

New Features

Granular Access Permissions for Queries

Note: In 2022.4.3, the Granular Access Permissions functionality for queries is not enabled by default. The feature can be made available to you for preview on demand. The general availability release of this functionality is planned for version 2023.1 (February, 2023).

Previously, Alation didn't have the capability to set granular access permissions for query objects that would allow only certain users to view, edit, or run them. Starting with this release, you can use the new **Sharing & Access** feature to configure access to queries in Compose and the catalog controlling who can view, edit, and run them. Levels of access supported for queries are: *No access*, *View Access*, *View & Run Access*, *View Run and Edit Access*, and *Owner*. The access rules are set on individual query objects and apply to both unpublished and published queries. For new queries, the default is *View & Run Access* for Everybody.

If you are using the `alation_conf` parameter `alation.catalog.unpublished_query_visibility_level=not_visible` to control visibility of unpublished queries, note that this setting will take precedence over individual access permissions on unpublished query objects. If the parameter `alation.catalog.unpublished_query_visibility_level` is currently set to `not_visible` (unpublished queries can only be accessed by Owners and users with whom they were explicitly shared), then after enabling the **Sharing & Access** feature, unpublished queries will remain invisible to users even if they are given *View Access*

in the user interface. We do not recommend changing the parameter `alation.catalog.unpublished_query_visibility_level` to visible before all protected unpublished queries have been assigned the correct access privileges.

User V2 Public API

A new public API **User V2** is now available to users with the Server Admin role. The API introduces two endpoints:

- GET `{protocol}://{base_URL}/integration/v2/user/`—Lists all users.
- GET `{protocol}://{base_URL}/integration/v2/user/{id}/`—Fetches an individual user.

Compared to the already existing API User V1, the User V2 API allows retrieving more user attributes (for example, *user created date* and *user last login date*) as part of the request.

Suspend Duplicate User Account(s) Using API

Two new endpoints were added to the **User V1** public API to provide an alternative way to remove duplicate user accounts. This API is available to users with the Server Admin role.

- GET `{protocol}://{base_URL}/integration/v1/generate_dup_users_accts_csv_file/`—Creates a file with a list of duplicate accounts.
- POST `{protocol}://{base_URL}/integration/v1/generate_dup_users_accts_csv_file/remove_dup_users_accts/`—Updates user accounts in Alation and eliminates duplicates.

Send Email Notifications About the Lineage V2 to V3 Migration Status

The Lineage V2 to V3 migration script now supports the ability to specify emails to send notifications to users about the lineage migration job status.

Improvements

- Added the ability to limit Bulk Utility to working with Articles and Terms only. This is controlled by the **alation_conf** parameter `alation.feature_flags.DEV_enable_bulk_utility_gating`. When this development flag is turned on:
 - If a user attempts to download saved searches that contain objects other than Articles or Terms, those objects will be ignored and not present in the downloaded payload.
 - If a user attempts to upload any files which contain non-Article non-Term rows, those rows will be ignored and no objects created or updated for those rows.
- For the custom API BulkMetadata V2 (available on demand), added flags `create_new` and `replace_values`. BulkMetadata V2 API will update article titles when `replace_values` is false and `create_new` is true.

Bug Fixes

- Fixed an issue where the headline color on the customized homepage was reset to the default dark gray after the upgrade to 2022.4. Now, the custom headline color will stay as was.
- Fixed an issue where for some queries ingested during QLI, the Alation SQL parser issued an error containing an encoded form of emoticons. This exception caused QLI to fail. Now, this exception will be handled gracefully so that the queries causing the error do not cause QLI to fail.
- Fixed a user interface issue with the Lineage diagram where column-level connectors were not highlighted when users selected the corresponding dataflow object node. Now, when users click a dataflow object, its upstream and downstream column-level connectors will be highlighted.
- Improved loading time for catalog pages that contain single picker or multi-picker custom fields. Previously, some customers reported noticeable slowness when opening such pages.
- Added logic for the column filter on catalog pages of column objects to search for characters % and _ in column names. Previously, they were interpreted as wildcard characters and the search rendered no result.
- Fixed an issue where users were unable to create a Change Request workflow for the **Type** field on policy objects. The approval of such change requests resulted in an error. Now, we added support for updating the **Type** field for policy objects using Change Request workflows.
- Fixed an issue where Bulk Utility might incorrectly associate a migrated term with an RDBMS object stored as a value in an Object Set field as such values do not contain unique identifiers. Added logic to handle non-unique RDBMS Object Set field values while using Bulk Utility.
- Fixed an issue where custom fields were not displayed in the left-side filter panel on the Search results page when **All** was selected from the **Object Type** filter.

1.5.10 RELEASE 2022.4.2

build 14.4.2.23237

Encryption Security Fix

Important: Action required after this Alation update on customer-managed (on-prem) instances.

Alation Cloud Service customers will receive this update as part of the routine maintenance, and no action is required on the part of instance administrators.

This patch includes a security fix removing a risk identified in the current encryption mechanism which makes Alation potentially susceptible to key reuse attacks. Information about this risk is available in [November 3, 2022 - Encryption Key Reuse Security Advisory](#) available on Alation Community (requires Community login).

The fix ensures that all new data will be properly encrypted using new, randomized initialization vectors appended to the ciphertext and eliminates the possibility to reuse the non-unique initialization vector.

On customer-managed instances, after the update to this patch, the existing data needs to be re-encrypted. Re-encryption will run as a background process and no separate downtime is required apart from that needed for the regular version upgrade. Users can work in Alation while the re-encryption job is in progress. For re-encryption steps, refer to [/installconfig/Update/ReencryptData](#).

Other Fixes

- The Impact Analysis report will now exclude temp objects by default. This is based on the **alation_conf** parameter `lineage-service.filters.exclude_temp`, which is now `True` by default. Previously, the default view of the Impact Analysis report included temp objects.
- Fixed an issue with the order of steps in change request workflows where the change requestor and reviewers saw a different order of steps. This fix ensures that the change requestor and reviewers see the same order of the workflow steps.
- Optimized the lineage service query that is used to enable the Lineage tab on catalog pages. The fix improves page load experience for catalog objects that support lineage.
- Improved page load time for catalog pages of view objects with long definition queries from minutes to a few seconds. Now, Alation will skip SQL highlighting for view definition queries with more than 1,000 characters to ensure faster page load time.

1.5.11 RELEASE 2022.4.1

build 14.4.1.22725

Note: An Alation Cloud Service release.

This build is not available for customer-managed (on-prem) instances.

- Updated the Connected Sheets add-on installation URL.
- Fixed an issue with publishing manual lineage and adding lineage via API, where new lineage did not appear on the diagram after publishing.
- When migrating a data source connection details to OCF, Alation reads the data from the `native_to_ocf_mapping` field in the manifest file. A connector manifest file will have this key when the connector implements `RdbmsNativeConfigurationMap` or `BINativeConfigurationMap`. This change provides a default value if the field does not exist in the manifest file.
- Fixed an issue of **pgBackRest** failing to initialize.
- Previously, in case of an OCF metadata extraction job failing, the object counts (for example, tables per schema and attributes per table) were incorrect and attribute breadcrumbs were missing the schema and table references. After this fix, the counts and the objects references should be populated correctly even when extraction fails, and the incorrect counts should be corrected during the next successful extraction.
- Fixed an issue where the `backup_all` action stalled at the Event Bus backup step, resulting in no backups getting generated.
- On the settings page of data sources, in the Migrate to OCF dialog, added a text box for validation requiring users to type the connector name and version. After the validation, if the text matches with the value in the dropdown list, the **Migrate** button becomes enabled.
- The Object Set values will now be properly handled by Bulk Utility (Article to Term migration).
- Fixed an issue with collection nodes not expanding on Lineage diagrams. An error occurred instead. Now, this action is available to users.
- Fixed an issue where invalid connection strings were successfully added to the **Connection Settings** dialog instead of Alation throwing an error during the Compose connection testing. This fix ensures that connection errors are properly handled.

- Previously, for Query Service-based connections, DDL and DML queries returned the update count (number of rows updated) in the `ResultSet` object instead of the `Statement` object as in JDBC-based connections. This caused the system to not process update counts properly. This fix ensures that the update counts are extracted from the `ResultSet` object for Query Service-based query executions.
- Fixed an issue where incremental MDE was not triggered due to a conditional check in Compose, disabling the incremental MDE or QLI flows for Query Service-based data sources. This change removes the conditional check.
- Fixed an issue where bulk actions failed for the object types `doc_schema`, `docstore_folder`, and `docstore_collection`. Now, bulk actions with these object types will succeed.
- Fixed an issue with uploading Lexicon terms from a CSV file where the upload resulted in inconsistent handling of existing confirmed titles.

1.5.12 RELEASE 2022.4.0 — General Availability

build 14.4.0.21319

Alation Anywhere

Alation Anywhere is a set of features that integrates Alation with other software that you use, so that you can access and use Alation data anywhere.

Alation Anywhere for Slack

With Alation Anywhere for Slack, rich contextual information about data from Alation is surfaced via previews in Slack. Data teams can search the catalog, preview helpful metadata at a glance, and navigate to the object in Alation to get more information. This could be anything from a business glossary term, to a table in a database, or a SQL query someone has published and described.

Alation Anywhere for Tableau

With Alation Anywhere for Tableau, users of Tableau can see certain metadata from Alation directly within Tableau. This enables Tableau users to see data descriptions, warnings, and more in Tableau itself. Alation serves as the single source of truth while supporting data users in their chosen context.

Connected Sheets

Connected sheets is a new spreadsheet tool using Google Sheets that helps users discover data sources in the catalog, download data using filter and sort criteria, and manage downloaded sheets including auto-refresh. Customers must be on the newly released multi-tenant Alation Cloud Service and a Creator or Explorer License is required.

Catalog

Term and Glossary Navigation Improvements

Term and Glossary Navigation improvements include:

- Added the ability to recursively add and remove terms in Glossary to a Domain.
- Added new optional columns to both Term tables (on the Hub as well as Glossary detail pages): **Last Updated**, **Created** and **Term Type**.
- Term Table column selection can now have a saved default view chosen by Catalog Admins.
- Added new custom fields search facet for Glossary Terms.
- Added breadcrumbs to Term catalog detail pages, and added **Last Updated** to Term and Glossary Properties panels.
- Added a new **Glossary Terms Total Count** feature flag (off by default), which controls whether the total Term count is calculated and displayed at the top of the Glossary Hub Terms table.

Terms Public API

The Terms Public API enhancements include:

- The GET endpoint of the Terms Public API now includes additional fields to match the CREATE endpoint plus the unique ID.
- Added an UPDATE endpoint to the Terms Public API.
- Indicate in the documentation that search, skip, limit, and deleted=true are supported query parameters in the Term bulk GET request.
- The Terms Public API List endpoint is now object-level permissions aware.

Bulk Utility

Bulk Utility is an admin-only feature that enables users to specify which Articles they would like to migrate to Terms easily by leveraging download of a saved search to csv. All critical Article fields will be contained within newly created Terms and associated with Glossaries and Term Types as desired. The Bulk Utility can also be used to create and edit multiple Terms and Articles at a time. Links to detailed instructions for use can be found on the feature page itself.

Customized Search User Interface (UI) Improvements

The **More Filters** drawer closes instead of refreshing when users switch otype filters. Flags and Star Watched History filters are compatible with custom filter settings, and move to the filter drawer when removed from the outer filters.

Persona-Based Homepage Preview Mode

Custom Homepages can now be previewed using the **Preview** button. The new homepage renders so that users can see their custom homepage in action before enabling it for viewing.

Propagated Data Quality Table on Downstream Data Objects

Alation's Data Quality service now includes built-in quality values being generated from ingestion from OCF data sources. On detection of deleted data objects, Data Quality creates a 'deleted' quality value on that given object which then propagates to downstream nodes on the lineage graph. This information will be displayed for downstream Tables, BIReports, and BIDataSources and is surfaced on the new Data Health > Upstream Issues tab for each respective catalog page.

Search Ranking Improvements

Catalog Admins can now influence the order of how search results appear through ranking by object type. The customization is available via *alation_conf* settings.

Elasticsearch Upgraded to 7.17.6

Upgraded the Elasticsearch version from 7.17.3 to 7.17.6. This was completed for both on-premise and cloud environments.

Other Catalog Improvements

Added a search flag `alation.search.enable_search_with_literals`. The search will return the exact match for a query with a period.

Governance App

Create Publicly Accessible APIs

Created a Public API to list Policy Groups.

Policy Center Template Support

Added support for various multi-template layout functionality for business policy templates, including displaying the business policy template name in the All policies and Policy group links table, registering a new built-in template for business policy, and migrating the existing built in template to a non-built in template.

Updated Database Configuration Settings Modal

Updated the Connection Settings dialog used in Add Data flow of Policy Center, bringing consistency and providing capabilities around managing multiple connections.

Enabled “Select All” for Stewardship Workbench Bulk Actions

One-Click approach to select up to 10K objects and update them all at once using bulk actions.

Improved Add/Remove Stewards Bulk Action

Merged the add or remove Stewards bulk action with the capability to add or remove people-sets as Stewards.

Object-level Permissions for Stewardship Workbench Actions

When a Steward is performing bulk actions in the Stewardship Workbench, the Steward who initiated a bulk action has object-level permissions including edit access to both the selected objects and the selected field.

Policy Center UI/UX Updates

Added the “Create Policy Group” and “Create Policy” buttons to replace the “Create New” button on the Policy Center homepage.

Compose

Query Service Integration with Compose

Compose now supports connections requiring Alation Agent. This is achieved through the Query Service and requires that an OCF Connector has been modified to support this. This support includes all query execution, manually in either Compose or query forms, using scheduled queries, or using Excel live reports and data upload. This feature requires Query Service enabled OCF Connectors. Oracle OCF Connector is the first supported one. Additional supported OCF Connectors will be announced as they become available.

Alation Analytics

Support for Deploying Alation Analytics V2 on Docker as Non-root User

Support is added for deploying Alation Analytics V2 on Docker daemon running as a non-root user. Alation Analytics V2 works in Docker rootless mode for the following platforms:

- Redhat 8
- Fedora 33, 34
- Ubuntu 18
- Debian 10
- Centos 8
- Oracle Enterprise Linux (OEL)

For installation of docker in rootless mode, see <https://docs.docker.com/engine/security/rootless/>

Connectors

Please note that connectors are released independent of Alation releases. The following connectors and connector changes are now available unless specified otherwise.

LA OCF Connectors

- Azure Blob Storage OCF Connector is LA. This connector enables cataloging objects/files in a file/folder heirarchy for Azure Blob Storage and ADLS Gen 2.
- Databricks Unity OCF Connector is LA. This connector provides connectivity to Databricks' Unity Catalog (AWS and Azure) for metadata extraction.
- SSRS OCF Connector is LA. The SSRS OCF connector does not support Sharepoint mode, and is verified with 2016, 2017 and 2019.
- Tibco Spotfire OCF Connector is LA. This connector catalogs Tibco Spotfire as a Business Intelligence source in Alation and catalogs Tibco Spotfire artifacts such as Folders, Reports, Report Fields, Data sources, and Connections.

GA Connectors

- [Amazon S3 OCF Connector](#)
- Denodo OCF Connector
- CDATA Azure Analysis Services CustomDB Connector
- Impala with CDH OCF Connector
- Google BigQuery OCF Connector
- MemSQL OCF Connector: All features are implemented as per the support matrix with the exception of LDAP.
- [Qlik Sense Connector](#)
- SAP Business Objects OCF Connector
- Tableau OCF BI Connector

Other Connector Updates

- CDATA SSAS OCF Connector: Added Expressions in the Catalog and an option to catalog hidden tables
- Google BigQuery (Native): Profiling timeout is behind a configuration flag to be increased as required.
- Looker OCF Connector: Upgraded the latest SDK, Fixed the BI connection and Reports preview issues.
- Power BI OCF Connector: Populate report owner field from scanner API response for Power BI Report
- Snowflake: Added support for Snowflake system tags for data classification and updated section headers.
- For BI object catalog pages, the **Source Comments** field is now renamed to **Source Descriptions**.

Alation Connector Manager

With Alation Connector Manager 1.2.1.1168, you can now configure Alation Connector Manager to use the host network so it appears to be coming from the same host as Alation. This may be needed if your network is configured with an Access Control List. See the Alation Connector Manager configuration instructions for more information.

Connector Platform Service

Public APIs for OCF Data Sources and Connectors

- APIs for creating data sources, listing all data sources, getting details for a given data source, and getting details for all OCF connectors are now supported for OCF data sources.
- Metadata extraction and selective extraction are now supported via Public APIs for OCF data sources.

Important: The existing Data Sources API (v1) will no longer support OCF data sources. It could previously be used to get limited data for OCF data sources. Use the new OCF Data Sources API (v2) for working with OCF data sources.

Lineage

Lineage V3 Is Now Enabled by Default

If the instance is for a new Alation customer, Lineage V3 is enabled by default. If the instance being upgraded to 2022.4 or later has Lineage V3 disabled, it will remain disabled after the upgrade to 2022.4. When Lineage V3 is enabled, the Lineage V3 functionality, such as manual lineage, can be used immediately without a Celery restart. All V2 functionality is also available in V3, using V3 just adds extra functionality and improves the lineage graph performance.

Column Level Lineage (CLL) for DB2 Is Available

Column Level Lineage (CLL) for DB2 is now generally available (GA).

Column Level Lineage (CLL) for Azure Synapse Is Available

Column Level Lineage for Azure Synapse is now generally available (GA). Enabled feature flag to turn on Automated Column Level Lineage for Azure Synapse data source.

Lineage Depth Based Filtering improvements

Ability to 'View More' lineage based on lineage-service.max_depth configuration. The default setting is to view all lineage. Change the max_depth value to set the initial view distance for both upstream and downstream.

Default View Settings Page

Default View Settings page adds the ability for server admins to set default settings for lineage graphs through the **Settings** page. The settings can be seen on the customize catalog section of the settings page, that is, **Settings > Customize Catalog > Lineage**.^{*} After the settings are modified, the lineage service needs to be restarted.

Lineage V2-V3 Migration Enhancements

When the Lineage V2 to Lineage V3 migration is run, email notifications are sent to instance admins to inform them about the outcome of the migration. This is for any outcome whether successful or failed. Occasional progress status notifications are also sent for long running jobs. This means monitoring the logs as the migration runs while still recommended is not strictly required unless something goes wrong.

General Platform

Upgrade Kafka for On-Premise Instances

Kafka upgraded to version 3.1.1.

Cloud

Cloud Console Dashboard

Created a multi-region cloud dashboard to manage cloud accounts.

Certifications

Connector Certifications

OCF PostgreSQL EC2 14.5 Version is certified.

Driver Updates

Driver updates include:

- Azure SQL DB OCF Connector is certified w/ 11.2.0 Driver
- SQL Server OCF Connector is certified w/ 11.2.0 Driver

Known Issues

Homepage Headline Color Turns to Dark Gray

After upgrading to 2022.4, any homepages that use the default headline color will default to a dark gray. You can fix this by editing your homepages to enter the desired headline color manually after upgrading. See /admins/CustomizableHomePage/CustomizeHomepage2022_3To2023_3_3 for help editing your homepage. **(FIXED in 2022.4.3)**

Domain Default Field Should Not Show “remove” Icon

When navigating to a custom template > Policy template, and hover over the Domain field, the “remove” icon should not display.

Glossary Membership Modal Does Not Show All Glossaries

When viewing the Glossary Membership modal, all glossaries are not shown because of a pagination issue, buttons are disabled, and it's possible selected glossaries are shown incorrectly.

Relational Update APIs Inconsistency

New Data Quality values added to Columns will be shown to be ingested successfully but will not be displayed on the column's table.

NoSQL User Interface Issue

The NoSQL catalog table data requires to return an indicator whether a schema has nested children with it or not. There is a known issue causing the user to be unable to see a child elements link when the child elements are more than the depth configured, and unable to see the Load More button to load elements in batches.

Workaround: Configure the depth flag “alation.feature_flags.docstore_tree_table_depth” to a higher number to see the child elements.

Fixed Issues

- Fixed an issue where the backup process failed at the Event Bus backup generation stage resulting in no backup file being created in the backup folder.
- Alation Analytics V2 no longer shows deleted flags.
- Fixed issue for proper handling of data values in scientific notation when unwrapping.
- Fixed issue to give access to admins other than server admin when updating OCF settings.
- Cleaned up older backup files failing because of a permission issue.
- Fixed issue with Alation's Relational Integration API with column calls on SAP BW Data Source.
- Fixed issue causing deleted terms through Public APIs to display on Search Page.
- Fixed issues where users cannot copy rows from results of Run Form.

- Postgres folders permissions update, secured by default.
- Fixed issues where search filters are not cleared when more than four options exist.
- Fixed issue causing high CPU load when backup run after upgrading to 2021.4.12.
- Added ability to run queries using script mode for Google Big Query (GBQ) datasource.
- Fixed issue causing Column profiling for Custom DB data sources to not work as expected when custom queries are used for sampling tables.
- Term-accepting ObjectSet fields that are hydrated with Terms will now properly allow the object types they are placed on to be indexed by search.
- Fixed issue causing backup status missing in health checklist.
- Fixed issue causing homepage to always load in default state, instead of custom homepage.
- Checkbox styling is now consistent across all browsers.
- Fixed issue causing tables to display governance count percent mismatch.
- FileSystems in the left pane are now sorted by Title.
- Restored the ability to add PDF and Excel files as types for upload via Rich Text Fields.
- After ingesting schemas via API for a NoSQL Virtual Data Source, record types should preserve capitalization style.

1.6 Release Notes 2022.3 (General Availability)

1.6.1 RELEASE 2022.3.15

build 13.4.15.29988

Note: This release is available for customer-managed (on-prem) instances.

- Previously, due to certain race conditions on the server, scheduled query executions sometimes failed with the following error: *Could not connect to the database. Please reconnect or reload.* The issue has now been fixed.

1.6.2 RELEASE 2022.3.14

build 13.4.14.26176

Note: This release is available for customer-managed (on-prem) instances.

- Fixed an issue where the metadata and data dictionary upload into a virtual data source from a CSV file failed with the *No correct data found* error if the file was more than 2.5 MB in size. Addressed by improving the logic of dialect detection for input files.
- Previously, the **Download Latest Usage Data** button on the **Admin Settings > Reporting > Upload Manually** page did not initiate the data download. This issue was fixed, and all usage data download steps will work as intended.

1.6.3 RELEASE 2022.3.13

build 13.4.13.23874

Encryption Security Fix

Important: Action required after this Alation update on customer-managed (on-prem) instances.

Alation Cloud Service customers will receive this update as part of the routine maintenance, and no action is required on the part of instance administrators.

This patch includes a security fix removing a risk identified in the current encryption mechanism which makes Alation potentially susceptible to key reuse attacks. Information about this risk is available in [November 3, 2022 - Encryption Key Reuse Security Advisory](#) available on Alation Community (requires Community login).

The fix ensures that all new data will be properly encrypted using new, randomized initialization vectors appended to the ciphertext and eliminates the possibility to reuse the non-unique initialization vector.

On customer-managed instances, after the update to this patch, the existing data needs to be re-encrypted. Re-encryption will run as a background process and no separate downtime is required apart from that needed for the regular version upgrade. Users can work in Alation while the re-encryption job is in progress. For re-encryption steps, refer to [/installconfig/Update/ReencryptData](#).

Other Fixes

- Improved loading time for catalog pages that contain single picker or multipicker custom fields. Previously, some customers reported noticeable slowness when opening such pages.
- Fixed an issue in the user interface where the hint under the automatic schedule sections in settings of data sources was **in UTC**. This was misleading as schedules are set in local time zones.

1.6.4 RELEASE 2022.3.12

build 13.4.12.21484

Note: This release is available for customer-managed (on-prem) instances.

- Fixed an issue where users could not create temp tables in Google Big Query using Compose. Added the ability to run queries using script mode for Google Big Query data sources for the native (built-in) connector. Temp tables can now be created by running the corresponding queries as scripts.
- Previously, when the metadata extraction job failed for an OCF data source, this resulted in incorrect object counts (for example, tables per schema or columns per table) for extracted objects and in missing schema and table references in column breadcrumbs. After this fix, object counts and object references will be populated correctly even if metadata extraction fails; or they will be corrected during the next successful metadata extraction job.
- Fixed an issue with uploading Lexicon terms from a CSV file where the upload resulted in inconsistent handling of existing confirmed titles.
- Fixed an issue where enabling pgbackrest as the backup tool resulted in an error.

1.6.5 RELEASE 2022.3.11

build 13.4.11.20480

- Fixed an issue where custom fields were not displayed in the left-side filter panel on the Search results page when **All** was selected from the **Object Type** filter.
- Fixed an issue where users with the Catalog Admin and Source Admin roles who were assigned as Data Source Admins to OCF-based data sources were unable to edit the data source settings. Now, Catalog Admins and Source Admins added as Data Source Admins can fully manage the data source settings (edit the **General Settings** tab, run MDE, sampling, and QLI).
- Fixed a bug where users were unable to copy multiple rows or columns from query form results. Now, they can use their mouse or touchpad to select a part of the expanded query results table and use copy-and-paste to copy the selection to a document.

1.6.6 RELEASE 2022.3.10

build 13.4.10.19871

Data Catalog

- Fixed an issue with term-accepting Object Set fields where the term objects were not indexed in the catalog and did not appear in search after adding terms as values to such Object Set fields. Now, the term objects that are added as values will be properly indexed by search.
- Fixed an issue where a homepage customized for a role was loading slowly exposing the default homepage in the first few seconds. Resolved by introducing a loading screen that covers the page while the customized homepage data is being loaded. The homepage data is then rendered on screen as expected.
- Fixed an issue where in some cases Compose showed *No data to display* message in the results pane, simultaneously displaying the *Retrieved N rows* message. The issue was caused by failure to unwrap values in scientific notation bigger than `Number.MAX_VALUE` and smaller than `Number.MIN_VALUE`. Addressed by adding logic to handle corner cases in the `unwrap_scientific_notation` method which caused this rendering failure.

Multi-Tenant Cloud Architecture

- Fixed an issue where the `alationfc-ha` container was not updating because we passed the wrong container name. After this fix, we'll dynamically pass the container name, which ensures that the update of `alationfc-ha` containers works correctly.
- Fixed an issue with pre-upgrade failing for the `alationfc` pod. The `alationfc` pod has multiple containers and all calls to `connect_get_namespaced_pod_exec` should specify the container name or pre-upgrade will fail. Now, the container names will be passed in the calls, ensuring that pre-upgrade is successful.
- Fixed an issue where search incremental indexing did not work for some of the cloud tenants. Fixed by removing aliases from stale indices while running search indexing scripts. This will prevent multiple indices from using the same alias.
- Fixed an issue with high availability set up on multi-tenant cloud architecture where both `alationfc` and `alationfc-ha` pods read the same configuration. Now, they read different groups of parameters in `alation_conf`, as expected.

- Fixed a check for the `live` and `live_large` index on the `alationfc` pod before backup restore.
- Temporarily disabled the active task monitoring user interface in **Admin Settings > Monitor** on cloud service versions of Alation to further investigate an issue with queuing of a number of task types.

1.6.7 RELEASE 2022.3.9

build 13.4.9.18541

- Fixed an issue with object sync tasks for lineage where they were reported to take a very long time to run. This fix optimizes the runtime of each object-sync task in the Celery component. As the result of the fix, on the **Active Tasks** page (Admin Settings > Monitor > Active Tasks), you will notice that the Celery tasks named `lineage_object_syncing_task` that used to run on the `Default` queue now run on the `Lineagepublishing` queue. The tasks will also be fewer in number than before. Previously, for every 10,000 objects ingested into Alation during MDE, there were ten of the lineage object-sync tasks queued or running at the same time. Now there will be one task queued or running for every 100,000 objects ingested during MDE. The fix ensures that the lineage object sync tasks do not block other high priority tasks that may be simultaneously running on the `Default` queue.

1.6.8 RELEASE 2022.3.8

build 13.4.8.17648

An Alation Cloud Service release.

This build is not available for customer-managed (on-prem) instances.

1.6.9 RELEASE 2022.3.7

build 13.4.7.17205

Note: This release is available for customer-managed (on-prem) instances.

- For the Alation Analytics application, upgraded the library versions for the `postgres` and `Transform` containers to eliminate potential security risks identified in the previous library versions.
- Fixed an issue with the backup cleanup process where old backup files were not automatically removed based on the value in `alation_conf` parameter `alation.backup.data_dump_versions`, causing the disk space to fill up. Now, old backups are cleaned up successfully based on the configuration in `alation_conf`.
- Fixed an issue where Alation previously did not allow console calls that would display certain internal API errors, which made error reporting more difficult. These calls are now allowed and can be caught in the developer console and reported. Additionally, fixed an issue with the error page where in certain cases, it did not differentiate between server-side and client-side errors. The error page now uses the unknown error code to generate the *unknown error* page instead of the *server error* page. This should make it more obvious to disambiguate between a client side issue vs. a server side issue.
- Fixed an issue with the Relational Integration API (`<your_Alation_URL>/openapi/integration_apis/`) where an invalid foreign key reference caused the creation of an invalid column. If an invalid column was referenced as a foreign key during the creation of another column with the Integration API, an invalid column was created. After this fix, the API will create a foreign key association only if the foreign key is valid.

- Added a new **alation_conf** parameter `alation.profiling.single_table_profiling_timeout` that allows setting the profiling timeout. Previously the timeout was hardcoded. The default is 600 seconds.
- When adding or removing terms using the public API, the Elasticsearch framework does not re-index these terms. Prior to this fix, users had to visit the catalog page of a term for newly created or deleted terms to show up in or be cleared from the search results. This fix will now correctly add or remove the terms created or deleted using the public API to or from the Alation search.
- Fixed a bug with the **Policy Center** page where the Stewards column was not displayed in the **Policy Groups** table. Now, the Stewards column is displayed in the user interface as expected.
- Fixed an issue with the user interface of the Compose web application where the UI elements on the top toolbar became misaligned when users reduced the size of the browser window. This fix adjusts the styles of the Compose toolbar to ensure a better responsiveness of the Compose page when the browser window is resized.
- Fixed a bug where the alphanumeric sorting of data sources in the left sidebar did not apply to data source names containing special characters, which caused a mismatch between the list in the left sidebar and the list on the **Sources** page. Now, the data source list in the left sidebar is sorted with the special characters taken into account.
- In version 2022.3.3 or earlier versions of Alation 2022.3, the upgrade of OCF RDBMS connectors to the latest available version could result in the following error:

Installation failed, changes not committed. Wrong data schema in manifest file: {"features": ["Field features can take one of the following values: [BI_GENERAL, BI_MDE_GBM_V2, BI_CERT_GBM_V2, [...]}. Invalid value : RDBMS_QUERY_SERVICE."}]}

This issue happens due to Alation not recognizing RDBMS_QUERY_SERVICE as a new feature added to connectors. The issue is fixed in this patch release. We recommend updating to 2022.3.7 or later, when available, to be able to upgrade to and use the latest versions of OCF RDBMS connectors.

1.6.10 RELEASE 2022.3.6

build 13.4.6.15702

An Alation Cloud Service release.

The build is not available for customer-managed (on-prem) instances.

1.6.11 RELEASE 2022.3.5

build 13.4.5.14981

This build is not available for customer installations.

1.6.12 RELEASE 2022.3.4

build 13.4.4.14662

An Alation Cloud Service release.

This build is not available for customer-managed (on-prem) instances.

1.6.13 RELEASE 2022.3.3

build 13.4.3.13911

- Fixed an issue when a bulk action using Stewardship Workbench failed if the objects selected for the action included at least one file or directory object. Now, bulk actions will successfully apply to file and directory objects.
- Fixed an issue with checking the OCF SDK version for BI connectors developed on the unified version of the SDK (Tableau OCF connector). Previously, this issue prevented an OCF BI source in Alation from establishing a successful connection to the BI system.
- Previously, when term objects were added to Object Set fields as values on the catalog pages of articles, the corresponding back references were not displayed on the catalog pages of terms. After this fix, back references to articles will be displayed under the **Relevant Article** section on the catalog pages of terms.
- Fixed an issue where the backup process changed the permissions on the Postgres folders on the Alation server to world writable. With this fix, the Postgres folders permissions are secured by default. Also, the Postgres folders permissions will be auto-corrected in case an admin, a privileged user, or a process changes them to a state that is different from the expected state.
- Previously, if the **Description** field included an at-mention or a link, this information was not exported with the API. The at-mentions and links were left out of the exported field value. After this fix, the full Description field will be exported.

Known Issue

Note: Action is required after updating your instance to 2022.3.3.

In this version, users may run into the error *Sorry, something went wrong* when they access the Policy Center page. The error occurs due to an issue with loading the policy objects when there are more than 9,000 policy objects to load. After updating to 2022.3.3, go to **Govern > Policy Center** and check if you experience this error. If this is the case on your instance, apply the following workaround:

1. From the Alation shell, change the value of the **alation_conf** parameter `alation.api.object_level_permissions.limit` to 9000.

```
alation_conf alation.api.object_level_permissions.limit - s 9000
```

2. Restart the Web component.

```
alation_supervisor restart web:*
```

3. In Alation, go to **Govern > Policy Center**, and verify that the error no longer appears.

1.6.14 RELEASE 2022.3.2

build 13.4.2.13511

An Alation Cloud Service release.

This build is not available for customer-managed (on-prem) instances.

1.6.15 RELEASE 2022.3.1

build 13.4.1.13309

An Alation Cloud Service release.

This build is not available for customer-managed (on-prem) instances.

1.6.16 RELEASE 2022.3.0 — General Availability

build 13.3.0.12402

Catalog

Role-based Homepage Customization

Server or Catalog Admins can now design separate homepages for each Alation user role, so every role can get targeted and relevant content. The new homepage design capabilities are available in **Admin Settings > Customize Homepage**. You can create homepages for one or multiple roles, adding the content that your users find the most engaging.

Alation Marketplaces

[Alation Marketplaces](#) is now live. Alation Marketplaces is a new Alation site where you can search for datasets like those on the AWS Data Exchange & Snowflake Data Marketplaces. A **Marketplaces** tile should now appear in your **Apps** drawer. See [/welcome/CatalogBasics/Marketplaces](#) for complete details.

Glossary and Glossary Term Permissions Management

Server and Catalog Admins can set access permissions for a glossary and choose if they want glossary terms to inherit those permissions or set them individually. To configure access permissions, use the **Access Settings** link available under the **Manage** menu on the catalog page of a glossary. The access can be set to:

- **Public** – Everyone can view and edit.
- **Editing Restricted** – Everyone can view and only specific users can edit.
- **Private** – Only the creator and Catalog and Server Admins can view and edit.

Editing includes the ability to edit any fields on the page, linking or unlinking glossary terms for a glossary, and the ability to delete the object.

Customizable Search Filters

Catalog and Server Admins can now customize the search filters to surface the most relevant filters on the Search page, tailoring the default Alation search experience to better suit their users' needs. In **Admin Settings > Customize Catalog > Customize Search Filters**, admins can select and deselect **Left Side Filters**, **Hidden Filters**, and **Object Type Filters** and rearrange the selected filters. The custom search filter selections will be reflected on the Search page for all users. The **Left Side** and **Hidden Filter** settings will be applied to the **Filters** panel on the left side of the Search page. The **Object Type Filter** settings will be reflected in the top object type filter panel of the Search page.

Adding Glossary Terms to Object Set Fields

Server and Catalog Admins can now add the glossary term object type as an option in the Object Set custom fields. As a result, users can link catalog pages to glossary terms using Object Sets.

Object and People Set Custom Fields in the Main Column of the Template

Server and Catalog Admins customizing the object templates in the catalog can now add or move the Object Set and People Set custom fields to the main column of the template. When added to the main column, an Object or a People Set field will be displayed as a table listing links to the catalog objects added as values. Hovering over a row of the Object or People Set table will reveal an action icon allowing users to add or remove values to or from the field.

Governance App

Change Management Workflow for Glossaries and Glossary Terms

Server and Catalog Admins now have the ability to configure change management workflows for glossary and glossary term catalog pages. They can enroll their glossaries and terms into a change management workflow, which enables users to suggest changes to the custom fields on the corresponding catalog pages. The changes proposed are made available for approval by the reviewers designated in the workflow configuration. The updated values become visible to all users only after the approval by the designated reviewers.

Change Management Workflow for Policy Groups

Change management workflow is now available for policy group objects. Catalog and Server Admins can enroll policy groups into workflows, which enables users to suggest changes to the custom fields on the corresponding catalog pages. The updated values become visible to all users only after the approval by the designated reviewers.

Bulk Updating Object and People Sets In Stewardship Workbench

Server and Catalog Admins can now bulk update Object and People Set custom fields using the corresponding bulk actions available in the Stewardship Workbench.

Policy Center Permissions Management

Server and Catalog Admins can set and manage the view and edit permissions for each policy group and choose if they want policies to inherit those permissions or set them individually. To configure the access permissions for a policy or a policy group, use the **Access Settings** link available under the **More** menu on the policy and policy group catalog pages. The access can be set to:

- **Public** – Everyone can view and edit.
- **Editing Restricted** – Everyone can view and only specific users can edit.
- **Private** – Only the creator and Catalog and Server Admins can view and edit.

Additionally, Stewards have been added to the list of roles that are able to create policies.

Compose

Advanced Query Scheduling

Compose users can now use Cron expressions to specify a custom schedule for a query. This ability is available in the **Schedule Settings** dialog in Compose under **Advanced Scheduler**. Alation will translate the Cron expressions into natural language for users to easily validate their correctness.

Support for Functions in SmartSuggest for PostgreSQL

For PostgreSQL data sources, the SmartSuggest feature in Compose now supports built-in and user-defined functions. As users type their query, they will see the list of available functions in the SmartSuggest dialog and will be able to add a function by clicking the corresponding prompt.

Deferred Query Execution Enabled by Default

Deferred Query Execution is now enabled by default.

Alation Analytics

Groups Filter for the Leaderboard

Added the ability to filter the Alation Analytics Leaderboard available on the **Built-in** tab by groups. Applying the **Groups** filter will narrow the leaderboard report dataset down to the included groups of users. The filter will also apply to the Leaderboard added to the homepage.

Flexible ETL Scheduling

The ETL schedule configurations using the `alation_conf` parameters `alation_analytics-v2.etl.schedule.hour` and `alation_analytics-v2.etl.schedule.minute` now allow more Cron expressions and enable users to schedule the ETL multiple times a day. After changing the values, the Celery Beat component requires a restart.

Note:

- To schedule for every 4 hours:

```
alation_conf alation_analytics-v2.etl.schedule.hour -s '*/4'
```

- To schedule for every 30 minutes:

```
alation_conf alation_analytics-v2.etl.schedule.minute -s '*/45'
```

Lineage

Ability to Delete Lineage Using Manual Lineage

Added the ability to delete lineage links and dataflow objects generated by metadata extraction, query log ingestion, and API using the Manual Lineage capabilities.

Dataflow Source Filter on the Lineage Diagram

Users can now filter the lineage data on Lineage diagrams using the new **Dataflow Source** filter available in the **Filters** menu on the top left of the Lineage diagram. By default, the Lineage diagram shows all dataflow objects and lineage links for a given data object. The Dataflow Source filter allows users to select specific lineage sources and view the lineage data that comes from those sources only.

Lineage Diagram Improvements

There are multiple user experience improvements to the lineage visualization:

- Users can now access the **Lineage** tab from the Column catalog pages.
- The sidebar of the Manual Lineage editor now automatically scrolls as users add more subpaths.
- Previously, when users scrolled down using the mouse wheel to the object preview panel under the Lineage diagram with the focus still on the diagram, the diagram was zoomed in. We removed the automatic zoom on scroll. With a mouse, users can now zoom in with pressing Control and scrolling. With the touchpad, users can zoom by pinching in and out.
- Server Admins can configure the lineage data depth using the **alation_conf** parameter `lineage-service.filters.max_depth`. By default, the parameter is set to -1, which means all the lineage data available for a data object is displayed. Setting the parameter to a positive number configures the number of upstream and downstream lineage levels to display on the lineage diagram.
- Server Admins can exclude all temporary objects from lineage diagrams using the **alation_conf** parameter `lineage-service.filters.exclude_temp`. By default, it is set to False, which means temporary objects are displayed. The parameter is universal and applies to all data sources.

Column-Level Lineage Parser Add-Ons GA

The column-level lineage parser add-ons for Oracle and MySQL are now in general availability (GA). In previous versions, these add-ons were released as beta features.

API

Glossary Terms API

The new Glossary Terms API is now available at `<your_Alation_URL>/openapi/terms`. This API allows Catalog and Server Admins to get, create, and delete glossary term objects. The API exposes new bulk public API endpoints for glossary terms:

- Bulk create using POST
- Bulk list using GET (currently, Title and Description are included)
- Bulk delete using DELETE

New Endpoints in Data Health API

The Data Health API `<your_Alation_URL>/openapi/data_quality/` now supports the following new methods and endpoints:

- Data Health DELETE — deleting of data quality fields and values through the new DELETE action.
- Data Health GET fields — a GET endpoint `integration/v1/data_quality/fields/` to retrieve the list of fields created by ingestion.
- Data Health GET values - a GET endpoint `integration/v1/data_quality/values/` to retrieve the list of values and the objects they have been applied to. The values can be filtered by `object_key` as well as `field_key` allowing users to find which values were applied to specified objects, as well as which objects have specified fields.

V2 Logical Metadata Field Value API

Added a new version of the Logical Metadata Field Value API that supports asynchronous processing with significantly improved processing speed for `write`. The documentation for this API is available at `<your_Alation_URL>/openapi/custom_field_value_async/`. The Logical Metadata Field Value API enables creating and updating of the Title, Description, and custom fields for the following object types:

- Data (source, schema, table, attribute (column))
- BI (server, folder, connection, source, report)
- Dataflow
- Domain
- Policy (business and data policies)
- File system (file, directory)
- Glossary term

Data Health (Open Data Quality Framework)

Improved the user experience for the Data Health flags:

- The Data Health flags, powered by the Open Data Quality Framework (ODQF) APIs, are now displayed in the new **Health** column next to the individual column previews on the table page, on the **Overview** tab and the Columns tab.
- The **Health** column can be sorted.

Connect Platform

Open Connector Framework SDK Enhancements

We implemented multiple enhancements in the OCF SDK and on the Alation server backend that enable connector developers to add more advanced features to the OCF connectors they create using the SDK. They include:

- Support for extracting tag, policy, and external table metadata from Snowflake.
- Support for computing the file system to table lineage from external tables for Snowflake.
- Support for lineage generation from extracted metadata for connectors that support direct lineage generation, such as SAP HANA or SAP ECC. The lineage job will be triggered as a downstream job for metadata extraction with the interim lineage data written to temporary files which will be then processed by the direct lineage creation (DLC) job to feed into the lineage service.
- Enhanced the `sdk_tests` module with examples and tooling to empower OCF connector developers to write integration tests.

Migrating Selective Extraction Settings from Native to OCF

Previously, during the migration of a native data source to the corresponding OCF connector, we did not preserve the filters specified for selective extraction. Now, all the schema filters for RDBMS connectors will be migrated and displayed on the data source Settings page.

Selective Extraction Filter UI Improvements

On the Settings > **Metadata Extraction** tab of OCF RDBMS data sources, added a new **Select Schemas** button and the ability to bulk-add the schemas that match the keyword used in the search filter. Previously, the schemas matching the keyword had to be added one by one.

Surfacing Source Comments for OCF BI

The BI object descriptions sent by BI OCF connectors are now displayed in the Alation user interface as source comments.

New OCF Connectors

- Snowflake (AWS, Azure, and GCP)
- Tableau Server and Tableau Online
- PostgreSQL Enterprise
- Greenplum
- SAP Hana
- Vertica
- Azure Databricks
- Teradata
- Azure Databricks on Azure Government Cloud
- Informatica PowerCenter
- SAP ECC

Proxy Support for the OCF Looker

Added proxy support for the OCF Looker connector. Additional fields have been added to the connector settings to provide proxy details like host, port, username, and password.

Data Source Certifications

- Certified support for Starburst Trino version 380-e LTS (Custom DB) with a newer version of the driver.
- Certified support for MongoDB version 5 (Custom DB) with the CData driver for MongoDB 2021.

Deprecation of Some of Built-in (Native) Data Source Types

From this release, some of the built-in (native) data source types are no longer available in the **Add a Data Source** list in the Alation user interface by default. We recommend using the corresponding OCF connectors instead. The deprecated built-in data source types are:

- MySQL
- PostgreSQL
- SQL Server
- Azure SQL Data Warehouse

This change does not affect your existing data sources of these types. They remain fully functional.

Some other data sources that could be added using the Custom DB data source type are also deprecated. We recommend using the corresponding OCF connectors instead. For the full list, see the dedicated blog on Alation Community: [OCF Migration and End of Life for Native Connectors](#) (requires Community login).

Alation Agent for the Cloud

Multiple Agent Installations on Cloud Instances

Alation now supports multiple Alation Agent installations per cloud instance. Multiple Agents may be required if data sources are in different geographical locations, network segments or security zones. A Server Admin for a cloud instance can now install Alation Agent multiple times to enable connection to the data sources from the catalog.

Agent Enhancements

- Users can now delete an Agent even when the certificate is revoked.
- The default port of **80** for web_proxy is now honored by the Agent. Previously port **80** was incorrectly ignored and defaulted to **3128**.

Enterprise Readiness

User Interface for Authentication Configuration

Server Admins can now configure data source authentication for metadata extraction, query log ingestion, sampling, and profiling in the Alation user interface. Previously, this configuration was available on the Alation server backend only. The user interface allows for configuring authentication with AWS IAM, OAuth, and Azure Key Vault.

Case-insensitive Usernames

Catalog now supports login with usernames that are case-insensitive. For example, such usernames as `user1@xyz.com` and `User1@XYZ.com`, differing only in the use of lowercase and uppercase letters are considered identical usernames. Logins using these usernames will be successful and tied to the case-insensitive account `user1@xyz.com`.

Postgres Upgrade

The internal server database (Postgres) was updated to version **13.6**.

Elasticsearch Upgrade

Upgraded the Elasticsearch component to a newer minor version for better code security.

Improvements

- Users with the Viewer role are now able to view shared query results if they have been allowed access to the underlying data source. Previously, the minimal role for viewing shared results was Composer.
- Google BigQuery struct objects are now captured in the catalog after one MDE. Previously, users had to run MDE twice to capture the struct data type. After the first MDE, struct data was not loaded and Alation issued the error *DocSchema could not be loaded* in the **Job History** table. The second MDE run loaded the struct data successfully. With this improvement, there is no need to run MDE several times in order to get struct columns into the catalog.
- The JDK on the Alation server was upgraded to a newer version of 11.x.x to take advantage of enhancements and security fixes in this version.
- The usage of time zones on the Alation server has been unified. Alation now uses UTC by default, and the **alation_conf** parameter `alation.install.timezone` defaults to UTC. Previously, the **alation-backup.log** file captured logs in PDT. As part of this improvement, UTC is now used to capture the backup logs.
- Improved sorting by Title in the left-side navigation panel. Relational data sources, non-relational data sources, and API resources are now ordered by special characters first, then numerically, and finally alphabetically. Sorting is case-insensitive.
- Previously, when downloading a data dictionary at the table level, the column order was not preserved in the resulting CSV file. Now, the downloaded data dictionary is sorted using the name field for Schema, Table, and Attribute (Column) object types.
- Previously, Stewards could select a maximum of 25 objects per page in the Stewardship Workbench for performing a bulk action. If the search results were on multiple pages, Stewards had to navigate to each page to select the next 25 objects. Now, they have the ability to select the number of rows to be displayed per page: 25, 50, 75, or 100, thus increasing the number of objects they can bulk-select in one click.
- If a bulk action in the Stewardship Workbench fails, Alation will now send an email notification to the admin performing the change that the task could not be completed successfully.
- The configuration-based Hive connector JDK was upgraded to a newer version to ensure application security.
- During the installation of Alation Analytics, the value of `umask` is now auto-set to the required default value. There is no requirement to check or set the value manually in the Alation shell at the time of installation.

Bug Fixes

- Fixed an issue where Elasticsearch did not start if the user `elasticsearch` existed on the system.
- Fixed an issue with searching for terms in the **Terms** table on a glossary page. Users can now successfully search for terms using the table search feature.
- Fixed a bug where the desktop Compose application did not load. Now, users can successfully open the application.
- Fixed an issue with the Advanced Query Scheduling dialog where the query search and the advanced schedule showed different next run times. We now convert the cron string to the server time as the next run function uses the server time, which aligns the next run times displayed in the user interface.
- Fixed an issue where invalid cron strings were allowed in query schedules. We added an internal endpoint that validates cron strings for queries that use advanced scheduling.

- Fixed an issue where the unexpected error was displayed when a user attempted to view a saved advanced schedule in Simple Scheduler. Now, users can successfully navigate from Advanced Scheduler to Simple Scheduler in the user interface.
- Fixed the user interface styles of the **Add/Remove Term** and **Add/Remove Policy** dialogs that allow users to add glossary terms to glossaries and policies to policy groups respectively.
- Previously, if there were more glossaries than displayed on one page and users used pagination to see the next batch of glossaries, then the list of glossaries was not displayed. Instead, it displayed the message *No glossaries to display*. This issue was fixed and now users can successfully use pagination on the **Glossaries** page.
- Previously, when @-mentions were used in the **Description** field of a glossary term object, the **Terms** table on the **Glossaries** page did not show the mentions in the **Description** column. This issue was fixed. Now @-mentions are correctly displayed in the **Description** column.
- Updated the prompt on the **Add** dropdown menu in the **Article Groups > Suggested Terms** table, changing *Glossaries* to *Article Groups* when the new Glossaries feature is enabled.
- Previously, on the glossary term object custom template, the **Save** button was always enabled, although users could not save the template until they added a Title. Now, the **Save** button only becomes enabled after users populate the Title field.
- Previously, when adding or removing glossary terms to or from an untitled glossary, the action status popover message did not display any glossary name. Now, this message will now show *Untitled Glossary [id]*.
- Previously, the glossary pages that a user visited were not shown in the **Recently Visited** section on the Alation homepage. This issue was fixed, and now users will see the glossary pages under **Recently Visited**.
- Previously, a Server or a Catalog Admin could accidentally delete a term template even when it was in use. Now, a template can only be deleted if it is not in use by any term type object.
- When accessing a deleted term using the URL with the object ID, admin users see the **Restore** action. Previously, this action was visible to Viewer users too. After this fix, the **Restore** action is available to admin users only, while Viewers only see the *object deleted* message.
- Fixed a user interface issue where on the glossary and glossary term object pages, the column organization list displayed a blank column. Now, the column organization list correctly displays the columns of the table.
- Fixed an issue where the **Documentation** search filter in the Object Types filter was not displayed when the New Glossaries feature flag was enabled. Now, the article, article group, glossary, and term object types are grouped under the Documentation filter.
- The **Delete** button on the new term type template now appears as disabled before the template is saved.
- Fixed an issue where the left-side filters were not displayed for glossary or term objects in full-page search. Now, for the glossary object type, the filter pane displays tags. For the term object type, it displays custom fields and tags.
- Previously, the **Date Range** filter was incorrectly applied to the Top Contributors report: a filtered report returned less data for the selected date range than the data available in the database. This fix ensures that the **Date Range** filter covers the correct date range from start to end of the selected month.
- Previously, custom fields with more than 25 characters were not loaded to the `public.customfield` table of Alation Analytics. We extended the length of some fields in the `public.customfield` table to the maximum length so that it can support larger values.

- Previously, custom template titles were not updated in the `public.article` table of Alation Analytics. This issue was due to the fact that the extraction query for articles didn't take into consideration the updates to custom templates. To fix this, we updated the extract query for articles to also identify the changes to the custom templates. As a result, the updates in custom templates will be reflected in Alation Analytics.
- Updated the transform job for the article data in Alation Analytics so that it updates the `ts_updated` field when a field that is not present in the `public.article` table is updated in Rosemeta. For example, the `ts_updated` field is now updated if a user favorites or un-favorites an article page, follows or unfollows an article, or edits the Description field so that the `ts_updated` in Alation Analytics and the corresponding value in Rosemeta will match. If you experienced this issue, contact Alation Support to reset the ETL checkpoint for article data on your Alation Analytics database.
- Added a fix to remove the Docker volumes that aren't associated with any containers on each update and upon uninstallation. Customers having older releases can do this manually following these steps:
 1. List all unassociated volumes.

```
sudo docker volume ls -f dangling=true -q
```

2. Remove these volumes.

```
sudo docker volume rm $(sudo docker volume ls -f dangling=true -q)
```

- Fixed an issue with Alation Sandbox where the pre-check test reported a failure based on available disk space when the available disk space was measured in TB. We modified the pre-check command to standardize the output disk space in GB.
- Fixed an issue where the **Health Checks** tab did not appear on the **Admin Settings > Monitor** page.
- Fixed an issue where the catalog object definitions kept changing for every metadata extraction run using the API. Now, the multipart schema value for Custom DB will be consistent with the catalog object format.
- Previously, the task queue could become saturated due to an error in processing certain objects in that queue. When the object at fault was processed, it appeared to stop the queue from processing entirely. For example, this route could be triggered when users created new articles, causing the new articles to not be immediately indexed and searchable. This fix allows the task queue to continue progressing even when this unexpected error occurs.
- Fixed an issue where deleted BI objects still appeared in search results. This was resolved by fixing the incremental and search reindexing jobs for BI objects.
- Fixed an issue with setting a password for the lineage DB. Previously, the parsing of passwords with special characters could cause errors. Now, an admin can successfully set a password that contains special characters.
- Fixed a bug where the search autocomplete functionality caused search issues after users used search criteria with tab delimiters. Resolved by adding a filter for the autocomplete entries that contain tabs.
- Fixed an issue with articles enrolled in Agile Approval. Previously, when a reviewer approved a version, this sometimes removed the existing values from custom fields. Users were also unable to restore the previous version of the article as the field history was not preserved. Now, when reviewers approve a version of an article, the field values stay intact, and version history is updated correctly.
- Fixed a bug where the filters on the left on the Search page were not populated for some object types selected in the **Object Types** filter (All, Data, Article, BI).
- While uploading a data dictionary from an ER studio file, some description values present in the file were absent from the preview and in the desc field after the upload. The issue occurred when the

description had a duplicate empty entry in the SmallString section of the ER file. To resolve, we now prioritize non-empty values over empty values.

- Fixed an issue with the Oracle OCF connector where the **Last Altered** field showed an incorrect value. Although an object was not modified, the **Last Altered** field had the same value as the **Created** timestamp. After the fix, the **Last Altered** field will only be displayed after an object was modified.
- Fixed an issue where the Alation Analytics Leaderboard report on the homepage did not load for Viewer and Steward users. Now, the Leaderboard on the homepage is visible to all users.
- Fixed an issue where the scheduled query executions did not close the database connection after they were completed. Now, the connection opened by a scheduled query will be closed after executing the query.
- Previously, the column profile histogram for a column only added new counts but never removed values that were removed from the database. After the fix, the report should correctly reflect the removal of a value.
- Fixed an issue where the lineage data resulting from the CREATE OR REPLACE TABLE USING statement could not be calculated for a Databricks on Azure data source. Now, such queries are supported by the Alation column-level lineage parser add-on for Databricks.
- Fixed an issue with Elasticsearch log rotation where the logs were not cleaned up based on rotation settings in **alation_conf**. Added a weekly task to clean up Elasticsearch logs.
- Fixed an issue where the ampersand symbol was rendered differently in the downloaded dictionary table Title and Description. The issue was fixed by updating the parser logic in the DataDictionaryExporter so that the ampersand symbol is parsed correctly.
- Fixed an issue where the query results table filter that opens in a popover on the query form page was still displayed after the user clicked the back button in the browser. Now, using the back button returns the user to the previous page without persisting popovers.
- Fixed an issue where only Server Admins were able to remove their cached database connection in **Account Settings > DB Connections**. Now all users are allowed to remove their cached database connections.
- Fixed an issue where Lexicon abbreviations confirmed by users disappeared from the Alation user interface after the Lexicon job was run.
- Fixed an issue where only the title was displayed for a Steward added using the **Add Stewards** bulk action in Stewardship Workbench, without displaying the Steward's name. Now, both the name and the title are displayed.
- Fixed an issue where the links to **Alerts** and **Branding** in the **Customize Catalog** sections in the catalog settings dropdown list did not work for users with the Catalog Admin role. Now, Catalog Admins can successfully open the corresponding pages of the settings.
- Fixed a bug with displaying the domain names that included the ampersand symbol. This symbol was displayed as & when the domain was @-mentioned in glossaries and articles. Now, ampersands in domain names are displayed correctly.
- For the Teradata data source, fixed an issue where single-line comments were removed from the SQL body when creating views in Compose. Single-line comments are no longer removed when users create views or execute queries in Compose.

Known Issues

- Back references to glossary term objects that result from the People Set custom fields are currently not displayed on the pages of groups. If the custom template of a term object contains a People Set custom field and a group is added as a value, the back reference to this term will not appear as a back reference on the catalog page of the corresponding group.
- Previously extracted column profiles are not refreshed when sampling is re-run on the data source level. This results in old data being displayed alongside the new data in column profiles after sampling the data source. Performing the column profiling again returns the correct profiling data.
- When a user selects file and/or directory objects in Stewardship Workbench and applies a bulk action (Add/Remove Stewards, Bulk Update People Sets, or Bulk Update Custom Fields), processing of the bulk action will fail. The issue happens when a bulk action is applied to a selection that only includes files and/or directories and when the selection includes at least one file and/or directory object among other object types. **(FIXED IN 2022.3.3)**
- Currently, only Server Admins and Catalog Admins can create new business policies using the **Create Policy** button in Policy Center. When a Source Admin, Steward, or Composer attempts this action, they get an error saying “Failed to Create New Business policy.” Users with these roles should contact a Server Admin or Catalog Admin to create policies.
- The job schedule for metadata extraction (MDE), query log ingestion (QLI), and profiling will be shifted eight hours earlier. See [Fix Job Schedules after 2022.3 Upgrade](#) for details on how to fix the issue.

1.7 Release Notes 2022.2 (General Availability)

1.7.1 RELEASE 2022.2.7

build 12.3.7.27641

- Previously, due to certain race conditions on the server, scheduled query executions sometimes failed with the following error: *Could not connect to the database. Please reconnect or reload.* The issue has now been fixed.
- Previously, the **Download Latest Usage Data** button on the **Admin Settings > Reporting > Upload Manually** page did not initiate the data download. This issue was fixed, and all usage data download steps will work as intended.

1.7.2 RELEASE 2022.2.6

build 12.3.6.16423

- Fixed an issue in the user interface of the Compose web application where the elements on the top toolbar became misaligned if users reduced the size of the browser window. This fix adjusts the styles of the Compose toolbar to ensure a better responsiveness of the Compose page when the browser window is resized.
- Fixed an issue where Alation previously did not allow console calls that would display certain internal errors, which made error reporting more difficult. These calls are now captured in the developer console and may be used to report issues. Additionally, fixed an issue with the error page where in certain cases, it did not differentiate between server-side and client-side errors. The error page now uses the *unknown error* code to display the *unknown error* message instead of the *server error* message to differentiate between client-side and server-side issues.

- In version 2022.2.5 or earlier versions of Alation 2022.2, the upgrade of OCF RDBMS connectors to the latest available version could result in the following error:

Installation failed, changes not committed. Wrong data schema in manifest file: {"features": ["Field features can take one of the following values: [BI_GENERAL, BI_MDE_GBM_V2, BI_CERT_GBM_V2, [...]}. Invalid value : RDBMS_QUERY_SERVICE."}]}

This issue happens due to Alation not recognizing RDBMS_QUERY_SERVICE as a new feature added to connectors. The issue is fixed in this patch release. We recommend updating to 2022.2.6 to be able to upgrade to and use the latest versions of OCF RDBMS connectors.

1.7.3 RELEASE 2022.2.5

build 12.3.5.13731

- Fixed an issue where users were unable to select and copy multiple rows from the expanded view of a query form results table. Now, users can use their mouse or touchpad to select a part of the result set and then copy and paste it to a document.
- Fixed an issue where Lexicon abbreviations confirmed by users disappeared from the Alation user interface after the Lexicon job was run.
- Fixed an issue where the connection opened for executing a scheduled query was closed too soon in a race condition preventing some code from completing. Now, once the execution of a scheduled query is complete, the database connection will be closed correctly and only after completing all necessary code.
- Fixed an issue where the backup process changed the permissions on the Postgres folders on the Alation server to world writable. With this fix, the Postgres folders permissions are secured by default. Also, the Postgres folders permissions will be auto-corrected in case an admin, a privileged user, or a process changes them to a state that is different from the expected state.
- Addressed an issue with high resource contention during the backup run. Previously, users reported performance degradation in Alation at specific hours of the day. Further investigation showed that the slowness coincided with the time of the `dump_all_data` stage of the Alation backup process and was due to resource contention on the host. This fix introduces an **alation_conf** parameter `alation.backup_v2.num_of_thread` that allows tuning the CPU usage down if needed.
- Previously, users reported an issue with uploading the user account certificate when authenticating to a Google BigQuery data source in Compose. The issue was caused by the deprecation of the previously used user account authentication method by Google. The issue has been fixed. Now, both new and old Google API credentials work.

1.7.4 RELEASE 2022.2.4

build 12.3.4.10641

- Updated the transform job for the article data in Alation Analytics so that it updates the `ts_updated` field when a field that is not present in the `public.article` table is updated in Rosemeta. For example, the `ts_updated` field is now updated if a user favorites or un-favorites an article page, follows or unfollows an article, or edits the Description field so that the `ts_updated` in Alation Analytics and the corresponding value in Rosemeta will match. If you experienced this issue, contact Alation Support to reset the ETL checkpoint for article data on your Alation Analytics database.
- Previously, users reported issues with highlighting SQL errors in Compose queries containing comments. Incorrect rows were highlighted, making it difficult to find the error. Now, if a query body includes

comments, the logic behind highlighting will correctly recognize them and highlight the row containing the error.

1.7.5 RELEASE 2022.2.3

build 12.3.3.9209

Improvements

For instances with a large amount of lineage data and large lineage diagrams (more than 100,000 lineage nodes), Server Admins can now fine-tune the amount of lineage data to be retrieved and displayed on the **Lineage** tab. Previously, users reported issues with loading such lineage diagrams in the user interface.

We added two new **alation_conf** parameters:

- `lineage-service.filters.max_depth` – Sets the maximum depth of lineage diagrams. The diagram depth is the number of lineage levels upstream and downstream from the current object that are displayed in the user interface. The default value is `-1`, which means by default, Alation will attempt to load all available lineage data into the lineage diagram. You can set this parameter to a desired number, starting with `1`. This will define how many upstream and downstream lineage levels will appear on lineage diagrams in the user interface. No restart is required when changing the value of this parameter.
- `lineage-service.filters.exclude_temp` – Controls the display of temporary objects on lineage diagrams. The default value is `False`. This means that by default, the temporary objects are displayed. Set the parameter to `True` to hide temporary objects from lineage diagrams. No restart is required when changing the value of this parameter.

The parameters apply to all data sources in the catalog. These settings cannot be applied to a specific data source only.

Bug Fixes

- Addressed issues in Event Bus by removing the potentially vulnerable classes of Log4j from the corresponding Log4j.jar file included into the Alation application. Upgrading to this patch will apply the fix automatically.
- Fixed an issue that prevents articles with duplicate titles to be created for a custom API.
- Previously, the Data Health public API could not process decimal values, for example `0.25` `.5` `1.5`. This bug has been addressed and decimal values are now processed.
- Fixed an issue causing the user profile page to load slowly when having many back references, resulting in a 502 error. The profile page now loads faster, even if there are many back references.
- Fixed a polling issue in Compose causing warning alerts in the Compose user interface during SQL execution. The issue was the result of long polling when trying to fetch Hive logs. This fix handles the log resulting from the polling separately.

Known Issues in This Patch

- On the lineage diagrams that contain column-level lineage, if a user selects a dataflow object, then only table-level paths going through this object are highlighted. The column-level lineage paths are not highlighted. It is expected that both table-level and column-level paths should be highlighted when users select a dataflow object on the diagram.
- After migrating Lineage from V2 to V3, some lineage objects are displayed as temporary on the lineage diagram for new lineage created after the migration. The issue manifests itself for the objects that were extracted into the catalog but were not referenced in any lineage data before the migration to V3.

1.7.6 RELEASE 2022.2.2

build 12.3.2.6877

New Features

Support for the DELETE method in the Data Health API

The Data Health API (<your_Alation_URL>/openapi/data_quality/) now supports the deleting of data quality fields and values using the new DELETE action.

Bug Fixes

- Previously, the migration from Lineage V2 to V3 could not complete and stopped midstream when processing data with the *file* object type. The issue was fixed, and the stopped migrations can be continued after updating Alation to this patch.
- Fixed an issue where the data upload from a CSV file failed for Hive data sources. Now, data can be successfully uploaded into Hive data sources from the Alation catalog.
- Previously, the task queue could become saturated due to an error in processing certain objects in that queue. When the object at fault was processed, it appeared to stop the queue from processing entirely. For example, this route could be triggered when users created new articles, causing the new articles to not be immediately indexed and searchable. This fix allows the task queue to continue progressing even when this unexpected error occurs.

1.7.7 RELEASE 2022.2.1

build 12.3.1.4339

- Fixed a user interface issue on the Alation Search page table view where the clear filters action did not correctly return to the column headers set for the **All** object type filter.
- Fixed a user interface issue on the Alation Search page table view where the **Edit Columns** dropdown list was not updated correctly when users switched between the object type filters.
- Fixed an issue on the Alation Search page where users were sometimes unable to select filter checkboxes on the **Filters** panel on the left. After the fix, users can select filters as is expected.
- Fixed an issue where the Alation Analytics database did not correctly reflect the changes to the picker and multi-select picker field options in **Customize Catalog > Custom Fields** that were associated with existing catalog objects. The issue was caused by the `ts_updated` field value in the internal server

database not capturing the event of the change. Now, the `ts_updated` field correctly reflects the time of the change and the corresponding data is transferred to the Alation Analytics database during the ETL.

- Previously, the Title and Description fields on domain, glossary, and glossary term catalog pages could not be restored using the **Restore** action in the field history dialog. As a result of this fix, it is now possible to restore the Title and Description values from field history.
- Fixed an issue with custom field history and version history for articles enrolled in Agile Approval, where approving an article sometimes cleared the field history. Now, the `bulk_metadata` endpoint of the internal API will correctly update the revision history of custom fields. When a user approves an article revision, the new revision will be in use instead of a blank revision, which was the cause of the issue.
- Previously, when authenticating to Snowflake from Compose using single sign-on (SSO), users sometimes got an Authorization Terminated Unexpectedly error message after the authentication tab closed. Despite the error message, the authorization was successful and the users could use Compose. This fix simplifies the authorization success and failure handling for Compose SSO to avoid such “false negative” scenarios.

1.7.8 RELEASE 2022.2 — General Availability

build 12.3.0.3415

Data Governance App

Ability to Add or Remove Policies to/from Groups on the Policy Center Main Page

Server and Catalog Admins now have the ability to add policies to groups and remove policies from groups directly on the main page of Policy Center. Actions that allow modifying group membership for policies are available in the **Policy Groups** and **Policies** tables.

Add or Remove Stewards Bulk Actions

Server Admins, Catalog Admins, and Stewards now have the ability to add or remove Stewards to or from multiple Catalog objects using Stewardship Workbench.

Update Custom Fields Bulk Action

Server Admins, Catalog Admins, and Stewards now have the ability to update custom fields for multiple Catalog objects using Stewardship Workbench. This applies to custom fields of the single picker and multipicker types.

Change Management Workflow for Policies

Server and Catalog Admins now have the ability to configure workflows to manage changes on the Catalog pages of policy and data policy objects. Catalog users can suggest a change to the Catalog fields for policies or data policies that are enrolled into a Change Management workflow.

Catalog

New Glossaries and Glossary Terms

The new glossaries are an alternative to business glossaries in the Alation Catalog. *Glossaries* and *glossary terms* are new Catalog object types with their own Catalog templates. They can be assigned to domains and discovered using Alation Search. Admins can create and configure new glossaries. Users can view, create, update, and delete glossary terms.

The new glossaries co-exist with the business glossary functionality.

Server Admins can enable new glossaries in **Admin Settings > Feature Configuration**. First, use the feature flag **Glossary Terms Setup**. This reveals the new glossary configuration page. After configuring glossaries, use the second feature flag **Glossary Terms** to make the new Glossaries available for all users and discoverable using Alation Search. When the new glossaries are enabled, the existing business glossaries are renamed to **article groups**.

Health Tab for Table Objects

Alation now has the ability to display data quality information pushed using the new **Data Health API**. When such data becomes available for a table object, the Catalog page of the table will display an additional tab **Health**, giving users an at-a-glance view into data quality rules, statuses, and descriptions, sourced from an external data quality tool.

Back References UI Enhancement

Added a new user interface component for back references on the pages of articles and RDBMS objects. This change leverages the Object Set custom field to create back references. Each back reference is presented as a separate field on the Catalog page of an object.

Homepage Enhancements

Admin Settings > Catalog Admin > Customize Homepage now allows for the following configurations:

- When customizing the homepage, admins can choose one of the three sizes for the Search area: small, medium, or large (default).
- Admins can now configure the homepage links to open in a new tab using the `openInNewTab: true` property in the homepage JSON file. By default, the links open in the same tab.
- Users can now dismiss alerts displayed on the homepage by clicking the **X** icon on the alert banner.
- Admins can set a custom accent color for the headline text and the top banner icons.

UI Refresh for Catalog Admin Settings

The Catalog Admin section of the Admin Settings now uses tabs, such as **Branding**, **Homepage Customization**, **Custom Fields**, and **Custom Templates** for easier navigation.

Search

Elasticsearch Upgrade to 7.17

Elasticsearch was upgraded to version 7.17, which resolves the security issues found in Log4j 2 in December 2021.

Alation Analytics

QLI Data Available in the Alation Analytics Schema

Alation Analytics now includes query ingestion data. Users can analyze such metrics as number of queries executed per day per data source, query computation time by data source, and total number of queries executed on a data source. The new tables that store QLI data are `aggr_qli_metrics` and `aggr_qli_mentions`.

New created_at Column for BI Objects

Added the `ts_created` timestamp column for the following BI objects: BI Folder, BI Report, BI Report Column, BI Datasources, BI Datasource Columns, BI Connection, BI Connection Column, BI Server.

Compose

Ability to Disable Export and Download per Data Source

A Data Source Admin can now toggle the setting **Allow Export and Download** to disable the export and download actions for an individual data source. This settings can be found:

- On the Data Source Settings > **General Settings** tab for native data sources and Custom DB
- On the Data Source Settings > **Compose** tab of OCF data sources.

Disabling the data source export will have the following effect on the user interface:

- Disable the **Export** dropdown on the Compose toolbar.
- Hide the option to download Live Excel Reports.
- Disable the **Download** menu in the Compose results pane.
- Hide the **Export** button when results are viewed on query form pages, Catalog page for query results, or on the Samples tab of a data object.
- Hide the **Download** link for query results.

Master Link to a Scheduled Query

At the time of configuring the sharing settings of a scheduled query, users will now get a master link pointing to the latest successful result of this scheduled query.

Ability to Share Query Results to Groups

Users can now share scheduled query results with a group of recipients instead of entering them one by one.

Edit, Delete, and Select Active Connections for Scheduled Queries

Users can set a schedule for a query in Compose using a new revamped Scheduling dialog. Users can set a new connection and credentials for the schedule or select a saved connection. When a saved connection is altered, the schedule automatically runs using the updated connection. The data from query schedules for existing queries is retained in the new Scheduling dialog after the update.

Edit, Delete, and Select Active Connections for Dynamic Profiling

Using the new Connection Settings dialog to run dynamic profiling on columns, users can add, edit, and delete connections and credentials. The new dialog always displays the active connection as selected when opened.

Enterprise Deployment

Catalog SSO with OpenID Connect (OIDC)

Server Admins can now configure Catalog authentication using OAuth and OIDC in the Alation UI. Users will be authenticated with Auth0 based on OAuth and OIDC protocols.

pgBackRest as the Backup Tool (Beta)

Server Admins can enable and use the pgBackRest tool for Alation backups.

- The full backup taken with pgBackRest has a smaller disk footprint.
- pgBackRest offers the ability to resume the backup process. If an admin encounters an issue, the backup process can be resumed after troubleshooting.
- pgBackRest provides more compression options, such as lz4, zst, gz, bz2.
- The ability to encrypt backups.

Reducing the Space Occupied by KVStore

Server Admins now have the ability to reduce the space occupied by KVStore by deleting the old parsed metadata.

Connect Platform

Alation Agent

build 1.2.0.815

Server Admins of the cloud instances can now install Alation Agent on their network. When IP whitelisting is not an option for connecting the Alation cloud instance to data sources, installing the Alation Agent on the customer's network solves the task of securely connecting to data sources from Alation. All connections will be initiated from the customer's VPC and target an Alation endpoint.

New OCF Connectors

- Power BI Scanner OCF connector

The Power BI Scanner OCF connector supports Power BI on Azure and catalogs workflows, apps, reports, dashboards, tiles, datasets, and dataset fields. The user experience is similar to the Power BI console: users can navigate to the Catalog objects, browsing the list of apps, workspaces, and datasets. Full lineage is supported for datasets, reports, tiles, and dashboards. Lineage between data sources and Power BI datasets has some limitations.

- AWS Glue OCF connector
- Azure Synapse Analytics OCF connector
- DB2 OCF connector
- SAP IQ (Sybase IQ) OCF connector
- SAS Base OCF connector
- Amazon S3 OCF connector (beta)
- Informatica PowerCenter OCF connector (beta)

Built-In Driver Upgrades

- The built-in driver for PostgreSQL was upgraded to version **42.3.3**.
- The built-in driver for Databricks was upgraded to version **2.6.21**.
- The driver version `denodo-vdp-jdbcdriver-8.0-update-20220126` was certified for Denodo data sources (Custom DB).

Column Level Lineage Parser Add-Ons

- Column Level Lineage for SQL Server sources (general availability)
- Column Level Lineage for PostgreSQL sources (beta)
- Column Level Lineage for DB2 sources (beta)
- Column Level Lineage for Oracle sources (beta)
- Column Level Lineage for MySQL sources (beta)

Dataflow Objects for BI Lineage

Added dataflow objects to existing table level lineage for BI sources.

Obfuscate Literals for OCF Data Sources

Added an option to obfuscate literals for OCF data sources. Enabling it will mask literals in queries and query statements.

Public APIs

- Enabled access for the Source Admin role to the Lineage API. Previously, this API was only accessible to Server Admins.
- Added a new Data Health API (<your_Alation_URL>/openapi/data_quality/) that powers the [Health Tab for Table Objects](#).

Improvements

- The Alation server now runs multiple threads for Hive query log extraction, reducing the total QLI time.
- Improved the performance of *write* requests to the LMS Public Field Value API. Bulk writes to the LMS Public Field Value API are now faster at scale. This is the first of many steps we plan to take to improve bulk API performance.
- Improved performance of the application start times. Previously, some customers reported a number of slow pages and poor user experience because of this slowness. As a result of this improvement, the homepage loading times should significantly decrease. We also optimized the amount of JSON encoding and serialization during the application start when Alation is loaded from a bookmark or after reloading a page. In some cases this should noticeably reduce page loading times. In addition, the page loading indicator is now displayed consistently and users can always see that their actions, such as clicking on links, have the appropriate effect.
- The legacy homepage has been removed from the application.
- Improved performance of the field level permission checks, which contributes to the decrease of the loading time of Catalog pages.
- The ability to select multiple domains in the **Domains** search filter is now enabled by default.
- For Alation Analytics, upgraded the Postgres, RabbitMQ and base-transform Docker images to newer versions to address security vulnerabilities.

- The OCF SDK now ensures that `tableType` is included in the `getTableMetadata()` method for incremental MDE.
- The **Archive processed QLI data** action now respects the archiving settings that an admin provides in the user interface on the Query Log Ingestion tab of a data source settings. Previously, the scheduled auto-archiving process always deleted QLI data older than six months.
- For MySQL, SQL Server, and PostgreSQL data sources, statements with CTE are now parsed to generate lineage and filter data.
- For SQL Server, PostgreSQL, and Oracle data sources, automatic column level lineage now correctly recognizes column names that contain spaces.
- For Compose desktop client, updated version of Electron to **3.1.13**.

Bug Fixes

- Previously, the Growth chart on the Governance Dashboard did not show the month if there was no data for this month. Now, the Growth chart displays labels for all months, including the months without data.
- Fixed an issue where the file type domain members were not extracted into Alation Analytics V2.
- Fixed an issue where the ETL failed to insert catalog sets into Alation Analytics V2 when the description was null. To resolve this, we have set the catalog set description to an empty string instead of using nulls.
- Previously, in the `public.user_group_membership` table of Alation Analytics V2, after updating a user with a new group, the previous group that was mapped to the user was not removed. As a result, analyst users could not tell which groups the user was active in. Now, inactive groups, or groups that were removed from the user, are also un-mapped from this user in Alation Analytics V2.
- For the Alation Analytics V2 ETL, increased the RabbitMQ timeout from 60 seconds to 3 hours so that the load from `stg` to `dim` tables does not get disrupted. As part of this change, identity will be restarted for the `total_replace` tables so that not to exceed the int upper limit of 2 billion rows.
- Custom visualizations on the Alation Analytics dashboard are now displayed to all admin roles (Catalog, Source, and Server Admins).
- In Alation Analytics V2, the ETL Manager is updated to generate the load ID sequence from the ETL checkpoint table instead of the Redis cache.
- Fixed an issue where the `aamanager` container did not start automatically after the update of Alation Analytics or after a system reboot. The `aamanager` service now starts automatically.
- Added a new prompt as part of the Alation Analytics database upgrade from **9.6** to **13**. The admin performing the upgrade can now provide a custom location to store the PostgreSQL dump. This helps avoid the disk space issue of the default `/tmp` folder.
- Fixed a bug where the **Customize Catalog** page sometimes displayed the 500 error when users were creating custom fields.
- Fixed an issue where users could accidentally add rich text fields to the right column of the custom template layout.
- Fixed a bug with the Metadata Change Log, where the changes were not reflected when schemas or tables were unchanged, but their child objects did change. After this fix, the correct number of changes is displayed.
- Fixed an issue with Compose parsing where records were passed as a set instead of a list. As a result, published queries were not displayed under the **Queries** tab on the Catalog pages of table objects,

which was accompanied by the error 'set' object is not subscriptable in the **celery-parsing_error.log** file. This fix ensures the correct Compose query parsing.

- Fixed an issue where users with the non-admin roles of Steward and Composer were unable to authenticate against a Google BigQuery data source in Compose when the Viewer role enforcement was enabled. Now, Stewards and Composers can connect to Google BigQuery data sources in Compose using their database accounts.
- Previously, when metadata was uploaded in bulk using the public API, values only appeared in filters in Advanced Search after users accessed the page of the corresponding Catalog object. After this fix, the metadata uploaded in bulk will be successfully indexed and available in Search filters.
- Previously, the changes to existing functions were not reflected in Alation after MDE. This issue has been fixed, and changes to existing functions are now reflected correctly in the Catalog.
- Fixed an issue where the server health check process sent multiple emails daily for jobs that were running on schedule and without issues. Now, alerts are only sent when the corresponding alert conditions are met.
- Resolved an issue where scheduled queries failed because they were interrupted by the service outage caused by excessive memory consumption by the lineage adjustment process. Added a fix that allows disabling lineage invalidation sync from Compose for the DROP queries. This is controlled by the **alation_conf** parameter `alation.lineage.sync_invalidation_from_compose` (True by default).
- Previously, when a null value was present in the **Steward** field in `logical_metadata_genericfieldvalue` table, the `sync_all_objects_curation_progress` async job was failing. This issue was fixed, and the `sync_all_objects_curation_progress` should now complete successfully.
- Fixed an issue where Compose users could not add new credentials for Google BigQuery sources in the Transient mode. After the fix, credentials can be successfully added to the **Connection Settings** dialog when the Transient mode is enabled.
- After loading the metadata into a virtual data source, the column count did not match the expected value. This fix adds a query to update the number of columns for all tables after loading metadata into virtual data sources.
- Previously, the Query API for retrieving the SQL text of a query did not return any data when the query was unpublished and the last saved date was null. The solution was to determine the current state of the query and return the appropriate content.
- Fixed an issue where the upload of data dictionaries to virtual data sources failed when there was a non-English character in the data objects names.
- Fixed an issue where sampling did not work for data sources when Compose SSO was enabled. Now, it works without issues when Compose SSO is configured for a data source.
- Fixed an issue where the Alation Analytics V2 ETL was failing to insert a catalog set in the database when the description was null. To resolve this, we have set the catalog set description to an empty string when it is null.
- Fixed an issue where functions were not reflected in Alation after MDE. Now, functions extracted during MDE or created in Compose are successfully displayed in the Catalog.
- Fixed an issue with uploading a data dictionary to virtual data sources when there were non-English characters in a schema.
- Fixed an issue with the Connection Settings dialog, where adding a new Google BigQuery credential resulted in the 500 error when the Transient mode was enabled. This affected features like Compose, Query Forms, and Dynamic Profiling.
- Fixed an issue where after upgrading from a previous version, the Alation Analytics app icon disappeared from the **Apps** menu in the Alation user interface.

Known Issues

Celery Job Scheduling

For instances on Debian-based operating systems, some scheduled Celery jobs may fail with the following error the **celery-beat_error.log**: `celery.platforms.LockFailed: [Errno 13] Permission denied: '/opt/alation/django/celerybeat.pid'`. If you notice that this is the case on your instance, perform these steps to resolve the issue:

1. On the Alation host, enter the Alation shell.

```
sudo /etc/init.d/alation shell
```

2. Change user to alation.

```
sudo su alation
```

3. Run the commands given below.

```
sudo chown -R alation:alation /opt/alation/django  
sudo chown -R alation:alation /opt/alation/ops
```

4. Exit the shell.

```
exit
```

New Glossaries

- Glossaries are not shown under the **Recently Visited** section on the Alation homepage.
- Links to term objects are not displayed on Catalog pages of groups as back references.
- Links to term objects are not displayed as back reference on Catalog pages of data source objects.
- On the Catalog page for a term or glossary that doesn't have a title, the title will not be displayed in the breadcrumbs.

Workflows

- A workflow with the Working state doesn't change the state after re-publishing.

Stewardship Workbench

- After a user assigns a Steward to objects using the **Add Stewards** bulk action and then performs a search for the objects with this Steward assignment, the search results only return a subset of the objects to which the Add Steward bulk action was applied.

Policy Center

- Duplicate Snowflake tags show up after deleting and re-adding a Snowflake data source.

Alation Agent

Admin users are currently unable to delete the Agent when it is in the *Disconnected* state.

Lineage V3

Lineage V3 for SQL Server: CREATE VIEW queries from Compose create an additional TMP node on the corresponding lineage graph.

Compose

- For Azure Synapse Analytics data sources, sampling does not work after enabling SSO for Compose.
- In the new Scheduled Query Settings dialog, when the connection URI of a saved connection is edited, the previously attached database credentials belonging to the old URI are still used to run the schedule.
- Sharing unpublished queries fails with the *permission denied* error when the **alation_conf** flag `alation.catalog.unpublished_query_visibility_level` is set to `not_visible`.
- The Saved field in Compose displays date instead of time.

Alation Analytics

- The Leadership dashboard on the homepage does not load for users with the Viewer and Steward roles. Instead, users see a page loading icon.

Authentication

- The **Use custom groups to assign roles** feature does not work when the OpenID Connect (OIDC) authentication is enabled.
- In case of the OpenID Connect (OIDC) authentication, the profile picture from the identity provider is not updated in Alation after a user is created.

1.8 Release Notes 2022.1 (General Availability)

1.8.1 RELEASE 2022.1.11

11.2.13.14725

- Fixed an issue where Lexicon abbreviations confirmed by users disappeared from the Alation user interface after the Lexicon job was run.

- Addressed an issue with high resource contention during the backup run. Previously, users observed slowness in Alation at specific hours of the day. Further investigation showed that the slowness coincided with the time of the `dump_all_data` stage of the Alation backup process and was due to resource contention on the host. This fix introduces an **alation_conf** parameter `alation.backup_v2.num_of_thread` that allows tuning the CPU usage down if needed.

1.8.2 RELEASE 2022.1.10

11.2.12.162488

- Updated the transform job for the article data in Alation Analytics so that it updates the `ts_updated` field when a field that is not present in the `public.article` table is updated in Rosemeta. For example, the `ts_updated` field is now updated if a user favorites or un-favorites an article page, follows or unfollows an article, or edits the Description field so that the `ts_updated` in Alation Analytics and the corresponding value in Rosemeta will match. If you experienced this issue, contact Alation Support to reset the ETL checkpoint for article data on your Alation Analytics database.
- Fixed an issue where the backup process changed the permissions of the some directories on the Alation server to world writable. Now, the permissions stay in the required format and the backup process does not affect them.

1.8.3 RELEASE 2022.1.9

11.2.11.162416

- Fixed an issue where the data upload from a CSV file failed for Hive data sources. Now, data can be successfully uploaded into Hive data sources from the Alation catalog.
- Previously, users reported issues with highlighting SQL errors in Compose queries containing comments. Incorrect rows were highlighted, making it difficult to find the error. Now, if a query body includes comments, the logic behind highlighting will correctly recognize them and highlight the row containing the error.
- Fixed an issue in the user interface of the full-page Search where the **More Types** filter was not reset after users left the Search page and then clicked on the Search icon in the Search bar to go back to the Search page. Both **All** and **More Types** filters were highlighted. Now, only the **All** filter is highlighted.

1.8.4 RELEASE 2022.1.8

build 11.2.10.162352

- Fixed an issue that prevents articles with duplicate titles to be created for a custom API.
- Fixed an issue where the Alation Analytics database did not correctly reflect the changes to the picker and multi-select picker field options in **Customize Catalog > Custom Fields** that were associated with existing catalog objects. The issue was caused by the `ts_updated` field value in the internal server database not capturing the event of the change. Now, the `ts_updated` field correctly reflects the time of the change and the corresponding data is transferred to the Alation Analytics database during the ETL.
- Fixed an issue causing the user profile page to load slowly when having many back references, resulting in a 502 error. The profile page now loads faster, even if there are many back references.
- Fixed an issue causing a user without access to a private data source, set as watcher, even if the user no longer has access to the data source, to keep receiving notification emails for changes done at the object level. To fix this, we removed watchers in the potential list of recipients that do not have permissions to view the data source that is the subject of the notification.

- Fixed a polling issue in Compose causing warning alerts in the Compose user interface during SQL execution, resulting from long polling when trying to fetch Hive logs. The fix handles the log resulting from the polling separately.

1.8.5 RELEASE 2022.1.7

build 11.2.9.162228

- Addressed issues in Event Bus by removing the potentially vulnerable classes of Log4j from the corresponding Log4j.jar file included into the Alation application.

1.8.6 RELEASE 2022.1.6

build 11.2.8.162150

- Fixed an issue where users with the non-admin roles of Steward and Composer were unable to authenticate against a Google BigQuery data source in Compose when the Viewer role enforcement was enabled. Now, Stewards and Composers can connect to Google BigQuery data sources in Compose using their database accounts.
- Resolved an issue where scheduled queries failed because they were interrupted by the service outage caused by excessive memory consumption by the lineage adjustment process. Added a fix that allows disabling lineage invalidation sync from Compose for the DROP queries. This is controlled by the **alation_conf** parameter `alation.lineage.sync_invalidation_from_compose` (True by default).
- Fixed an issue where after upgrading from a previous version, the Alation Analytics app icon disappeared from the **Apps** menu in the Alation user interface.

Alation Cloud

- Fixed an issue where the Alation Analytics V2 data source did not get updated in EKS upon restoring. Now, the `aav2_downstream_jobs` updates the Alation Analytics V2 password, URI, and username correctly.

1.8.7 RELEASE 2022.1.5

build 11.2.7.162035

- Fixed an issue where the left-side filters on the Search page were not populated when users selected specific object types in the **Object Types** filter (**Articles, Data, BI, All Objects**). The issue manifested itself if the Catalog instance had permissions on custom fields set up using people sets (for example, *Stewards*). After the fix, all filter facets on the Search results page load successfully.
- OCF connector logs are now limited in size and automatically rotated.
- Previously, we introduced log rotation for OCF connectors to prevent disk space issues. However, the change only applied to newly installed connectors and did not apply to existing connectors. Now, the existing OCF connectors will have their logs rotated too.
- Fixed an issue with Compose parsing where records were passed as a set instead of a list. As a result, published queries were not displayed under the **Queries** tab on the Catalog pages of table objects, which was accompanied by the error `'set' object is not subscriptable in the celery-parsing_error.log file`. This fix ensures the correct Compose query parsing.

- Previously, users reported an issue with the built-in connector for SAP HANA data sources where in some cases it extracted columns as tables. To address this issue, we introduced an **alation_conf** parameter `alation.feature_flags.enable_filtering_of_column_hierarchy_views_in_SAP_HANA`. For filtering out the auto-generated column hierarchy views from the `_SYS_BIC` schema during MDE for SAP HANA, set this parameter to `True`.
- For the Alation Analytics V2 ETL, increased the RabbitMQ timeout from 60 seconds to 3 hours so that the load from the `stg` to `dim` tables does not get disrupted. As part of this change, identity will be restarted for the `total_replace` tables so that not to exceed the `int` upper limit of 2 billion rows.
- Previously, when metadata was uploaded in bulk using the public API, values only appeared in filters in Advanced Search after users accessed the page of the corresponding Catalog object. After this fix, the metadata uploaded in bulk will be successfully indexed and available in Search filters.
- Added a new prompt as part of the Alation Analytics V2 Postgres upgrade from version **9.6** to **13**, where an admin can provide a location where they want to store their SQL dump before proceeding with the upgrade.
- Improved the performance of Articles when field-level permissions and multi-picker custom fields are in use in the article templates.
- Improved the performance of the application start times. Previously, some customers reported a number of slow pages and poor user experience because of this slowness. As a result of this improvement, the homepage loading times should significantly decrease. We also optimized the amount of JSON encoding and serialization during the application start when Alation is loaded from a bookmark or after reloading a page. In some cases this should noticeably reduce page loading times. In addition, the page loading indicator is now displayed consistently and users can always see that their actions, such as clicking on links, have the appropriate effect.
- In Alation Analytics V2, the ETL Manager is updated to generate the load ID sequence from the ETL checkpoint table instead of the Redis cache.
- Fixed an issue where the server health check process sent multiple emails daily for jobs that were running on schedule and without issues. Now, alerts are only sent when the corresponding alert conditions are met.
- Fixed an issue with email notifications for the data dictionary download: the download data dictionary email task failed for objects of large size with about 100k or more child objects. After the fix, this should no longer be the case.

Cloud

The Alation Analytics V2 database in EKS is externalized to AWS RDS. If RDS is enabled, the Postgres pod will not be deployed and all the database requests will be sent to the RDS host that is provisioned by the Unified Installer.

1.8.8 RELEASE 2022.1.4

build 11.2.6.161723

- For Stewardship Workbench (bulk-actions on the Search page), added ability to bulk-add objects to and bulk-remove objects from policies (business policies).
- For Stewardship Workbench, added success and failure messages to the actions of adding objects to and removing objects from policies and data policies. Users adding objects to or removing objects from a policy are now notified in the user interface if their action was successful.

- Fixed an issue with Alation Analytics V2 where the `user_type` value was not fetched during ETL for users associated with both a built-in and a custom group. This issue was resolved by regrouping the timestamp filters of the extract query. Now, if a user belongs to a custom group, this information is extracted to the Alation Analytics V2 correctly. If a user does not belong to a custom group, the name of the built-in group will be extracted as the `user_type` value for this user.

Important: Action Required

After the update, admins need to reset the ETL checkpoint for user data to reconcile the existing records. See [Resetting ETL Checkpoint for User Data](#).

- The following parameters were enabled on the internal Postgres database:

- `async_archiving`
- `wal_compress`
- multi-threaded backup

This change applies to the Backup V2 tool. It does not affect the steps of the backup process. Enabling these parameters should improve the backup performance and help avoid high I/O contention from the backup process.

- Fixed an issue where the information about users and groups who were given access to data sources on the **Access** tab of the data source settings page was not extracted into the Alation Analytics database during the ETL. Added logic to update the `ts_updated` column for a data source when a user or group is added or removed from data source access settings.
- Some workbooks were not extracted from the Tableau source due to a gateway error from the Tableau Rest API after making a certain number of calls. In order to resolve this, a delay is set when Alation receives this error from the Tableau Rest API before retrying.
- Added ability to run multiple threads for Hive query log extraction, reducing the total QLI time for Hive data sources.
- This patch release provides an updated version of the Upload Logical Metadata API. This modified API version allows updating multiple articles with the same key (article title) using the `article_id` parameter in the API payload. This API version is available on demand only.

Alation Cloud

Fixed an issue with the update process on the Alation Cloud to ensure that external Postgres, if configured, is not upgraded as part of the Alation update.

1.8.9 RELEASE 2022.1.2

build 11.2.4.161105

- Improved the process of migrating lineage data from Lineage V2 to Lineage V3. This change fixes a data discrepancy issue where there were fewer links in the lineage database than in the Rosemeta database after the migration. This change also improves the performance of the lineage migration process and decreases the time required for the migration.
- Upgraded the built-in driver for Databricks from **2.6.17** to **2.6.21**. This addresses the Databricks driver security vulnerability described in *January 24, 2022 - Log4j 2 Update and Databricks Connector Security Advisory*.

Alation Cloud Service

- Alation Analytics V2 database in EKS is externalized to AWS RDS. If RDS is enabled, the Postgres pod will not be deployed and all the database requests will be sent to the RDS host that is provisioned by the unified installer.

1.8.10 RELEASE 2022.1.1

build 11.2.3.159997

This patch is a security update for the built-in PostgreSQL driver vulnerability described in *March 5, 2022 - JDBC Driver Security Advisory*. In this patch, the built-in PostgreSQL driver was upgraded to a newer version where the vulnerability CVE-2022-21724 is not present. This upgrade addresses the PostgreSQL driver vulnerability for PostgreSQL, Greenplum, and Tableau sources in the Alation Catalog.

1.8.11 RELEASE 2022.1 - General Availability

build 11.2.2.159547

Data Governance App

Policy Groups

Policy Center users can now categorize policy objects by grouping them together into **Policy Groups**. Policy Groups can help organize data and business policies. One policy can be added to multiple Policy Groups. Policy Groups have a dedicated Catalog template and can be curated using custom fields associated with the template. They are discoverable using Alation Search and Search filters. The UI of the Policy Center page was enhanced to support Policy Groups: they appear in the **Policy Groups** table on the top of the page, while policy objects are listed in the **All Policies** table under the Policy Groups table.

Snowflake Tags Synchronization

Alation has added support for Snowflake tags, increasing the number of data governance activities that users can perform for the Snowflake sources in the Catalog. Server Admins can enable Snowflake tags synchronization in **Admin Settings > Feature Configuration**. During metadata extraction, Alation synchronizes Snowflake tags with the picker custom fields in **Admin Settings > Customize Catalog > Custom Fields**. Catalog Admins can associate tags with Snowflake data objects, such as Data Source, Schema, Table, and Column. Catalog Stewards can update tag values on the Snowflake objects. All users with access to Catalog pages of Snowflake objects can view the associated tag values.

Stewardship Workbench: Data Policy Bulk Actions

Stewards can now perform bulk curation of Snowflake table and column objects by adding or removing Snowflake data policies in bulk directly on the Search results page. These capabilities fall under **Stewardship Workbench** in the Governance App. Stewardship Workbench can be accessed from the Governance access page **Apps > Govern** by users with the Steward role or higher. The **Perform Bulk Actions** button is also available on the Search results page to users with appropriate roles. In this release, Alation adds the ability to bulk add or remove Data Masking and Row Access policies: Alation validates the selected objects for relevance and sends an API request to Snowflake to apply the changes. Specific types of policies will be

applied to relevant objects only: Data Masking policies to column objects, and Row Access policies to table objects. The user will receive an email notification when policies are successfully applied.

Workflow Center Improvements

Ability to Use People Sets to Assign Reviewers

When assigning reviewers for a workflow, users are now able to select People Sets.

Ability to Include or Exclude Top Level RDBMS Objects from Workflows

When configuring a workflow, users can include or exclude the top level RDBMS object from the workflow using the **Top level source included/excluded** toggle on the workflow configuration page. For example, users can select a Schema as the Source Object, but exclude this schema object from the workflow, only enrolling the child Tables and Columns of this schema into the review process.

Curation Growth Report

The new **Curation Growth** report shows the curation trends for Catalog objects (articles, datasources, schemas, tables, columns) over the last year from the current date. The Curation Growth report is a tool for Catalog Admins and Stewards to further drive the Catalog adoption, curation effort, and governance activities.

Governance Dashboard Accessibility

The Governance Dashboard of the Governance App is now accessible to users with all roles.

Search and Discovery

Elasticsearch Upgrade

The Elasticsearch component has been upgraded to version **7.16**, which removes Log4j 2 vulnerabilities discovered in December 2021.

Alation Full-Page Search Improvement

The Alation Full-Page Search interface now allows for selecting multiple object types from the **Object Types** Search filter.

Homepage Improvements

- Catalog Admins can now configure the Alation Homepage. In Admin Settings, the **Customize Homepage** and **Branding** links are now located in the **Catalog Admins** section. Catalog Admins will see the **Customize Catalog**, **Branding**, and **Alerts** links when they click on the Settings icon on the top right.
- Server and Catalog Admins can now customize the background image of the Homepage header. They can upload a custom image in **Admin Settings > Catalog Admins > Customize Homepage > Homepage Search Banner**.
- Server and Catalog Admins now have the ability to add alerts to the Homepage. A message and style for an alert can be set up in **Admin Settings > Catalog Admins > Alerts**.

Back References for RDBMS to RDBMS Object Sets

Users can now add back references between Alation object types, leveraging the Object Set custom field. Back referencing is now supported for RDBMS objects. On Catalog pages of RDBMS objects, you can add references to other RDBMS objects as values of Object Set fields. In such a case, back references will be added under the **Referenced By** section on the page of the linked RDBMS object.

Public API Enhancements

Search API

Users can now browse the contents of Domains using the GET method of the Search API `<your_Alation_URL>/openapi/search/>`. The GET endpoint now supports the `domain_ids` parameter, providing the ability to list member objects of a specific Domain.

Domains API

- The Domains API `<your_Alation_URL>/openapi/domain/>` now provides an endpoint to programmatically view which membership rules are applied to domains and subdomains: **POST /domain/membership/**.
- Server and Catalog Admins can now delete Domain objects in bulk using the DELETE endpoint of the Domains API.

Data Source API

The PUT endpoint of the Data Source API `<your_Alation_URL>/openapi/datasources/>` now has a new parameter `compose_default_uri` that allows users to update the Compose Default Connection URI field for a data source. This does not impact the Service Account URI of the data source.

Compose

Select, Edit, and Delete Connections in Compose

Users can now select, edit, and delete data source connections and user credentials in Compose and the Catalog. This applies to:

- Connection settings in Compose
- Connection setting for running query forms
- Connection settings for Excel Live Reports

Previously, users did not have the ability to modify or remove existing connection strings in Compose or the Catalog.

Connect Platform

Extracting Descriptions for Struct Data Type for Google BigQuery

MDE from the Google BigQuery data source was extended to extract descriptions of the struct data type. The extracted descriptions are not editable in the Catalog. They are displayed in the Source Comments field on the Catalog page of a struct column object.

Column-Level Lineage Parser Add-Ons

- Query parser has been enhanced to automatically extract column-level lineage from common queries for **Databricks** and **Azure Synapse** (SQL Server) data sources. Column-level lineage parser add-ons for Databricks and SQL Server are available as beta versions.
- Column-level lineage extraction from **Google BigQuery** and **Amazon Redshift** sources is now general availability. In the previous release, these add-ons were available as beta versions.

Manual Lineage Enhancements (Beta)

- The lineage links can now be created, edited, and removed by all admin roles: Server Admin, Catalog Admin, and Source Admin. Previously, creation of manual lineage was only available to users with the Source Admin role.
- Users with the Server and Source Admin roles can now create column-level lineage links. Previously, only table-level lineage could be created.

Ability to Turn Off Automatic Lineage for RDBMS Data Sources

Data Source Admins now have the ability to disable automatic lineage generation in the settings of RDBMS data sources. When disabled, lineage will not be created automatically during query log ingestion, metadata extraction, and from Compose queries. Lineage can be updated via API or manually. The **Disable Automatic Lineage Generation** setting is located on the data source settings page > **General Settings** tab > **Lineage Creation** section (bottom of the page).

Key Pair Authentication for Snowflake

Previously, when using key pair authentication for extraction from Snowflake, an admin had to specify sensitive authentication credentials in the JDBC connection URI. From this release, Key Pair Authentication is supported in a secure way where no credentials information is exposed in the Alation UI, URI, internal APIs, or logs.

OCF Connector for SAP BW and SAP BW/4HANA (Beta)

Users can now catalog their SAP BW sources in Alation and discover, search, and enrich SAP BW metadata for better analysis and governance:

- BW Classic and BW on HANA support InfoObject, DSOs, InfoCubes, MultiProvider, and Query (MDE only)
- BW/4HANA supports InfoObject, ADSOs, Composite Provider, and Query (MDE only)

Built-In Driver Upgrades

- The built-in driver for DB2 was upgraded to version **11.5**
- The built-in driver for Presto has been updated to version **0.265.1**

Enterprise Deployment

Viewer Role Enforced by Default

From 2022.1, the **Viewer** role enforcement is enabled by default. Previously, the default value of the feature flag `alation.feature_flags.enable_permissions_middleware_feature` was `False`, which meant that the Viewer role restrictions were not enforced and Composer, Steward, and Viewer roles had the same level of access. Now, with the default value of this parameter being `True`, new installations of Alation 2022.1 will feature a more restricted access for Viewers by default. This change also affects instances where the Viewer role was not enforced before updating to 2022.1: after the update, the access level of the Viewer role will change to the scope permitted for this role.

Manage Service Account Credentials in Azure Key Vault

Alation now provides a way to securely store the service account credentials in the Azure Key Vault at a customer's organization instead of the Alation database. The Catalog will retrieve the account credentials from Azure Key Vault for MDE, QLI, Profiling, and Sampling. This guarantees greater security and allows IT admins to rotate the credentials based on their organization's security policies.

Use MS Modern Email Authentication for Outgoing Email

Alation now supports authentication of the SMTP server used for outgoing email with Microsoft modern email authentication. Email server configuration options are available in **Admin Settings > Server Admin > Email Server** page.

Improvements

- As part of the Log4j 2 vulnerability mitigation effort, the Hive Connector was patched to use a newer version of Log4j 2. This Log4j 2 version specifically addresses CVE-2021-44832.
- To further enhance security of the Alation application, the OpenJDK version used in the Hive adapter has been updated to *8u312*.
- Query Forms in Compose now allow for multi-selection of filter values. The user experience with Query Forms in the Catalog and Compose has also been improved with better aligning the UI elements of the filter and enhancements to the handling of default values, cached variables, and quotation marks displayed in the fields.
- Previously, the scope parameter in the OAuth token response was treated as a mandatory field by Alation; however, this parameter is not provided by some of the SSO applications, for example, Pingfederate, which caused an error during authentication. The scope parameter is now treated as optional by the authentication logic in Alation.
- In order to configure the Deferred Query Execution, you can now use the **alation_conf** flags `alation.connector.managed_execution.max_query_concurrency` and `alation.connector.managed_execution.max_export_concurrency`. In the previous release, these parameters did not work. This is an alternative to setting the JVM flags directly using the parameter **alation_conf** `connector.extra_flags`. Changes to parameters `alation.connector.managed_execution.max_query_concurrency` and `alation.connector.managed_execution.max_export_concurrency` require the action `alation_action_deploy_conf_all` and a restart of the Connector component.
- It is now possible to use SSL connections from SSL-enabled databases when whitelisting IP addresses for the Alation Analytics V2 data source. Previously, non-SSL connections were added for whitelisted IPs even when SSL was enabled on the Postgres server.
- The results of the `scan_postgres` action are now stored in the internal server database tables `jobs_scanpostgresstatus` and `jobs_corruptedtables`. The `scan_postgres` script will no longer fail after detecting the first corruption; instead, it will continue to verify the database until all tables are checked. It is also possible to add the `scan_postgres` alert to the list of health alerts for administrators: admins will get emails in case a Postgres scan finds a case of Postgres corruption.
- Added the ability to configure the format for the **Name ID** property of SAML authentication requests. The Name ID format can be selected in **Admin Settings > Authentication > SAML** settings. Users can set Name ID to persistent or emailAddress.
- The **Archive processed QLI data now** action in **Admin Settings > Server Admins > Miscellaneous** will archive the session data along with the events and mentions data.

Bug Fixes

Note: 2022.1 GA includes all bug fixes listed for the 2022.1 LA version.

- Fixed multiple minor CSS issues that were due to the CSS library upgrade in the LA version.
- Fixed an issue where names of query authors were not displayed in the Table view on the search results page.
- Fixed an issue where the Popular Articles report on the Alation Analytics Dashboard was not rendered correctly. After this fix, the report will display the top articles data, as expected.
- Fixed a bug where since the update to 2021.4.x, users experienced issues with SSO connections from Compose to AWS sources and Azure Synapse (SQL DB). As a result, Data Source Admins had to reconfigure those SSO connections in Alation. SSO connections for AWS and Azure data sources will continue to work as expected after the update to 2022.1.
- Fixed a bug where the object preview page of the selected table or column object did not load in the object preview section under the Lineage diagram, and the loading icon was displayed indefinitely. Now, the previews for the table or column objects selected on the diagram will be displayed as expected.
- Fixed an issue where users were unable to set a Compose connection as an active connection until they reloaded the Compose page. Now, users can create a connection and select it as an active connection without the page reload.
- Fixed an issue where users with the Composer role received the *permission denied* error when they attempted to perform a data upload. Now, users with the Composer or Steward roles can successfully upload data for a table.
- Fixed an issue where users with the Composer or Steward roles could not manage the credentials of Compose connections. After this fix, all Compose users can manage their database credentials.
- Fixed an issue where performing a bulk action after selecting all objects in the search results did not produce the expected outcome: actions were applied to a subset of the selected objects. Now, bulk actions apply to all selected objects.
- Fixed a bug where the **Submit for Review** button was disabled when changes were proposed for the Title field. Now, users are able to successfully submit changes to object titles.
- Fixed an issue where the Snowflake key pair authentication checkbox was displayed on the settings page of all data source types.
- Fixed an issue where the restore process for a backup taken with the Backup V1 tool issued an error *Failed to access directory "/data2/backup/pgbackup/backups/alation": No such file or directory ERROR: Instance 'alation' does not exist in this backup catalog.*

1.8.12 RELEASE 2022.1 - Limited Availability

build 11.2.0.158074

For feature list, see the list of features in the General Availability section.

Bug Fixes In the LA Version

- Fixed an issue where export to CSV failed with a generic error *Something went wrong while processing your CSV file for <table_name>* for the Impact Analysis report with column-level lineage data. After this fix, exports to CSV should be successful.
- Fixed an issue with Search results when only the name of a schema was displayed in the suggested results list instead of the full name of the type database.schema. Now, the full name is displayed for schemas.
- Fixed a time zone issue in query forms for dates selected from the custom date picker widget. Previously, users reported that after selecting a date, the date which was one day earlier was populated in the Date filter. Now, the date that is selected will be displayed.
- Fixed an issue where redundant error messages related to `pg_last_xlog_receive_location` were written to the Postgres log, creating logging noise and preventing an effective search for useful error messages. Added logic to clean up WAL logs regularly.
- Fixed an issue where users were unable to reject a Lexicon expansion due to a specific network security setup. The issue was resolved by improving the handling of payloads within the DELETE request: the request information is now sent in the request URL parameter instead of the request body.
- Fixed an issue where table aliases were represented as temporary tables on the Lineage diagram. Only real table objects are now plotted on the Lineage diagram.
- Fixed an issue where extraction from Tableau failed due to a large log message size. To avoid the issue, the log message is currently truncated to the specific number of characters that is configurable. This is configurable using the parameter `alation.job.maxMessageLen` that defaults to 1000.
- Previously, the SAML assertion encryption feature did not work with ADFS. Now the encrypted assertions from ADFS are correctly decrypted in Alation.
- Fixed an issue where users querying the Alation Analytics V2 database for search queries retrieved empty results although search was in heavy use on the instance. The problem was caused by the column `num_results` of the table `search_queries` being populated incorrectly. Now, Alation Analytics queries about search will return results.
- Fixed an issue where the values of rich text custom fields with an @-mention were not correctly stored in the Alation Analytics database. After the fix, the full value of such fields can be retrieved when querying Alation Analytics.
- Previously, if Alation Analytics V2 had no data or had not been enabled on an instance, the Governance Dashboard displayed loading icons on all charts. Now, the dashboard will display empty charts with the *No data* message.
- Fixed a bug where the `assigned_to` field in the `public.conversations` table of the Alation Analytics V2 database was not populated and stored the NULL value instead of the actual value. The issue was resolved by improving the JOIN logic for joining the required Rosemeta tables in order to populate the `public.conversations` table.
- Previously, when the Users table was paginated and an admin attempted to approve a user by bypassing email confirmation or to change the role of a user that was located on pages after page 10, the page was refreshed displaying a list of different users. Now, actions on paginated User tables do not affect the list of users on the current page.
- Fixed an issue where extraction from Tableau failed with the error *out of memory DETAIL: Cannot enlarge string buffer*. This was resolved by adding logic that handles extraction of a large volume of BI permission information.
- Fixed an issue where Compose users could not connect to the Presto data source with an error *Connector Error: value too long for type character varying(214)*. The bug was due to a length limitation

on the field used to store the properties of `db_username` and `db_password`. This was resolved by changing the associated field type.

- Fixed an issue where QLI from Hive failed attempting to extract query history outside of the configured date range. Now, QLI will only pull in query history that falls within the specified range.
- Fixed an issue where republishing a query erased previous edits to the Description field of the query. Now, the Description field stays intact, as is expected.

1.8.13 Known Issues

Alation Analytics

The `aamanager` container may not start automatically after the update of Alation Analytics or after a system reboot. In this case, the admin performing the update needs to start it manually, on the Alation Analytics host, with the following command: `sudo service aamanager start`. To check the status of AA Manager: `sudo service aamanager status`.

Sources

- Dynamic table sampling does not work for Hive data sources added on the default Hive framework.
- When a user enters incorrect credentials for the Snowflake key pair authentication on the settings page > **General Settings** tab, the error message will prompt to check username and password. This wording is incorrect and users should check the key pair authentication credentials that they entered.
- For the Snowflake data source, on the settings page > **General Settings** > **Service Account**, the **Save** button for key pair authentication configuration does not automatically trigger the connection test. To test the connection, users have to click the **Test** button as a separate action.

UI

- You may encounter minor CSS issues in the LA version due to the most recent CSS library updates. On some pages and dialog windows, font styles may appear inconsistent, input fields may be misaligned, and the styles may look off. **FIXED IN GA**

Authentication

After the upgrade from 2021.4 to 2022.1 LA, Compose SSO for Azure Synapse with Azure AD may fail with the following error: *Error: Redirect handler: The scope must include an "openid" value.* **FIXED IN GA**

Compose

- Users are unable to set a newly added connection and credentials as an active connection in Compose until after the Compose page is reloaded. **FIXED IN GA**
- For SQL Server and Hive data sources, a new connection created in Compose may result in the error *Error verifying credentials. Please check your username and password.*
- For users with the Composer role, the **Upload Data** action returns a *Permission Denied* (403) error. This action should be accessible to Composers. **FIXED IN GA**
- New user credentials cannot be added to the **Connect as (Select user)** dropdown for Google BigQuery in the *transient* authentication mode. Instead, a black field is added to the list of users.

- Query forms don't prompt for credentials with Google BigQuery in the *transient* authentication mode.
- Users with non-admin roles (Composer and Steward) are unable to add new user credentials in the connection dialog. **FIXED IN GA**

Search

- Performing bulk actions after selecting all objects in the search results does not produce the expected outcome: actions are applied to a subset of the selected objects. **FIXED IN GA**

Workflows

- The **Submit for Review** action appears disabled when edits are suggested for the Title field. **FIXED IN GA**

Backup and Restore

- When Backup V2 is disabled and then a backup is created using Backup V1, the restore process from this backup generates the following error message in the log file although the restore is successful: *WARNING: Failed to access directory "/data2/backup/pgbackup/backups/alation": No such file or directory 2ERROR: Instance 'alation' does not exist in this backup catalog.* **FIXED IN GA**

1.9 Release Notes 2021.4 (General Availability)

Note:

1. If you are skipping patch 2021.4.4, please make sure to review the release information in [RELEASE 2021.4.4](#) below and the information in [/installconfig/Update/UpdateFrom20214To20214Patch](#). The 2021.4.4 patch includes important security updates which require changes to **alation_conf**.
 2. Review the information in [RELEASE 2021.4.10](#) as an additional action may be required after updating Alation Analytics V2 if you haven't performed it on a previous release.
-

1.9.1 RELEASE 2021.4.16

build 10.2.17.162478

- Updated the transform job for the article data in Alation Analytics so that it updates the `ts_updated` field when a field that is not present in the `public.article` table is updated in Rosemeta. For example, the `ts_updated` field is now updated if a user favorites or un-favorites an article page, follows or unfollows an article, or edits the Description field so that the `ts_updated` in Alation Analytics and the corresponding value in Rosemeta will match. If you experienced this issue, contact Alation Support to reset the ETL checkpoint for article data on your Alation Analytics database.
- Addressed an issue with high resource contention during the backup run. Previously, users observed slowness in Alation at specific hours of the day. Further investigation showed that the slowness coincided with the time of the `dump_all_data` stage of the Alation backup process and was due to resource contention on the host. This fix introduces an **alation_conf** parameter `alation.backup_v2.num_of_thread` that allows tuning the CPU usage down if needed.

1.9.2 RELEASE 2021.4.15

build 10.2.16.162361

- Fixed an issue where the data upload from a CSV file failed for Hive data sources. Now, data can be successfully uploaded into Hive data sources from the Alation catalog.
- Fixed an issue where after upgrading from a previous version, the Alation Analytics app icon disappeared from the **Apps** menu in the Alation user interface. The fix ensures that the Alation Analytics app icon is always available in the **Apps** menu when the Alation Analytics app is enabled in the catalog.

1.9.3 RELEASE 2021.4.14

build 10.2.15.162327

This patch is a security update providing an important fix to mitigate the Log4j 2 vulnerability discussed in *December 10, 2021 - Log4j2 Security Advisory*, and also provides other fixes as defined below.

Upgrading to this security patch will apply the fixes automatically. The patch includes:

- Addressed issues in Event Bus by removing the potentially vulnerable classes of Log4j from the corresponding Log4j.jar file included into the Alation application.
- Fixed an issue that prevents articles with duplicate titles to be created for a custom API.
- Fixed an issue where the Alation Analytics database did not correctly reflect the changes to the picker and multi-select picker field options in **Customize Catalog > Custom Fields** that were associated with existing catalog objects. The issue was caused by the `ts_updated` field value in the internal server database not capturing the event of the change. Now, the `ts_updated` field correctly reflects the time of the change and the corresponding data is transferred to the Alation Analytics database during the ETL.
- Fixed an issue causing the user profile page to load slowly when having many back references, resulting in a 502 error. The profile page now loads faster, even if there are many back references.
- Fixed an issue causing a user without access to a private data source, set as watcher, even if the user no longer has access to the data source, to keep receiving notification emails for changes done at the object level. To fix this, we removed watchers in the potential list of recipients that do not have permissions to view the data source that is the subject of the notification.
- Fixed a polling issue in Compose causing warning alerts shown on Compose UI during SQL execution, resulting from long polling when trying to fetch Hive logs. The fix handles the log resulting from the polling separately.

1.9.4 RELEASE 2021.4.13

build 10.2.14.162082

- Users can now see the default values in the input fields of query forms. Previously, the default values were not displayed.
- Added more improvements for query forms, such as support for multi-selection, aliases, single quotes, and the NOT IN clause. Fixed some minor UI issues with field alignment.
- Added a prompt to the Alation Analytics V2 installer as part of the upgrade from Postgres version 9.6 to version 13, where the admin can provide a custom location to store the Postgres dump before proceeding with the upgrade.

1.9.5 RELEASE 2021.4.12

build 10.2.13.161694

- Upgraded the built-in driver for Databricks from **2.6.17** to **2.6.21**. This addresses the Databricks driver security vulnerability described in *January 24, 2022 - Log4j 2 Update and Databricks Connector Security Advisory*.
- Fixed an issue with Compose authentication for Google BigQuery data sources after the Viewer role was enabled in Alation. Non-admin users received the permission denied error when connecting to their Google BigQuery data source in Compose. After the fix, non-admin users can successfully connect to Google BigQuery data sources with their database credentials.
- Previously, when metadata was uploaded in bulk using the public API, values only appeared in filters in Advanced Search after users accessed the page of the corresponding Catalog object. After this fix, the metadata uploaded in bulk will be successfully indexed and available in Search filters.
- In the Compose connection dialog, all user roles, both admin and non-admin, can now manage database credentials for the selected connection. Previously, non-admin users could not add new credentials or verify the connection status due to the permission denied error.

1.9.6 RELEASE 2021.4.11

build 10.2.12.161303

- This patch release provides an updated version of the Upload Logical Metadata API. This modified API version allows updating multiple articles with the same key (article title) using the `article_id` parameter in the API payload. This API version is available on demand only.
- The following parameters were enabled on the internal Postgres database:
 - `async_archiving`
 - `wal_compress`
 - multi-threaded backup

This change applies to the Backup V2 tool. It does not affect the steps of the backup process. Enabling these parameters should improve the backup performance and help avoid high I/O contention from the backup process.

- Some workbooks were not extracted from the Tableau source due to a gateway error from the Tableau Rest API after making a certain number of calls. In order to resolve this, a delay is set when Alation receives this error from the Tableau Rest API before retrying.
- Improved the performance of a number of pages for which users reported slowness. Loading time should significantly decrease for the homepage and Search. Page loading time when Alation is opened using a bookmark, has also been improved. The page loading indicator in the Alation user interface is now displayed consistently so that users see that their actions, such as clicking a link, have the appropriate effect.
- Fixed an issue with email notifications for the data dictionary download: the download data dictionary email task failed for objects of large size with about 100k or more child objects. After the fix, this should no longer be the case.

Alation Cloud

- Fixed an issue with the update process on the Alation Cloud to ensure that external Postgres, if configured, is not upgraded as part of the Alation update.

1.9.7 RELEASE 2021.4.10

build 10.2.11.160653

- Fixed an issue with Alation Analytics V2 where the `user_type` value was not fetched during ETL for users associated with both a built-in and a custom group. This issue was resolved by regrouping the timestamp filters of the extract query. Now, if a user belongs to a custom group, this information is extracted to the Alation Analytics V2 correctly. If a user does not belong to a custom group, the name of the built-in group will be extracted as the `user_type` value for this user.

Important: Action Required

After the update, admins need to reset the ETL checkpoint for user data to reconcile the existing records. See [Resetting ETL Checkpoint for User Data](#).

- Added the domains information to the ERD of the Alation Analytics V2 database. Using the ERD, analysts can explore how the tables with the domain data can be joined with other tables of the database. The domain tables in Alation Analytics provide the ability to retrieve data about page visits for domain objects, tags attached to a domain, custom fields of a domain object, and the Catalog objects associated with a domain. A Server Admin needs to **initiate the Alation Analytics database** in the Settings page to apply the changes.
- In Alation Analytics V2, the ETL Manager is updated to generate the load ID sequence from the ETL checkpoint table instead of the Redis cache.

1.9.8 RELEASE 2021.4.9

build 10.2.10.160030

This patch is a security update for the built-in PostgreSQL driver vulnerability described in *March 5, 2022 - JDBC Driver Security Advisory*. In this patch, the built-in PostgreSQL driver was upgraded to a newer version where the vulnerability CVE-2022-21724 is not present. This upgrade addresses the PostgreSQL driver vulnerability for PostgreSQL, Greenplum, and Tableau sources in the Alation Catalog.

1.9.9 RELEASE 2021.4.8

build 10.2.9.159647

- Updated Elasticsearch version to 7.16.
- Now you can set the **alation_conf** flags `alation.connector.managed_execution.max_query_concurrency` and `alation.connector.managed_execution.max_export_concurrency` in order to configure query queueing. Changes to these flags take effect after redeploying **alation_conf** using the `alation_action deploy_conf_all` command and restarting the Connector component. This is an alternative to setting the JVM flag `alation_conf.connector.extra_flags` directly.
- Fixed a bug where since the update to 2021.4, users experienced issues with SSO connections from Compose to AWS sources and Azure Synapse (SQL DB). As a result, Data Source Admins had to reconfigure those SSO connections in Alation.

- Fixed an issue where the information about users and groups who were given access to data sources on the **Access** tab of the data source settings page was not extracted into the Alation Analytics database during the ETL. Added logic to update the `ts_updated` column for a data source when a user or group is added or removed from data source access settings.
- In Workflow Center, Change Requests on people and object set fields that contained more than one object type were incorrectly marked as the last selected object type. This issue has been fixed.

1.9.10 RELEASE 2021.4.7

build 10.2.8.158918

- Fixed an issue where the Lineage diagram sometimes showed the loading cue indefinitely when loading large graphs. Now, large Lineage diagrams load successfully.
- Fixed a time zone issue in query forms for dates selected from the custom date picker widget. Previously, after selecting a date, the date which was one day earlier was populated in the Date filter. Now, the date that is selected will be displayed.
- Addressed an issue with the search index after restoring Alation from a backup when some Catalog objects became unavailable in Search results. Added a restart of the Celery services before the `rebuild_es_index` action is initiated by the `destructive_restore_all` script. This resolves potential indexing issues on restored instances.
- Fixed an issue where a backup resulted in an error if a new backup was started before the first one was completed. Several Alation backups cannot run simultaneously. A lock file `/var/tmp/alation_backup.lock` was added to the Alation server to ensure that a new backup cannot be started when another backup is already in progress. Now, when a backup process begins, it acquires the lock. The lock is released when the backup is complete.
- Fixed an issue where the restore process generated the following error message in the log file although the restore was successful: *WARNING: Failed to access directory "/data2/backup/pgbackup/backups/alation": No such file or directory 2ERROR: Instance 'alation' does not exist in this backup catalog.*

1.9.11 RELEASE 2021.4.6

build 10.2.7.157270

- Fixed an issue with updating Alation from 2021.3 to 2021.4 where the Elasticsearch upgrade stage of the update took a significant amount of time. This fix improves the logic behind copying the contents of the Elasticsearch directory, significantly decreasing the time of the Elasticsearch upgrade.
- Fixed an issue where the Alation Analytics V2 Leaderboard did not display the Showcase/Table View and Open Query controls. This fix restores the switch between the views and the **Open New Query in Compose** button for the Leaderboard.
- Fixed an issue where the data dictionary download failed with the error *unknown error while exporting data* for specific data sources. This error, caused by the "object not found" value in Object Set or People Set custom fields, represents an edge case. Fixed by adding logic for bypassing deleted objects when generating data dictionary downloads.
- Fixed an issue where the update of Alation Analytics V2 to version 13 failed with the error *exit status 1: initdb: warning: enabling "trust" authentication for local connections*. This fix improves the handling of intermittently running Postgres processes which were the cause for the error.
- Corrected an issue that could cause CPU over-utilization in certain error handling scenarios.

- Fixed an issue where the removal of the **Steward** value from objects like Data Sources, Schemas, Tables, and Columns was not correctly reflected on the Governance Dashboard and the removed values were still taken into account in reports. Added validation for cases when the Stewards Object Set field is empty. When the Stewards Object Set is empty, it won't be included into the `complete_curation` count.
- Added support for the SHA-256 algorithm when signing SAML authentication requests. There are two new **alation_conf** settings which can be used to configure the *SignatureAlgorithm* and *DigestAlgorithm* for the SAML authentication requests generated by Alation instance:
 - `alation.authentication.saml.authn_requests_sign_algo` (defaults to sha1; can be set to sha256)
 - `alation.authentication.saml.authn_requests_digest_algo` (defaults to sha1; can be set to sha256)
- Previously, when creating a manual catalog set at the Table level, users had to use the mouse wheel to scroll through the list of results. This fix improves the scrolling experience for the Quick Search in the manual catalog set editor by enabling the keyboard up and down actions and the ability to drag the scroll bar.
- Fixed an issue with the [Table Metadata Sync API](#) where using it resulted in the 500 error. The API now works as expected.

1.9.12 RELEASE 2021.4.5

build 10.2.6.155870

- The Manual Lineage links can now be created, edited, and removed by all admin roles: *Server Admin*, *Catalog Admin*, and *Source Admin*. Previously, creation of manual lineage was only available to users with the *Source Admin* role.
- Manual Lineage now enables Admins to create Column-Level lineage links. Previously, only table-level lineage could be created.
- For Snowflake External OAuth, the scope parameter in the OAuth token response is now treated as optional. Previously, this parameter was treated as a mandatory field by Alation; however, some SSO providers, for example, Pingfederate, do not provide this parameter, which caused an authentication error in Compose. This fix ensures that Snowflake External OAuth will now work with PingFederate.
- Disk space usage email alert is now enabled by default. The alert is set in the new **alation_conf** parameter `alation.healthcheck.default_list` using the value `['disk']`. This check validates if any of the root or data volumes are filled above the threshold and sends email notifications to *Server Admins* if the threshold is hit. The alerts set in the parameter `alation.healthcheck.default_list` are triggered if the parameter `alation.healthcheck.enable_admin_alert_checks` is set to False (disabled). If `alation.healthcheck.enable_admin_alert_checks` is set to True (enabled), then it supersedes the parameter `alation.healthcheck.default_list`. When Health Check Alerts are enabled, then the checks set in the parameter `alation.healthcheck.enable_admin_alert_checks` will run instead of the checks set in `alation.healthcheck.default_list`.

Note: The threshold is set using the **alation_conf** parameter `alation.healthcheck.disk_usage_threshold`. More about Health Check Alerts: [Configuring Server Health Alerts for Administrators](#).

- Retention for the tables `public.metrics_metric` and `public.celery_job_monitor_celeryjoblog` in the internal server database has been reduced from 365 days to 30 days. Jobs to clear records older than 30 days are scheduled to run on every Sunday. There is no option to change this schedule.

- Removed redundant checks from the `scan_postgres` script. Previously, these checks could generate redundant error messages in the `scan_postgres` output.
- Fixed an issue where the upgrade or a new installation of Alation generated the following error message in the installer log file: `ModuleNotFoundError: No module named 'alation_crypt'`. This message did not point to any upgrade or installation exception and was redundant. The issue that created this logging noise was fixed and the message no longer appears in the log file.

1.9.13 RELEASE 2021.4.4

build 10.2.5.155508

This patch is a security update and also deprecates the prior two patches, 2021.4.2 and 2021.4.3. This patch is cumulative and contains all the security fixes from the prior two patches. This patch also addresses a bug with the prior two patches that would cause the update to fail when updating from Alation 2021.3.8 and later.

Warning: Patches 2021.4.2 and 2021.4.3 have been deprecated. Use this patch 2021.4.4 or a later patch, for an upgrade from 2021.3 and other previous major releases.

Update information:

- [Update Alation from 2021.4.x to 2021.4.4 or Later Patches](#)
- [Update Alation to 2021.4](#)

This patch provides fixes to mitigate the Log4j 2 vulnerabilities discussed in these security advisories:

- *December 22, 2021 - Log4j2 Security Advisory*
- *December 17, 2021 - Update to December 15 log4j2 CVE-2021-45046 Security Advisory*
- *December 10, 2021 - log4j2 Security Advisory*

For additional information about the security vulnerabilities, see the following NIST publications:

- [CVE-2021-45105](#)
- [CVE-2021-45046](#)
- [CVE-2021-44228](#)

Upgrading to this security patch will apply the fixes automatically. This patch includes:

Fix for upgrades from 2021.3.5 and later

This patch fixes a bug in the prior two patches, 2021.4.2 and 2021.4.3. The bug would cause the update to fail when updating from Alation 2021.3.5 and later, resulting in Alation being inaccessible.

Security fixes for Hive connector

The Apache Hive connector is updated with a newer version of Apache Log4j 2. This addresses the security vulnerabilities listed above.

In addition to updating the Log4j version, all invocations of the Apache Hive connector will run with the `-Dlog4j2.formatMsgNoLookups=true` system property.

Security fixes for Elasticsearch

Elasticsearch is updated with a newer version of Apache Log4j 2. This addresses the security vulnerabilities listed above.

1.9.14 RELEASE 2021.4.3

build 10.2.4.155181

This patch is a security update providing a fix to mitigate the Log4j 2 vulnerability discussed in *December 17, 2021 - Update to December 15 log4j2 CVE-2021-45046 Security Advisory*.

Security fix for the Hive connector

The Hive connector is updated to use a newer version of Apache Log4j 2, which addresses a possible DoS vulnerability described in [CVE-2021-45046](#).

1.9.15 RELEASE 2021.4.2

build 10.2.3.154697

This patch is a security update providing two important fixes to mitigate the Log4j 2 vulnerability discussed in *December 10, 2021 - log4j2 Security Advisory*.

Upgrading to this security patch will apply the fixes automatically. The patch includes:

Security fix for the Hive connector

The Apache Hive connector is updated with a newer version of Apache Log4j 2 as per NIST publication [NVD - CVE-2021-44228](#). In addition to updating the Log4j 2, all invocations of the Apache Hive connector will run with the `-Dlog4j2.formatMsgNoLookups=true` system property.

With this security update, the Alation Hive connector is not susceptible to remote code execution based on the Log4j 2 vulnerability CVE-2021-44228.

Security fix for Elasticsearch

As per recommendation from Elastic found at [Apache Log4j2 Remote Code Execution \(RCE\) Vulnerability - CVE-2021-44228 - ESA-2021-31](#), this fix disables the message lookup substitution for Elasticsearch on the Alation server. With this security update, the **alation_conf** parameter `elasticsearch.env.es_java_opts`, which sets Java options for Log4j 2, will be set to `-Dlog4j2.formatMsgNoLookups=true` by default. This change ensures that the Log4j 2 Java options have the recommended setting and that the Elasticsearch component of the Alation server is not susceptible to remote code execution based on the Log4j 2 vulnerability CVE-2021-44228.

1.9.16 RELEASE 2021.4 - General Availability

build 10.2.1.153064

New Features

Data Governance App

Workflow Center and the Change Management Workflow

The Data Governance App now includes **Workflow Center** that is a centralized location for creation and management of data governance workflows. Catalog admins can create and configure multiple workflows, enroll Catalog objects into them, and appoint users to be Reviewers and Approvers.

For Schema, Table, and Column Catalog objects, users can initiate **Change Management and Approval** workflows that take changes to the fields on Catalog pages through a review process to approval. Approved changes become visible to all users who can access the Catalog page that was edited.

Data Governance Page

Admins can now enable a new Data Governance page that allows access to various parts of the Data Governance App, such as Policy Center, Governance Dashboard, and Workflow Center. The Data Governance page can be enabled in **Admin Settings > Feature Configuration**. When enabled, the **Apps** menu will include a new icon **Govern** that opens the Data Governance access page.

Alation Analytics Governance Dashboard

The new **Governance Dashboard** pulls curation data from the Alation Analytics database to provide insights into curation progress, Catalog growth, and Steward assignments. The Governance dashboard can be enabled in **Admin Settings > Feature Configuration** and accessed from the new Data Governance page. The data will be available on instances using both Alation Analytics and the Data Governance App.

Business Policies

A new type of Policy object has been added to Policy Center: a *Business Policy*. It is a generic Policy object that can be created manually and linked to other Catalog objects using Object Set fields. Business policies are discoverable using Search. They can be managed in Policy Center.

Snowflake Policy Enhancements

Ability to Apply Row Level Access Policies to Tables

It is now possible to apply Row Level Access policies to tables in Snowflake data sources in the Policy Center UI.

UI for Removing Data Masking or Row Level Access Policies from Objects

It is now possible to remove Column objects from Data Masking and Row Level Access policies in Snowflake data sources in the Policy Center user interface.

Search and Discovery

Domains Public API

Alation has added a public API for managing Domain objects. API users can *shallow add* or recursively add objects to Domains, create Domains in bulk, bulk-edit the Custom Fields, and remove objects from Domains using this API. The Domain API can be accessed by adding **/openapi/domain/** to your instance URL.

Ability to *Shallow Add* Objects to Domains

When adding an object to a Domain, users can now choose to add this object and all its child objects or add a single object only (*shallow add*), without adding the child objects.

Article Authors Can Add Articles to Domains

Article authors can now assign a Domain to the articles they create.

Object Set Back References

When users add values to Object Set fields on Catalog pages, Alation will add back references between the objects. The following bi-directional references are supported:

- Article to Article
- RDBMS objects to Articles
- Policies to RDBMS objects

Elasticsearch Upgrade

The Elasticsearch component of the Alation server was updated from version 7.4 to version **7.14** to take advantage of security enhancements and new features.

Alation Analytics

Alation Analytics V2 Database Version Updated to 13.1

The Postgres database of Alation Analytics V2 has been upgraded to version **13.1**. On instances that are updated from previous versions of Alation, the Alation Analytics V2 database version upgrade happens automatically during the update of Alation Analytics V2 to its 2021.4-compatible version.

Domains Data in Alation Analytics V2

The Alation Analytics database now includes data about Domains and Domain object membership.

Ability to Add the Alation Analytics Leaderboard to the Homepage

Admins configuring the Homepage for the Alation Catalog now have the ability to add the Alation Analytics **Top Contributors** report to the Homepage content. The Top Contributors report can be added by customizing the Homepage configuration JSON file by adding a new JSON object: { "type": "analytics_leaderboard_row", "objects": ["top_contributors"] }. There is no UI support for this feature in 2021.4 yet.

Ability to Hide the Leaderboard from the Alation Analytics Dashboard

It is now possible to control the display of the Alation Analytics **Top Contributors** report on the Alation Analytics Dashboard with the dedicated feature flag in **Admin Settings > Feature Configuration**. To hide or display the report, use the feature flag **Alation Analytics V2 Leaderboard - List of top ten contributors**.

Compose

For Compose, the focus in 2021.4 has been on performance and scalability.

Compose Query Execution Queue Enhancements

Query execution tasks will now be queued and throttled when the system constraints are hit. Queries that are the result of:

- Manual execution from Compose (all execution options for **Run and Export**)
- Execution of query forms
- Excel live reports
- Scheduled execution

will be throttled when Compose is out of execution queue capacity. Users will see appropriate error messages, indicating that they need to try again later due to system constraints.

Query Execution queueing is disabled by default and can be enabled in **Admin Settings > Compose Settings** using the **Allow Query Execution Deferral** toggle.

Monitoring of Query Results Table Size

Alation admins can now monitor the Query Results table size in Postgres relative to the available disk space. Disk space consumed by a table should never approach the remaining disk space. There is now a background job that checks the total space available to Postgres and monitors Postgres table sizes. This job is configurable using **alation_conf** and the `monitor_pg_disk_usage.*` group of parameters. The monitoring job sends email alerts to Server Admins when the set space threshold is hit.

Reduced Query Results Retention

The default value of the **alation_conf** parameter `alation.query_exec.results_auto_persist_bytes` now defaults to zero. This means that no query results are going to be stored longer than seven days (default) and that queries of any size are subject to the same seven-day retention period. Previously, the default value of this parameter was 1 MB, which led to small query results accumulating over time. Note that this change will affect results that will be created after the update to 2021.4 and will not change the retention of existing small query results.

alation_action for Clearing Stored Execution Results

Admins can now use the Alation action `alation_action purge_execution_results` to remove the old unused execution results and reclaim disk space.

Improved Ability to Reuse Query Parameters

Analyst users can now define a query variable and parameters at the beginning of the query and reference it every time they need it in the query without repeating the parameter values.

Connect Platform

New Connector for the Open Connector Framework

Out-of-the-box OCF connectors for MySQL and Postgres data sources are now available.

Extraction of External Tables from Snowflake and Google BigQuery

Alation now extracts external table metadata from the Snowflake and Google BigQuery data sources. External tables have dedicated Catalog pages and are available for Sampling, Profiling, and querying in Compose. Data location will also be cataloged.

Lineage for external tables is currently not supported.

Support for Multiple Active Directory or LDAP Domains for Tableau BI Sources

Prior to 2021.4, Alation supported access to Tableau objects in the Catalog for users from a single AD or LDAP domain only. Now, users from multiple domains will be able to access catalogued Tableau sources.

Support of the Complex Data Types for Impala Data Sources

MDE from the Impala data source now supports complex data types: Struct and Array.

Built-In Driver Upgrades

Built-in drivers have been upgraded for:

- SQLServer (driver version 9.4.0)
- Snowflake (driver version 3.13.6)
- Google BigQuery (driver version 1.2.19.1023)
- MySQL (driver version 8.0.26)
- Databricks (driver version 2.6.17)

Data Source Certifications

- AWS Databricks and Azure Databricks 8.4
- A number of sources supported with the Custom DB connector have been certified with newer versions of CData drivers:
 - Starburst
 - Cassandra
 - Cosmos DB
 - Amazon DynamoDB
 - MongoDB
 - MongoDB Atlas
 - Salesforce DB
 - ServiceNow

Full information is available in the Support Matrix version 2021.4.

Lineage

Automatic Column-Level Lineage for Cloud Data Sources

Query parser has been enhanced to automatically extract column-level lineage from common queries.

Column-level lineage can be auto-generated for:

- Snowflake data sources (GA)
 - Requires the *Snowflake Column-Level Lineage Parser Add-On* to be purchased
- Google BigQuery data sources (Beta)
- Amazon Redshift data sources (Beta)

Lineage Diagram UI Refresh

The Lineage diagram received a UI refresh and is now using a new dark blue and grey theme for the lineage elements and links.

Creation and Management of Lineage in the Alation UI (Beta)

Data Source admins can create or update lineage data directly from the Alation UI. Most of the Lineage functionality, including API compatibility, deprecation propagation, and impact analysis reports will be available for the manually created Lineage too. This feature requires that the existing lineage data should be migrated to the *Lineage V3* service.

Lineage V3 (Beta)

Lineage V3 (Beta) is a new service for managing lineage data and generating Lineage diagrams. To begin using Lineage V3, you will need to migrate the existing lineage data from Lineage V2 to Lineage V3. The Lineage V3 framework includes a separate internal Lineage database that stores lineage data and internal server APIs to manage the lineage information. Lineage V3 uses the *Event Bus* component to store and retrieve lineage data.

Enterprise Deployment

UI for Configuring SAML Authentication

The SAML authentication can now be configured in the Alation UI on the **Admin Settings > Authentication** page. Previously, SAML authentication had to be configured on the backend of the Alation server. The new UI allows configuring the required SAML parameters as well as uploading the *idp_metadata.xml* file to the Alation server. No restart of the server is required to apply the configuration.

Compose SSO with Azure AD for Azure Synapse

Data source admins can now set up SSO authentication for Compose and Dynamic Profiling with Azure Active Directory for Azure SQL DB and Azure SQL Serverless and Dedicated Pools. This configuration uses the Alation AuthService framework and requires access to the backend of the Alation server.

Synchronization of Azure AD Groups into Alation Using SCIM

When the authentication method in Alation is set to SAML and the IdP provider is Azure Active Directory, Server Admins can set up User and Group synchronization from Azure AD to Alation over the SCIM protocol. This eliminates the need to manually manage group assignments and changes to group membership as group management will be provisioned by Azure AD.

CentOS Removal from Chroot

CentOS 8 was replaced with Red Hat UBI Init 8.4 in the Alation Chroot.

Event Bus

The Alation server includes a new component: *Event Bus*. The Event Bus is a middle layer in the application stack of the Alation server that facilitates asynchronous messaging between server components. In 2021.4, the Event Bus is only used by the new Lineage V3 service. For example, during metadata extraction, messages related to lineage will be stored in the Event Bus. The Lineage V3 service can then use those stored messages at a later time to create lineage graphs.

Improvements

- The Search interfaces - **Advanced Search** and **Query Search** are turned off by default in 2021.4. These pages are behind **alation_conf** parameters `alation.feature_flags.enable_query_search` and `alation.feature_flags.enable_advanced_search`.
- Open Connector Framework is now enabled by default (**alation_conf** parameter `alation.feature_flags.enable_gbm_v2_connector_strategy` is set to `True` by default).
- The React library on the Alation server has been upgraded to version 17.
- Added an **alation_conf** parameter `elasticsearch.version` to record the active Elasticsearch version.
- Previously, when the SCIM sync process removed groups and the Group Role Assignment feature was enabled, Alation allowed the last admin group to be removed. Now, Alation will check for remaining admin users in SCIM-synchronized groups and prevent the last admin group from being deleted.
- Whenever the Domains list on a Catalog page is closed, it will be reset to its initial state, which is the list of top-level Domains.
- The NGINX component on the Alation server has been upgraded to version 1.21.0.
- The Secret key that is stored by the AWS IAM plug-in is now stored in the encrypted format.
- Added the columns `suspended` and `ts_suspended` to the `users` table in Alation Analytics V2 to make these details available in reports about users.

- Improved the UI experience and loading times for large collections with multiple nesting levels in NoSQL data sources. Previously, large collections (for example, containing ~ 100 schemas with ~ 100 attributes each on various nested levels) took a long time to load and sometimes timed out. Now, they are going to be loaded in batches of 10 items. Users can click **Load More** to load more elements of the collection.

Bug Fixes

Note: 2021.4 GA includes all bug fixes listed for the 2021.4 LA version.

- Fixed an issue with Compose performance where users experienced Compose instability when executing multiple large queries using objects that have Data Quality flags. The presence of Data Quality Flags on objects used in queries no longer affects Compose stability.
- Fixed an issue where the highlighting and Trust Check prompts were displayed with a delay or not displayed at all in Compose query editor for queries that used objects with Data Quality flags.
- Fixed an issue where built-in groups were not listed in the Quick Search list of potential Reviewers in the **Select Reviewers** dialog. Now, Quick Search includes all groups.
- Users can no longer make trivial “no change” suggestions. Previously (2021.4 LA), there was no validation of whether or not suggestions contained meaningful changes. This could lead to confusion either on the part of the editor or Reviewers as the visual diff for the suggested change was blank. Now, when a user attempts to submit a “no change” suggestion, a modal dialog will pop-up directing them to change the suggestion.
- Previously (2021.4 LA), Excel Live reports did not run, showing only a *No results to display* message if the connector component was configured to limit the number concurrent query executions or query exports. Now, Live reports will run with such a connector configuration. Users will see the message *Query cannot run due to server resource limitations. Please try again later* if the maximum number of concurrent query executions or query exports has been exceeded.

1.9.17 RELEASE 2021.4 - Limited Availability

build 10.2.0.151763

For the feature list, see [New Features](#)

Bugs Fixed in 2021.4 - LA

- Fixed an issue where the Search filter for schemas limited the number of schemas available for filtering to 100. The facet size for schemas was increased to 500. It will allow up to 500 facets in the search filter.
- Fixed a bug where the Reply-By-Email feature did not work for users of the Outlook client app. This fix is applicable for Outlook client apps on OS X and Windows. For email apps on OS X, there is an additional **alation_conf** parameter that can be used when troubleshooting Reply-By-Email for Outlook: `alation.email.original_html_message_patterns -s "[{'attrs': {'class': '[]', 'name': 'div'}}]"`.
- Fixed an issue where LDAP synchronization stopped if one of the LDAP groups had a name containing parentheses. Added logic to skip such problem groups, log the error, and continue the synchronization.

- Fixed an issue where users without groups were not loaded into the Alation Analytics V2 database during ETL.
- Fixed an issue where users did not see the message *You are logged out* after they logged out of Alation but saw either the 401 or 403 error page instead. The problem was discovered on instances using HTTP/2. After the fix, users see the required logout page.
- Fixed an issue with Lineage from AWS Databricks where Lineage calculation failed with an error in the **celery-ingestion_error.log** file that was related to the PRIMARY_KEY usage.
- Fixed an issue where the Function objects still persisted in Search results after the parent data source was deleted.
- Fixed an issue where statements with USING were not parsed for a Teradata data source and as a result, lineage was not generated. Alation now supports queries starting with USING for Teradata. Lineage, filters, and joins information will also be available for queries that start with the USING keyword.
- Fixed a bug where the text of an @-mention was shown twice when the content was underlined, bolded, or italicized.
- Fixed an issue with QLI from Presto data sources where it failed if the entire batch of the *ExecutionEvent* objects to be processed had `cancelled=true`. This created a *Null* value set. Now, query batches with the `cancelled=true` property are parsed correctly.
- Previously, when users navigated to a page in Lexicon and clicked the **Confirmed** filter, pagination stayed on the same page, although there were fewer pages of confirmed terms. As a result, the current page appeared empty and users had to navigate to previous pages to see the terms. After the fix, the Lexicon page resets pagination when a filter is changed, and users will be automatically navigated to the first page of the filtered list of terms.
- Previously, when user profiles were uploaded from a CSV file, Alation converted usernames to lowercase. The case of the username in the source CSV file was not respected. However, the SAML response is case-sensitive. As a result, Alation created a second user with the username spelling from the SAML response. This issue has been fixed, and user profile upload is now case-sensitive. Also added case-matching logic to prevent creation of duplicate accounts for users already loaded into the system. Updated the UI with a message explaining this change.
- Fixed an issue with QLI where previously ingested Query History information was not restored after KV Store was reset during a troubleshooting session. Now, the Query Template text for the Query History section will be populated from the internal Postgres database if KV Store does not have the required Template text.
- Fixed an issue where Postgres restart failed in certain cases after the backup restore. Now, the Postgres restart waits for the recovery to be complete during the restart.
- Fixed an issue where the Job API did not return the job status if the bulk upload payload included special characters.
- Fixed an issue where the RDBMS table count done on Alation Analytics V2 sometimes did not reflect the real number of RDBMS tables in the Catalog. The issue manifested itself for tables containing columns with complex data types and for tables with long names. After the fix, users should be able to see `rdbms_columns` with names larger than 1,000. Column data type for `data_type_details` in the `rdbms_columns` table has been updated. Data with length more than 1,000 is now properly synced during the ETL.
- Fixed an issue where the Alation Analytics objects became starred after the upgrade from 2020.3 to 2021.1. Only objects that were starred by users before the upgrade remain starred after upgrade. The fix does not remove starring from objects that became starred after the update.

- Fixed an issue where the Alation Analytics V2 database did not correctly reflect the count of column objects in the Catalog. The issue existed for columns with the Struct data type or columns with data type length $\geq 1,000$. Now, columns with data type length $\geq 1,000$ are properly ingested during ETL and the `rdbms_columns` table contains the correct number of column objects.
- Fixed an issue where users received empty Search results when they selected the Tag object type in the Object Type filter. Now, the Tag objects are searchable as expected.
- Fixed an issue where the Data Sources API returned an error when users attempted to update the **Description** field for an existing Data Source. Now, the Description field can be updated successfully.
- Fixed an issue where the tables `bi_folder` and `bi_report` in Alation Analytics V2 were not in sync with the BI server metadata cataloged in Alation, and Alation Analytics still contained objects removed from the Catalog during extraction. The fix introduces the `deleted` field in the tables `bi_report`, `bi_folder`, `bi_column`, `bi_report_column`, `bi_datasource_columns`, `bi_connection_column`, and `bi_connection`.
- The name of the filter link on the Alation Analytics **Built In** tab was changed from **Select Users/Groups** to **Select Users** as the Top Contributors report cannot be filtered by Groups.
- Cyrillic symbols in the Source Comments field within columns were not displayed for Hive data sources. This issue has been fixed.

1.9.18 Known Issues

Compose

- Compose users may encounter some Compose performance instability when executing multiple large queries using objects that have Data Quality flags. Queries may be taking a long time to run or the Compose window may become temporarily unresponsive. If users report such issues, admins can consider the following workaround: using the `alation_conf` command, set the value of the parameter `alation.feature_flags.enable_semantics_check` to `False` and restart the Alation server. (**FIXED IN GA**)
- Highlighting and the Trust Check prompts may be displayed with a delay or not displayed at all for queries that use objects with the Data Quality flags. (**FIXED IN GA**)
- Republishing a query erases previous edits to the **Description** field of the query.
- Currently, the value of the `alation_conf` parameter `alation.connector.managed_execution.max_query_concurrency` does not take effect. Until this issue is addressed, it is recommended to use the parameter `connector.extra_flags` instead to set the concurrency values.
- Excel Live Reports and Query Explain are not throttled when Deferred Query Execution is enabled. (**FIXED IN GA**)

Alation Analytics V2

- The Alation Analytics V2 Leaderboard does not display **Showcase/Table View** and **Open Query** buttons. (**FIXED in 2021.4.6**)
- The Alation Analytics V2 Leaderboard remains displayed on the Homepage after the Alation Analytics V2 data source is removed from the Alation instance. The report can be removed by editing the Homepage configuration JSON file and deleting the corresponding JSON object.
- The **Date Range** filter is incorrectly applied to the Top Contributors report: a filtered report may return less data for the selected date range than the data available in the database.

- On the Governance Dashboard, the Download icon is not displayed for the Catalog Objects and Total Curation Progress charts.
- When users click the **Apps** menu while staying on the Governance Dashboard page, the **Alation Analytics** icon will be selected under Apps. As the Governance Dashboard is part of Governance App, the **Govern** icon should be selected instead.
- A minor UI issue exists on the Catalog Objects and Curation Progress by Data Object charts: the metric for the Table object should be *Tables* (plural). Currently, it is *Table* (singular).
- The PNG files downloaded for the Curation Progress by Data Object and Stewardship charts have differences in font style and size and color.

Sources

- The CREATE EXTERNAL TABLE and DROP EXTERNAL TABLE IF EXISTS statements are not parsed during QLI from Google BigQuery data sources. As a result, such queries are not listed under the **Query History** tab on the Catalog page of the data source.
- The CREATE OR REPLACE EXTERNAL TABLE and ALTER EXTERNAL TABLE statements are not parsed during QLI from Snowflake data sources. As a result, such queries are not listed under the **Query History** tab on the Catalog page of the data source.
- The lineage data created by running the *parse_all_view* script from the Alation shell is incomplete. This issue may manifest itself when troubleshooting issues with lineage using steps provided by Alation with the *parse_all_view* script included into the steps.

Workflows

- Built-in groups are currently not listed in the Quick Search list of potential Reviewers in the **Select Reviewers** dialog. (**FIXED IN GA**)
- Users can submit an empty value as a suggested change for the Object Set field type without selecting any object from the list. (**FIXED IN GA**)

Backend Operations

- The Alation server Health Check sends daily emails to server admins that specific tasks missed a scheduled run even though the jobs that are monitored by the Check run on schedule and without issues. The emails include the message: *The following tasks missed their latest scheduled run* and can list metadata extraction, Query Log Ingestion, and some other tasks. If you do not wish to receive these notifications, please remove celery from alation.healthcheck.alert_list: [Configuring Server Health Alerts for Administrators](#)
- The Backup process fails with a “KV Store error” when another backup is initiated before the completion of the first one. The Alation Backup process does not currently support multiple simultaneous backup runs. (**FIXED IN 2021.4.7**)

Authentication

- The **Test Connection** function on the **Settings** page of a Tableau source fails if the service account used for extraction is an Active Directory user with Unicode characters in the username.
- After upgrading to 2021.2.5 from a previous release, Compose users cannot connect to a Presto data source and get an error *Connector Error: value too long for type character varying(214)*. Contact Alation Support to apply a workaround if you encounter this issue.
- On instances where the value of the **alation_conf** parameter `alation.authentication.saml.assertion_decrypt_id_attr` was changed in order to set a custom decryption id attribute for encrypted SAML assertions, after upgrading from version V R7 to 2021.2.3, SAML encryption stops working. This causes login problems for SAML users.

1.10 Release Notes 2021.3 (General Availability)

1.10.1 RELEASE 2021.3.11

build 9.2.12.161428

- The following parameters were enabled on the internal Postgres database:
 - `async_archiving`
 - `wal_compress`
 - multi-threaded backup

This change applies to the Backup V2 tool. It does not affect the steps of the backup process. Enabling these parameters should improve the backup performance and help avoid high I/O contention from the backup process.

Alation Cloud

- Fixed an issue with the update process on the Alation Cloud to ensure that external Postgres, if configured, is not updated as part of the Alation update.

1.10.2 RELEASE 2021.3.10

build 9.2.11.159996

This patch is a security update for the built-in PostgreSQL driver vulnerability described in *March 5, 2022 - JDBC Driver Security Advisory*. In this patch, the built-in PostgreSQL driver was upgraded to a newer version where the vulnerability CVE-2022-21724 is not present. This upgrade addresses the PostgreSQL driver vulnerability for PostgreSQL, Greenplum, and Tableau sources in the Alation Catalog.

1.10.3 RELEASE 2021.3.9

build 9.2.10.157611

- Fixed a time zone issue in query forms for dates selected from the custom date picker widget. Previously, after selecting a date, the date which was one day earlier was populated in the Date filter. Now, the date that is selected will be displayed.
- Fixed an issue with the [Table Metadata Sync API](#) where using it resulted in the 500 error. The API now works as expected.

1.10.4 RELEASE 2021.3.8

build 9.2.9.155519

This patch is a security update providing fixes to mitigate the Log4j 2 vulnerability discussed in the *December 22, 2021 - Log4j2 Security Advisory*. This addresses the security vulnerability described in [CVE-2021-45105](#).

This patch is cumulative and fixes all the security issues from the prior two patches, 2021.3.6 and 2021.3.7.

Upgrading to this security patch will apply the fixes automatically. This patch includes:

Security fix for the Hive connector

The Apache Hive connector was updated to address the security vulnerability [CVE-2021-45105](#).

Security fix for Elasticsearch

Elasticsearch is updated to addresses the security vulnerability [CVE-2021-45105](#).

1.10.5 RELEASE 2021.3.7

build 9.2.8.155180

This patch is a security update providing a fix to mitigate the Log4j 2 vulnerability discussed in *December 17, 2021 - Update to December 15 log4j2 CVE-2021-45046 Security Advisory*.

Security fix for the Hive connector

The Hive connector is updated to address a possible DoS vulnerability described in [CVE-2021-45046](#).

1.10.6 RELEASE 2021.3.6

build 9.2.7.154704

This patch is a security update providing two important fixes to mitigate the Log4j 2 vulnerability discussed in *December 10, 2021 - log4j2 Security Advisory*.

Upgrading to this security patch will apply the fixes automatically. The patch includes:

Security fix for the Hive connector

The Apache Hive connector is updated with a newer Apache Log4j 2 version as per NIST publication [NVD - CVE-2021-44228](#). In addition to updating the Log4j 2 version, all invocations of the Apache Hive connector will run with the `-Dlog4j2.formatMsgNoLookups=true` system property.

With this security update, the Alation Hive connector is not susceptible to remote code execution based on the Log4j 2 vulnerability CVE-2021-44228.

Security fix for Elasticsearch

As per recommendation from Elastic found at [Apache Log4j2 Remote Code Execution \(RCE\) Vulnerability - CVE-2021-44228 - ESA-2021-31](#), this fix disables the message lookup substitution for Elasticsearch on the Alation server. With this security update, the **alation_conf** parameter `elasticsearch.env.es_java_opts`, which sets Java options for Log4j 2, will be set to `-Dlog4j2.formatMsgNoLookups=true` by default. This change ensures that the Log4j 2 Java options have the recommended setting and that the Elasticsearch component of the Alation server is not susceptible to remote code execution based on the Log4j 2 vulnerability CVE-2021-44228.

1.10.7 RELEASE 2021.3.5

build 9.2.6.153371

- Fixed an issue where the PUT request of the Data Source API returned the *Status 400 (Bad Request)* error when users attempted to update the **Description** field for an existing Data Source. Now, the Description field can be updated successfully.
- Fixed an issue where the tables `bi_folder` and `bi_report` in Alation Analytics V2 were not in sync with the BI server metadata cataloged in Alation, and Alation Analytics still contained objects removed from the Catalog during extraction. The fix introduces the `deleted` field in the tables `bi_report`, `bi_folder`, `bi_column`, `bi_report_column`, `bi_datasource_columns`, `bi_connection_column`, and `bi_connection`.
- Fixed an issue with QLI where previously ingested Query History information was not restored after KV Store was reset during a troubleshooting session. Now, the Query Template text for the Query History section will be populated from the internal Postgres database if KV Store does not have the required Template text.
- Fixed a bug where the Reply-By-Email feature did not work for users of the Outlook client app. This fix is applicable for Outlook client apps on OS X and Windows. For email apps on OS X, there is an additional **alation_conf** parameter that can be used when troubleshooting Reply-By-Email for Outlook: `alation.email.original_html_message_patterns -s "[{'attrs': {'class': '[]', 'name': 'div'}}]"`.
- Fixed an issue where the names or titles of some object types were not displayed in the Search Autocomplete dialog when users were typing the search keyword.

1.10.8 RELEASE 2021.3.4

build 9.2.5.152222

- Fixed an issue with Compose where highlighting and the Trust Check prompts were displayed with a delay or not displayed at all for queries using objects with the Data Quality flags. After this fix, Compose is also more stable for queries with multiple statements.
- Fixed an issue where the update from version 2021.1.10 to version 2021.2.8 failed with the error *Failed to execute "/opt/alation/alation/opt/alation/ops/actions/unjailed_update_alation"*. After this fix, updates between Alation versions are expected to be successful.

1.10.9 RELEASE 2021.3.3

build 9.2.4.148937

- Updated the Alation Analytics V2 ERD article with the ERD diagram that corresponds to the database model in 2021.3.
- Fixed an issue with QLI from the SQL Server data source where it failed with a `java.lang.StackOverflowError` exception. The issue was previously caused by the recursive call `getProbableCopies` in the V2 Parser. The Parser logic has been enhanced for this call, which eliminates the error.
- Fixed an issue with QLI from the Snowflake data source where it failed with a `java.lang.StackOverflowError` exception. After the fix, QLI should successfully parse the query that previously resulted in this error.
- Fixed an issue where the parsing of OCF Connector logs failed with an error *This is a programming error* when there were invalid UTF-8 characters in the logs. Now, the Connector tail view on the OCF Admin Dashboard renders proper log output even if invalid UTF-8 characters are present in the Connector logs.
- Added a new flag to skip PostgreSQL validation while performing a backup.

1.10.10 RELEASE 2021.3.2

build 9.2.3.147624

- Fixed an issue with updating the Alation search index after a large number of objects have changed in the Catalog during metadata extraction. Previously, after metadata extraction removed a large number of objects from the Catalog, the search reindexing queue could become backed up while processing the delete messages. This could stall the reindexing process for a long period of time. This fix improves the performance of the search index queue and decreases the time of search reindexing after a large number of objects are removed by metadata extraction.
- Fixed an issue where user profiles were not loading on the **Advanced Search** page after searching for users. Now, when a user is selected in the Advanced Search results pane, the corresponding user profile will load in the right-hand view pane.
- Fixed an issue where the Alation updater validated if encryption key rotation has been performed on versions 2021.1.x and newer. Encryption key rotation is no longer required before updating Alation to 2021.2.x and 2021.3.x.
- Fixed an issue where Alation could not load the user profile image for LDAP users and returned a 500 error. Now, user profile images can be uploaded as expected.

1.10.11 RELEASE 2021.3.1

build 9.2.2.146480

- Fixed an issue where the BI folders were extracted as empty containers even though no objects from these folders were selected for extraction from a BI source. Now, sites, projects, and folders that are not selected for extraction will not be displayed in the Catalog.
- For Legacy Hive users, added a feature flag that allows specifying from where the query string is read during QLI. Enabling `alation.feature_flags.enable_usage_of_map_reduce_name_for_query_location` will cause Hive V1 QLI to read `mapreduce.workflow.name` for the query string. By default, Alation reads `hive.query.string` to get the query string. Note that Alation recommends using the Hive V2 framework for Hive sources.
- Fixed an issue with search filters behavior on the Search results page. Previously, when running a search, users attempted to select filters in the left pane but they were not selectable until the results were fully loaded. This created an impression that the search facets were not working. After this fix, the filters in the left sidebar will appear in a loading state until all the objects in the result set are fully loaded and available.
- Fixed an issue with Sampling from an Oracle source added with the OCF connector for Oracle where Sampling failed with the error Protocol message CustomTableRequest has no "configuration" field.
- Fixed an issue where on Alation instances with no Domains, the Search bar still displayed the **All Domains** dropdown for a moment while loading data.

1.10.12 RELEASE 2021.3 - General Availability

build 9.2.1.145216

New Features

Data Governance App: Policy Center

Important: With the 2021.3 release, Alation introduces the **Data Governance App**, which is priced and packaged separately from the Alation Data Catalog. The Data Governance App includes the **Policy Center** now, with more governance features to come in subsequent releases. Contact your account manager to learn more about pricing and licensing for the Data Governance App.

Policy Center is a centralized location in the Alation Data Catalog for managing data policies that apply to the cataloged data objects. Alation surfaces relevant policies at the time of metadata extraction, ensuring that data consumers are aware of their existence and fostering a more ethical data culture.

With the introduction of Policy Center, Alation also extends the capabilities of the native Snowflake connector. For Snowflake data sources, Alation can now extract Row Access and Data Masking policies. Extracted policies are listed in the new Policy Center. Users can add data objects found under the Snowflake data source to Data Masking policies directly in the Alation Catalog. Extracted policies are enforced in the Catalog for all users viewing the Catalog pages of the associated data objects. In addition to policy extraction and application, analyst users can create new data policies from Compose.

Alation Analytics V2

Alation Analytics V2 History

Alation Analytics V2 adds the ability to perform analytics on data object history. It has 2 new columns in the `curation_history` table to analyze metadata changes on built-in fields, custom fields, queries, and more. The `action` field stores the type of history change that has occurred. The `value` field contains the change that occurred, if applicable. For example, if the Description field was changed for an object, then the new Description will be available. The previous Description value will be visible on a separate row.

Alation Analytics V2 Top Contributors Report

A new tab **Built-In** has been added to the **Alation Analytics** page that features a visualization of top contributors to the Alation Catalog. The **Top Contributors** report has two views: *Showcase* and *Table*. The Showcase view displays the top 3 Alation Catalog users based on the action metrics total. The Table view displays a number of action metrics for top 10 users, such as *number of articles created* or *number of queries published* and the summed total of all actions for each user. The users are displayed in the descending order of the total. Dashboard users can switch between the *Table* and *Showcase* views of the report.

Compose and Query Forms

Compose Improvements

- Added ability to export to MS Excel to ensure compatibility of Alation exports with this application
- Users can now limit the number of rows that can be previewed in query results in Compose. Previously it was a predefined value (1,000) but now users can change it in Compose Settings. The configurable number of rows is displayed on the Results pane to enable browsing through more rows of query results.
- Previously, column aliases in SQL queries were not reflected in Compose exports. The export process has been modified to validate the presence of an alias first before using the column name.
- Admin can now disable the public sharing of query result in Compose: added the toggle **Allow Public Sharing Of Compose Results** to **Admin Settings > Compose Settings**. Disabling this toggle will disable the Sharing radio button in Compose.
- The Auto-Suggest pop-over box in Compose now shows the full name of the selected suggestion at the bottom of the box. Users can now change the size of the Auto-Suggest box.

Query Form Improvements

- Users can now cancel query form execution on the Query Form page. Clicking on the **Cancel** button will mark the query as cancelled. The **Cancel** button does not appear in the UI when query form execution is complete or cancelled.
- Added ability to select multiple WHERE filters for a query form. Users can now either select filter conditions from a multi-selection list or choose to type them as before.

Search and Discovery

Enhanced experience for adding objects to Domains

Subject-matter experts assigned as Stewards to Catalog objects can now add those objects to Domains or remove them from Domains even when they do not have the roles of the Catalog or Server Admin. If the object has recursive support, its child objects will be added to or removed from a Domain automatically. Query Authors can now add the queries they author to Domains too.

Improved UI for assigning Domains on Catalog pages

When adding an object to a Domain, users will now see the available Domains and Sub-Domains as a hierarchy instead of a list. This UI element provides a visual representation of the Domains structure and a unified experience of selecting Domains across all parts of the Catalog.

Search Autocomplete

As users start typing search keywords in the Search box, Alation now automatically suggests 3 auto-completed keyword prompts and up to 4 live search results. Keyword suggestions are based on:

- Keyword prefix matching
- Prediction ranking using past searches and page views

See [Search](#).

Saved Searches

Users now have the ability to save searches. Saved searches can be edited, re-saved, shared, and deleted. All saved searches currently available to a user are displayed in the left bottom corner of the Search page.

See [Search](#).

More Search Improvements

- Two new sorting options have been added to Full-Page Search: **Last Updated** and **Last Created**. In previous releases, search results were sorted by **Relevance** only. Now, users can choose and apply the sorting option they prefer.
- Conversation objects in Search results can now be filtered by **Resolve Status**: *Open* or *Resolved*
- For Article objects, Search results will now include information about the parent Article as breadcrumbs
- It is now possible to hide the legacy Advanced Search and/or Query search pages from the Alation UI using **alation_conf**. To hide Advanced Search, set `alation.feature_flags.enable_advanced_search` to False. To hide Query search, set `alation.feature_flags.enable_query_search` to False. Changing the values of these flags requires a restart of the Django component

Connect Platform

Open Connector Framework for RDBMS Data Sources

Open Connector Framework (OCF) and OCF SDK can now be used to create connectors for RDBMS data sources. The SDK supports full Alation capabilities, such as connection, authentication, extraction, deep profiling, dynamic profiling, query log extraction, and lineage. The SDK is unified for RDBMS and BI connector development. In addition to the SDK, Alation releases an out-of-the-box OCF connector for Oracle data sources.

OCF Oracle Connector

The OCF Oracle connector supports a full range of Alation functionality for RDBMS data sources: basic, Kerberos-based, and LDAP; metadata extraction, sampling and profiling (including dynamic profiling), query log ingestion, lineage, and Compose. This connector will extract a full range of metadata from schemas, tables, and columns to synonyms, primary and foreign keys and data types.

OCF Admin Dashboard

Server Admins can now use the OCF **Admin Dashboard** to easily install and manage OCF connectors. This is a UI-based method to install, manage, and delete OCF connectors that eliminates the need to perform these actions using shell commands.

Snowflake Connector Improvements

- Added support for extracting Materialized Views as Catalog objects, their view definition SQL, and and broader Lineage representation
- Now extracting the *last_altered* field for data objects and displaying it in the **Properties** panel of the Catalog page
- For QLI and Compose, added support for parsing statements that contain keywords:
 - COPY INTO
 - QUALIFY
 - MERGE INTO
 - TOP

GBQ Connector Improvements

- Added support for extracting Materialized Views as Catalog objects and their view definition SQL
- Added automatic Source Comments extraction for Table and Schema objects

Data Source Certifications

- Databricks on Google Cloud Platform
- Databricks 8.3
- Splunk
- Confluent Kafka
- GitHub (Beta)
- Parquet data files on Amazon S3
- Latest versions of supported sources (Tableau, Snowflake, Azure Synapse Analytics, Amazon Athena, AWS EMR, Amazon S3, AWS Glue)

Enterprise Deployment

Postgres Upgrade to 13.1

Internal Postgres database has been upgraded to version 13.1.

Automated User and Group Provisioning with SAML Authentication

Alation now supports automatic user and group provisioning with SAML authentication using the System for Cross-Domain Identity Management (SCIM) protocol 2.0. The following data will be propagated to Alation from the external IdP based on SCIM:

- Users
- Groups
- User-group membership

Previously, the Groups sync capability was only available for LDAP authentication and now becomes available for the SAML authentication method as well. An IdP admin can use the SCIM integration to control Alation user and group management including creating users, updating user profiles, deactivating users, creating groups, updating group, syncing group members, and deleting groups. This feature helps maintain consistency between users and groups in the IdP and in Alation. After the SCIM integration is set up both on the Alation server and on the IdP side, user and group management in Alation will be handed over to the IdP.

AWS IAM Authentication for Metadata Extraction (Beta)

Authentication for MDE, Profiling and Sampling, and QLI can now leverage AWS IAM Roles & Security Token Service (STS) which generates short-lived, limited-privilege credentials (similar to Compose SSO for AWS IAM). MDE and Profiling with the IAM role are supported for Amazon Athena, Amazon Redshift and Amazon S3 sources. This functionality is released as a **Beta** version in 2021.3 as there is a 12-hour limitation on the duration of the MDE run. MDE fails if extraction takes more than 12 hours, although this restriction can be overcome by using instance profile-based policies.

Alation Sandbox

- Integration with the Alation Cloud Service
- Added a sub-command for encryption key rotation
- Added support for downloading only the required Alation package
- Added support for HA setup
- Improved user experience for admins:
 - Disk space check as part of pre-checks
 - Checkpoint file check as part of pre-checks
 - Base URL displayed at the end of a successful run
- Improved stability and robustness

Improvements

- Added ability to turn off automatic Lineage generation for BI sources using the **Disable Automatic Lineage Generation** checkbox on the Settings page of a BI source.
- Lineage will now be created from MERGE INTO queries for Snowflake data sources.
- The `excluded column` was added to the `rdbms_schema` model Alation Analytics V2.
- Non-root users now can be allowed to access the Alation Analytics V2 configuration and installation directories.
- Previously, all Alation Analytics containers ran under the `root` user. Now they are updated to run as non-root users. The Postgres container runs as user `postgres`, RabbitMQ container runs as user `rabbitmq`, and all Transform containers run as the `transform` user.
- Added support for enabling or disabling Agile Approval for Articles using the Articles API. Updated the Articles API documentation to include information about changing this setting.
- Enabled support for databricks as one of the supported parser dbtypes for Custom DB data sources.
- The Sampling query for the Oracle data source has been updated, with the inner query removed so that a full table scan does not happen during Sampling.
- Code comments will no longer be removed from the SQL code of queries when queries are parsed.
- Added support for Presto over Hive with SSL and LDAP, a configuration which was previously unsupported.
- An email notification about a change to an article will now only go out to Stewards and watchers when the user editing the Article clicks on **Save**. Added the functionality that auto-saves Article drafts in the local storage, but those auto-saves will not trigger email notifications.
- Backup V2 log file location changed from `/var/log/` to `/opt/alation/site/logs/`. The change was made to group it with other log files for better accessibility.

Bug Fixes

Note: 2021.3 GA also includes all bug fixes listed for the 2021.3 LA version.

- Fixed the issue of Backup V1 restore failing after upgrading the internal PostgreSQL database to 13.1. The Backup V1 tool can now be used in version 2021.3 for creating backups.
- Fixed an issue with Backup V2 restore where the restore process failed with the error `pg_probackup archive-get failed to deliver WAL file` after several consecutive backup and restore operations.
- The public API to retrieve queries failed if any queries in the result set had a schedule; this was due to a bug introduced in 2020.4 and has been now fixed.
- Previously, handling non-ASCII data caused failures in several features like Metadata Extraction, attaching files to Articles, Virtual Data source Upload, and some others. Now, Alation handles the Unicode characters by using UTF-8 as the default encoding/decoding scheme, which resolves this issue.
- Added extra filters to generate the `synonyms.txt` file from Lexicon expansion pairs to fix an issue with Alation search.
- For the Query Forms UI, fixed a number of UX issues:
 - Now displaying the dropdown with example values only for clauses containing at least one “e.g.” value
 - Improved the UI for entering filter values
 - Added tooltips next to the freeform text box
 - Introduced date-picker element for variable values of the date type
 - Added the **Clear Form** button to enable clearing variable values
 - Not displaying the filter results pane when there are no variable items in the form
- Miscellaneous other bug fixes to stabilize the release.

1.10.13 RELEASE 2021.3 - Limited Availability

build 9.2.0.143424

For the feature list, see [New Features](#)

Bugs Fixed in 2021.3 - LA

- Updated the type of the `field_id` property of the object `custom_field` from `string` to `integer` in the specification of the Integration API.
- Ampersand, “Greater Than”, and “Less Than” symbols are now rendered properly on Catalog pages for Domain objects in the **Title**, **Preview**, and other fields.
- Ampersand, “Greater Than”, and “Less Than” symbols are now rendered properly on Catalog pages for Glossary objects in the **Title**, **Preview**, and other fields.
- Previously, the ampersand symbol in data source names was not rendered correctly. This problem has been addressed.

- Previously, the Users and Groups combined filter did not pull correct results on the Alation Analytics V2 dashboard. This issue has been fixed to ensure that the Users and Groups combined filtering works as expected.
- Previously, an error and a warning occurred when an admin tried to whitelist a new 3rd-party IP address with the Alation Analytics V2 data source on an instance where Alation Analytics was installed and removed multiple times. This issue has been fixed.
- Fixed an issue where metrics for Query Search logged object id's as integers, which caused the searchclickextract ETL job of Alation Analytics V2 to fail. Now, they will be logged as object type names.
- Queries containing the DATE() function and the FORCE keyword will now be parsed for Snowflake data sources.
- Alation Connector Manager no longer requires that IPv6 should be enabled in the kernel configurations.
- Previously, if all users were removed from a group in the active directory, the corresponding LDAP group in Alation would not be cleared after synchronization, still displaying the previously synchronized LDAP group members. Now, external LDAP groups that are empty will be correctly synchronized with the Catalog, resulting in the corresponding group without members on the Alation side too.
- The **Groups** page used to take a long time to load as it was scanning through all existing groups and their members before loading in UI. Now, it loads group metadata and 10 members per group, which should significantly improve the speed of page rendering.
- Previously, when users exported search results that contained NoSQL attribute objects, they received an email notification with the error "Something went wrong while processing your search results export. Please contact your help desk for further assistance". This issue has been fixed and now Alation will successfully generate a search results export that includes NoSQL attributes.
- Fixed an issue with QLI from Presto with Hive and AWS Glue as Custom DB data source. Previously, there was an inconsistency in how QLI jobs pulled query history into Alation: initial QLI worked but consecutive QLI jobs omitted queries. After this fix, Alation should correctly process query history according to the QLI settings of the data source.
- Fixed an issue with metadata extraction from Google BigQuery data sources behind a proxy.
- Improved performance for large queries in Compose: page reloading, query editing response, selecting auto-suggestions, and scrolling now work faster.
- Fixed repeated calls to the `results.caching.enabled` endpoint of the Compose API that used to happen at the start of query execution and could significantly increase execution time for large queries. After the fix, query execution begins immediately after a user triggers the query run.
- The **Compare User Roles** dialog was not scrollable, and users who have smaller screens could not see this dialog in full. Added scrolling to ensure that the full content of the dialog can be viewed.
- The CSV export data was truncated in Excel on Windows due to the presence of the "equals" symbol "=" in front of the string, which caused the string to be treated as a formula; and Excel formulas have a limit of 255 characters. "=" is added by Alation when the "wrap text in quotes" option is enabled on the Alation server. This issue has now been fixed, and full values are transferred to Excel from Alation exports.
- Previously, the SQL content in cloned queries could not be updated or saved in the Catalog unless these queries were published. Now, updates to cloned queries can be saved whenever they are edited without having to publish them first.
- Data quality flag pop-overs will now be shown for Table, Attribute, and Schema objects in Compose.

- Previously, the 500 error appeared intermittently after the update to 2021.2 when users clicked the **Schedule** button in Compose. This issue has been fixed, and opening the Schedule dialog does not cause server errors.
- Alation users reported a problem that they could not copy the SQL content of a query from the bottom half of the Compose screen. This issue has been fixed: the **SQL** button on the bottom toolbar of Compose will now open a tooltip and users will be able to copy the SQL content of the query by copying the text in the tooltip.
- Previously, the Catalog became unresponsive after several consecutive runs of a query form that returned a large result set. After the fix, query forms will only fetch the last 25 batches to the frontend.
- Previously, for MicroStrategy sources, full extraction took place even though the **all projects except** option was selected in the Selective Extraction filter on the BI source Settings page. This issue has been fixed, and Alation now selectively extracts the Projects specified in the extraction filter.
- Fixed an issue with metadata extraction from an Oracle data source where it did not extract any objects as Alation tried to access `ALL_` objects instead of `DBA_` objects system tables. Extraction from Oracle will now use correct system tables based on the service account role.
- Added logic to handle duplicate queries with different sequence numbers in the query log table in order to address a query parsing issue with QLI from Amazon Redshift.
- Column aliases in SQL queries were not reflected in results export from Compose. Now, the export logic has been improved to validate the presence of an alias first before using the column name in the export.
- Fixed an issue where self-referential lineage was created when the created view and the source table had identical names. If the schema name is specified in the query, then Alation resolves the object to the specified schema only. If the schema name is not specified, then Alation resolves the object to the schema where the target view exists.
- After upgrading to 2021.2, duplicate accounts were created for usernames that were previously truncated to 30 characters. This issue has been fixed. When users with usernames that are longer than 30 characters log in to Alation after the update, their names will be complete, and truncation will no longer happen.

1.10.14 Known Issues

- The Search by Lexicon Abbreviations feature is turned off in 2021.3. Previously, this feature was enabled by default. It is temporarily disabled as Alation has discovered an issue with handling special characters in Lexicon tokens. The issue will be fixed in one of the coming patch releases.
- Sampling from an Oracle source added with the OCF connector for Oracle fails with the error `Protocol message CustomTableRequest has no "configuration" field` (**FIXED IN 2021.3.1**)
- The RabbitMQ container of Alation Analytics V2 may restart due to the following permissions issue: `docker logs rabbitmq mkdir: cannot create directory '/var/lib/rabbitmq/mnesia': Permission denied`.
- An Alation error message ("Sorry, there was an error in Alation") pops up when clicking on a BI connector on the Connector Manager Dashboard, although the connector status is *Running*. The pop-up can be closed by clicking **OK** at the bottom of the dialog.
- BI site, project, and folder objects that are not selected in the Selective Extraction filter are still extracted but empty. (**FIXED IN 2021.3.1**)
- When an MDE job exceeds 12 hours on an Amazon S3 file system source or Amazon Athena data source, the STS token does not get refreshed and MDE fails.

- Scheduled queries may fail when Alation is processing bulk lineage operations.
- A limitation with Google BigQuery Materialized Views: Query History, Popularity, and Lineage are currently not supported.
- Snowflake ALTER, DROP, and CLONE queries are not shown in Query History.
- Snowflake Joins and Filters for queries on Materialized Views are currently not shown on the corresponding tab on the Catalog page.
- Snowflake Materialized Views created using aggregate functions have known limitations for column extraction, Query History, and Lineage.
- Backup V1 restore fails after upgrading the internal PostgreSQL database to 13.1 (**FIXED IN GA**)
- The Policy object ingestion sets the Policy title to “Untitled” instead of the Policy name (**FIXED IN GA**)
- Saved Search fails with a 403 error if the the **alation_conf** parameter `alation.feature_flags.enable_permissions_middleware_feature` is set to True (**FIXED IN GA**)
- Excel Live Reports display “No results to display” even though the query returns results (**FIXED IN GA**)
- Encryption key rotation is required as part of updating Alation to 2021.3.0 and 2021.3.1. This requirement applies regardless of whether the instance had been upgraded from a previous installation or had an initial installation of 2021.1.x or above. This issue is specific to the GA build of 2021.3 (9.2.1.145216) and to version 2021.3.1 (9.2.2.146480) and is **FIXED IN 2021.3.2** (9.2.3.147624)

1.11 Release Notes 2021.2 (General Availability)

1.11.1 RELEASE 2021.2.14

build 8.2.15.160970

The following parameters were enabled on the internal Postgres database:

- `async_archiving`
- `wal_compress`
- multi-threaded backup

This change applies to the Backup V2 tool. It does not affect the steps of the backup process. Enabling these parameters should improve the backup performance and help avoid high I/O contention from the backup process.

1.11.2 RELEASE 2021.2.13

build 8.2.14.159995

This patch is a security update for the built-in PostgreSQL driver vulnerability described in *March 5, 2022 - JDBC Driver Security Advisory*. In this patch, the built-in PostgreSQL driver was upgraded to a newer version where the vulnerability CVE-2022-21724 is not present. This upgrade addresses the PostgreSQL driver vulnerability for PostgreSQL, Greenplum, and Tableau sources in the Alation Catalog.

1.11.3 RELEASE 2021.2.12

build 8.2.13.155510

This patch is a security update providing fixes to mitigate the Log4j 2 vulnerability discussed in the *December 22, 2021 - Log4j2 Security Advisory*. This addresses the security vulnerability described in [CVE-2021-45105](#).

This patch is cumulative and fixes all the security issues from the prior two patches, 2021.2.10 and 2021.2.11.

Upgrading to this security patch will apply the fixes automatically. This patch includes:

Security fix for the Hive connector

The Apache Hive connector is updated with a newer version of Apache Log4j 2. This addresses security vulnerability [CVE-2021-45105](#).

Security fix for Elasticsearch

Elasticsearch is updated with a newer version of Apache Log4j 2. This addresses security vulnerability [CVE-2021-45105](#).

1.11.4 RELEASE 2021.2.11

build 8.2.12.155192

This patch is a security update providing a fix to mitigate the Log4j 2 vulnerability discussed in *December 17, 2021 - Update to December 15 log4j2 CVE-2021-45046 Security Advisory*.

Security fix for the Hive connector

The Hive connector is updated to use a newer version of Apache Log4j 2, which addresses a possible DoS vulnerability described in [CVE-2021-45046](#).

1.11.5 RELEASE 2021.2.10

build 8.2.11.154693

This patch is a security update providing two important fixes to mitigate the Log4j 2 vulnerability discussed in *December 10, 2021 - log4j2 Security Advisory*.

Upgrading to this security patch will apply the fixes automatically. The patch includes:

Security fix for the Hive connector

The Apache Hive connector is updated with a newer version of Apache Log4j 2 as per NIST publication [NVD - CVE-2021-44228](#). In addition to updating the Log4j 2 version, all invocations of the Apache Hive connector will run with the `-Dlog4j2.formatMsgNoLookups=true` system property.

With this security update, the Alation Hive connector is not susceptible to remote code execution based on the Log4j 2 vulnerability [CVE-2021-44228](#).

Security fix for Elasticsearch

As per recommendation from Elastic found at [Apache Log4j2 Remote Code Execution \(RCE\) Vulnerability - CVE-2021-44228 - ESA-2021-31](#), this fix disables the message lookup substitution for Elasticsearch on the Alation server. With this security update, the **alation_conf** parameter `elasticsearch.env.es_java_opts`, which sets Java options for Log4j 2, will be set to `-Dlog4j2.formatMsgNoLookups=true` by default. This change ensures that the Log4j 2 Java options have the recommended setting and that the Elasticsearch component of the Alation server is not susceptible to remote code execution based on the Log4j 2 vulnerability CVE-2021-44228.

Additional fix

- Fixed an issue with the [Table Metadata Sync API](#) where using it resulted in the 500 error. The API now works as expected.

1.11.6 RELEASE 2021.2.9

build 8.2.10.151952

- Fixed an issue with QLI where previously ingested Query History information was not restored as expected after resetting KV Store during a troubleshooting session. Now, the Query Template text for the Query History section will be populated from the internal Postgres database if KV Store does not have the required Template text.
- Previously, when there was a large Search result set, the left-side filters were loaded in an unexpected order, creating an error state when clicked. This fix adds logic for correctly handling filters returned in any order.

1.11.7 RELEASE 2021.2.8

build 8.2.9.150274

- Fixed an issue where upgrade from version 2021.1 (7.2.11) to 2021.2 failed due to an error during the Django migration stage. Now, the upgrade process is expected to complete successfully.
- Fixed a UI bug where on instances with no Domains, the **All Domains** dropdown was still visible on the Search page for a short time while the search results were loading. Now, the **All Domains** list is only displayed if Domain objects exist in the Catalog.
- The Search by Lexicon Abbreviations feature is now turned off by default. It is temporarily disabled as Alation has discovered an issue with handling special characters in Lexicon tokens.
- For a specific customer case, added a new **alation_conf** flag to skip Postgres validation during the backup process.
- Fixed an issue where user profiles were not loading on the **Advanced Search** page after searching for users. Now, when a user is selected in the Advanced Search results pane, the corresponding user profile will load in the right-hand view pane.
- Fixed an issue where the Alation updater validated if encryption key rotation has been performed on versions 2021.1.x and newer. Encryption key rotation is no longer required before updating Alation to 2021.2.x and 2021.3.x.

1.11.8 RELEASE 2021.2.7

build 8.2.8.145756

- Fixed an issue with updating the Alation search index after a large number of objects have changed in the Catalog during metadata extraction. Previously, after metadata extraction removed a large number of objects from the Catalog, the search reindexing queue could become backed up while processing the delete messages. This could stall the reindexing process for a long period of time. This fix improves the performance of the search index queue and decreases the time of search reindexing after a large number of objects are removed by metadata extraction.

1.11.9 RELEASE 2021.2.6

build 8.2.7.145401

- Changed backup files permissions. Alation backup files are now readable by user `alation` only. Change affects both backup tools, Backup V1 and V2.
- Fixed an issue with exporting search results where export failed to be created if results included the Object Type filter *NoSQL Attributes*.
- Previously, there was an issue with logging of the metric *search.click* that captures the usage of the Full-Page search. This fix corrects the logging problem for the *search.click* metric in the *metrics_metric* table of the internal database.
- Previously, an error and a warning occurred when an admin tried to whitelist a 3rd-party IP address with the Alation Analytics V2 data source on an instance where the Alation Analytics data source was installed and removed multiple times. This issue has now been fixed.
- Fixed an issue with Backup V2 restore where the restore process failed with the error `pg_probackup archive-get failed to deliver WAL file` after several consecutive backup and restore operations.
- Previously, Databricks queries that had self-referencing CTE caused a `java.lang.StackOverflowError` error when they were parsed during QLI. The internal query parsing logic of the Alation server has now been enhanced to parse self-referencing CTE queries for Databricks.
- For Legacy Hive users, added a feature flag that allows specifying from where the query string is read during QLI. Enabling `alation.feature_flags.enable_usage_of_map_reduce_name_for_query_location` will cause Hive V1 QLI to read `mapreduce.workflow.name` for the query string. By default, Alation reads `hive.query.string` to get the query string. Note that Alation recommends using the Hive V2 framework for Hive sources.

1.11.10 RELEASE 2021.2.5

build 8.2.6.143824

- Removed a vulnerability that could potentially allow executing script tags after uploading files to the Alation application.
- Fixed an issue with QLI from Presto with Hive and AWS Glue as Custom DB data source. Previously, there was an inconsistency in how QLI jobs pulled query history into Alation: initial QLI worked but consecutive QLI jobs omitted queries. After this fix, Alation should correctly process query history according to the QLI settings of the data source.
- Fixed an issue with duplicate usernames: after 30-character limitation on usernames was removed in 2021.2, duplicate accounts were created upon login for users with usernames longer than 30 characters.

To cleanup the duplicate accounts that were already created, run the following script from the Alation shell:

```
cd /opt/alation/django
sudo /opt/alation/env/bin/python -m rosemeta.one_off_scripts.rename_saml_duplicate_users
```

1.11.11 RELEASE 2021.2.4

build 8.2.5.141629

- Previously while performing metadata extraction or any other action on a connector and the connector is unresponsive, the connection will be terminated. Now, the connection timeout time has been increased.
- Previously the CSV exports in Compose were having few blank rows in between the rows. The CSV exports now should not have any blank rows in between the rows.
- Fixed an issue where the metrics in query search logged the Object ID (oid) as integers which fail the searchclickextract job. Now, the metrics in query search logged oid as otype for each search result click.

1.11.12 RELEASE 2021.2.3

build 8.2.4.140617

- Added server-side XSS protection for custom rich text fields within Articles. HTML elements are wiped and script tags are escaped. This change should prevent any malicious 3rd-party API requests. Custom rich text fields in Articles can be used as before, with no changes to user experience.
- For the Hive connector, fixed an issue with metadata extraction when it failed if the column data type value exceeded 100,000. Now, MDE supports struct schema definitions which have more than 100,000 characters.
- Fixed an issue with QLI from Presto on Hive when it failed with “UnicodeDecodeError” if column data had Unicode characters. Now, Alation defaults to UTF-8 encoding when parsing query log files.
- Fixed an issue with Lineage V2 migration where it failed due to database field limit breach caused by a large volume of Lineage V1 data to be migrated.
- Fixed an issue with MDE from a Greenplum data source when extraction failed with the error message *a bytes-like object is required, not 'str'*.

1.11.13 RELEASE 2021.2.2

build 8.2.3.140018

- Previously, when **Use Group to assign roles** was enabled, new users with the Viewer role did not appear in the Viewer built-in group. Those users received the “Permission Denied” error on numerous pages if the Viewer role was enforced on the instance. Now all new users are assigned to the correct built-in groups that correspond to their roles.
- Query execution and export failed when multibyte Unicode characters were part of the query result retrieved from a data source. This issue has been fixed by using a different decoder mechanism.
- Previously, metadata extraction from SQL Server sources failed to extract tables from a database when the database collation type was different from the collation type on the SQL Server level. This issue has been fixed.

- Fixed an issue with metadata extraction from Azure Data Warehouse where it failed with the 0 schemas found error. This was resolved by adding support for case sensitive collation.
- Now, when no Domains have been created on an instance, the **Domains** custom field will not be displayed on the Properties panel of Catalog pages.
- The **Explore Domain** button on the top right of the Domain object Catalog page has been replaced with the **Search Domain** button that is located on the top right of the **Browse This Domain** table.

1.11.14 RELEASE 2021.2.1

build 8.2.2.138904

- Fixed an issue when Article Template information was not displayed in the **Template** column of the Table view for Search results.
- Previously, when users navigated to the Full-Page Search from another page in Alation, the Back button of the browser did not function correctly. This issue has been fixed by correctly populating browser history.
- Previously, long labels in the Search filter overflowed the left sidebar. This fix introduces truncation and text wrapping to resolve the issue.
- Fixed a bug in processing incoming emails from the Reply-by-Email feature when Content-Type was application/ms-tnef. Added support for processing emails with Content-Type application/ms-tnef generated from Office 365 and some Outlook clients.

1.11.15 RELEASE 2021.2 - General Availability

build 8.2.1.137867

Note: To familiarize yourself with the new features in 2021.2, watch the [What's New in 2021.2](#) course in Alation University (*this link will require you log in to Alation University*)

Alation Cloud Service

Alation Cloud Service has become a reality! It is a fully managed Alation cluster for customers looking to minimize their infrastructure and maintenance efforts. Customers also benefit from faster time to value as their instance is regularly patched and upgraded with quarterly releases.

There is no difference between Alation Cloud Service and Alation on-premise or VPC deployment. It offers the same features: the newest Alation release and 100% of Alation functionality, including support for connectors, but without the need to host and manage the infrastructure and updates.

Alation uses Amazon Web Services (AWS) where Alation Cloud Service can be launched in different AWS regions. Alation also offers the services of migrating an on-prem deployment to Alation Cloud.

More information can be found on the Alation company site: [Alation Cloud Service](#)

New Features

Search and Discovery

More Catalog metadata fields have become searchable

More catalog metadata fields are now processed by Search:

- **Source Comments** field
- Descriptions of the Trust flags
- Custom SQL statements for datasource Connections under Tableau BI sources
- **BI Owner** field from the properties of BI objects (a new facet in the Filters pane)
- **Partition Keys** for table objects (a new facet in the Filters pane)

Search results are extended with confirmed Lexicon terms

Lexicon is an algorithm that generates comprehensive titles for Catalog objects using expansions for each abbreviated word that it discovers in catalog object names. When an *abbreviation-expansion pair* is confirmed by a human user (a Catalog or Server Admin), this pair is assigned the highest level of confidence and is given preference when Lexicon generates auto-titles. Alation Search will now consider such **confirmed** abbreviation-expansion pairs as search keyword synonyms. When a user enters a search word that is an abbreviation with a confirmed expansion in Lexicon, the search results will return objects that match both the abbreviation the user entered and its confirmed expansion from the Lexicon. For example, if a user enters an abbreviated keyword *dgt* which has a confirmed expansion *digit* in Lexicon, Search will return objects that contain *dgt*, *digit*, or both the words in searchable metadata.

Improved Search Ranking

The Endorse and Deprecate Trust flags now participate in search ranking of an object. *Endorsing* a Catalog object by adding the **Endorse** Trust flag will result in boosting its search rank. An endorsed Catalog object will be promoted to the top of search results enabling catalog users to find trusted data more easily. However, *deprecating* an object by adding the **Deprecate** flag will penalize its Search ranking and move this object down in the search results list. The extent of boosting or penalizing of the Search rank of flagged objects is controlled with the `alation_conf` parameters `alation.search.flag_boost_ranking_factor` and `alation.search.flag_penalty_ranking_factor`. The default values of these parameters can be customized by an Alation admin with backend access to the Alation host.

Search Table View

Search results can now be viewed in the tabular format: users can switch between the List and Table views on the full-page Search results page. The Table view displays results in a structured way with ability to sort, filter, and re-order the content of the table. The Table view includes descriptive metadata for each result, such as object type, Trust flags, path to the object in the Catalog, Author, query SQL, and other information. Users can adjust the view by adding or removing table columns. It is also possible to collapse the Filters pane on the left in order to expand the search results area.

Domains

- Domains are now enabled by default. The feature switch for Domains has been removed from **Admin Settings > Feature Configuration**.
- Object type coverage for Domains has been extended to:
 - Glossaries
 - Dataflow objects
 - File Systems
 - NoSQL objects
 - API Resources

These objects can now be added to Domains. Note that the child objects of these object types are not automatically added to Domains and need to be added separately.

Improved UI of the Query Form Page

The Query Form page gets a refresh, which includes the following changes:

- Users can now select a connection and enter their credentials to run the query form right on the query form page without leaving the page to go to Compose for authentication
- Users can see the **Description** field, the **Properties** sidebar, and the *Endorse*, *Warn*, or *Deprecate* Trust flags of the corresponding Query object on the Query Form page.

Governance and Curation

File Systems Data in Alation Analytics V2

File System data has been added to the Alation Analytics schema and ETL pipeline, allowing users to perform analytics on the File System objects in the Catalog. 3 new tables were added to provide valuable insights into how users interact with the File Systems cataloged in Alation:

- filesystems
- directories
- files

Ability to Embed 3rd-Party Visualizations

The Alation Analytics V2 dashboard can now display custom visualizations from 3rd-party BI tools, such as Tableau, on the **Custom** tab. A new **Custom Visualizations** section was added to the Alation Analytics V2 Settings page where an admin can configure custom views for the dashboard. The **Custom** tab will appear on the Alation Analytics V2 dashboard when at least one custom visualization is configured in the Settings.

Improved Alation Analytics V2 Experience

- A new visualization **Top Search Terms** on the Alation Analytics dashboard
- A new filter by **Group** on the Alation Analytics dashboard
- Alation Analytics V2 ETL time cut-off has been optimized. Previously, the daily ETL started at 3:59 AM (default schedule) and loaded the delta of the data between 1:00 AM of the previous day and 1:00 AM of the present day. Now, the ETL process will extract data up to the ETL start time. This makes the Alation Analytics V2 data as close to the current state of the Catalog as possible; and analysts do not need to wait an additional day for the data to flow into the Analytics database.
- Information about deleted RDBMS objects is now correctly extracted into Alation Analytics V2. The *deleted* attribute on a Schema, Table, or Column object now inherits the value from the corresponding attribute of the parent data source. If a data source is deleted, then all its child schemas, tables, and columns will be marked as *deleted* in the Alation Analytics V2 data.
- Admins can start the Alation Analytics V2 ETL process from the Alation UI: the Alation Analytics V2 Settings page now has the **Execute ETL** button that triggers ETL on demand.
- In 2020.4, 3 new columns `thread_resolved_at`, `thread_resolved_by`, and `thread_status` were added to the conversation table to keep track of the status of Conversations. However, these columns were not updated after the ETL for Conversations that had been resolved before this change to the Alation Analytics V2 model. Now, the columns will be updated during ETL for all conversations without any intervention on the admins' part.

Connect Platform

Lineage V2 Feature Flag Is Enabled by Default

Lineage V2 is now enabled by default on new installations of Alation. On existing installations, the Lineage V2 feature flag will become enabled after the update to 2021.2 unless an admin intervenes before performing the update. It is highly recommended to migrate your Lineage data to the Lineage V2 framework if you have not done so yet: </installconfig/LineageV3/EnableLineageV2>

Open APIs and Swagger UI Are Enabled by Default

From 2021.2, the integration with Swagger is enabled by default. API documentation in the OAS 3.0 format becomes available on Alation instances at **`http(s)://<your_Alation_URL>/openapi/<openapi_endpoint>`** without the need to set the corresponding feature flag.

RDBMS Integration Open API

Added the **Integration API** that brings together physical and logical catalog metadata for the Schema, Table, and Column objects, allowing users to retrieve and create these objects in Virtual or non-virtual RDBMS data sources in Alation. This API can be used by Source, Catalog, and Server Admins who are also assigned as *Data Source Admins* to the data source they are working with.

Using the RDBMS Integration API, you can:

- Get a list of schemas, tables, and columns and their custom fields from virtual and non-virtual data sources;

- Create physical (schemas, tables, columns) and logical metadata (custom field values) in non-virtual data sources.
- Upload new and update existing physical (schemas, tables, columns) and logical (custom field values) metadata into a Virtual Data Source at the same time;
- Get a list of schemas, tables, and columns characterized by specific custom fields.

The RDBMS Integration API documentation is available in the OAS 3.0 format at **[http\(s\)://<your_Alation_URL>/openapi/integration_apis/](http(s)://<your_Alation_URL>/openapi/integration_apis/)**.

Support for Recovery of Curated Data for BI Sources

Enhanced GBM V2 and OCF frameworks to support “Soft Delete” from the Alation Catalog in order to protect curated data from permanent deletion. Soft Delete is enabled by default. The checkbox **Disable Hard Sync** has been removed from the Settings page of BI Sources on GBM and OCF.

New Connector Certifications and Driver Upgrades

- **MarkLogic** is now in the list of supported data sources
- Certified newer versions of supported data sources, such as SQL Server 2019 on RDS, Salesforce DB Winter 2021 edition, EMR 6.2 with Spark SQL, Databricks 8.0, and PostgreSQL 13 on RDS.
- Updated a number of built-in drivers to newer versions:
 - **Google BigQuery**: Simba 1.2.12.1015
 - **Amazon Redshift**: JDBC42-1.2.43.1067
 - **Snowflake**: 3.12.16

Open Connector Framework

- Alation 2021.2 comes with a pre-packaged Docker installer allowing to install Docker on any OS supported by Alation without the need to additionally configure installation repositories or download Docker packages from available sources.
- The OCF connector zip packages are now included into the Alation backup.
- Added a new Ypireti command that can be run to reinstall OCF connectors after a backup restore or a failover to Secondary on the HA Pair.

MicroStrategy GBM V2 Extraction Stabilization

The MicroStrategy connector for GBM V2 received multiple stabilization and improvement fixes:

- Fixed an issue where only the first project from the list of projects selected for extraction was extracted into the Catalog. Now, all selected projects should be extracted as expected
- Fixed a number of NPEs reported during metadata extraction
- All document, dossier, and dashboard previews are now extracted using the MicroStrategy REST APIs
- Alation now determines if a cube is published or if it is not using isPublished API during extraction
- Made changes to populate sub-reports of a document using search-based API. This works for both published and unpublished cubes

- Cube data sources are now linked with reports and documents using search API, which enables Alation to calculate lineage
- All data source cubes in the project are extracted using the search API
- Simple reports are extracted using the NoAction execution flag. This opens the reports in the design mode before extraction and works for both published and unpublished cubes
- Alation now retrieves object prompts using the search API and extracts the corresponding field metadata
- Alation can now extract the document and lineage information irrespective of whether the underlying cube is published or not. Fields will not be extracted if the document depends on an unpublished cube
- Report metadata is extracted asynchronously. If a large number of reports are extracted, Alation creates multiple connections to extract report data. The MicroStrategy server has different types of governing job limits set to restrict the number of connections per user account, user session, and per project. Alation will respect the limits set on the MicroStrategy server and only use the server resources allowed to be used for extraction
- Preview generation will now be skipped if a report or document is dependent on an unpublished cube.
- Fixed multiple issues with prompt validation and extraction
- Added a fix for the VALUE prompt to answer all prompts and only fail for required prompts that do not have any answers
- Alation will now extract only valid metric types. The auto-generated field row count that used to cause the MicroStrategy server crash is no longer extracted

Enterprise Deployment

Backup V2 Becomes Default

Backup V2 becomes the default backup tool. The switch-over to V2 happens automatically when the Alation application is updated to 2021.2. Backup V2 will also be the default backup tool on new installations of Alation from 2021.2 and onward.

Improved Group Profile Page

The Group Profile page received a UI refresh and a significant boost in performance.

Support for Microsoft Modern Authentication for the Alation Email Server

It is now possible to configure inbound email interactions for Microsoft-based email servers on the **Admin Settings > Email Server** tab by leveraging *Microsoft Modern Authentication*. The 2021.2 release introduces this configuration capability for the *Reply By Email* feature of Conversations.

Note: Modern Authentication is the term Microsoft uses to refer to their implementation of the OAuth 2.0 authorization framework for client/server authentication. Alation uses MS Graph APIs for accessing application mailboxes and OAuth 2.0 endpoints for authentication.

Data Source Admins Can Add Sources

Users with the **Source Admin** role now gain the ability to add new sources to the Catalog. This feature needs to be enabled using the parameter **alation_conf alation.roles.allow_source_and_catalog_admins_to_create_ds** in **alation_conf**.

Role Selector for AWS Data Sources

For AWS data sources configured to use SSO in Compose, users can now select an AWS role from a list if multiple roles are returned for this user by the IdP.

Ability to Disable Offline AuthService Connections

An Alation admin can enable or disable offline connections (used for scheduling of queries and executing Excel Live reports) for connections that use the AuthService component of Alation, for example, SSO-based connections for AWS data sources. When offline connections are disabled (default for SSO-based connections to AWS data sources), the Compose UI will prevent users from scheduling a query or creating an Excel live report. Execution of a scheduled query or an Excel Live report that use such a connection will not succeed.

Improvements

- The **limitation of 30 characters maximum for usernames in Alation has been removed**. Usernames up to 150 characters are now allowed by default. Usernames that were previously truncated will be migrated to their full form when the corresponding users sign in to Alation after the update to 2021.2.
- Previously, the SQL query content had a limitation of 10,000 characters per line. This resulted in truncation of some of the query content in the **Dataflow Content** field of Dataflow objects, making lineage information incomplete. Now, the limitation has been removed and the full SQL query content of an ingested SQL query will be displayed.
- Data Source and Schema names were added as **Source** and **Schema** columns to CSV exports of Lineage Impact Analysis reports.
- Suspended users will no longer appear in the **Top Users** field of data sources.
- Internal Java component was updated from version 11.0.7 to 11.0.9.
- In case of SAML authentication, Alation will no longer allow replaying a SAML response when the parameter **alation.authentication.saml.allow_unsolicited** is set to **True**. Alation will deny access when a consumer assertion request contains a SAML Response with an assertion ID that has already been authenticated.
- In case of SAML authentication, Alation admins can now choose to use the **nameID** attribute instead of the **uid** attribute as the username in Alation. This feature needs to be enabled in **alation_conf** by setting the parameter **alation.authentication.saml.use_name_id_as_username** to **True**.
- The data type length will now be displayed for SAP HANA data sources.
- Improved handling of the tasks queue so that Search index updates receive priority over bulk-profiling tasks. This will allow changes to the Catalog objects that happened after the bulk-profiling tasks were queued to be reflected in Search results sooner and before bulk-profiling tasks are completed.
- Added a new a property **alation.qli.fs.defaultFS** to the **aha-conf.xml** file to explicitly set the file system to be used as the **fs.defaultFS** for QLI. This property is used when there are multiple active *Nameservices*.

Bug Fixes

Note: 2021.2 GA also includes all bug fixes listed for the 2021.2 LA version.

- Fixed the issue of the system health alerts not sent to admins. Now works as expected.
- Fixed a problem where new schedules for existing queries could not be created after the Alation update. This fix also updates the partially created schedules so that they apply to the queries as intended.
- Fixed a bug when information about Connections based on custom SQL queries was missing from the Tableau report pages after extraction from Tableau BI sources.
- Fixed an issue of authorization redirects not appearing for default connections and data source users created before the update to release 2020.3.x. Resolved by updating connections on blank `rosemeta_userpreferredconnection` rows and by removing duplicate rows in the `rosemeta_userpreferredconnection` table.
- Fixed the issue of duplicate result tables on the Query Form page that occurred when running the query form.
- Fixed the slowness of the query forms page loading in case of large result sets.
- Fixed an issue with permission mirroring not working for SSRS sources. Permission mirroring can now be enabled or disabled as expected.
- Fixed the issue of Postgres auto-scan failing with the “permission denied” error by moving the Postgres scan logs from `/var/log/scan-postgres.log` to `/opt/alation/site/logs/scan-postgres.log`. Auto-scan now works as expected.

1.11.16 RELEASE 2021.2 - Limited Availability

build 8.2.0.136682

For new features in 2021.2, see [New Features](#).

Bug Fixes

- Previously, in the Edit mode for Articles, the rich text editor re-formatted ordered lists with lowercase letter list markers to numbered lists. This issue was fixed, and the editor now saves ordered lists in the required format.
- For the user profile upload from a CSV file, changed the error message for invalid characters in column names to be clear and user-friendly.
- Paginated tables on the **Admin Settings > Users** tab displayed the same user list on the first page and other pages when the page length was set to be below 100 rows. This issue has been fixed, and table pagination is applied correctly for all table page lengths.
- Fixed a bug where the values of a multi-picker field type were incorrectly sorted in tables in the Alation UI. Now, table sorting applies to multi-picker fields as expected.
- For Alation Connector Manager, extended the dependency for the installer for Debian-based systems to both `docker.io` and `docker-ce`. This allows to install either of the packages on Debian-based operating systems.

- Fixed an issue where the dialog that confirms addition of child objects to a Domain was displayed for objects without child objects. Now, this dialog will not be displayed if an object does not have child objects.
- In a new instance of Alation or in an existing instance where the QLI archiving job had never been run, an attempt to schedule the QLI archiving job caused the Alation server to crash. This issue has been fixed, and QLI archiving works as expected.
- For Presto data sources, column names enclosed within single quotes are obfuscated and column names within double quotes are not.
- The **History** section of the **Queries** tab of the Catalog page for Table objects used to display queries that were not executed. Clicking on the statement template for such a query caused the 404 error in the Alation UI as execution details were unavailable. Now, queries that have not been executed are filtered out from the History section.
- Fixed an issue with query-based QLI when queries containing strings `False AS cancelled`, `' ' as defaultDatabase` or `null as defaultDatabase` were not parsed correctly. After the fix, such queries are successfully parsed.
- When LDAP was selected as the authentication method for Alation, sometimes user authentication failed even though the user credentials are valid. The issue was identified and fixed so that authentication with valid credentials is always successful.
- Previously, when a Custom Group name was updated on the Group Profile page, the action did not update the Search index and the old group name was still displayed in the Search results. Now, updating a Custom Group name causes the Search index to update and the updated group name is displayed in Search results as expected.
- Previously, if a user assigned as a *Data Source Admin* did not have the *Server Admin* role, they could not see some sections of the data source Settings page, for example **Network Connection** and **Service Account**. This issue has been fixed and all Data Source Admins can see full information on the Settings page of the data sources they are admins of.
- Previously, metadata extraction from MySQL data sources did not extract table source comments. This issue has been fixed, and Table comments are now extracted as **Source Comments** on the Table catalog page.
- Fixed a bug where the Data Sources tab on the **Sources > Manage Sources** page displayed incorrect statuses for Metadata Extraction, Profiling and Query Log Ingestion.
- Now, Server Admins can add query parameters to the **Contact your Help Desk** link when they configure it using **alation_conf** and they will be processed correctly. Previously, including special symbols, such as an ampersand, resulted in redirecting to an incorrect page.
- Fixed a bug when the exporting of Search results failed if the *Queries* option was selected in the object type filter on the Search page.
- Server Admins can now configure the timeout period for the Advanced Search page using **alation_conf** and the parameters **alation.search.timeout** and **alation.search.max_retries**. These parameters can be modified if Search encounters timeout issues.
- Fixed a performance issue with the **Groups** page. The Groups page loading time has been significantly improved.
- For Alation Analytics V2, fixed an issue where ETL could not load data if a custom field created in Alation had the name `id`, which is the name of the primary key column in the Alation Analytics V2 tables. During ETL, fields named `id` are now renamed as `id_[field_id]` in the Alation Analytics V2 database, for example `id_10007`. This will affect all objects that have this custom field associated with them as the field will appear as `id_[field_id]` in the corresponding object table.

- Fixed an issue during Query Log Ingestion from Hive data sources that was due to the presence of multiple AM_STARTED attribute values in query logs. The last attribute in the log will now be used in ingestion.
- Fixed a bug with Query Log Ingestion from SQL Server 16 with NTLM authentication which caused the Lineage information to be incomplete for View DDL statements containing the COALESCE keyword.
- Fixed an issue with high CPU usage by Taskserver during extraction, QLI, profiling, search index update, and other processes. The issue was caused by Hive grammar projection grammar rule: antlr went into an infinite loop while parsing a query. As a solution, the antlr version and the timeout for execution of parser threads were updated.
- Fixed an issue with Query Log Ingestion for Custom DB where it stalled with the get() returned more than one AlationVersion error. Added a unique constraint to the tree_hash field of AlationVersion. Now duplicate entries are deleted if any exist and new duplicate entries will not be created.
- Fixed an extraction bug for Tableau BI sources where some datasource columns were not extracted into Alation.
- Fixed a bug with DB2 data sources where primary key information was not displayed in the Alation UI after metadata extraction.
- Fixed an issue in Compose where the *Loading...* message was displayed for an excessively long time in the **Results** section after a query was cancelled. Now, when a query is cancelled by a user clicking the **Cancel** button, the query results will either be displayed immediately if Alation managed to retrieve some results before cancellation or no results will be displayed if the cancelled execution has returned no results.
- Fixed a UI bug with Compose results column headings. Now, the headings will be properly aligned in supported versions of Firefox, Chrome, and Edge. The fix does not apply to Internet Explorer.
- Fixed an issue when the uploaded SQL statement was not displayed correctly in UI when a SQL query was uploaded to a data source Catalog page.
- The SAML request generated to authenticate a Compose user while performing SSO for AWS sources was using an invalid format for the IssueInstant attribute. This resulted in an error in Compose when users tried to authenticate using SSO. The issue was fixed by introducing the ISO format based on the SAML 2.0 specification for this attribute.
- Fixed a bug which prevented Alation Analytics V2 to be installed when the server FQDN was specified during the installation. Now, admin can specify either the host IP or FQDN and successfully install Alation Analytics V2.
- Resolved an issue with executing query forms whose SQL contains single-line comments (starting with --). Previously, if the comment line was at the end of the query, Alation attempted to execute the comment as a separate statement. If the comment appeared immediately above a statement, Alation issued an error regarding the statement being closed. This is no longer the case as Alation correctly handles comments in SQL statements in query forms.
- Miscellaneous other bug fixes to improve application quality.

1.11.17 Known Issues

- If the initial installation of Alation was done on release 2020.4.x or 2021.1.x, then the update to 2021.2 will include a full reindexing of the internal application database tables. Reindexing may increase the time required for the update on instances with large internal databases. **(Fixed in 2021.3)**
- Unable to create or update query schedules for new queries or for queries without a schedule prior to upgrade to 2021.1. Still present in 2021.2 LA. Query objects with a schedule created in a prior release are not affected; in these cases, schedules can be updated successfully. **(FIXED IN GA)**
- Admin Health alerts are not sent. **(FIXED IN GA)**
- Update to 2021.2 will fail if there is a password set on the internal Postgres database. In order to upgrade successfully, the Postgres password must be cleared before the update and set back after the update is completed.
- Query form Catalog page displays 5 duplicate tables upon clicking **Get Results**. **(FIXED IN GA)**
- During the Re-key job run, if the re-key operation fails for a single row in data_storage_blobaccesspostgres, the resultset data is updated with empty values instead of throwing an exception, which causes the encryption key rotation process to fail.
- In Compose, authorization redirects do not appear for the default connection and data source users created in Alation before version 2020.3. **(FIXED IN GA)**
- Automatic Postgres scan fails with a “permission denied” error. **(FIXED IN GA)**
- For MicroStrategy BI sources, full extraction takes place even though **all projects except** is selected during selective extraction. **(Fixed in 2021.3)**
- In Compose, **Run and Export** returns different file sizes on two successive attempts. **(Fixed in 2021.3)**
- Permission mirroring cannot be disabled for SSRS BI sources. **(FIXED IN GA)**
- MDE fails with the “out of memory” exception on large datasets (~ 3,000 tables per schema, 100 columns each) when using the Oracle native connector.
- Query form page stops loading when a large result set is returned, which causes Alation UI to freeze. **(FIXED IN GA)**
- The **Execute ETL** button on the Alation Analytics V2 Settings page flashes if Alation Analytics needs to be updated.
- The Alation Analytics V2 data source page shows an incorrect number of extracted tables in the schema Public.
- When both **User** and **Group** filters are applied to the Alation Analytics V2 dashboard at the same time, the filtered result is incorrect across all dashboard visualizations. **(Fixed in 2021.3)**
- When restarting the Alation server, you may see a warning message `python3: can't open file '/opt/alation/event_bus/kafka_topic_manager.py': [Errno 2] No such file or directory`. Please ignore this warning.
- Alation does not load the Active Directory user picture and returns a 500 error.
- Scheduling a discarded query in 2021.1 results in leaving anomalous records in `django_celery_beat_periodictask`.
- **Remove Projects that are not captured by the list above checkbox** appears on the Settings page of BI sources on GBM V1 and appears disabled. This checkbox should not be present for sources added on GBM V1.
- Metadata extraction from Hive data sources fails MDE if the column type length is more than 100,000. **(FIXED in 2021.2.3)**

- Snowflake DATE() function and FORCE keyword are not supported by the current GSP version. (**Fixed in 2021.3**)
- MSTR 2021: Fields are not displayed in catalog if Dossiers contain Shapes
- Extraction from an Azure Data Warehouse data source may issue the error “Found zero schemas” (**Fixed in 2021.2 - patch 8.2.3**)
- The 500 error is displayed when opening the query scheduling dialog in 2021.2. (**Fixed in 2021.3**)
- After upgrade to 2021.1.x, users are facing a decode error when running queries in Compose. (**FIXED in 2021.2.2**)

1.12 Release Notes 2021.1 (General Availability)

1.12.1 RELEASE 2021.1.13

build 7.2.14.155515

This patch is a security update providing fixes to mitigate the Log4j 2 vulnerability discussed in *December 22, 2021 - Log4j2 Security Advisory*. This addresses the security vulnerability described in [CVE-2021-45105](#).

This patch is cumulative and fixes all the security issues from the prior two patches, 2021.1.11 and 2021.1.12.

Upgrading to this security patch will apply the fixes automatically. This patch includes:

Security fix for the Hive connector

The Apache Hive connector is updated with a newer version of Apache Log4j 2. This addresses security vulnerability [CVE-2021-45105](#).

Security fix for Elasticsearch

Elasticsearch is updated with a newer version of Apache Log4j 2. This addresses security vulnerability [CVE-2021-45105](#).

1.12.2 RELEASE 2021.1.12

build 7.2.13.155179

This patch is a security update providing a fix to mitigate the Log4j 2 vulnerability discussed in *December 17, 2021 - Update to December 15 log4j2 CVE-2021-45046 Security Advisory*.

Security fix for the Hive connector

The Hive connector is updated to use a newer version of Apache Log4j 2, which addresses a possible DoS vulnerability described in [CVE-2021-45046](#).

1.12.3 RELEASE 2021.1.11

build 7.2.12.154834

This patch is a security update providing two important fixes to mitigate the Log4j 2 vulnerability discussed in [December 10, 2021 - log4j2 Security Advisory](#).

Upgrading to this security patch will apply the fixes automatically. The patch includes:

Security fix for the Hive connector

The Apache Hive connector is updated with a newer version of Apache Log4j 2 as per NIST publication [NVD - CVE-2021-44228](#). In addition to updating the Log4j 2 version, all invocations of the Apache Hive connector will run with the `-Dlog4j2.formatMsgNoLookups=true` system property.

With this security update, the Alation Hive connector is not susceptible to remote code execution based on the Log4j 2 vulnerability CVE-2021-44228.

Security fix for Elasticsearch

As per recommendation from Elastic found at [Apache Log4j2 Remote Code Execution \(RCE\) Vulnerability - CVE-2021-44228 - ESA-2021-31](#), this fix disables the message lookup substitution for Elasticsearch on the Alation server. With this security update, the `alation_conf` parameter `elasticsearch.env.es_java_opts`, which sets Java options for Log4j 2, will be set to `-Dlog4j2.formatMsgNoLookups=true` by default. This change ensures that the Log4j 2 Java options have the recommended setting and that the Elasticsearch component of the Alation server is not susceptible to remote code execution based on the Log4j 2 vulnerability CVE-2021-44228.

1.12.4 RELEASE 2021.1.10

build 7.2.11.146829

- Fixed an issue where user profiles were not loading on the **Advanced Search** page after searching for users. Now, when a user is selected in the Advanced Search results pane, the corresponding user profile will load in the right-hand view pane.
- Fixed an issue with Backup V2 restore where the restore process failed with the error `pg_probackup archive-get failed to deliver WAL file` after several consecutive backup and restore operations.

1.12.5 RELEASE 2021.1.9

build 7.2.10.144381

- Previously, there was an issue with logging of the metric `search.click` that captures the usage of the Full-Page search. This fix corrects the logging problem for `search.click` in the `metrics_metric` table of the internal database.
- For Legacy Hive users, added a feature flag that allows specifying from where the query string is read during QLI. Enabling `alation.feature_flags.enable_usage_of_map_reduce_name_for_query_location` will cause Hive V1 QLI to read `mapreduce.workflow.name` for the query string. By default, Alation reads `hive.query.string` to get the query string. Note that Alation recommends using the Hive V2 framework for Hive sources.

1.12.6 RELEASE 2021.1.8

build 7.2.9.142382

- Fixed an issue in Compose where the *Loading...* message was displayed for an excessively long time in the **Results** section after a query was cancelled. Now, when a query is cancelled by a user clicking the **Cancel** button, the query results will either be displayed immediately if Alation managed to retrieve some results before cancellation or no results will be displayed if the cancelled execution has returned no results.
- Fixed the issue of duplicate result tables on the Query Form page that occurred when running the query form.
- Previously, metadata extraction from SQL Server sources failed to extract tables from a database when the database collation type was different from the collation type on the SQL Server level. This issue has been fixed.
- Fixed the issue with QLI from Presto with Hive and AWS Glue as *Custom DB* data source. Previously, there was inconsistency in how QLI jobs pulled query history into Alation: initial QLI worked but consecutive QLI jobs omitted queries. After this fix, Alation should correctly process query history according to the QLI settings of the data source.

1.12.7 RELEASE 2021.1.7

build 7.2.8.141049

- Presto on Hive QLI fails with `UnicodeDecodeError` when column data has unicode characters. This issue has been fixed by defaulting to UTF-8 encoding for reading QLI log files.
- Fixed an issue with Greenplum metadata extraction where it fails for unicode characters in functions.
- Previously while performing metadata extraction or any other action on a connector and the connector is unresponsive, the connection will be terminated. Now, the connection timeout time has been increased.
- Fixed an issue where the metrics in query search logged the Object ID (OID) as integers which fail the searchclickextract job. Now, the metrics in query search logged oid as otype for each search result click.
- Previously the CSV exports in Compose were having few blank rows in between the rows. The CSV exports now should not have any blank rows in between the rows.

1.12.8 RELEASE 2021.1.6

build 7.2.7.139384

- Previously, when **Use Group to assign roles** was enabled, new users with the Viewer role did not appear in the Viewer built-in group. Those users received the “Permission Denied” error on numerous pages if the Viewer role was enforced on the instance. Now all new users are assigned to the correct built-in groups that correspond to their roles.
- Query execution and export failed when multibyte Unicode characters were part of the query result retrieved from a data source. This issue has been fixed by using a different decoder mechanism.
- Fixed the issue of the system health alerts not being sent to admins. System health alerts are now triggered and sent as expected.
- Fixed a problem where new schedules for existing queries could not be created after the Alation update to 2021.1. This fix also updates the partially created schedules and applies them to the queries as was intended.

- The search group strategy for LDAP groups uses a query that pulls out all the users. Previously, in case of very large LDAP groups, this query could time out or reach the LDAP server limit on how many results were permitted to be retrieved. This fix introduces pagination for the LDAP group search query and an **alation_conf** setting to control it: **alation.authentication.ldap.search_page_size** (defaults to 500). The parameter should be set to a value that is less than the group result size allowed by the LDAP server.

1.12.9 RELEASE 2021.1.5

build 7.2.6.138178

- Fixed an issue with metadata extraction from Azure Data Warehouse where it failed with the `0 schemas found` error. This was resolved by adding support for case sensitive collation.
- Fixed an issue with high CPU usage by Taskserver during extraction, QLI, profiling, search index update, and some other processes. The issue was caused by Hive grammar projection grammar rule: `antlr` went into an infinite loop while parsing a query. As a solution, the `antlr` version and the timeout for execution of parser threads were updated.
- Fixed an issue when information about Connections based on custom SQL queries was missing from the Tableau report pages after extraction from the Tableau server.
- Fixed an issue with query log ingestion from Hive data sources when it failed with the `QueryLog.setStartTime was called more than once` error when multiple `AM_STARTED` attributes were present in the query log. The last `AM_STARTED` attribute in the log will now be the value used in ingestion.

1.12.10 RELEASE 2021.1.4

build 7.2.5.136994

- Fixed a performance issue when a large number of Groups with multiple members in the system caused the **Customize Catalog** page to load slowly. Now, the page loading time has significantly decreased.
- Fixed a problem when sometimes a subset of web server workers on the Alation server hung and did not serve requests. This fix improves throughput and availability of the Alation web service.
- Previously, QLI could become stalled at the `StatementParsingJob` stage with `get()` returned more than one `AlationVersion` error message. Fixed by adding a unique constraint to the `tree_hash` field of the `AlationVersion` model so that no duplicate entries can be created due to race condition.
- Added two **alation_conf** parameters to mitigate search timeout issues. Admins can change the values in case they observe timeout issues on the **Advanced Search** page or when using the public Search API:
 - **alation.search.timeout = 10.0**
 - **alation.search.max_retries = 1**
- SAML requests generated to authenticate Compose users to SSO-enabled AWS data sources used an invalid format for the `IssueInstant` attribute in the request, which resulted in a connection error. This problem is now fixed by assigning the `IssueInstant` attribute the correct ISO format based on the SAML 2.0 specification.
- Fixed an issue of the manual upload of usage data to Alation failing due to changes in how model data are written to files in Python 3. After the fix, manual uploading of usage data completes successfully.

1.12.11 RELEASE 2021.1.3

build 7.2.4.134495

- Fixed an issue with lineage calculation where Table information was not available on the **Connections** tab of the Tableau BI source and the lineage relationships with Tableau objects were not displayed on the **Lineage** tab of the Amazon Redshift data source when there were multiple clusters with the same URI and schema and table names.
- Fixed an issue of the lineage diagram for the Tableau source not displaying the upstream objects from a Presto on Hive data source when a number of lines of the connection query SQL were commented out. Added support for array expressions to Alation Presto grammar.
- Fixed a bug related to open stream writer that caused excessive memory usage during metadata extraction from a file system source.
- Fixed a QLI issue where it failed with the following error: *A string literal cannot contain NUL (0x00) characters*. Now, queries containing null characters will not cause the QLI job to fail.

MicroStrategy GBM V2 Extraction Stabilization

- Fixed an issue where only the first project from the list of projects selected for extraction was extracted into the Catalog. Now, all selected projects should be extracted as expected
- Fixed a number of NPEs reported during metadata extraction
- All document, dossier, and dashboard previews are now extracted using the MicroStrategy REST APIs
- Alation now determines if a cube is published or if it is not using isPublished API during extraction
- Made changes to populate sub-reports of a document using search-based API. This works for both published and unpublished cubes
- Cube data sources are now linked with reports and documents using search API, which enables Alation to calculate lineage
- All data source cubes in the project are extracted using the search API
- Simple reports are extracted using the NoAction execution flag. This opens the reports in the design mode before extraction and works for both published and unpublished cubes
- Alation now retrieves object prompts using the search API and extracts the corresponding field metadata
- Alation can now extract the document and lineage information irrespective of whether the underlying cube is published or not. Fields will not be extracted if the document depends on an unpublished cube
- Report metadata is extracted asynchronously. If a large number of reports are extracted, Alation creates multiple connections to extract report data. The MicroStrategy server has different types of governing job limits set to restrict the number of connections per user account, user session, and per project. Alation will respect the limits set on the MicroStrategy server and only use the server resources allowed to be used for extraction
- Preview generation will now be skipped if a report or document is dependent on an unpublished cube.
- Fixed multiple issues with prompt validation and extraction
- Added a fix for the VALUE prompt to answer all prompts and only fail for required prompts that do not have any answers
- Alation will now extract only valid metric types. The auto-generated field row count that used to cause the MicroStrategy server crash is no longer extracted

1.12.12 RELEASE 2021.1.2

build 7.2.3.133419

- Fixed an issue where long-running extraction jobs were given the *Failed* status in the Job History table in Alation UI before they were completed. The job status will now be properly tracked for the entire duration of the job lifecycle.
- Added support for the back quote character ` in the Query Log Ingestion parser logic, which should ensure that QLI is successful for the data sources where query SQL includes this character.
- Improved handling of recursive objects when they were added to a Data Domain, removed, and then re-added.
- Added a description of the Search API to the API Documentation on Customer Portal.
- Fixed an issue where the **Articles** section of the left-hand navigation panel was unavailable to users who logged in for the first time, requiring a refresh of the browser page. The **Articles** panel is now always displayed in the left-hand navigation panel.

1.12.13 RELEASE 2021.1.1

build 7.2.2.132324

- Added an **alation_conf** flag **alation.backup_v2.pgbackup_compression** to exclude the Postgres backup from Alation backup tarball creation. This reduces backup tarball creation time. Postgres backups should be handled manually and restored separately. For details, refer to [How to Exclude Postgres Backup from Alation Backup](#)
- Search using the * wildcard is now case insensitive.
- Fixed an issue with SSO authentication for AWS data sources where it did not work in Chrome without disabling the SameSite flags. This change introduces the use of the RelayState parameter in SAML requests to track asynchronous communication with identity providers and deprecates the use of “user session” to track such asynchronous communication, rendering the SSO solution browser-agnostic.

1.12.14 RELEASE 2021.1 - General Availability

build 7.2.1.131666

New Features

Search and Discovery

Data Domains

A Data Domain is a new catalog object type that can be used to categorize other catalog objects. With the help of Data Domains, you can organize your catalog data by business areas, departments or any other unifying criteria. Data Domains can be created, edited, deleted, and organized hierarchically. Like other Catalog object types, Data Domains have a dedicated customizable Catalog page template that allows for the adding, removing, and re-arranging of custom fields. Data Domains increase discoverability of other Catalog objects: by selecting the relevant Domain in the Search filter, you can retrieve the list of all catalog objects residing in this Domain. The Data Domains feature can be enabled on the Admin Settings > **Feature Configuration** page (former **Labs**).

Search API

2021.1 adds a new public API for retrieving Search data from the Catalog: you can now send requests to the GET Search API endpoint and retrieve search results. GET Search API documentation is available in the OAS 3.0 format at **`http(s)://<your_Alation_URL>/openapi/search/`** (requires the Swagger integration to be enabled: `alation_conf alation.feature_flags.enable_swagger = True`). The Search API offers the capability to interact with other applications used at your organization and to integrate the Alation Catalog into the existing application workflows.

Full-Page Search Improvements

Alation's full-page Search now supports the asterisk wildcard syntax. The asterisk symbol can be placed at the beginning, at the end, and in the middle of the search string: "ab*", "*yz", "ab*yz", and "*mn*". The search will retrieve relevant Catalog objects where the search string with the wildcard symbol matches the title or the name of the object.

Governance and Curation

Alation Analytics V2 Improvements

- Added search data to the Alation Analytics V2 schema surfacing the information on:
 - the number of searches performed by users
 - search queries
 - search timestamps
 - the number of searches that resulted in a click-through
 - the pages from which search was initiated;
- Added dashboard-level filters to the Alation Analytics V2 Dashboard that allow filtering the dashboard data by User and Date Range;
- Added the ability to remove the Alation Analytics V2 data source from the Catalog;
- Alation Analytics V2 database password is now stored in the encrypted format in the `alation-analytics.env` file.

Lineage Impact Analysis

The Lineage page now includes 2 new tabular Lineage reports: **Impact Analysis** and **Upstream Audit**. These reports list data objects downstream or upstream from a selected object in the tabular format. The tables can be filtered by relevant fields, such as *Steward*, *Object Type*, *Trust Flag*, and others. The reports can be exported into the CSV format. For details, see [Lineage Impact Analysis](#).

Custom DB Profiling Improvements

- Column catalog page load time is significantly decreased
- Custom Profiling information has been added to the Profiling preview in Compose: [Custom DB Column Profile Preview](#)
- Improved the UI of the Frequency Distribution chart: Frequency Distribution Chart
- The names of UI elements and text messages for deep Column Profiling (Profiling V2) and Sampling (default) features were updated for a more consistent user experience

Connect Platform

Data Source Certifications

Certified support for more data sources and new versions of supported sources:

- Starburst Enterprise Trino
- Teradata Vantage
- Denodo
- MongoDB Atlas and others.

Enterprise Deployment

Federated Compose User Login For AWS Sources

Alation Compose now supports SSO for data sources on AWS. Alation integrates with AWS IAM to direct the authentication flow through SAML 2.0-supporting IdP providers used at the Organization (such as Azure AD, Okta, ADFS, PingOne, and others).

User Role Assignment Using LDAP Groups

Server Admins can now map Alation roles to Custom Groups, and users in this group are automatically assigned the mapped role. Admins also have an option to auto-suspend LDAP users who do not belong to any LDAP group: [Use Custom Groups to Assign User Roles](#).

Alation Sandbox Improvements

- Added the `--resume` flag for the `run` command
- Download progress status is now displayed as the build is being downloaded
- After a successful run, users are prompted to use the `alation-sandbox-admin` user for UI login
- Improved error messages
- Key rotation warning is displayed when upgrading to the 2021.1 Alation build
- Improved handling of the Sandbox license
- Numerous other bug fixes and improvements to enhance user experience, stability, and security.

Alation Server

- Upgraded internal codebase components to Python 3
- Expanded Operating Systems support for new installations of the Alation server by adding:
 - Fedora
 - Oracle Linux with RHEL compatible kernel
 - Debian

Improvements

- An admin can now disable the query export feature in Compose.
- The **Labs** page was renamed to **Feature Configuration**.
- Previously, some of the queries from Presto were not parsed by the Presto parser component because they contained built-in functions with no arguments. Now, the Presto parser will properly handle such queries and they will be displayed on the **Query History** tab.
- The built-in driver for Snowflake was updated to version 3.12.14.
- Added support for the NULL value in the IS_NULLABLE column in the Columns query for the Query-Based MDE.
- Previously, making numerous changes to fields on a custom template could cause a spike in the number of background tasks processed on the Alation server. Now, this action will be processed in parallel to other running tasks and not cause the queued tasks to back up.
- The temporary folder **pg_backup** created under **/data2/backup/** during backup restore is deleted automatically after the backup restore process completes successfully.
- Tableau BI source logs will now capture BI type, framework version, and error code.
- For Databricks on Azure as Custom DB, the complex data types *Array* and *Struct* can be displayed as a structured JSON object in Alation UI.

Bug Fixes

Note: 2021.1 GA also includes all bug fixes listed for the 2021.1 LA version.

- Fixed an issue where Table and Column samples were not loading for users with the Viewer role when the Viewer role was enforced. After the fix:
 - When **Dynamic Profiling** is OFF, only *Server Admins* can run samples and profiles.
 - When **Dynamic Profiling** is ON, all roles, including Viewers, can run samples and profiles.
- Fixed an issue where after update to 2021.1 LA Tableau certificates had to be re-uploaded on the Tableau source Settings page. No action is required after updating to 2021.1 GA from a previous major release.
- For the Amazon Redshift data sources, decimal, real, float, and double data types will now display the correct length and precision.
- Miscellaneous other fixes.

1.12.15 RELEASE 2021.1 - Limited Availability

build 7.2.0.130782

For new features in 2021.1, see [New Features](#).

Bug Fixes

- Fixed an issue with MDE from Amazon Redshift data sources when tables with column comments were not extracted. Now, table extraction works as expected.
- Fixed an issue when the Monitor > **Active Tasks** tab displayed a pop-up error message when the number of active tasks in queues was too large to be displayed. After the fix, Active Tasks no longer issues the error. If the task list for a specific queue is too large to be displayed (> 500 items), the Active Tasks list will show a collective record “other tasks in this queue” with the number of overflowing tasks.
- Fixed a bug where Alation Analytics V2 did not reflect the updated Glossary article counts correctly after Glossary rules were updated.
- Fixed an issue when the Popularity values were not updated correctly after QLI in a specific case of QLI.
- Previously, table aliases from merge statements from Teradata were incorrectly recognized as temporary tables, which generated incorrect lineage. Now the aliases are properly ignored.
- Added support for Hive data sources with the Oracle metastore on the configuration-based Hive framework (Hive V2).
- For MySQL data sources, fixed an issue with the MDE property **Remove schemas from the catalog that are not captured by the lists below** when a previously extracted schema was not removed during the next MDE after being placed into the Exclude Schemas list.
- The Compare User Roles matrix on the Admin Settings > **Users** page now appropriately reflects the current access rules. For example, a Steward can access Compose, and the Data Source Management tooltip no longer includes “wizard deployment” as Catalog and Source Admins do not have access to this feature.
- Fixed the text overflow issues in the Quick Links section of the Homepage and the Apps menu in Internet Explorer 11.
- Fixed a bug where embedded videos with the allowfullscreen HTML attribute could not be viewed in full-screen mode.
- Fixed the issue of QLI archiving not working after enabling the Lineage V2 feature.
- Restored the ability to upload data to a table using Excel files.
- Fixed a performance issue that manifested itself when loading batches of execution results while a Compose query tab was being loaded (or when adding results for a new execution). This fix provides a dramatic improvement in execution time and will free up database resources as a result.
- When OAuth is enabled for a data source but the default Compose URI does not include the OAuth parameter, using the Export Connection dialog in Query Forms and Data Upload with manual URI and without saved credentials will use the manually entered credentials from the username and password fields.
- Removed the Alation Analytics V2 staging tables from the UI on the Per-Object Parameters tab as these tables are not displayed under the Public schema and are not meant to be profiled.
- Addressed a potential Stored XSS issue where the “title” parameter of a query would be stored and rendered without any sanitization.

- Since release V R7 (5.12.x), the script to convert usernames did not work. This has now been fixed, and admin users can continue to use the instructions provided in [How to Convert Usernames](#).
- Fixed a bug in Compose when clicking on **Run Current Statement** executed the statement found in the next line. Now, the statement on the line which currently has the cursor will be executed.
- Fixed a UI issue with table formatting in the Dusk theme in Compose: the row shading of the table in the Preview pane now uses 2 alternating colors from the navy blue palette.
- Fixed the incorrect Lexicon file size limit value in the error message that is displayed when the file size limit is exceeded. The value in the message now correctly reflects the file size limit set in the Alation application.
- Previously, the Open API specification of the Custom Field Values API incorrectly rendered the example for the PUT endpoint. This issue has now been fixed.
- Fixed an issue with installing Alation Connector Manager on SELinux-enabled systems. Alation Connector Manager is now configured to automatically start on system startup.

1.12.16 Known Issues

- If you have a password set on the internal PostgreSQL database (Rosemeta), the update to 2021.1 will result in an error: Password for user alation: psql: fe_sendauth: no password supplied ERROR: Cannot upgrade alation. Update to 2021.1 requires that the password on the internal PostgreSQL database should be cleared for the time of the update. The password needs to be set again after the update is completed.
- The command `alation_action scan_postgres` will fail if a password is set on the internal PostgreSQL database.
- The search using the wildcard symbol with uppercase symbols (for example, TEST*) only returns results from object titles and disregards object names.
- With Backup V2 enabled, automatic backup process cannot write to `alation_backup.log` due to a file permissions error.
- Only users with the Server Admin permission can see reports extracted from an SSRS BI source. Contact Alation Support to apply a known workaround if you run into this issue.
- If during the upgrade you encounter the error that points to the corrupted index on the `djcelery_periodictask` table, please contact Alation Support to fix the corrupted indexes. The error message will include the following text: (duplicate key value violates unique constraint \"django_celery_beat_periodictask_pkey\").
- CREATE VIEW queries in Denodo do not support the `catalog.table` format, and as a result, the CREATE VIEW queries are not parsed during QLI and Lineage cannot be built.
- Alation saves the timestamps of the Catalog events in the UTC time. As a result, the events created after 12 am UTC will be loaded into the Alation Analytics V2 database with a day's delay.
- With the Viewer role enforcement enabled, the Viewer users cannot run Table and Column profiles on data sources they have access to **[FIXED IN 2021.1 GA]**
- For AWS Data Sources with SSO enabled in Compose, when the STS token expires, Compose issues an error in the query results page instead of displaying the Authentication dialog.
- In Chrome, Compose SSO for AWS data sources may not work correctly if Alation uses HTTP and the security flags are enabled in the browser settings. As a workaround, disable the SameSite setting (set SameSite=None) in your Chrome settings if Alation uses HTTP.

- The `alation_repair` action does not fix permissions on the **extra_config** folder, which may cause the Alation upgrades to fail. The recommendation to check permissions on **extra_config** has been included into the upgrade steps for 2021.1.
- Using the up or down arrow keys on the keyboard to select a search item from the Quick Search list closes the Quick Search **[FIXED IN 2021.1 GA]**
- Upgrades from 2021.1 LA to a newer build of 2021.1 do not work. **[FIXED IN 2021.1 GA]**
- Backups of instances on release 2021.1 LA cannot be restored on clean 2021.1 systems. **[FIXED IN 2021.1 GA]**
- New configuration-based Hive connections (Hive V2) cannot be created as the Hive configuration file upload does not work in 2021.1 LA. Existing configuration-based Hive data sources continue to work as expected. **[FIXED IN 2021.1 GA]**
- If you update an unpublished query created by cloning a published query, the changes are not displayed on the corresponding Catalog page.
- For extraction jobs that take a long time, the Job History UI may mark the job as “failed” before it actually completes. The job will eventually complete, and the Job History UI will be updated to reflect this; however, initially, the job is marked as “failed” and false alerts are issued to admins.
- Compose does not work for Salesforce DB with the custom CData driver.
- Customers upgrading their 2021.1 LA test environments to 2021.1 GA still need to re-upload the SSL certificates on the Tableau source Settings page after the update. This step is required only for customers upgrading from 2021.1 LA to 2021.1 GA and does not affect updates from previous major releases to 2021.1 GA.

1.13 Release Notes 2020.4 (General Availability)

1.13.1 RELEASE 2020.4.15

build 6.2.17.155524

This patch is a security update providing fixes to mitigate the Log4j 2 vulnerability discussed in the *December 22, 2021 - Log4j2 Security Advisory*. This addresses the security vulnerability described in [CVE-2021-45105](#).

This patch is cumulative and fixes all the security issues from the prior two patches, 2021.4.13 and 2021.4.14.

Upgrading to this security patch will apply the fixes automatically. This patch includes:

Security fix for the Hive connector

The Apache Hive connector is updated with a newer version of Apache Log4j 2. This addresses security vulnerability [CVE-2021-45105](#).

Security fix for Elasticsearch

Elasticsearch is updated with a newer version of Apache Log4j 2. This addresses security vulnerability [CVE-2021-45105](#).

1.13.2 RELEASE 2020.4.14

build 6.2.16.155178

This patch is a security update providing a fix to mitigate the Log4j 2 vulnerability discussed in *December 17, 2021 - Update to December 15 log4j2 CVE-2021-45046 Security Advisory*.

Security fix for the Hive connector

The Hive connector is updated to use a newer version of Apache Log4j 2, which addresses a possible DoS vulnerability described in [CVE-2021-45046](#).

1.13.3 RELEASE 2020.4.13

build 6.2.15.154832

This patch is a security update providing two important fixes to mitigate the Log4j 2 vulnerability discussed in *December 10, 2021 - log4j2 Security Advisory*.

Upgrading to this security patch will apply the fixes automatically. The patch includes:

Security fix for the Hive connector

The Apache Hive connector is updated with a newer version of Apache Log4j 2 as per NIST publication [NVD - CVE-2021-44228](#). In addition to updating the Log4j 2 version, all invocations of the Apache Hive connector will run with the `-Dlog4j2.formatMsgNoLookups=true` system property.

With this security update, the Alation Hive connector is not susceptible to remote code execution based on the Log4j 2 vulnerability [CVE-2021-44228](#).

Security fix for Elasticsearch

As per recommendation from Elastic found at [Apache Log4j2 Remote Code Execution \(RCE\) Vulnerability - CVE-2021-44228 - ESA-2021-31](#), this fix disables the message lookup substitution for Elasticsearch on the Alation server. With this security update, the `alation_conf` parameter `elasticsearch.env.es_java_opts`, which sets Java options for Log4j 2, will be set to `-Dlog4j2.formatMsgNoLookups=true` by default. This change ensures that the Log4j 2 Java options have the recommended setting and that the Elasticsearch component of the Alation server is not susceptible to remote code execution based on the Log4j 2 vulnerability [CVE-2021-44228](#).

1.13.4 RELEASE 2020.4.12

build 6.2.14.145722

- Fixed an issue with Backup V2 restore where the restore process failed with the error `pg_probackup archive-get` failed to deliver WAL file.

1.13.5 RELEASE 2020.4.11

build 6.2.13.143825

- Previously, metadata extraction from SQL Server sources failed to extract tables from a database when the database collation type was different from the collation type on the SQL Server level. This issue has been fixed.
- Fixed an issue with metadata extraction from Azure Data Warehouse where it failed with the `0 schemas found` error. This was resolved by adding support for case sensitive collation.

1.13.6 RELEASE 2020.4.10

build 6.2.12.138408

- The search group strategy for LDAP groups uses a query that pulls out all the users. Previously, in case of very large LDAP groups, this query could time out or reach the LDAP server limit on how many results were permitted to be retrieved. This fix introduces pagination for the LDAP group search query and an **alation_conf** setting to control it: **alation.authentication.ldap.search_page_size** (defaults to 500). The parameter should be set to a value that is less than the group result size allowed by the LDAP server.

1.13.7 RELEASE 2020.4.9

build 6.2.11.137415

- Fixed an issue with high CPU usage by Taskserver during extraction, QLI, profiling, search index update, and some other processes. The issue was caused by Hive grammar projection grammar rule: `antlr` went into an infinite loop while parsing a query. As a solution, the `antlr` version and the timeout for execution of parser threads were updated.
- Previously, QLI could become stalled at the `StatmentParsingJob` stage with `get()` returned more than one `AlationVersion` error message. Fixed by adding a unique constraint to the `tree_hash` field of the `AlationVersion` model so that no duplicate entries can be created due to race condition.
- Fixed a problem when sometimes a subset of web server workers on the Alation server hung and did not serve requests. This fix improves throughput and availability of the Alation web service.
- Added two **alation_conf** parameters to mitigate search timeout issues. Admins can change the values in case they observe timeout issues on the **Advanced Search** page or when using the public Search API:
 - **alation.search.timeout = 10.0**
 - **alation.search.max_retries = 1**

1.13.8 RELEASE 2020.4.8

build 6.2.10.135810

- Fixed a QLI issue where it failed with the following error: *A string literal cannot contain NUL (0x00) characters*. Now, queries containing nulls will not cause the QLI job to fail.
- Fixed a performance issue when a large number of Groups with multiple members in the system caused the **Customize Catalog** page to load slowly. Now, the page loading time has significantly decreased.

1.13.9 RELEASE 2020.4.7

build 6.2.9.134079

- Fixed an issue where the **Articles** section of the left-hand navigation panel was unavailable to users who logged in for the first time, requiring a refresh of the browser page. The **Articles** panel is now always available in the left-hand navigation panel.

1.13.10 RELEASE 2020.4.6

build 6.2.8.132865

- Added support for the back quote character ` in the Query Log Ingestion parser logic, which should ensure that QLI is successful for the data sources where query SQL includes this character.
- Fixed an issue where long-running extraction jobs were given the *Failed* status in the Job History table in Alation UI before they were completed. The job status will now be properly tracked for the entire duration of the job lifecycle.

1.13.11 RELEASE 2020.4.5

build 6.2.7.132077

- Added an **alation_conf** flag **alation.backup_v2.pgbackup_compression** to exclude the Postgres backup from Alation backup tarball creation. This reduces backup tarball creation time. Postgres backups should be handled manually and restored separately.

For details, refer to [How to Exclude Postgres Backup from Alation Backup](#)

1.13.12 RELEASE 2020.4.4

build 6.2.6.131177

- Fixed a performance issue with search indexing for file objects. This change is expected to make file system objects indexing 3 to 5 times faster.
- Previously, text in tables in the Description field of file system objects was not searchable. This issue was fixed, and the Description field of file system objects is now correctly parsed. The file system objects where the text in tables in the Description field matches the search string will now appear in search results.
- Added support for user impersonation for EMR Presto as Custom DB: users can use impersonation for running queries in Compose.

1.13.13 RELEASE 2020.4.3

build 6.2.5.129768

- Fixed an issue when Unicode characters in View definition SQL caused MDE to fail. MDE will now correctly parse View definitions that contain Unicode characters.

1.13.14 RELEASE 2020.4.2

build 6.2.4.128370

New Features in 2020.4.2

- For Custom DB, Alation now allows Data Source Admins to enable concurrent queries for Compose. This can be done on the **Custom Settings** tab of the Custom DB data source Settings page with the help of the **Enable Concurrent Queries** checkbox.
- For Custom DB, Alation now allows Data Source Admins to enable **User Impersonation**. This can be used when configuring Kerberos and Keytab authentication for Compose for the data sources that support this setup. User Impersonation can be turned **On** on the **General Settings** tab of the Custom DB data source Settings page.

Bug Fixes

- Fixed the issue of Lexicon overflow error by capping the approve/reject count to 7,000. In the Lexicon job, if the approve/reject count exceeds 7,000, the exceeded count will be capped at 7,000.
- Fixed an issue with delay in loading the **Glossaries** page and glossary pages.
- Previously, metadata extraction from Snowflake extracted the view names in upper case when the original view name was in mixed case. Now, metadata extraction of views extracts the view names in mixed case.

1.13.15 RELEASE 2020.4.1

build 6.2.3.127279

- For Amazon S3 users, fixed an issue where Alation would not extract the last batch of files when extracting more than 1,000 files. For example, if a filesystem had 1,234 files, 1,000 would be extracted and 234 would be ignored.
- For BI users, fixed an issue where setting the automated extraction job from the UI would not make any changes.
- For Hive users, fixed an issue where query logs generated by Hive 2.x would not be extracted when they were URL encoded. Users facing this issue need to include a configuration file named **aha-conf.xml** in their configuration tarball for the datasource. That configuration file should contain the following lines:

```
<configuration>
  <property>
    <name>alation.qli.check-url-decode</name>
    <value>true</value>
  </property>
</configuration>
```

- Upgrading HA systems to 2020.4 GA with preserving the cluster required an extra step of creating the `reindex_rosemeta_success` file on the Secondary server. When upgrading HA systems to 2020.4.1, this extra step is no longer required.

1.13.16 RELEASE 2020.4 - General Availability

build 6.2.2.125680

New Features

Alation Search Improvements

- BI Field objects are now included into the BI filter to allow for searching and finding the full range of BI object types;
- The quotation mark search operator is now supported: by including your search keywords in quotation marks, you can look for an exact match in the Catalog data.
- Misspelled search keywords are now auto-corrected to speed your search.
- A number of improvements were made to the search matching algorithm to enhance matching accuracy and search comprehensiveness.
- **Elasticsearch version 1 has been removed** from the Alation server. During the update to 2020.4, Alation will determine if the Elasticsearch 1 index data is still in use. If it is no longer in use, it will automatically be removed, as well as the Elasticsearch 1.4 directory `/data1/elasticsearch`.
- Breadcrumbs for BI objects have been added to the Search results page.
- BI Objects in full-page Search can now be filtered by Source.

New Lexicon UI

The new UI of the Lexicon page makes Lexicon management an easy and logical experience. Lexicon abbreviations and expansions are presented in a tabular format with multiple filters. You can easily find and update suggestions, add new abbreviations and expansions, or go through the unconfirmed suggestions and curate them.

Brand Color Consistency

The custom branding color configured for the Customizable Homepage is now extended to the rest of the Catalog: icons, tabs, font highlighting, and other UI elements - to create a consistent color theme across all Catalog pages.

New Features for Alation Analytics V2

In 2020.4:

- Alation Analytics V2 schema and ETL are extended with BI object and Catalog Set data;
- The **Insights** page has been renamed to **Alation Analytics Dashboard**;
- 2 new visualizations have been added to the Alation Analytics Dashboard: **Popular Tables** and **Popular Queries**;
- Alation Analytics V2 becomes upgradeable to newer versions. The upgrade process is separate from the Alation application upgrade;
- The capability to remove the Alation Analytics V2 data source from the Catalog has been added;
- On the Alation Analytics **Settings** page, there is now an ability to download the Alation Analytics V2 package using a Curl command;
- Alation Analytics V2 can be installed into a custom directory that can be created automatically during installation or specified during installation if created in advance;
- The Alation Analytics **Settings** page will now display the currently installed version of Alation Analytics V2;
- There is now an ability to remove the Alation Analytics V1 data source after enabling Alation Analytics V2.

Table Object Privacy Management

In previous releases, Table object access permissions were inherited from the parent data source. From 2020.4, a *Server Admin* or a *Data Source Admin* can decide to make a particular table an exception by configuring access settings on the Table object level. For Table objects, there is now a **Settings** page which allows for selecting a Privacy option: *Public* or *Private*. For Private Table objects, the admin can choose individual users and groups who can view this table. This feature is not enabled by default and can be activated using the **alation_conf** command and the dedicated feature flag `alation.granular_object_privacy.enabled`

[Table Privacy Settings](#)

Custom DB: Query-Based Column Profiling, Value Distribution Chart, and Profiling Stats

In order to provide better and faster data insights from the Catalog, Custom DB sources have been enhanced with column-level Profiling capabilities. More than that, Profiling is now customizable for each individual column. You can:

- write a custom query for each column to be profiled, for both numeric and non-numeric data types;
- run a full profile of an entire column;
- view the Value Distribution chart and Profiling stats for numeric data types.

This feature can be enabled using **alation_conf** and the parameter `alation.feature_flags.enable_profiling_v2`.

[Sampling and Profiling](#)

GET Lineage API

Added a new API allowing users to get objects in lineage paths and traverse lineage graphs for root cause analysis. GET Lineage API documentation is available in the OAS 3.0 format at **`http(s)://<your_Alation_URL>/openapi/lineage/`** (requires the Swagger UI to be enabled on the Alation instance: `alation_conf.alation.feature_flags.enable_swagger = True`).

Datasources API

The Datasources API allows for integrating with CI/CD data pipelines when creating, updating, or deleting data sources and for creating automation to get schemas, run MDE, do configuration checks, or update data source credentials. This API is available in the OAS 3.0 format at **`http(s)://<your_Alation_URL>/openapi/datasources/`** (requires the Swagger UI to be enabled on the Alation instance: `alation_conf.alation.feature_flags.enable_swagger = True`). Additional data sources have been certified and now the following database types are covered:

- Snowflake
- Amazon Redshift
- PostgreSQL
- Teradata
- SQL Server
- MySQL
- Oracle

Lineage Diagram Enhancements

The Lineage diagram now has the ability to filter out *TMP* objects and a full-screen view mode. General usability of the diagram has also been improved:

- GBM V2 catalog pages now have a link icon to their respective catalog pages
- Untitled Dataflow nodes now have a default title in the tooltip instead of just the Dataflow icon.
- Lineage nodes on the diagram now support multiple badges. For example, the TMP or EXT badge will now be displayed next to the Deprecation badge if both are applied to a node.
- Action buttons on the diagram are now spaced evenly
- The Preview section shows a default visual cue to select a node when no preview is displayed
- UI performance of the Lineage diagram has been improved

Open Connector Framework (OCF) for BI Sources

Open Connector Framework (OCF) allows Alation to support external connectors developed by partnering teams in order to add specific sources to the Catalog. An OCF connector can be viewed as a plugged-in microservice that connects to an external source (such as a BI server) in order to extract raw metadata. Connectors created on the basis of OCF can be installed in addition to the Alation application and maintained separately from the Data Catalog. In addition to OCF, Alation provides OCF-based out-of-the-box connectors for the following BI tools:

- Power BI
- Looker

Alation will provide OOB connector packages for downloading upon request. For details, refer to section [Open Connector Framework](#).

Pass-Through Credentials Compose Authentication

A new authentication mode - **Transient** - is added to Compose when user database credentials are not stored and require to be provided every time connection to the database is established. In the *Transient* mode, users establish a connection by entering their database username and password in Compose, but their credentials are not stored on the Alation server side. This applies to Compose query execution, query forms, data upload, and Dynamic Profiling, and other Compose functionality. The credentials storage mode is controlled by the **alation_conf** parameter `alation.datasource_auth.credentials.storage_mode`. The default is *Persistent* mode (parameter value = **0**) when credentials are stored on the Alation server. To enable the *Transient* mode, set this parameter to **1**.

[Transient User Credentials](#)

OAuth in Compose for Databricks Sources

Support for connections using OAuth 2.0 has been added for Azure Databricks with AAD as the authorization server. OAuth can be configured for Databricks data sources and OAuth connections can be made with such data sources for query execution, query scheduling, query forms, and Dynamic Profiling, Excel live reports, and data upload. OAuth can be set up on the General Settings tab of the Databricks data sources, under Compose Connections.

[Databricks Azure OAuth for Compose](#)

Data Source Certifications

Alation certifies support for new data sources and new versions of previously certified sources:

- [Salesforce](#)
- [ServiceNow](#)
- [EMR Presto on Hive with AWS Glue as Metastore](#)
- newer versions of supported sources (PostgreSQL, Greenplum, Tableau, and others).

Data Source Built-In Driver Updates

In 2020.4, a number of native (built-in) connectors were updated with newer driver versions to keep up with the driver bug fixes and improved functionality. During the update to 2020.4, the data sources using built-in drivers will be automatically switched to the newer driver versions.

- **Databricks:** added Simba JDBC driver for Databricks, driver version 2.6.16:
 - *AWS Databricks:* the new driver is auto-switched during update. The URI is auto-updated to conform to the format required by the driver;
 - *Azure Databricks* using Custom DB: no changes to the driver during update.
- **Google BigQuery:** updated JDBC driver for Google BigQuery to driver version 1.2.11.1014:
 - no changes to the URI. The driver is auto-switched during the update.
- **MySQL:** added the JDBC driver for MySQL, driver version 8.0.21:
 - no changes to the URI. The driver is auto-switched during the update.
- **PostgreSQL:** updated the JDBC driver for PostgreSQL to driver version 42.2.14:
 - no changes to the URI. The driver is auto-switched during the update.
- **Snowflake:** changed the JDBC driver for Snowflake to driver version 3.12.9:
 - no changes to the URI. The driver is auto-switched during the update.
- **SQL Server:** added Microsoft JDBC driver for SQL Server, driver version 8.2.2:
 - SQL Server sources using older built-in driver versions 6.2 or 7.4.1 are auto-switched to the newer driver version 8.2.2;
 - SQL Server sources using the NTLM authentication and the built-in jTDS driver are NOT switched over to the new driver during the update. It is recommended to manually change the jTDS driver to 8.2.2 and update the URI after updating to 2020.4.

CentOS 8 for Alation Chroot

Alation Chroot has been updated to use CentOS 8.2 to take advantage of the new features and bug fixes on CentOS.

CentOS 8 and RHEL 8 for Alation Host OS

The Alation application can now be installed on hosts running on CentOS 8.x and RHEL 8.x. Please note that the older host OS versions, such as CentOS 6.x, RHEL 6.x, and Ubuntu 12 and 14 are no longer supported in 2020.4.

Internal Database Performance Optimization

To avoid issues during Alation upgrades, new detection tools have been implemented to find and troubleshoot Rosemeta indexing issues.

Alation Sandbox

The upgrade testing framework - *Scarlet* - has been renamed to **Alation Sandbox** and considerably improved with the focus on security, reliability, and efficiency. As a result:

- Alation Sandbox can work in a completely offline mode without requiring a connection to Production;
- There is a new CLI tool for easily modifying the configuration parameters;
- There are new configuration parameters to verify the stability of application after upgrade;
- The ability to set up the connection over SSL has been added;
- Alation Sandbox logging has been improved and now includes the PostgreSQL logs.

[Alation Sandbox](#)

Improvements

Catalog UI

- A lot of UI improvements have been done to the Alation Catalog UI as team Alation strives to make Catalog usage an easy and pleasant task.
- Improved the scrolling experience of the Customizable Homepage.

Sources

- Users now have an option **Enable default schema extraction** for Oracle data sources which explicitly enables extractions from default schemas.
- Alation now supports extracting **Impala** query logs from an Amazon S3 bucket. Previously, queries were ingested from a local file system. To extract from S3, the Impala logs should be placed in an S3 bucket in the same manner they would normally be placed on HDFS. A Server Admin should then provide the location and credentials in Alation to access the S3 bucket.
- Alation now ships with the Hive driver version 1.1.0 that is capable of using Zookeeper in a Hive 1.X environment. When setting up Hive with Zookeeper, select the driver **com.alation.drivers.hive.one.cdh** on your Hive datasource on the **General Settings** tab.
- Alation will now extract SQL Server synonyms only for tables and views as synonym base objects. Synonyms for other base objects, such as functions or stored procedures, will be skipped during metadata extraction.
- Optimized performance for extraction from MicroStrategy BI sources on GBM V2, reducing extraction time for extraction with previews.

Platform

- Upgraded the internal NGINX component to version 1.19.2.

Bug Fixes

Note: 2020.4 GA also includes all bug fixes listed for the 2020.4 LA version.

- Badging and Workbook Promotion for Tableau data sources did not work after upgrading from V R7 to 2020.3. This issue has been fixed: Badging and Workbook Promotion now work as expected.
- Fixed an issue with table UI where the table would not display the dataset when table pagination was used.
- Fixed an issue with table filtering where applying filters to paginated tables only returned filtered results from the current table page. After the fix, the whole dataset in the table is filtered and displayed.
- Fixed a bug where for MySQL data sources as Custom DB, during Custom Query-Based MDE, the data types that should have been normalized as INT or FLOAT were returned as non-numeric. As a result, the Query-based Column Profiling did not work for such columns.
- Fixed an issue where for Snowflake data sources with OAuth configured, the OAuth authentication dialog was only displayed for users with the Server Admin role.
- Fixed a bug that caused the Data Objects Without Stewards and Curation Progress reports to fail to load the data due to missing information about certain catalog objects.
- Previously, Alation users with the Steward, Composer, Source Admin, and Viewer roles were not able to update their email notification preferences when the Viewer role was enabled. This issue has been fixed and users with any role can update their email notification preferences.
- Fixed a bug where users with the Viewer role were unable to view the Statement Template object catalog pages when the Viewer role was enabled. Now, users with any role can view the Statement Template catalog page if they are provided access to the data source.
- Fixed an issue where after upgrading Alation Analytics V2, the migration status and logs were not displayed on the **Alation Analytics Settings** page.
- Miscellaneous other bug fixes to improve application quality.

1.13.17 RELEASE 2020.4 - Limited Availability

build 6.2.0.124044

For new features in 2020.4, see [New Features](#).

Bug Fixes

- When the internal Postgres password was set in **alation_conf**, the Backup V2 failed as it was waiting for the password prompt resolution. This issue has been fixed.
- Previously, identical API requests to update the same value for the custom field of Picker type using the Upload Logical Metadata API would result in multiple field history entries in Alation. After this fix, when the same value update is requested multiple times via the API, the history will only store one update record.
- Fixed an issue with customized Glossary views. When an admin customizes the order of columns for a Glossary page and saves the view, then refreshing the page will not change the order back to default. The user's preferences are now preserved.
- Nulls and array values are now parsed as expected for Doc schemas for MongoDB and NoSQL sources.
- Views created in the Alation Analytics V1 database are no longer cleared during the ETL run.
- Previously, the resolution job in the QLI pipeline might fail if SQL queries in the query log contained numerous predicates (more than 50k predicates in 100k SQL statements). This issue has been fixed.
- Fixed lineage computation in case of custom query-based extraction for Amazon Athena data sources.
- Disabled auto MDE and profiling for the Alation Analytics database as default behavior to avoid false alarm emails from Datadog monitoring.
- Previously, for ADLS Gen2, the Access token was not refreshed after an hour, which caused metadata extraction on large datasets to fail. After this fix, the Access token is refreshed as expected, and MDE can complete successfully.
- Fixed an issue with saving API requests that contain both Description and other custom fields and their values.
- Previously, @-mentioning a synonym in an article did not show up in the Relevant Article section on the page of the synonym. Now, @-mentioning of a synonym will resolve the synonym type as table in order to retrieve the corresponding article for this synonym in the Relevant Article section. Note that for existing articles, it is required to re-add the @-mention of the synonym to the article for re-calculating the @-mention of synonym type and for the Relevant Article section to be updated correctly.
- Previously, the GET data sources endpoint would request extra data for each row, which caused long response times on the Sources page. This issue has now been fixed as the data sources API was optimized to reduce the load time on the Sources page.
- For Google BigQuery sources, fixed an issue with passing an incorrect username to Compose. Now, Alation uses the authentication parameters provided for the new connection before running the queries and not the user authentication parameters saved in the MDE settings. This also applies to Dynamic Profiling, data upload, query forms, scheduled queries, Excel Live reports, and Desktop Compose.
- Alation Analytics V2 admins ran into an issue during installation when multiple AA manager entries were made in the crontab when the installer was run multiple times. This has now been fixed.
- Azure Databricks data source would throw an error when using the Simba driver because batch insertion of data was not supported for data upload. Made changes to ensure the Databricks sources follow a sequential execution of Insert queries.
- Previously, with the Viewer Role Enforcement feature enabled, users with the Composer and Steward roles were able to see **Upload Data** option in the **More** dropdown menu on a data source catalog page. They would get the 403 error when they clicked on the **Upload Data** link. Now, **Upload Data** option in the **More** dropdown menu is no longer accessible to Steward and Composer roles.

- Previously, when users ran queries in Compose with results exceeding 40MB, they observed a number of UI issues, such as the *Completed* and *Warning* icons not appearing on the results tab or the *File too large* banner not popping up. This bug has now been fixed.
- For OCF, the internal state is now kept consistent even if connector authors produce multiple different connectors with the exact same name and version. Improved compatibility with middleboxes that do not support TLS 1.3.
- Fixed a UI issue for the Firefox browser when Object names were wrapped and cut off on the full-page Search results page.
- Fixed an issue with permission mirroring for MicroStrategy BI Sources on GBM V2. After the fix, enforcing view permissions when performing extraction is supported.
- Fixed an issue with the *New Query* tab in Compose where it would not open in Firefox.
- Fixed a bug where query search on the **Shared with Me** tab in Compose Query Browser resulted in an error in the browser console and no search results.
- Previously, extraction from Tableau 2020.1 would miss some workbooks when permission mirroring was enabled. This issue has now been fixed.
- For Oracle data sources, fixed an issue where Synonym Columns were not extracted when the *Exclude Schema* MDE filter was used.
- Previously, catalog sets that had filter rules on the Data Source with unicode characters embedded in their name or Description would cause an error during indexing. This is now fixed.
- Previously, Stewards of Virtual Data Sources were not included into query results when querying for Steward information in the Alation Analytics V1. After this fix, the `value_fp` field is updated correctly and the Alation Analytics V1 queries that retrieve Stewards return Steward information for Virtual Data Sources too.
- Previously, Alation Analytics V1 tables still showed objects as existing even after tables were dropped from data sources. To fix this, the `ts_updated` field is also updated when `ts_deleted` field is updated.
- Data source configuration check API calls have resulted in an error code 500 since V R7. The issue has now been resolved.
- Previously, intermittent login issues and connection issues during long query runs in Compose were observed by a customer after upgrading Alation from V R2 to V R5. The root cause of this issue was identified and fixed.
- Fixed an issue where lineage calculation failed in some scenarios involving connections between “temp” columns to actual columns in the graph.
- Fixed an issue with data migration during upgrade to 2020.3 or higher where the upgrade failed if no authenticated user existed in the system with a primary key of 1. This has been changed to assign ownership of the saved connections to the data source creator.
- Previously, even if the *Signup Moderation* feature was enabled, all new users logging into Alation through LDAP would have their user accounts active before the Admins can approve their account. After this fix, Admin users can approve/reject user accounts of LDAP users before they can access Alation if signup moderation is enabled.
- – Miscellaneous other bug fixes.

1.13.18 Known Issues

- **Sampling in Query Log Ingestion** is not supported for Google BigQuery data sources. This feature is located on the **General Settings** page of the data source Settings page and is disabled by default. Leave this feature in the **OFF** state for Google BigQuery sources.
- For Snowflake data sources with OAuth configured, the OAuth authentication dialog is only displayed for users with the Server Admin role. No workaround. This will affect all instances with Snowflake sources and OAuth enabled. **FIXED in 2020.4 GA.**
- Alation Analytics V2 ETL fails after updating Alation to 2020.4 from a previous release and then updating Alation Analytics V2. This can be remedied by editing the Alation Analytics V2 env file is not edited first. **FIXED in 2020.4 GA.** *The steps below have limited application and are relevant to 2020.4 LA (build 6.2.0.124044) only*

Note: Before upgrading Alation Analytics V2 in 2020.4, perform the following steps:

1. Each Alation Analytics V2 installation **.tar** will has its own installer bundled with the updated files. Extract this tar outside of the installation folder of Alation Analytics V2. Running the update command will copy the bundled files to the installed Analytics folder and proceed to update Alation Analytics to the newer version.

```
mkdir /tmp/update-analytics
sudo tar -C /tmp/update-analytics -xzf ./alation-analytics-
package.tar
```

2. Open the **analytics.env** file:

```
vi /etc/default/alation-analytics.env
```

3. Add the following line at the end of the file:

```
INSTALL_DIRECTORY=/opt/alation-analytics
```

4. Run the update Analytics script:

```
sudo /tmp/update-analytics/alation-analytics-x.x.x.xxxxxx/alation-
analytics-installer-v-x.x.x -u
```

5. Once the update is successful, go to the Alation Analytics V2 Settings page and click on **Run Database Migration**.
6. Refresh the Alation Analytics V2 Settings page to view the migration status.

- After upgrading Alation Analytics V2, when the admin clicks on the **Run Database Migration** button, the migration status and logs are not displayed. As a workaround, refresh the page for the logs to be displayed. Additional validation of the migration status: check **Active/Completed Tasks** for Alation Analytics V2 MDE/ETL tasks. **FIXED in 2020.4 GA**
- The Lexicon job may fail because of index corruption. This manifests itself with the following error during the Lexicon job run: "IntegrityError True duplicate key value violates unique constraint "lexicon_token_token_key" DETAIL: Key (token)=(metroarea) already exists". As a temporary workaround, the table `lexicon_token` has to be reindexed from the Alation Postgres shell on the instance. **FIXED in 2020.4 GA.**
- Tableau Server integrations - Data Source Badging and Workbook Promotion - does not work since 2020.3. **FIXED in 2020.4 GA.**

- Dynamic column profiling for Custom DB results in the loading error when Transient storage mode is enabled for Compose. In LA, it is not recommended to enable these features at the same time. **FIXED in 2020.4 GA**
- Data upload fails for Custom DB in transient mode when Transient storage mode is enabled for Compose. **FIXED in 2020.4 GA**
- The “Use COP” checkbox cannot be selected for the Teradata data source during data source setup. **FIXED in 2020.4 GA**
- For MySQL data sources as Custom DB, during Custom Query-Based MDE, the data types that should be normalized as INT or FLOAT are returned as non-numeric. As a result, the Query-based Column Profiling does not work for such columns. **FIXED in 2020.4 GA.**
- Data Objects Without Stewards and Curation Progress pages do not load data. **FIXED in 2020.4.**
- Oracle synonym extraction query performance may be slow due to sequential scans for filters.
- After changing the flag `ALATION_CACHE_DB_CREDENTIAL_IN_MEMORY` to True in `settings_default.py`, the query tab in Compose fails with the “query does not exist” error.
- After enabling **User Impersonation** for a data source, and attempt to run a query in Compose in the **Transient** mode displays an Authentication error message.
- Previously extracted SSRS reports are deleted when new reports are ingested.
- SSRS attributes of a previously extracted report are deleted after extraction of a new report.
- For Teradata data sources, Table aliases in the Select queries are identified as temp tables.
- Incremental restore fails on a clean instance if the `/tmp` folder and the restore folder under `/data2` have not been manually created before restoring.
- During update from V R7 to 2020.4 on RHEL 7.x, 8.x, and CentOS 8.x, there may be a non-critical error “Failed to clear cacheops”.
- Filters and Joins do not work as expected for Amazon Athena - Presto grammar: queries shows up in Query History but not on the Filter and Join tab.
- Alation Connector Manager cannot be installed on Debian 9 and CentOS 8 with dependencies enabled: cannot auto-install Docker components. For these systems, install Docker from any accessible repo and then install Alation Connector Manager without installing the dependencies. On Debian 9, you can install Docker 18.09.1, although the recommended version is 18.09.7.
- After a deprecated object is deleted, the deprecation propagated to other objects from the deleted object persists.
- When performing full extraction after selective extraction, aggregate attribute pages show an “unexpected error”.
- Automated upgrades from Alation configurable on **Admin Settings > Software Updates** will not work in 2020.4 GA. Follow the regular manual update steps to update Alation.
- The automatic ETL is not triggered by the system after running **Initiate Analytics Database** on the Alation Analytics Settings page after installing or upgrading Alation Analytics V2. The scheduled nightly ETL works as expected.

1.14 Release Notes 2020.3 (General Availability)

1.14.1 RELEASE 2020.3.9

build 5.17.10.138502

- Fixed a problem when sometimes a subset of web server workers on the Alation server hung and did not serve requests. This fix improves throughput and availability of the Alation web service.

1.14.2 RELEASE 2020.3.8

build 5.17.9.130830

- Previously, metadata extraction from Snowflake extracted the view names in upper case when the original view name was in mixed case. After this fix, metadata extraction of views extracts the view names in mixed case.
- Fixed a performance issue with search indexing for file objects. This change is expected to make file system objects indexing 3 to 5 times faster.
- Added support for user impersonation for EMR Presto as Custom DB: users can use impersonation for running queries in Compose.

1.14.3 RELEASE 2020.3.7

build 5.17.8.127669

New features

- For Custom DB, Alation now allows Data Source Admins to enable concurrent queries for Compose. This can be done on the **Custom Settings** tab of the Custom DB data source Settings page with the help of the **Enable Concurrent Queries** checkbox.
- For Custom DB, Alation now allows Data Source Admins to enable **User Impersonation**. This can be used when configuring Kerberos and Keytab authentication for Compose for the data sources that support this setup. User Impersonation can be turned **On** on the **General Settings** tab of the Custom DB data source Settings page.

Bug Fix

- Fixed the issue of Lexicon overflow error by capping the approve/reject count to 7,000. In the Lexicon job, if the approve/reject count exceeds 7,000, the exceeded count will be capped at 7,000.

1.14.4 RELEASE 2020.3.6

build 5.17.7.125879

- For Amazon S3 users, fixed an issue where Alation would not extract the last batch of files when extracting more than 1,000 files. For example, if a filesystem had 1,234 files, 1,000 would be extracted and 234 would be ignored.
- For BI users, fixed an issue where setting the automated extraction job from the UI would not make any changes.
- For Hive users, fixed an issue where query logs generated by Hive 2.x would not be extracted when they were URL encoded. Users facing this issue need to include a configuration file named **aha-conf.xml** in their configuration tarball for the datasource. That configuration file should contain the following lines:

```
<configuration>
  <property>
    <name>alation.qli.check-url-decode</name>
    <value>true</value>
  </property>
</configuration>
```

1.14.5 RELEASE 2020.3.5

build 5.17.6.124318

- Fixed an issue with data migration during upgrade to 2020.3 or higher where the upgrade failed if no authenticated user existed in the system with a primary key of 1. This has been changed to assign ownership of the saved connections to the data source creator.
- Previously, extraction from Tableau 2020.1 skipped some workbooks when permission mirroring was enabled. This issue has now been fixed.
- Previously, even if the **Signup Moderation** feature was enabled, all new users logging into Alation using LDAP would have their user accounts active before the Admins could approve them. After this fix, if Signup Moderation is enabled, admin users can approve/reject user accounts of LDAP users before those users can log in to Alation.
- For Google BigQuery sources, fixed an issue with MDE and Profiling not working after changes to the service account setup. When a user account certificate was uploaded in addition to the service account and then deleted, MDE and Profiling became unavailable. This has now been fixed.
- Fixed an issue where lineage calculation failed in some scenarios involving connections between “temp” columns to actual columns in the graph.

1.14.6 RELEASE 2020.3.3

build 5.17.4.121765

- Fixed an issue with rich text formatting where text formatted as bold in the Edit mode did not appear as bold after the field was saved. Now, emphasizing with bold works as expected.
- Previously, when one database was cataloged in Alation as multiple Data Sources, each of these sources containing a different schema from the database, and one BI object was sourced from database objects in different Data Sources, only one lineage diagram was displayed correctly, while others would display source data objects as TMP objects. This issue has now been fixed.

- Fixed an issue of the Export feature not working for data samples when exported values contained a mix of HTML-escaped and unescaped characters.
- Previously, triggering a search on the Compose page would sometimes result in a 500 error message on the page but no network error, particularly when the search result has been highlighted. This search results error has now been fixed.
- Previously, certain BI and File System objects were incorrectly shown as “Untitled [Object Type]” in Full Search results. This fix resolves the issue. Relevant object titles will now appear for BI and File System objects in search results.

1.14.7 RELEASE 2020.3.2

build 5.17.3.120743

- Added the “Inactive” state for Trust Flags in the Alation UI. If a user does not have permission to edit the Trust Flags for an object, then to this user, they will appear as disabled on the Catalog page of the object.
- Fixed an issue when Oracle synonyms were extracted even if the synonym extraction was disabled. Oracle synonym extraction can be enabled or disabled on the Admin Settings > **Labs** page.
- Fixed an issue with LDAP authentication when users with usernames containing Cyrillic characters could not log in to Alation.
- Improved initial page loading performance by deferring the loading of conversations in the left-hand navigation panel.
- Optimized the data sources API to reduce the loading time of the **Sources** page.
- Added subject indexes to Flags to improve Flag lookup performance for a given catalog object. This is intended to decrease page loading time when Flag propagation is enabled.
- Alation Connector Manager of the Open Connector Framework is no longer dependent upon a particular behavior of “pidof” in order to query the run status of the Docker daemon.

1.14.8 RELEASE 2020.3.1

build 5.17.2.118938

- Fixed the issue of email notifications not being sent to all users @-mentioned in a Conversation if multiple users were @-mentioned in one post.
- Alation now supports Azure Blob Storage extraction for AzureCN access endpoint in the file system.
- Fixes several issues reported in 2020.3 LA (5.17.0) **related to the Snowflake OAuth** functionality and the related default Compose connection URI editing feature:
 - For non-Snowflake data sources or Snowflake data sources without OAuth enabled, data upload would use the service account URI to connect rather than the default Compose connection URI. Also, for OAuth-enabled Snowflake data sources, data upload would not run with the selected and/or authorized Snowflake user if there was already a connection open on the server for the specified URI that had not yet been closed due to being idle. These issues with data upload have been fixed.
 - For non-Snowflake data sources or Snowflake data sources without OAuth enabled, a query form would use the service account URI to connect rather than the default Compose connection URI. Also, for OAuth-enabled Snowflake data sources, query forms would not run with the selected and/or authorized Snowflake user if there was already a connection open on the server for the

specified URI and for that particular query that had not yet been closed due to being idle. These issues with query forms have been fixed.

- For non-Snowflake data sources or Snowflake data sources without OAuth enabled, dynamic profiling would use the service account URI to connect rather than the default Compose connection URI. This issue has been fixed.
- The OAuth client secret was serialized and sent back to the Alation UI. This has been corrected by changing the serializer to treat the client secret field as write-only. After the fix, the client id and secret are not updated on each call.
- Issues with authorization redirect in Desktop Compose have been fixed; it would erroneously display “Authorization terminated unexpectedly” when authorization was successful; similarly, when the browser URL differs from the base URL or hostname of the Alation server (if the browser is using the Alation server’s IP address), the same error message would be displayed. When the browser URL doesn’t specify the correct Alation base URL hostname, “Authorization terminated unexpectedly” will be displayed only if there is an authorization error; with Desktop Compose, it will try to make a connection regardless of whether authorization had succeeded and report whether there are any subsequent connection issues due to expired or missing authorization tokens.
- When OAuth is enabled for a Datasource and the “Test Authorization” button is clicked, the username associated with the authorization is not displayed in the “Authorized as:” section. Have made changes to ensure that the username is updated and displayed appropriately.
- Added auto-trimming extra spaces in Compose for the fixed length char type and the nchar type.
- Fetching of published queries for Schema and Table catalog pages is now done on demand to prevent slowness of the catalog page loading and to reduce the loading time for these pages.
- Added `alation.feature_flags.enabled_flash_plugin` flag to the **alation_conf** utility. Setting it to `False` hides the Adobe Flash-dependent **Copy** and **Download** buttons from QLI Preview actions.
- Previously, Snowflake CREATE VIEW queries with COPY GRANTS keywords did not show up in the **View SQL** area and the Lineage diagram. Now, such queries will be displayed correctly.

1.14.9 RELEASE 2020.3 - General Availability

build 5.17.1.118021

Full-Page Unified Search

Introducing the **New Search** experience with

- a prominent search box on the front center or top center of the page;
- the full-page search results list;
- comprehensive multi-dimensional filters;
- detailed, descriptive result previews.

The New Search unifies two separate search features - the Quick Search (or Left-Hand Navigation) and Advanced Search - into one tool for every search need. It includes the query auto-complete feature; relevancy ranking; ability to browse all objects, objects by category or search within a specific object type; and ability to download the search results list. Trust flags are now displayed in search results next to the object title to direct users to trusted data. Note that the previous search tools are still available to allow for a smooth transition from the old to the new experience.

See [Search](#) for an overview of the New Search.

New Catalog UI

This release reveals the updated Catalog UI, with changes to color scheme, UI element design, typography, and iconography. Most of the UI changes appear purely visual and aesthetic, but there are also a number of subtle changes to the layout on pages of Catalog objects. Team Alation worked to make this change unobtrusive so that your knowledge of the Catalog actions and workflows remains intact. The work on the new UI will continue in the future releases.

New Homepage UI and Configuration page

In 2020.3, the customizable Homepage (first available in V R4 (5.9.x)) has been restyled with:

- an improved layout,
- new custom branding options,
- a dedicated configuration page in the Admin Settings.

The New Homepage is enabled by default when your instance is upgraded to 2020.3. If you have never enabled a customizable homepage, then your instance will inherit the default Homepage version. If you have previously enabled the customizable Homepage, your customizations will transition seamlessly into the new style.

Starting with 2020.3, Server Admins can customize the Homepage from a dedicated configuration tab **Customize Homepage** in the Admin Settings.

Improved Global Navigation

The main toolbar now features an **Apps** menu and a Search bar instead of multiple links to sections of the Catalog. The new **Apps** menu provides access to the various Catalog applications, such as Governance Tracking Dashboard, Glossaries, or Compose. The central component of the top navigation panel is now the Search bar for users to find the data they need from anywhere in the Catalog, anytime.

Snowflake SSO for Compose

Alation now supports SSO using OAuth in Compose for Snowflake data sources. Both Snowflake Built-in OAuth and External OAuth with Okta are supported. Data Source Admins can enable and configure OAuth connection for Compose on the Snowflake data source Settings > **General Settings** page. When a Snowflake data source is OAuth-enabled, Compose users will connect to it by authenticating on the OAuth authorization server instead of connecting directly with their Snowflake user account. When scheduling a query, users have an option to select an OAuth connection to be used for the scheduled query run. To run a query, to schedule it, or to run a query form, users may need to re-authorize with the source if their access token has expired.

Custom DB Enhancements

Added the ability:

- to set the **Limit Query Template** for Profiling that limits the number of rows selected for Profiling to 10,000: Custom Settings > Limit Query Template
- to select the **Catalog Object Definition** to correctly represent the extracted metadata in the Catalog: Set the Catalog Object Definition
- to define a custom subset of data to be sampled with custom Profiling queries for Table objects that can be set up in Per-Object Parameters

MicroStrategy on GBM V2

It is now possible to add MicroStrategy sources to Alation using the GBM V2 framework.

Generic BI Model V2 for MicroStrategy Sources

Backup and Restore V2

Introducing an improved backup and restore functionality that is designed to reduce the time of the backup process, increases the backup and restore reliability, and adds backup integrity validation, as well as backup logging and monitoring capabilities. With Backup and Restore V2, you can choose to create incremental backups saving disk space and considerably decreasing the backup time. This feature must be enabled with a dedicated *alation_action* command and requires a number of configuration flags to be set before usage.

</installconfig/BackupandRestore/index>

Alation Analytics V2

Alation Analytics Version 2 (V2) is the new version of Alation Analytics that gives you the data you need to accelerate and simplify adoption of the Catalog. Alation Analytics V2 addresses the known performance and data accessibility issues of the existing Alation Analytics database (Version 1). It features a completely redesigned data model with the focus on the scope and accessibility of the data available for analysis. Ready-to-use visualizations of popular catalog adoption metrics have been added to the catalog as part of V2. The database ERD and data dictionary are installed as part of the package. V2 also meets the need for scalability as it makes it possible to remotely deploy the Alation Analytics server, with its components installed using Docker containers.

[Alation Analytics V2](#)

API Access Token Management

[Generate Tokens for the Alation API](#)

Access Token Management API

Adding new public APIs for managing API access token generation and authentication to use the Alation APIs. The main difference from the previous releases is that the generated tokens now have an expiration value. This API introduces two types of Tokens:

- **Refresh Token:** A long-lived token that can be used to manage and create API AccessTokens.
- **API Access Token:** A short-lived access token to be used to interact with the Alation APIs. This token must be passed as the value for the 'TOKEN' header in API requests to the Alation server.

The lifetime of tokens is configurable using the **alation_conf**. The specification for this API is available in the OAS 3.0 format at `{AlationInstanceURL}/openapi/api_authentication/` (requires the Swagger UI flag `alation.feature_flags.enable_swagger` to be set to `True` in **alation_conf**).

The Get Token API from previous releases remains fully functional but can be disabled using the **alation_conf**.

Access Token Management UI

Users can generate Refresh and Access tokens in the Alation UI on the User Account Settings > **Authentication** page.

Lineage V2 API Enhancement: Dataflow Objects API

Introducing a new API for managing Dataflow Objects for Lineage V2: POST, DELETE, PATCH, and GET requests for Dataflow objects have now been added for users to be able to build detailed lineage in Alation. You can bulk-create, update, and read the metadata for Dataflow objects, including names, descriptions, and paths. The specification for this API is available in the OAS 3.0 format in Swagger UI at `{AlationInstanceURL}/openapi/lineage/` (requires the Swagger UI flag `alation.feature_flags.enable_swagger` to be set to `True` in **alation_conf**).

Custom Field API

Added the Custom Field API (read-only) that lets you get all custom fields available on the Admin Settings > Catalog Admin > **Customize Catalog** page. The Custom Field API is accessible to users with the Server Admin role only. The specification for this API is available in the OAS 3.0 format in Swagger UI at `{AlationInstanceURL}/openapi/custom_field/` (requires the Swagger UI flag `alation.feature_flags.enable_swagger` to be set to `True` in **alation_conf**).

Custom Field Value API V2

Added the Custom Field Value API V2 that lets you get and update Custom Field values for the DataFlow objects and the BI objects (GBM V2 only). The API requires the Server Admin role. The specification for this API is available in the OAS 3.0 format in Swagger UI at `{AlationInstanceURL}/openapi/custom_field_value/` (requires the Swagger UI flag `alation.feature_flags.enable_swagger` to be set to `True` in **alation_conf**).

User API

Added a new User API (read-only) that allows a Server Admin user to get basic User metadata from the Alation instance. The specification for this API is available in the OAS 3.0 format in Swagger UI at `{AlationInstanceURL}/openapi/user/` (requires the Swagger UI flag `alation.feature_flags.enable_swagger` to be set to `True` in **alation_conf**).

Group API

Added the new Group API (read-only) that allows a Server Admin user to get basic Group metadata from the Alation instance. The specification for this API is available in the OAS 3.0 format in Swagger UI at `{AlationInstanceURL}/openapi/group/` (requires the Swagger UI flag `alation.feature_flags.enable_swagger` to be set to `True` in **alation_conf**).

API Throttling

API throttling is an optional feature for the public APIs that gives admin users the ability to set rate limits for their users. Note that the feature is available to all APIs using the Public API functionality but is only enabled for the users if they are registered in the `api_throttles.yaml` configuration file. Execute the following command in the Alation shell to see the list of available configurations: `alation_conf throttling`.

Scarlet

Scarlet is complementary software developed by Alation as a system to test upgrades to patch releases or new major releases before updates are applied to the production environment. Its other goal is to reduce the feedback cycle for new features and bug fixes by automating the collection of all troubleshooting information. Scarlet automatically creates a copy of the Alation production instance and upgrades it to a newer Alation version allowing the admin user to uncover any update or post-update issues that would potentially affect Production. In 2020.3 Alation releases Scarlet on a case-by-case basis. Please contact your account manager if you would like to learn more about the Scarlet program or would like to try out Scarlet.

[Scarlet Installation and Configuration](#)

1.14.10 Data Source Certifications

Confirmed support for:

- Amazon Athena with the Simba driver
- Amazon DynamoDB with the CData driver
- Azure Cosmos DB with the CData driver
- Databricks with query logs on ADLS Gen 2
- Databricks 6.3 on Azure as Custom DB with Simba driver
- Greenplum 6
- MicroStrategy Cloud 104,11.1 2019 and 2020
- MongoDB Atlas
- Oracle EBS 12.2.5
- Oracle v18 and v19 on RDS

- Snowflake on GCP, AWS, and Azure with the Snowflake JDBC driver version 3.12.2
- PostgreSQL 12 on RDS

1.14.11 Improvements

Sources

- Updated the built-in driver for **Snowflake** data sources to version 3.12.2.
- Updated the built-in driver for **Google BigQuery** data sources to version 1.2.3.1006.
- **Google BigQuery** data sources now support the Profiling for Record and Struct data types. These data types are represented in the Catalog using the DocStore format.
- For **Custom DB** sources, the Profiling job now successfully extracts the array data type.
- Alation now has the ability to ingest query logs from ADLS Gen 1, Gen 2, and Azure BLOB storage for the Azure Databricks data sources.
- **GBM V2 enabled by default:** The GBM V2 feature flag is now TRUE by default. The Tableau Legacy framework is no longer available. Tableau sources can only be added using the GBM V2 BI model.

The Catalog

- **Search index:** Improved performance of search re-indexing for the Column type of objects.
- **Hats removed:** the Hats selection feature was removed. Users will no longer have to select a Hat when they log in. Access to functionality in Alation is controlled by Alation roles assigned to users in the Admin interface.
- **Default Alation role:** Added the ability to configure the default Alation role. Server Admins can now set a default role on the Admin Settings > **Authentication** tab. The default role will be assigned to all new users who join Alation. The default role setting works in all authentication modes: built-in, SAML, or LDAP.
- **Branding:** The Branding tab in the Admin Settings > **Server Admin** section is now enabled by default.

1.14.12 Bug Fixes

- Fixed an issue where bulk-upload of usernames could not complete if the CSV file contained usernames with more than 30 characters. Now, admins can successfully bulk-upload users from CSV. Note that usernames longer than 30 characters are still truncated to 30 characters in the Alation user account.
- Previously, users whose account was inactive (suspended or pending activation) became caught in a redirection loop between their IdP and the Alation Server when they tried to log in to Alation. Now, users with inactive accounts will be redirected to the **Inactive Users** page.
- In V R7 (5.12.x), the temporary files created during SAML authentication were not cleared from the Alation server and accumulated over time. Now, the SAML authentication temp files are auto-deleted as soon as authentication is done.
- Fixed a number of priority issues with the New UI.
- Fixed an issue where the Top Banner would show the **Powered by** text as a result of specific customization settings.
- Improved the scrolling experience on Homepage 2020.3.

- Fixed a number of Compose styling issues that were reported for the 3 Compose UI modes. For the Dusk mode, the following is fixed:

- Current database account name was too dark for the background;
- The Schedule Details title in the scheduling tooltip was too dark and unreadable;

A few more similar UI bugs were fixed for all 3 modes.

- Previously, the license expiration time on the Admin Settings > **License** page was inaccurate in some scenarios. This issue was fixed, and now the license expiration date information in Alation always matches the date in the license file.
- Previously, Unicode characters in user display names caused an internal error in search indexing; conversations would not be indexed and attempts to navigate to them would result in errors. This has now been addressed.
- Fixed a bug where no synonyms were extracted if the user had permissions on some but not all tables with synonyms. Now, the synonyms are extracted if users have permissions on the base table.
- Fixed an issue where the use of the usage reporting feature (Admin Settings > Reporting) resulted in an incomplete usage data package and an error. After the fix, an admin can successfully upload usage metrics from the Alation instance.
- Previously, on the last day before the license expiration, all Alation users would have been locked out of the UI. After this fix, users can continue to work in their Alation Catalog during the last day before license expiration. It is recommended to renew the Alation license in timely fashion.
- Fixed an issue where the first attempt to upload a custom logo to the Customizable Homepage did not succeed.
- Fixed a bug where an attempt to use the User and Group APIs with the feature flag *enable_permissions_middleware_feature* set to *True* results in the 500 error.
- Previously, server admins could not revoke all API Access tokens for a user in bulk. This is now fixed.
- Fixed an issue with the *alation_chmod_permissions* script which did not work after Django upgrade in V R7 and 2020.3 LA.
- Fixed an issue with deleting Dataflow objects: when an Dataflow object was deleted but a user tried to access it by directly typing the URL, Alation UI would show a loading icon indefinitely instead of the 404 error page.
- Fixed an issue where the Lexicon job run fails if the column name has more than 200 characters.
- Miscellaneous other bug fixes.

1.14.13 RELEASE 2020.3 - Limited Availability

build 5.17.0.116556

For new features in 2020.3, see the GA release notes.

Bug Fixes

- Previously, when SAML authentication was used with the option to moderate user sign-up requests, a server admin was unable to reject any sign-up requests that appeared in Alation. After this fix, the admins can reject sign-up requests as expected.
- Fixed a bug where the server URL for Tableau data sources was extracted incorrectly.
- For Alation Analytics, changed the condition in *ObjectFieldLogicalETLJob* to prevent the processing of records when *value_otype* or *value_oid* is null, which previously caused this ETL job to fail with errors and caused performance issues.
- Previously, metadata extraction from Amazon Redshift failed with an NPE due to NULL value types in the distribution node extraction query result. To fix the issue, distribution node metadata extraction was disabled for Redshift as it is not required to be displayed on the Catalog pages.
- MicroStrategy sources no longer require to add the HTTP or HTTPS protocol prefix to the URL when setting up the connection.
- Previously, the start time for active tasks was displayed in the UTC time zone in Admin Settings > **Monitor** > **Active Tasks**. Now, the start time is shown in the local time zone that the user logs in to Alation from.
- Previously, in Compose, when users tried to run the first query for a new data source, clicking **Run** or **Run Full Query** would result in no response until the page was refreshed. This also caused the auto-suggest to not work for this first query. This issue has been fixed, and the Run functionality and auto-suggest will work for all queries as expected.
- Fixed an issue where scheduled extraction from Tableau would run a day after it was scheduled. Now, scheduled extraction will be performed as expected.
- The jobs API response has been updated to correctly reflect the status of a running job for programmatic handling.
- Fixed an issue with metadata extraction from MongoDB data sources where it failed if DBRef objects were present. After the fix, MDE will complete successfully.
- Extraction from a MicroStrategy source could cause exceptions on the MicroStrategy server when the auto-generated Row Count field (metric type = "Unknown") was included into the extracted project. This fix adds validation for metrics of such type, which will correctly extract their metadata into Alation.
- Fixed an issue where the query title disappeared from the Queries tab of a data source. After the fix, query titles and descriptions are displayed as expected in the list of queries.
- Fixed an issue where MDE failed with a *UnicodeDecodeException* error when a Function/SP was ingested from the yaml files that contained specific characters. After the fix, MDE will interpret the symbols correctly and proceed.
- Fixed the ability to create *BIRReport* and *BIRReportColumn* objects with the field *description_at_source* using the GBM V2 API.
- Fixed a bug where users received a 404 error when they clicked on the link for an object with the camel-cased name on the search results page.
- Fixed a UI bug where Custom Field titles were truncated on catalog pages of some objects. Now, the full titles will be displayed as expected.
- Fixed a bug where the Server, Workbook, and Reports links from a MicroStrategy source redirected to an incorrect URL. The fix applies to sources on the GBM V2 framework.

- Previously, when the Viewer role was enabled, the users with the Catalog Admin, Source Admin, Composer, or Steward roles could not discard an auto suggested title. Now, all the users except Viewers can discard auto-suggested titles.
- Fixed a bug on the Customizable Homepage when users with the Viewer role could see the Write Articles and Write SQL cards although the Viewer role was enforced. After the fix, these cards are only visible to users with enough permissions to create articles and queries.
- Fixed a bug where the **Go to Access** dialog would throw a Permission denied error when a server admin tried to access the settings of the Alation Analytics data source. After the fix, Go to Access button opens the Access tab of the Settings page.
- Fixed a bug which allowed logged-in users to make API requests without first having verified their email.
- API response for the NoSQL job API was changed to the JSON format.
- Fixed a bug where scale and precision were not displayed for the data types in the Vertica data sources.

Known Issues in LA

- New Customizable Homepage: the first attempt to upload a custom logo does not succeed. Admin Settings > **Customize Homepage**: when a valid logo file is uploaded, the Alation UI does not display the uploaded file. To troubleshoot, refresh the browser page and re-upload the logo a second time. The second attempt should be successful. **FIXED IN GA**
- An attempt to use the User and Group APIs with the feature flag *enable_permissions_middleware_feature* set to *True* results in the 500 error. **FIXED IN GA**
- Currently Server Admins cannot revoke all API Access tokens for a user at once. **FIXED IN GA**
- *alation_chmod_permissions* script does not work after Django upgrade (in R7 and above). **FIXED IN GA**
- For Databricks on Azure as Custom DB, an attempt to export the Profiling results for a table may end in an error when values contain complex data types (struct, record, and array), and/or mixed Unicode and ASCII characters.
- When an Dataflow object is deleted, but a user tries to access it by directly typing the URL, Alation UI shows a loading icon indefinitely instead of the 404 error page. **FIXED IN GA**
- For Google BigQuery sources, if selective extraction is followed by full extraction, the catalog pages of complex data types display an *Unexpected error* message.
- MDE from Azure ADLS Gen 2 may fail if it takes longer than the access token lifetime. This issue only applies to large systems where MDE takes a long time due to the size of the metadata to be extracted.
- For Oracle 11g using the SSL with Kerberos connection setup, the Service Account validation fails when the built-in driver *oracle.jdbc.OracleDriver.ojdbc7.12.1.0.2-tlsPatch* is used. As a workaround, use the built-in driver *oracle.jdbc.OracleDriver.ojdbc7.12.1.0.2*.
- For Greenplum 6, the CREATE queries run from Compose are stored as empty values in the *query_text* field in the *queries_history* table.
- After a deprecated Dataflow object is deleted from the catalog, the propagated deprecation remains on the downstream objects and cannot be removed.
- When the Alation server is configured to use LDAP authentication, and at the same time the User Signup Moderation Preference is enabled, all new user accounts created are active by default and do go through the Admin moderation as expected.

SUPPORT MATRICES

The support matrix tables in this section provide information on supported database, file system, and BI tool types and versions.

Contact your account manager with additional questions about databases supported by a specific Alation version and possible limitations.

For more information, see the [Alation Support Matrices](#).

END OF SUPPORT

Alation provides technical support for current and previous versions of Alation Software as documented in this [Support Policy](#) page.

This section provides information on end of support for database versions, connectors, drivers, browser versions, operating systems, and features.

3.1 Alation End of Support Matrix

This document provides information about the planned end of support (EOS) by Alation for data sources, connectors, drivers, browsers, operating systems, and features. Alation ends support of the listed versions from the release given in the **End of Support** column.

The EOS versions provided in this document are intended for general planning purposes and do not supersede any of Alation's contractual commitments.

Note: The Alation versions included into the **End of Support** columns reflect the current end of support schedule and may change due to changes in product requirements or release scoping. Most of the older versions of data sources are retired based on the end-of-support or end-of-life schedules announced by the respective companies.

3.1.1 Data Source Versions

Datasources	Versions	Release Ver- sion	End of Support
AWS Databricks	6.4, 6.5, 6.6, 7.0, 7.1	VR2+	2021.2
	7.3 and below	VR5+	2021.3
	8.2 and below	VR5+	2021.4
	8.4 and below	VR5+	2022.1
	10.1	VR5+	2022.3
Azure Databricks	7.2 and below	VR5+	2021.2
	7.3 and below	VR5+	2021.3
	8.2 and below	VR5+	2021.4
	8.4 and below	VR5+	2022.1
	10.1	VR5+	2022.3
Azure HDInsights	3.5	VR5+	2021.3
Azure HDP-Hive2 [Config Based]	3.x	VR2+	2021.3

continues on next page

Table 1 – continued from previous page

Datasources	Versions	Release Version	End of Support
Cassandra	2.1	VR1+	2022.1
CDH	5.14 and below	VR1+	2021.3
	6.0 and below	VR1+	2021.4
	6.1 and below	VR1+	2022.1
EMR Presto	5.3 and below	VR1+	2021.2
GCP Databricks	8.4 and below	2021.3	2022.1
	10.1	2021.3	2022.3
Greenplum	4	VR1+	2022.1
Greenplum OCF Connector	5, 6.3, 6.9	2022.3	2024.1
HDFS - CDH	5.2 to 5.16, 6.0, 6.1	VR1+	2022.1
HDFS - HDP	All Versions	VR1+	2022.1
HDP	2.5 and below	VR1+	2021.3
	3.0 and below	VR1+	2021.4
	All versions	VR1+	2022.1
HDI Insights	3.5	VR5+	2021.3
	3.6	VR5+	2022.4
Hive on HDP	All Versions	VR1+	2022.1
Hive 2 [Config-Based]	CDH 5.14 and below, EMR <5.3 and HDP 2.5.x versions	VR2+	2021.3
Hive 3 on HDP with Tez	All Versions	VR1+	2022.1
Impala [Default Hive]	CDH 5.14 and below	VR1+	2021.3
	CDH 5.8 to 5.16, 6.0, 6.1	VR1+	2022.1
MicroStrategy	10.4	VR1+	2021.2
	11.x	VR1+	2021.3
	2019 Cloud	VR1+	2022.1
MongoDB	3.4 and below	VR4+	2021.2
	3.6 and below	VR4+	2021.3
	4	VR4+	2022.2
MySQL EC2	5.6	VR1+	2021.2
MySQL RDS	5.6	VR1+	2021.2
MySQL OCF Connector EC2	5.7	2021.4	2024.1
MySQL OCF Connector RDS	5.7	2021.4	2024.1
Oracle EC2	11g	VR1+	2021.2
	18C	VR1+	2021.4
	12.1C	VR1+	2022.3
Oracle RDS	11g	VR1+	2021.2
	18	VR1+	2021.3
	12.2c	VR1+	2022.2
	12.1C	VR1+	2022.3
Oracle OCF Connector (E-Business Suite)	12.2.5 on 11g	2023.1	2023.3
PostgreSQL EC2	9.5	VR1+	2021.2
	9.6 and below	VR1+	2022.1
	10	VR1+	2023.1
PostgreSQL Enterprise	9.6 and below	VR1+	2022.1
	10	VR1+	2023.1
PostgreSQL RDS	9.5	VR1+	2021.2
	10	VR1+	2023.1

continues on next page

Table 1 – continued from previous page

Datasources	Versions	Release Version	End of Support
PostgreSQL OCF Connector EC2	11.4, 11.5	2021.4	2024.1
PostgreSQL OCF Connector RDS	11.4, 11.5	2021.4	2024.1
PostgreSQL OCF Connector Enterprise	11	2022.3	2024.1
SQL Server	2012	VR1+	2022.3
SSAS	2012	2022.2	2022.3
SSRS (with Sharepoint)	2012	2020.3	2022.3
SSRS (without Sharepoint)	2012	VR2+	2022.3
Straburst Enterprise (Trino)	345-e LTS and below	2021.1	2021.4
	350-e LTS	2021.2	2022.2
	360-e LTS	2021.2	2022.3
	380-e LTS	2021.2	2024.1
Straburst Enterprise (Trino) OCF Connector	380-e LTS	2023.1	2024.1
Tableau	10.5.x, 2018.x and below	VR5+	2021.2
	2019.1 and below	VR5+	2021.4
	2020.3 and below	VR5+	2022.1
Teradata	15	VR1+	2022.1
Teradata Aster	4.6	VR1+	2022.1
Tez Engine [Config-Based]	CDH 5.14 and below, EMR <5.3 and HDP 2.5.x versions	VR1+	2021.3

3.1.2 Deprecated Data Sources

The data sources listed below are no longer supported by Alation.

Connector	Versions	End Of Support
Cassandra (Native)	2.1	EOL in 2022.1
Cognos	All versions	not supported
MicroStrategy GBM	V1	2021.2
Netezza	7.2.1.0	EOL
SFDC Einstein Analytics	All versions	not supported
Tableau Legacy	All versions	2020.3
Teradata Aster	4.6	EOL in 2022.1
Vertica (Native)	7.2.3	EOL in 2022.3

3.1.3 Deprecated Native Connectors

The data sources listed below are no longer supported with the native (built-in) connectors. Use the corresponding OCF connectors instead.

Connector	Status	End Of Support
Azure SQL Data Warehouse/ Azure Synapse Analytics	Deprecated	15-Dec-2022

continues on next page

Table 2 – continued from previous page

Connector	Status	End Of Support
MySQL EC2	Deprecated	15-Dec-2022
MySQL RDS	Deprecated	15-Dec-2022
PostgreSQL (enterprise)	Deprecated	15-Dec-2022
PostgreSQL EC2	Deprecated	15-Dec-2022
PostgreSQL RDS	Deprecated	15-Dec-2022
SQL Server	Deprecated	15-Dec-2022
SQL Server RDS	Deprecated	15-Dec-2022
DB2	Deprecated	01-Dec-2023
SAP HANA	Deprecated	01-Dec-2023
SAS Base (Data files in SAS) server	Deprecated	01-Dec-2023
Sybase IQ	Deprecated	01-Dec-2023
Sybase ASE	Deprecated	01-Dec-2023
Teradata	Deprecated	01-Dec-2023
Google BigQuery	Deprecated	01-Dec-2023
AWS Databricks	Deprecated	01-Dec-2023
Greenplum	Deprecated	01-Dec-2023
MemSQL	Deprecated	01-Dec-2023
Snowflake (AWS, Azure, GCP)	Deprecated	01-Dec-2023
EMR Presto on Hive	Deprecated	01-Dec-2023
Hive on CDH	Deprecated	01-Dec-2023
Hive with AWS Glue as Metastore	Deprecated	01-Dec-2023
Hive 2 (Config Based)	Deprecated	01-Dec-2023
Hive 3 on EMR with Tez	Deprecated	01-Dec-2023
Hive 3 on CDP with Tez	Deprecated	01-Dec-2023
Impala	Deprecated	01-Dec-2023
Hive 3 on Azure HDInsight	Deprecated	01-Dec-2023
AWS Glue	Deprecated	01-Dec-2023
Amazon Redshift	Deprecated	01-Dec-2023
Amazon Redshift Spectrum	Deprecated	01-Dec-2023
Oracle	Deprecated	01-Dec-2023
Oracle E-Business Suite	Deprecated	01-Dec-2023

3.1.4 Features

Features	End Of Support
Alation Analytics V1	2021.3
Lineage V1	2021.3
Backup V1	Tentative 2021.4
Advanced Search	2021.4
Query Search	2021.4

3.1.5 Drivers

Drivers	Versions	End of Support
jTDS Driver for SQLServer	All versions and editions	2021.3
Progress driver support for Azure Databricks	All versions	2020.3

3.1.6 Browsers

Browsers	Versions	End of Support
Internet Explorer	IE 11	2021.3
Internet Explorer	IE 10	2020.3
Safari	11.1.2 on OS X 10.11 El Capitan and 12.1.2 on macOS 10.12 Sierra	2021.1
Windows 7 Client for Browsers	All Browsers and Versions	2021.2

3.1.7 Operating Systems

Operating Systems	Versions	End of Support
CentOS	6.x	2020.4
	8.x	2022.1
	7.x	2024.1
Fedora	31, 32	2021.2
	35	2024.1
RedHat	6.x	2020.4
Ubuntu	20.x	2024.1
Debian	9.x	2023.3
	10.x	2024.1
Oracle Linux on Red Hat-compatible Kernel	7.x	2024.1

SNOWFLAKE COLUMN-LEVEL LINEAGE PARSER ADD-ON

Snowflake column-level lineage add-on is a parser enhancement that allows the Alation Snowflake connector to parse and extract column-level lineage information from Snowflake query logs.

Refer to Support Matrices in this section for information about query grammar supported by the add-on.

4.1 Snowflake Column-Level Lineage 2021.4

Refer to the table below for details about query grammar supported by the Snowflake column-level lineage parser add-on.

Grammar	Can Be Parsed by QLI	Parent Level Lineage (Tables and Files)	Column-Level Lineage Parsing
CREATE TABLE AS SELECT (Basic)	✓	✓	✓
CREATE VIEW AS SELECT (Basic)	✓	✓	✓
CREATE MATERIALIZED VIEW (Basic)	✓	✓	✓
INSERT	✓	✓	✓
MERGE	✓	✓	✓
UPDATE	✓	✓	✓
CTAS with JOIN and USING keywords	✓	✓	✓
CTAS with SET operations	✓	✓	✓
CTAS with CTE	✓	✓	✓
CTAS with recursive CTE	✓	✓	✓
CTAS with CASE and WHEN	✓	✓	✓
CTAS with subquery in result column	✓	✓	✓
CREATE AS CLONE	✓	✓	✓
CREATE AS LIKE	✓	✓	✓
COPY INTO	✓	✓	x*
CREATE TABLE	✓	✓	<i>not applicable</i>
CREATE EXTERNAL TABLE	x*	x**	<i>not applicable</i>

continues on next page

Table 1 – continued from previous page

Grammar	Can Be Parsed by QLI	Parent Level Lineage (Tables and Files)	Column-Level Lineage Parsing
CREATE FUNCTION	x*	x	<i>not applicable</i>
CREATE SCHEMA	✓	✓	<i>not applicable</i>
CREATE STAGE	x*	x	<i>not applicable</i>
CREATE WAREHOUSE	x	<i>not applicable</i>	<i>not applicable</i>
CREATE PROCEDURE	x*	x	<i>not applicable</i>
DELETE	✓	✓	<i>not applicable</i>
DROP	✓	✓	<i>not applicable</i>
DROP EXTERNAL TABLE	✓	✓	<i>not applicable</i>
DROP DATABASE	✓	<i>not applicable</i>	<i>not applicable</i>
DROP PIPE	x*	x	<i>not applicable</i>
DROP SCHEMA	✓	✓	<i>not applicable</i>
DROP STAGE	x*	x	<i>not applicable</i>
DROP STREAM	✓	x*	<i>not applicable</i>
DROP TABLE	✓	✓	<i>not applicable</i>
DROP VIEW	✓	✓	<i>not applicable</i>
TRUNCATE TABLE	✓	✓	<i>not applicable</i>
ALTER TABLE	✓	<i>not applicable</i>	<i>not applicable</i>
ALTER VIEW	✓	<i>not applicable</i>	<i>not applicable</i>
ALTER WAREHOUSE	✓	<i>not applicable</i>	<i>not applicable</i>
SELECT	✓	<i>not applicable</i>	<i>not applicable</i>

* - Under consideration as a future improvement

** - Alation supports extracting external tables. Lineage will be added in the future.

IMPORTANT NOTICES AND SECURITY ALERTS

Security Advisories are available on the Alation Community. Viewing this information requires a Community login: [Security Advisories](#).

Alation has issued the following advisories:

- June 17, 2023 - Oracle Java SE Multiple Vulnerabilities
- November 3, 2022 - Encryption Key Reuse Security Advisory
- October 28, 2022 - Update to Log4j 1.x Security Advisory
- August 12, 2022 - Log4j 1.x Security Advisory
- April 26, 2022 - Hive Log4j 1.x Advisory
- April 8, 2022 - Spring4Shell Advisory
- March 5, 2022 - JDBC Driver Security Advisory
- January 24, 2022 - Log4j 2 Update and Databricks Connector Security Advisory
- December 22, 2021 - Log4j 2 CVE-2021-45105 Security Advisory
- December 17, 2021 - Update to December 15 Log4j2 CVE-2021-45046 Security Advisory
- December 15, 2021 - Log4j 2 CVE-2021-45046 Security Advisory
- December 10, 2021 - Log4j 2 Security Advisory
- Migration to New Alation Support Help Desk on July 3, 2020
- November 11, 2019 - Java JDK Security Advisory
- September 25, 2019 - Update to April 17 Important Security Alert
- June 4, 2019 - Patch Alert Security Bulletin - Data Source Credentials in taskserver.log Files
- April 23, 2019 - Important Security Update - Potential Transmission of Access Credentials to Alation
- April 17, 2019 - Important Security Alert - PII on Customer Servers
- April 16, 2019 - Important Update to April 15, 2019 - Patch Alert Security Bulletin
- April 15, 2019 - Patch Alert Security Bulletin - Potential Internal LDAP-Related Access to Alation Usernames and Passwords
- January 5, 2018 - Meltdown and Spectre Security Flaws
- Old Backups Are Not Cleaned Up in V R3 (Builds 5.5.0 Through 5.6.5)